

Tutorial

LNCS 2900

Michel Bidoit  
Peter D. Mosses

# CASL User Manual

Introduction to Using the  
Common Algebraic Specification Language



Springer



# Lecture Notes in Computer Science

Edited by G. Goos, J. Hartmanis, and J. van Leeuwen

2900

**Springer**

*Berlin*

*Heidelberg*

*New York*

*Hong Kong*

*London*

*Milan*

*Paris*

*Tokyo*

Michel Bidoit Peter D. Mosses

# CASL

## User Manual

Introduction to Using the  
Common Algebraic Specification Language

With chapters by

Till Mossakowski, Donald Sannella,  
and Andrzej Tarlecki



Springer

## Series Editors

Gerhard Goos, Karlsruhe University, Germany  
Juris Hartmanis, Cornell University, NY, USA  
Jan van Leeuwen, Utrecht University, The Netherlands

## Authors

Michel Bidoit  
Laboratoire Spécification et Vérification, CNRS UMR 8643  
École Normale Supérieure de Cachan  
61, Avenue du Président Wilson, 94235 Cachan Cedex, France  
E-mail: bidoit@lsv.ens-cachan.fr

Peter D. Mosses  
University of Aarhus, BRICS and Department of Computer Science  
Aabogade 34, 8200 Aarhus N, Denmark  
E-mail: pdmosses@brics.dk

## Cataloging-in-Publication Data applied for

A catalog record for this book is available from the Library of Congress.

Bibliographic information published by Die Deutsche Bibliothek  
Die Deutsche Bibliothek lists this publication in the Deutsche Nationalbibliografie;  
detailed bibliographic data is available in the Internet at <<http://dnb.ddb.de>>.

CR Subject Classification (1998): D.2.1, D.3.1, D.2, D.3, F.3

ISSN 0302-9743

ISBN 3-540-20766-X Springer-Verlag Berlin Heidelberg New York

This work is subject to copyright. All rights are reserved, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, re-use of illustrations, recitation, broadcasting, reproduction on microfilms or in any other way, and storage in data banks. Duplication of this publication or parts thereof is permitted only under the provisions of the German Copyright Law of September 9, 1965, in its current version, and permission for use must always be obtained from Springer-Verlag. Violations are liable for prosecution under the German Copyright Law.

Springer-Verlag is a part of Springer Science+Business Media  
[springeronline.com](http://springeronline.com)

© 2004 IFIP International Federation for Information Processing, Hofstrasse 3, A-2361 Laxenburg, Austria  
Printed in Germany

Typesetting: Camera-ready by author, data conversion by Olgun Computergrafik  
Printed on acid-free paper SPIN: 10931011 06/3142 5 4 3 2 1 0

---

## Preface

CASL, the *Common Algebraic Specification Language*, has been designed by CoFI, the *Common Framework Initiative* for algebraic specification and development. CASL is an expressive language for specifying requirements and design for conventional software. It is algebraic in the sense that models of CASL specifications are algebras; the axioms can be arbitrary first-order formulas.

*This User Manual illustrates and discusses how to write CASL specifications.*

CASL is a major new algebraic specification language. It has been carefully designed by a large group of experts as a general-purpose language for practical use in software development – in particular, for specifying both requirements and design. CASL includes carefully-selected features from many previous specification languages, as well as some novel features that allow algebraic specifications to be written much more concisely and perspicuously than hitherto. It may ultimately replace most of the previous languages, and provide a common basis for future research and development.

CASL has already attracted widespread interest within the algebraic specification community, and is generally regarded as a de facto standard. Various sublanguages of CASL are available – primarily for use in connection with existing tools that were developed in connection with previous languages. Extensions of CASL provide languages oriented toward development of particular kinds of software (reactive, concurrent, etc.).

Major libraries of validated CASL specifications are freely available on the Internet, and the specifications can be reused simply by referring to their names. Tools are provided to support practical use of CASL: checking the correctness of specifications, proving facts about them, and managing the formal software development process.

The companion *CASL Reference Manual* [20] provides full details of the CASL design, including its formal semantics.

After briefly reviewing the background of CoFI and CASL, and the underlying concepts of algebraic specification languages, this book introduces the potential user to the features of CASL mainly by means of illustrative examples. It presents and discusses the typical ways in which the language concepts and constructs are expected to be used in the course of building system specifications. Thus, the presentation focuses on *what* the constructs and concepts of CASL are *for*, and *how* they should (and should not) be used. These points are made as clear as possible by referring to simple examples, and by discussing both the general ideas and some details of CASL specifications.

Further chapters introduce the reader to the CASL Reference Manual, to some of the currently available CASL support tools, and to a couple of the CASL libraries of basic datatypes. A substantial case study of the practical use of CASL in an industrially-relevant context completes the material. Appendices provide a quick reference of CASL constructs, a list of the main points to bear in mind when using CASL, and the original informal requirements for the case study.

## Structure

### Part I: Background

Chapter 1 describes the origins of CASL: how CoFI was formed in response to the proliferation of algebraic specification languages in the preceding two decades, and the aims and scope that were formulated for this international initiative.

For the benefit of readers not already familiar with other algebraic specification languages, Chap. 2 reviews the main concepts of algebraic specification, explaining standard terminology regarding specification language constructs and models (i.e., algebras).

### Part II: Writing CASL Specifications

Chapter 3 shows how some familiar datatypes involving total functions are specified in CASL, essentially as in many other algebraic specification languages. Loose, generated, and free specifications are discussed in turn, with illustrative examples and advice on the use of different specification styles.

Partial functions arise naturally. Chapter 4 explains how CASL supports specification of partial functions, drawing attention to where special care is needed compared to specifications involving only total functions.

Subsorts and supersorts are often useful in CASL specifications. Chapter 5 illustrates how they can be declared and defined, and that they can sometimes be used to avoid the need for partial functions.

The examples given so far make use of named and structured specifications in a simple and natural way. Chapter 6 takes a much closer look at the constructs CASL provides for structuring specifications, explaining how large and complex specifications are easily built out of simpler ones by means of a small number of specification-building operations.

Chapter 7 shows how making a specification generic (when appropriate) improves its reusability, allowing it to be instantiated with different arguments; compound identifiers avoid the need for explicit renaming when combining the results of different instantiations. It also introduces the constructs for expressing so-called views between specifications.

While specification-building operations are useful to structure the text of large specifications, architectural specifications are meant for imposing structure on implementations. Chapter 8 discusses and illustrates the role of architectural specifications, and shows how to express them in CASL.

Chapter 9 explains and illustrates how libraries of named specifications can be formed, and made available over the Internet, to encourage widespread reuse and evolution of specifications. Version control is of crucial importance here.

### **Part III: Carrying On**

Chapter 10 gives a detailed overview of the foundations of CASL, which are established in the accompanying CASL Reference Manual.

Tool support is vital for efficient use of formal specifications in connection with practical software design and development. Chapter 11 presents the main tools that have been implemented so far; several of them allow use of CASL specifications in connection with tools that were originally developed for other specification languages, showing how CASL provides tool interoperability.

Chapter 12 introduces a few of the many specifications that are available in the CASL libraries of basic datatypes.

Finally, Chap. 13 gives a realistic case-study of the use of CASL in practice, in connection with the design of software for a Steam-Boiler Control System. This particular example is one of the standard bench-marks for comparing specification frameworks [1].

### **Appendices and Indexes**

This volume is completed by three appendices: App. A provides a compact overview of all CASL constructs, for quick reference; App. B lists all the main points to bear in mind when using CASL; and App. C reproduces the informal requirements specification for the case study.

The names of all the specifications given in this book are listed at the back, together with an index of concepts and a list of references to the literature. (A comprehensive annotated bibliography of publications involving CASL is provided in the Reference Manual.)



An accompanying CD-ROM contains the source files for all the illustrative specifications, and a copy of the libraries of specifications of basic datatypes.

## Organization

*All the main points are highlighted like this.*

The material in this book is organized in a tutorial fashion. Each main point is usually accompanied by an illustrative example of a complete CASL specification; the names of these specifications are listed (both in order of presentation and alphabetically) at the end of the book. Moreover, the points themselves are repeated (in order of presentation) in App. B.

Readers who are familiar with previous algebraic specification languages, and especially those who have been following or participating in the design and development of CASL, may prefer to skip lightly through Chaps. 1 and 2. Chapter 3, however, is mandatory, since it is there that many CASL features needed to understand the subsequent chapters are introduced.

In contrast, Chaps. 4 and 5 can be skipped at first reading if the reader is not so much interested in partial functions, resp. subsorting (with the proviso though that there are some references to the examples given in these chapters from later chapters).

Chapters 6 and 7 present mainstream material, and until one feels comfortable with all the main points and examples, it is advisable to wait with proceeding to Chaps. 8 and 9.

Chapter 10 is primarily for those who will want to follow up on this book with a more detailed study of CASL, based on the Reference Manual. Part of Chap. 11 assumes familiarity with concepts introduced in Chap. 7. In Chap. 12, Sect. 12.1 assumes Chaps. 4 and 5, whereas Sect. 12.2 assumes also Chaps. 6 and 7. Finally, most of Chap. 13 can be studied after Chap. 7, but Sect. 13.10 requires Chap. 8.

*Acknowledgement.* Chapter 10 was written by Donald Sannella and Andrzej Tarlecki. Chapter 11 was written by Till Mossakowski, with contributions from Mark van den Brand and Markus Roggenbach. Till Mossakowski also provided Chap. 12, based on libraries of CASL specifications developed at the Bremen Institute for Secure Systems in joint work with Markus Roggenbach and Lutz Schröder; and he checked the well-formedness of all the specifications in this book, using CATS, the CASL Tool Set.

Public drafts of this book were released in July and October 2003. The many insightful comments from CoFI participants and other readers were very helpful to the authors during the preparation of the final version. Detailed comments on all or part of the public drafts were received from Jørgen Iversen, Christian Maeder,

Guillem Marpons, Narciso Martí-Oliet, Till Mossakowski, Don Sannella, Giuseppe Scollo, Andrzej Tarlecki, Frédéric Voisin, and Alexandre Zamulin. Responsibility for any mistakes in the final version belongs, of course, to the authors.

Michel Bidoit gratefully acknowledges support from LSV,<sup>1</sup> CNRS, and École normale supérieure de Cachan. Peter Mosses gratefully acknowledges support from BRICS<sup>2</sup> and the Department of Computer Science, University of Aarhus. Much of the material on which this book is based was developed in connection with activities of CoFI-WG (ESPRIT Working Group 29432) and IFIP WG 1.3 (Foundations of System Specification).

Finally, special thanks to Springer, and in particular to Alfred Hofmann as Executive Editor, for their willingness to publish this book, and for helpful advice concerning its preparation.

Michel Bidoit and Peter D. Mosses

October, 2003

---

<sup>1</sup> Laboratoire Spécification et Vérification ([www.lsv.ens-cachan.fr](http://www.lsv.ens-cachan.fr)).

<sup>2</sup> Basic Research in Computer Science ([www.brics.dk](http://www.brics.dk)), funded by the Danish National Research Foundation.

---

# Contents

---

## Part I Background

---

<b>1</b>	<b>Introduction</b> .....	3
1.1	CoFI .....	3
1.2	CASL .....	7
<b>2</b>	<b>Underlying Concepts</b> .....	11
2.1	Basic Specifications .....	11
2.2	Structured Specifications .....	15
2.3	Architectural Specifications .....	19
2.4	Libraries of Specifications .....	20

---

## Part II CASL Specifications

---

<b>3</b>	<b>Getting Started</b> .....	23
3.1	Loose Specifications .....	24
3.2	Generated Specifications .....	33
3.3	Free Specifications .....	36
<b>4</b>	<b>Partial Functions</b> .....	47
4.1	Declaring Partial Functions .....	47
4.2	Specifying Domains of Definition .....	50
4.3	Partial Selectors and Constructors .....	54
4.4	Existential Equality .....	55
<b>5</b>	<b>Subsorting</b> .....	57
5.1	Subsort Declarations and Definitions .....	57
5.2	Subsorts and Overloading .....	61
5.3	Subsorts and Partiality .....	62

<b>6</b>	<b>Structuring Specifications</b>	67
6.1	Union and Extension	67
6.2	Renaming	69
6.3	Hiding	71
6.4	Local Specifications	73
6.5	Named Specifications	75
<b>7</b>	<b>Generic Specifications</b>	77
7.1	Parameters and Instantiation	78
7.2	Compound Symbols	85
7.3	Generic Specifications with Imports	88
7.4	Views	90
<b>8</b>	<b>Specifying the Architecture of Implementations</b>	93
8.1	Architectural Specifications	95
8.2	Generic Components	100
8.3	Writing Meaningful Architectural Specifications	106
<b>9</b>	<b>Libraries</b>	111
9.1	Local Libraries	112
9.2	Distributed Libraries	116
9.3	Version Control	120

---

## Part III Carrying On

---

<b>10</b>	<b>Foundations</b>	125
<b>11</b>	<b>Tools</b>	131
11.1	The Heterogeneous Tool Set (HETS)	132
11.2	HOL-CASL	138
11.3	ASF+SDF Parser and Syntax-Directed Editor	139
11.4	Other Tools	140
<b>12</b>	<b>Basic Libraries</b>	143
12.1	Library BASIC/NUMBERS	144
12.2	Library BASIC/STRUCTURED DATATYPES	147
<b>13</b>	<b>Case Study: The Steam-Boiler Control System</b>	155
13.1	Introduction	156
13.2	Getting Started	157
13.3	Carrying On	161
13.4	Specifying the Mode of Operation	163
13.5	Specifying the Detection of Equipment Failures	167
13.6	Predicting the Behavior of the Steam-Boiler	176
13.7	Specifying the Messages to Send	182

13.8	The Steam-Boiler Control System Specification . . . . .	183
13.9	Validation of the CASL Requirements Specification . . . . .	184
13.10	Designing the Architecture . . . . .	186

---

## Appendices

---

<b>A</b>	<b>CASL Quick Reference . . . . .</b>	<b>193</b>
A.1	Basic Specifications . . . . .	194
A.2	Structured Specifications . . . . .	199
A.3	Architectural Specifications . . . . .	200
A.4	Libraries . . . . .	201
<b>B</b>	<b>Points to Bear in Mind . . . . .</b>	<b>203</b>
B.1	Introduction . . . . .	203
B.2	Underlying Concepts . . . . .	203
B.3	Getting Started . . . . .	204
B.4	Partial Functions . . . . .	205
B.5	Subsorting . . . . .	206
B.6	Structuring Specifications . . . . .	206
B.7	Generic Specifications . . . . .	207
B.8	Specifying the Architecture of Implementations . . . . .	207
B.9	Libraries . . . . .	208
B.10	Foundations . . . . .	209
B.11	Tools . . . . .	209
B.12	Basic Libraries . . . . .	210
<b>C</b>	<b>The Steam-Boiler Control Specification Problem . . . . .</b>	<b>211</b>
C.1	Introduction . . . . .	211
C.2	Physical Environment . . . . .	212
C.3	The Overall Operation of the Program . . . . .	214
C.4	Operation Modes of the Program . . . . .	214
C.5	Messages Sent by the Program . . . . .	216
C.6	Messages Received by the Program . . . . .	217
C.7	Detection of Equipment Failures . . . . .	219
	<b>References . . . . .</b>	<b>221</b>
	<b>List of Named Specifications . . . . .</b>	<b>225</b>
	<b>Index of Library and Specification Names . . . . .</b>	<b>231</b>
	<b>Concept Index . . . . .</b>	<b>235</b>

## Part I

---

### Background

# Introduction

This chapter first explains the background and aims of CoFI, the Common Framework Initiative for algebraic specification and development of software. It then gives an overview of the main features of CASL, the Common Algebraic Specification Language.

## 1.1 CoFI

In 1995, an open collaborative effort was initiated: to design a common framework for algebraic specification and development of software. It is referred to as *The Common Framework Initiative*, CoFI.<sup>1</sup>

*There was an urgent need for a common framework.*

The rationale behind this initiative was that the lack of such a common framework was a major hindrance for the dissemination and application of algebraic specification techniques. In particular, there was a proliferation of languages – some differing in only quite minor ways from each other. The major languages included ACT ONE/ACT TWO [19], ASF [6], ASL [36], CLEAR [15], EXTENDED ML [26], LARCH [25], OBJ3 [24], PLUSS [9], and SPECTRUM [14]. This abundance of languages was an obstacle for the adoption of algebraic methods for use in industrial contexts, making it difficult to exploit standard examples, case studies and training material. A common framework, with widespread support at least throughout the research community, was urgently needed.

---

<sup>1</sup> CoFI is pronounced like ‘coffee’.

The aim of CoFI was to base the common framework as much as possible on a critical selection of features that had already been explored in previous research and applications (see the IFIP State-of-the-Art Report on *Algebraic Foundations of Systems Specification* [3] for the background). The collective experience and expertise of the CoFI participants provided a unique opportunity to achieve this aim within a reasonably short time-span.

The various groups working on algebraic specification frameworks had already had ample opportunity to develop their own particular variations on the theme of algebraic specification [16], yet no clear ‘winner’ had emerged (although there were several strong contenders).

*CoFI aims at establishing a wide consensus.*

The aim of CoFI was to design a framework incorporating just those features for which there would be a wide consensus regarding their appropriateness. This framework should be able to subsume many of the existing frameworks, and be seen as an attractive common basis for future research and development – with high potential for strong collaboration between the various groups.

The initial overall aims of CoFI were formulated as follows:

- A common framework for algebraic specification and software development is to be designed, developed, and disseminated.
- The production of the common framework is to be a collaborative effort, involving a large number of experts (30–50) from many different groups (20–30) working on algebraic specification.
- In the short term, the common framework is to become accepted as an appropriate basis for a significant proportion of the research and development in algebraic specification.
- Specifications in the common framework are to have a uniform, user-friendly syntax and straightforward semantics.
- The common framework is to be able to replace many existing algebraic specification frameworks.
- The common framework is to be supported by a concise reference manual, user’s guide, libraries of specifications, tools, and educational materials.
- In the longer term, the common framework is to be made attractive for use in industrial contexts.
- The common framework is to be available free of charge, both to academic institutions and to industrial companies. It is to be protected against appropriation.



*The focus of CoFI is on algebraic techniques.*

The functionality of the common framework is to allow and be useful for:

- algebraic specification of the functional requirements of software systems, for some significant class of software systems;
- formal development of design specifications from requirements specifications, using some particular methods;
- documenting the relation between informal statements of requirements and formal specifications;
- verification of correctness of development steps from (formal) requirements to design specifications;
- documenting the relation between design specifications and implementations in software;
- exploration of the (logical) consequences of specifications: e.g., rewriting, theorem-proving, prototyping;
- reuse of parts of specifications;
- adjustment of specifications and developments to changes in requirements;
- providing a library of useful specification modules; and
- providing a workbench of tools supporting the above.

*CoFI has already achieved its main aims.*

The first major achievement of CoFI was the completion of the design of CASL, the Common Algebraic Specification Language. The CASL design effort started in September 1995. An initial design was proposed in May 1997 (with a language summary, abstract syntax, and formal semantics, but no concrete syntax) and tentatively approved by IFIP WG1.3. The report of the IFIP referees on the initial CASL design proposal suggested reconsideration of several points in the language design, and requested some improvements to the documents describing the design. Apart from a few details, the design was finalized in April 1998, and CASL version 1.0 was released in October 1998. IFIP WG 1.3 was asked to review the final design of CASL version 1.0.1 in May 2000, and subsequently approved that design in April 2001. The current version (1.0.2) was adopted in October 2003; it incorporates adjustments to some minor details of the concrete syntax and the semantics. No further revisions of the CASL design are anticipated.

The *CASL Reference Manual* [20], published as a companion volume to the present book, includes a detailed yet concise (60 pages) summary of the CASL design; the rest of it is concerned mainly with the formal syntax and semantics of CASL, and with the libraries of basic datatype specifications. An introduction to the Reference Manual is given in Chap. 10 of this book.

In parallel with the design of CASL, CoFI has developed tool support for the use of CASL (see Chap. 11), and substantial libraries of CASL specifications (see Chap. 12).

Despite the previous lack of a CASL User Manual, there is already much evidence that CASL is now accepted as an appropriate basis for research and development in algebraic specification. Reference [33] gives an overview of what was achieved in the period 1998–2001, and the annotated bibliography in the CASL Reference Manual lists a significant number of further publications that involve CASL. At the time of writing, it remains to be seen whether significant industrial take-up will follow.

*CoFI is an open, voluntary initiative.*

CoFI was started by COMPASS (ESPRIT Basic Research WG 3264/6112, 1989–96), in cooperation with IFIP WG 1.3 (Foundations of System Specification, founded 1992), on the basis of proposals made during their meetings in 1994 (Santa Margherita Ligure, Italy) and 1995 (Oslo, Norway); participation in CoFI was, however, never confined to members of those working groups. The active participants have included some 30 leading researchers in algebraic specification, with representatives from almost all the European groups working in this area. (Ideally, representatives from non-European groups would have been involved too, but logistic difficulties prevented this.)

Originally, CoFI had separate task groups concerning language design, semantics, tools, methodology, and reactive systems. There was a substantial amount of interaction between the task groups, which was facilitated by many of the CoFI participants being involved in more than one task group. The overall coordination of these task groups was managed by Peter Mosses from the start of CoFI in September 1995 until August 1998, and subsequently by Don Sannella. In 2003, the CoFI task groups were replaced by a looser coordination mechanism, with a steering committee chaired by Don Sannella.

*CoFI has received funding as an ESPRIT Working Group, and is sponsored by IFIP WG 1.3.*

The European Commission provided funding for the European component of CoFI as ESPRIT Working Group 29432 from 1998 to 2001 [33]. The partners were the coordinating sites of the various CoFI task groups (University of Aarhus, University of Bremen, École normale supérieure de Cachan, University of Genova, INRIA Lorraine, Warsaw University) with the University of Edinburgh as overall coordinator. The goals of the working group were to coordinate the completion of and disseminate the Common Framework, to

demonstrate its practical applicability in industrial contexts, and to establish the infrastructure needed for future European collaborative research in algebraic techniques. Apart from this period of funding, and support for meetings from the COMPASS Working Group until its termination in 1996, CoFI has relied entirely on unfunded efforts by its participants. Participation in the frequent working meetings was often supported by generous subsidies from the local organizers.

IFIP WG 1.3 sponsors CoFI by reviewing proposals for changes to the design of CASL, and proposals for extensions of CASL. Moreover, a considerable number of members of IFIP WG 1.3 have been (and, at the time of writing, still are) active participants of CoFI.

*New participants are welcome!*

CoFI is an open collaboration, and new participants are always welcome. Current information about CoFI activities is available at the main CoFI web site: <http://www.cofi.info>. The low-volume mailing list [cofi@cofi.info](mailto:cofi@cofi.info) is reserved for CoFI announcements, and discussions about CoFI activities generally take place on the mailing list [cofi-discuss@cofi.info](mailto:cofi-discuss@cofi.info); see the CoFI web site for how to subscribe, and for access to the archives.

## 1.2 CASL

*CASL has been designed as a general-purpose algebraic specification language, subsuming many existing languages.*

The primary specification language developed by CoFI is called CASL: the Common Algebraic Specification Language. Its main features are:

- The design of CASL is based on a critical selection of the concepts and constructs found in existing algebraic specification frameworks.
- CASL is an expressive specification language with simple semantics and good pragmatics.
- CASL is appropriate for specifying requirements and design of conventional software packages.
- CASL is at the heart of a coherent family of languages that are obtained as sublanguages or extensions of CASL.


CASL subsumes many previous languages for the formal specification of functional requirements and modular software design. Tools for CASL are inter-operable, i.e., capable of being used in combination rather than in isolation. CASL interfaces to existing tools extend this inter-operability (see Chap. 11).

The intention was to base the design of CASL on a critical selection of concepts and constructs from existing specification languages. However, it was not easy to reach a consensus on a coherent language design. A great deal of careful consideration was given to the effect that the constructs available in the language would have on such aspects as the methodology and tools. A complete formal semantics for CASL was produced in parallel with the later stages of the language design, and the desire for a relatively straightforward semantics was one factor in the choice between various alternatives in the design.

CASL represents a consolidation of past work on the design of algebraic specification languages. With a few minor exceptions, all its features are present in some form in other languages, but there is no language that comes close to subsuming it. Designing a language with this particular novel collection of features required solutions to a number of subtle problems in the interaction between features. An overview of the CASL design is presented in [2], and full details are provided in the CASL Reference Manual [20].

*CASL is at the center of a family of languages.*

It was clear from the start that no single language could suit all purposes. On the one hand, sophisticated features are required to deal with specific programming paradigms and special applications. On the other hand, important methods for prototyping and reasoning about specifications only work in the *absence* of certain features: for instance, term rewriting requires specifications with equational or conditional equational axioms.



**Fig. 1.1.** The CASL Family of Languages

CASL is therefore at the center of a *family* of languages, see Fig. 1.1. Some tools will make use of well-delineated *sublanguages* of CASL, obtained by syntactic or semantic restrictions [29], while *extensions* of CASL are generally

designed to support various paradigms and applications. The design of CASL took into account the need to define sublanguages and extensions.

*CASL itself has several major parts.*

The major parts of CASL are concerned with *basic* specifications, *structured* specifications, *architectural* specifications, and *libraries* of specifications. They have been designed to be used together: basic specifications can be used in structured specifications, which in turn can be used in architectural specifications; structured and/or architectural specifications can be collected into libraries. However, these parts of CASL are quite independent, and may be understood separately, as we shall see in Part II.

---

## Underlying Concepts

*CASL is based on standard concepts of algebraic specification.*

This chapter reviews the main concepts of algebraic specification. It briefly explains and illustrates standard terminology regarding specification language constructs and models of specifications (i.e., algebras), and indicates the differences between basic, structured, and architectural specifications.

The focus here is on concepts that are relevant to CASL, and which will be needed in later chapters. For comprehensive presentations of concepts and results concerning algebraic specification, see [3, 10, 16, 27, 34, 35, 37]; for an overview of the design of CASL, see [2]; and for full details of CASL, see the CASL Reference Manual [20].

The reader is assumed to be familiar with basic mathematical notions (sets, relations, and total and partial functions) and with the use of logical formulas as axioms.

### 2.1 Basic Specifications

*A basic specification declares symbols, and gives axioms and constraints.*

A *basic specification* in an algebraic specification language generally consists of a set of *declarations* of *symbols*, and a set of *axioms* and *constraints*, which restrict the *interpretations* of the declared symbols. CASL allows basic specifications to include also items which simultaneously declare symbols and restrict their interpretations.

*The semantics of a basic specification is a signature and a class of models.*

The *meaning* or *semantics* of a basic specification  $SP$  generally has two parts:

- a *signature*  $\Sigma$ , corresponding to the *symbols* introduced by the specification, and
- a *class of  $\Sigma$ -models*,<sup>1</sup> corresponding to those *interpretations* of the signature  $\Sigma$  that *satisfy* the *axioms* and *constraints* of the specification.

When a model  $M$  satisfies a specification  $SP$ , we write  $M \models SP$  and say that  $M$  is a *model of  $SP$* . (Formalizing this within the theory of so-called *institutions* involves categorical structure on the set of signatures and on the class of models, and a natural condition on the *satisfaction relation*. We need not bother with the details here – but see however the concept of a signature morphism in Sect. 2.2.) A specification is said to be *consistent* when its class of models is non-empty, and otherwise *inconsistent*.

*CASL specifications may declare sorts, subsorts, operations, and predicates.*

A CASL *signature* represents declarations of sorts, subsorts, operations, and predicates. The signature is called *many-sorted* when there are no subsort declarations, and otherwise *subsorted*; it is called *algebraic* when there are no predicate declarations.

*Sorts are interpreted as carrier sets.*

A *sort* is a symbol which is interpreted as a set, called a *carrier set*. The elements of a carrier set are generally abstract representations of the data processed by software: numbers, characters, lists, etc. Thus a sort declared by a specification corresponds to a type in a programming language. Sort symbols are usually chosen to be strongly suggestive of their intended interpretations, e.g., *Int* for a sort to be interpreted as the set of integers, *List* for a set of lists. CASL allows also *compound sort* symbols, such as *List[Int]* for lists of integers.

<sup>1</sup> Readers who are not interested in foundational aspects may treat the word ‘class’ as a synonym for ‘set’. In general, the models of an algebraic specification in CASL constitute a *proper class*, because there is no restriction on the elements of the carrier sets.

*Subsorts declarations are interpreted as embeddings.*

A sort may be declared to be a *subsort* or a *supersort* of other sorts. The subsort relation between two sorts could be interpreted as *set inclusion*. Its interpretation in CASL is however more general: it is interpreted as an *embedding*, i.e., a 1-1 function from the carrier set of the subsort to that of the supersort. For example, if *ASCII* is specified to be a subsort of *ISO\_Latin1* in CASL, the carrier set for *ASCII* could be simply a subset of that for *ISO\_Latin1*. If *Char* were to be declared as a subsort of *String*, however, the carrier sets for *Char* and *String* could be disjoint, with the embedding mapping each character to the corresponding single-character string. (See also the concept of overloading, below.)

*Operations may be declared as total or partial.*

An *operation* symbol consists of the name of the operation together with its *profile*, which indicates the number and sorts of the arguments, and the *result sort*. In CASL, a declared operation symbol is interpreted as either a *total* or a *partial function* from the Cartesian product of the carrier sets of the argument sorts to the carrier set of the result sort; the subset of the argument tuples for which the result of a function is defined is called its *domain of definition*. The declaration indicates whether the function is total or partial.<sup>2</sup> For example, integer addition would be declared as total, but integer division as partial. The result of applying an operation is undefined whenever any of its arguments is undefined (regardless of whether the operation itself is total or partial).

When there are no arguments, the operation is called a *constant*. A constant is interpreted simply as an element of the result sort.

*Predicates are different from boolean-valued operations.*

A *predicate* symbol consists of the name of the predicate together with its *profile*, which indicates the number and sorts of the arguments but no result sort: predicates are syntactically different from boolean-valued operations, and are used to form atomic formulas rather than terms. In CASL, a declared predicate symbol is interpreted as a relation on (i.e., a subset of) the Cartesian product of the carrier sets of the argument sorts. An application of a predicate is said to *hold* when the tuple of arguments is in the relation. For example, a

<sup>2</sup> A partial function might just happen to be everywhere defined, of course.



symbol ‘ $<$ ’ to be interpreted as the less-than relation could be declared as a binary predicate on integers.

An application of a predicate simply fails to hold when any of its arguments is undefined: there is no undefinedness about holding or not. This allows the logic to remain two-valued, and the logical connectives to have their familiar interpretations.

In contrast, the result of evaluating an application of even a total boolean-valued operation could be true, false, or undefined: the last case arises when any argument of the application is undefined. Thus boolean-valued operations corresponding to logical connectives (conjunction, implication, etc.) have to take account of undefinedness, which leads to a three-valued logic.

A further significant difference between predicates and boolean-valued operations shows up in connection with the concept of initiality, see Sect. 2.2. (Predicates of two-valued logic can be represented accurately by *partial* operations with a *single-valued* result sort, holding being represented by definedness.)

*Operation symbols and predicate symbols may be overloaded.*

An operation or predicate name can be declared with different profiles in the same specification. This is called *overloading*. For example, the constant ‘*empty*’ could be overloaded, being interpreted as (unrelated) elements of the sorts *List* and *Set*, according to the context of its use. Similarly, a predicate name such as ‘ $<$ ’ could be overloaded on unrelated sorts such as *Char* and *Int*.

In CASL, overloading is required to be *compatible* with embeddings between subsorts. For example, the sort *Nat*, interpreted as the set of natural numbers, might be a subsort of *Int*, interpreted as the set of all integers; then when the operation name ‘ $+$ ’ and the predicate name ‘ $<$ ’ are declared both on *Nat* and on *Int*, their interpretations are required to be such that it makes no difference whether the embedding from *Nat* to *Int* is applied to the arguments or to the result of the operation, and whether it is applied to the arguments of the predicate or not.

*Axioms are formulas of first-order logic.*

The interpretation of quantification (universal, existential, and unique-existential) and of the usual logical connectives (negation, conjunction, disjunction, implication, and equivalence) in CASL axioms is completely standard. Variables in formulas range over the carrier sets of specified sorts.

Apart from the usual predicate applications, the atomic formulas in CASL axioms are equations (strong or existential), definedness assertions, and subsort membership assertions. An *existential* equation holds when the values of its terms are defined and equal; a *strong* equation holds moreover when the values of the terms are both undefined.

Regardless of whether the values of the terms occurring in an axiom are defined, the axiom either holds or it does not hold in a particular model: the logic is two-valued, there is no “maybe” or undefinedness about the holding of axioms. Recall that when the value of any argument term is undefined, an application of a predicate never holds; similarly, definedness and subsort membership assertions never hold when their arguments are undefined.

*Sort generation constraints eliminate ‘junk’ from specific carrier sets.*

In general, the carrier sets of the models of a specification may contain ‘junk’ elements, i.e., elements which cannot be obtained by any composition of the operations declared by the signature of the specification.

A *sort generation constraint* in CASL concerns specific sorts and operations, and is satisfied in a model when no elements of the indicated carrier sets are junk with respect to the indicated operations – i.e., all the elements of those sets can be obtained by consecutively applying those operations to elements of the carrier sets of the remaining sorts. For example, the carrier set for the sort *Container* might be constrained to be generated from that for the sort *Elem* by the following operations:

- a constant ‘*empty*’ of sort *Container*, and
- a binary operation ‘*insert*’ with argument sorts *Elem* and *Container*, and result sort *Container*.

This constraint would ensure that the only elements of the *Container* carrier are those obtained by a finite number of successive applications of the *insert* operation to elements of sort *Elem*, starting with the *empty* value of sort *Container*.

## 2.2 Structured Specifications

Structured specifications are formed from basic specifications, *references* to *named specifications*, and *instantiations* of *generic specifications*, using various constructs for composing specifications.

*The semantics of a structured specification is simply a signature and a class of models.*

The semantics of a structured specification is of the same form as that of a basic specification: a signature, together with a class of models. Thus the structure of a specification is *not* reflected in its models: it is used only to present the specification in a modular style. (Specification of the *architecture* of implementations is addressed in Sect. 2.3.) The symbols in the signature are called the *exported* symbols of the specification.

The interpretation of structured specification constructs involves mappings between signatures  $\Sigma$ , called *signature morphisms*, and corresponding mappings between models  $M$ , called *reducts along morphisms*. In CASL, a signature morphism  $\sigma$  from  $\Sigma$  to  $\Sigma'$  consists of a mapping which gives:

- for each sort of  $\Sigma$ , a corresponding sort of  $\Sigma'$ , preserving any subsort relationships, and
- for each operation or predicate symbol whose profile has sorts in  $\Sigma$ , a corresponding symbol in  $\Sigma'$  whose profile has the corresponding sorts, preserving any overloading between symbols whose profiles are related by subsorting. A partial operation may be mapped to a total operation, but not vice versa.

Let  $M'$  be any  $\Sigma'$ -model. We can define its *reduct* along the signature morphism  $\sigma$  to be the  $\Sigma$ -model  $M$  obtained as follows: each symbol of  $\Sigma$  is interpreted in  $M$  in exactly the same way as the corresponding symbol in  $\Sigma'$  is interpreted in  $M'$ . Conversely, given  $M$ , a model  $M'$  is said to be an *expansion* of  $M$  when the reduct of  $M'$  is  $M$ .

Suppose that a specification  $SP$  has signature  $\Sigma$  and  $SP'$  has signature  $\Sigma'$ . A signature morphism  $\sigma$  from  $\Sigma$  to  $\Sigma'$  is said to be a *specification morphism* from  $SP$  to  $SP'$  when the reduct along  $\sigma$  of each model of  $SP'$  is a model of  $SP$ .

For two  $\Sigma$ -models  $M_1, M_2$ , a (weak) *homomorphism* from  $M_1$  to  $M_2$  maps the elements of the carrier sets of  $M_1$  to the elements of the corresponding carrier sets of  $M_2$ , preserving the embeddings between subsorts, the values (and definedness) of operations, and the holding of predicates. A homomorphism is an *isomorphism* when it has an inverse homomorphism.

A model  $M$  is *initial* in a class of  $\Sigma$ -models if there is a unique homomorphism from  $M$  to each model in the class. When a class of models has an initial model (which need not be the case in CASL) it is unique, up to isomorphism.

With reference to the above concepts of signature morphism, model reduct, and homomorphism, we can now proceed to explain the constructs involved with structured specifications.

*A translation merely renames symbols.*

Translating a sort symbol requires translating the profiles of all operation and predicate symbols involving that sort; translating an operation or predicate symbol has to respect overloading between symbols whose profiles are related by subsorting. The translation of sort, operation, and predicate names in CASL determines a signature morphism  $\sigma$  mapping the signature  $\Sigma$  of a specification  $SP$  onto a new signature  $\Sigma'$ . The models of the translation specification are all those models interpreting  $\Sigma'$  whose reducts along  $\sigma$  are models of  $SP$ .

*Hiding symbols removes parts of models.*

Hiding a sort symbol implies hiding also all operation and predicate symbols whose profiles involve that sort; hiding an operation or predicate symbol, however, does not have further implications. Hiding a set of symbols that occur in the signature  $\Sigma$  of a specification  $SP$  to give a subsignature  $\Sigma'$  determines a signature morphism  $\sigma$  which simply includes  $\Sigma'$  in  $\Sigma$ . The models of the hiding specification are the reducts of the models of  $SP$  along  $\sigma$ .

For example, the operation *suc* might be introduced purely to facilitate the specification of the natural numbers, with sort *Nat*, constants *0* and *1*, and the usual arithmetic operations. Hiding *suc* removes the interpretation of *suc* from the models of the specification<sup>3</sup> but leaves the carrier set for *Nat* unchanged.

*Union of specifications identifies common symbols.*

The signature of a union of specifications  $SP_1, SP_2$  is simply the union of their respective signatures  $\Sigma_1, \Sigma_2$ . The models of the union are those models of the union signature whose reducts to  $\Sigma_1$  and  $\Sigma_2$  along the signature inclusions satisfy  $SP_1$  and  $SP_2$  respectively. Thus each symbol that the two signatures have in common has a single interpretation in any model of the union specification. This is known as the ‘same name, same thing’ principle.

<sup>3</sup> The successor of a number can of course still be obtained, using addition and *1*.

*Extension of specifications identifies common symbols too.*

The signature of the extension of a specification  $SP$  by further specification items (declarations, axioms, and constraints) is simply the extension of the signature  $\Sigma$  of  $SP$  with the symbols of the new declarations. The models of the extension are those models of the extended signature which satisfy the axioms and constraints specified by the extension and whose reducts to  $\Sigma$  satisfy  $SP$ . If the extension redeclares a symbol of  $SP$ , there is still only one occurrence of that symbol in the signature of the extension, and hence only one interpretation of it – again the ‘same name, same thing’ principle.

In CASL, unions, extensions, and other kinds of structured specification can be formed from specification fragments that determine only signature *extensions*, not necessarily complete signatures.

*Free specifications restrict models to being free, with initiality as a special case.*

When a specification is freely extended by additional specification items, the interpretations of the additional declarations are required to satisfy the axioms, but nothing more: that is, properties that are not consequences of the axioms should not hold. In particular, the domains of definition of partial operations – and the sets of arguments for which predicates hold – are as small as possible. The carriers for the original sorts are left unchanged; any new carriers are no larger than is required to provide interpretations for the operations, without unnecessary junk elements. This restriction of the models is referred to as a *freeness constraint*. In the degenerate case where the specification being enriched is empty, the models of the free extension are just the initial models.

The difference between predicates and boolean-valued operations is particularly apparent in free specifications: with predicates, it is only required to specify when they hold, since not holding is the default; with boolean-valued operations, however, the true and false values are treated symmetrically, and it is necessary to specify both cases, since neither is the default.

*Generic specifications have parameters, and have to be instantiated when referenced.*

A named specification may declare some *parameters*, the union of which is extended by a *body*; it is then called *generic*. The purpose of a generic

specification is to reuse the body in different contexts; hence a reference to a generic specification should *instantiate* it by providing, for each parameter, an *argument specification* together with a *fitting morphism* from the parameter to the argument specification. Fitting may also be achieved by use of named *views* between the parameter and argument specifications. The instantiation of the generic specification gives the union of the arguments, together with the translation of the generic specification by an expansion of the fitting morphism. This corresponds to a so-called push-out construction – taking into account any explicit *imports* of the generic specification.

## 2.3 Architectural Specifications

*The semantics of an architectural specification reflects its modular structure.*

The intention with architectural specifications is primarily to impose structure on implementations, expressing their *composition* from component units – and thereby also a *decomposition* of the task of developing such implementations, from requirements specifications. This is in contrast to the structured specifications considered in Sect. 2.2, where the specified models have no more structure than do those of the basic specifications considered in Sect. 2.1.

*Architectural specifications involve the notions of persistent function and conservative extension.*

A function  $F$  mapping  $\Sigma$ -models to  $\Sigma'$ -models, where the signature  $\Sigma'$  extends  $\Sigma$ , is said to be *persistent* when for each  $\Sigma$ -model  $M$ , the reduct of  $F(M)$  to a  $\Sigma$ -model is exactly  $M$ .

A specification extension  $SP'$  of  $SP$  is said to be *conservative* when each model of  $SP$  can be expanded to a model of  $SP'$ .

A persistent function mapping models of  $SP$  to models of  $SP'$  exists only if  $SP'$  is a conservative extension of  $SP$ .

## 2.4 Libraries of Specifications

*The semantics of a library of specifications is a mapping from the names of the specifications to their semantics.*

The specification of a library gives also a library name, and determines a version number.

## CASL Specifications



---

## Getting Started

*Simple specifications may be written in CASL essentially as in many other algebraic specification languages.*

The simplest kind of algebraic specification is when each specified operation is to be interpreted as an ordinary *total* mathematical function: it takes values of particular types as arguments, and always returns a well-defined value. Total functions correspond to software whose execution always terminates normally. The types of values are named by simple symbols called *sorts*.

In practice, a realistic software specification involves *partial* as well as total functions. However, it may well be formed from simpler specifications, some of which involve only total functions. This chapter explains how to express such simple specifications in CASL, illustrating various features of the language.

The simple specifications discussed in this chapter can also be expressed in many previous specification languages; it is usually straightforward to reformulate them in CASL. Readers who know other specification languages will probably recognize some familiar examples among the illustrations given in this chapter.

*CASL provides also useful abbreviations.*

The technique of algebraic specification by axioms is generally well-suited to expressing properties of functions. However, when functions have commonly-occurring mathematical properties, it can be tedious to give the corresponding axioms explicitly. CASL provides some useful abbreviations for such cases. Similarly, so-called free datatype declarations allow sorts and value constructors to be specified much as in functional programming languages, using a concise and suggestive notation.

*CASL allows loose, generated and free specifications.*

The models of a loose specification include all those where the declared functions have the specified properties, without any restrictions on the sets of values corresponding to the various sorts. In models of a generated specification, in contrast, it is required that all values can be expressed by terms formed from the specified constructors, i.e. unreachable values are prohibited. In models of free specifications, it is required that values of terms are distinct except when their equality follows from the specified axioms: the possibility of unintended coincidence between them is prohibited.

Section 3.1 below focuses on loose specifications; Sect. 3.2 discusses the use of generated specifications, and Sect. 3.3 does the same for free specifications. Loose, generated, and free specifications are often used together in CASL: each style has its advantages in particular circumstances, as explained below in connection with the illustrative examples.

### 3.1 Loose Specifications

*CASL syntax for declarations and axioms involves familiar notation, and is mostly self-explanatory.*

```
spec STRICT_PARTIAL_ORDER =
  %% Let's start with a simple example !
  sort Elem
  pred -- < -- : Elem × Elem %% pred abbreviates predicate
  ∀x, y, z : Elem
    • ¬(x < x)                                %(strict)%
    • x < y ⇒ ¬(y < x)                        %(asymmetric)%
    • x < y ∧ y < z ⇒ x < z                    %(transitive)%
  %%{ Note that there may exist x, y such that
    neither x < y nor y < x. }%
end
```

The above (basic) specification, named `STRICT_PARTIAL_ORDER`, introduces a sort `Elem` and a binary infix predicate symbol '`<`'. In the declaration of a predicate symbol, argument sorts are separated by the sign '`×`', which can be input directly as such in ISO Latin-1 or as '`*`' in ASCII. Note that CASL allows so-called *mixfix notation*, i.e., the specifier is free to indicate, using '`--`' (pairs of underscores) as *place-holders*, how to place arguments when building

terms (single underscores are treated as letters in identifiers).<sup>1</sup> Using mixfix notation generally allows the use of familiar mathematical and programming notations, which contributes substantially to the readability of specifications.

The interpretation of the binary predicate symbol ‘<’ is then constrained by three axioms. A set of axioms is generally presented as a ‘bulleted’ list of formulas, preceded by the universally quantified declaration of the relevant variables, together with their respective sorts, as shown in the above example. In CASL, axioms are written in first-order logic with equality, using quantifiers and the usual logical connectives. The universal quantification preceding a list of axioms applies to the entire list. Axioms can be annotated by labels written `%(...)%`, which is convenient for proper reference in explanations or by tools.

Note that ‘ $\forall$ ’ is input as ‘`forall`’, and that ‘ $\bullet$ ’ is input as ‘`.`’ or ‘`·`’. The usual logical connectives ‘ $\Rightarrow$ ’, ‘ $\Leftrightarrow$ ’, ‘ $\wedge$ ’, ‘ $\vee$ ’, and ‘ $\neg$ ’, are input as ‘`=>`’, ‘`<=>`’, ‘`/\`’, ‘`\V`’, and ‘`not`’, respectively; ‘ $\neg$ ’ can also be input directly as an ISO Latin-1 character. The existential quantifier ‘ $\exists$ ’ is input as ‘`exists`’, and ‘ $\exists!$ ’ is input as ‘`exists!`’.

It is advisable to comment as appropriate the various elements introduced in a specification. The syntax for end-of-line and grouped multi-line comments is illustrated in the above example. The ‘`end`’ keyword ending a specification is optional.

The above `STRICT_PARTIAL_ORDER` specification is loose in the sense that it has many (non-isomorphic) models, among which models where ‘<’ is interpreted by a total ordering relation and models where it is interpreted by a partial one.

*Specifications can easily be extended by new declarations and axioms.*

```
spec TOTAL_ORDER =
    STRICT_PARTIAL_ORDER
then  $\forall x, y : Elem \bullet x < y \vee y < x \vee x = y$       %(total)%
end
```

Extensions, introduced by the keyword ‘**then**’, may specify new symbols, possibly constrained by some axioms, or merely require further properties of old ones, as in the above `TOTAL_ORDER` example, or more generally do both at the same time. In `TOTAL_ORDER`, we further constrain the interpretation of the predicate symbol ‘<’ by requiring it to be a total ordering relation.

All symbols introduced in a specification are by default exported by it and visible in its extensions. This is for instance the case here for the sort *Elem* and

<sup>1</sup> Mixfix notation is so-called because it generalizes infix, prefix, and postfix notation to allow arbitrary mixing of argument positions and identifier tokens.

the predicate symbol ' $<$ ', which are introduced in `STRICT_PARTIAL_ORDER`, exported by it, and therefore available in `TOTAL_ORDER`.<sup>2</sup>

*In simple cases, an operation (or a predicate) symbol may be declared and its intended interpretation defined at the same time.*

```
spec TOTAL_ORDER_WITH_MINMAX =
  TOTAL_ORDER
then ops min(x, y : Elem) : Elem = x when x < y else y;
      max(x, y : Elem) : Elem = y when min(x, y) = x else x
end
```

`TOTAL_ORDER_WITH_MINMAX` extends `TOTAL_ORDER` by introducing two binary operation symbols *min* and *max*, for which a functional notation is to be used, so no place-holders are given. The intended interpretation of the symbol *min* is defined simultaneously with its declaration, and the same is done for *max*. For instance:

**op** *min*(*x*, *y* : *Elem*) : *Elem* = *x* when *x* < *y* else *y*

abbreviates:

**op** *min* : *Elem* × *Elem* → *Elem*  
 $\forall x, y : Elem \bullet min(x, y) = x \text{ when } x < y \text{ else } y$

(and similarly for *max*). As for predicate symbol declarations, in an operation symbol declaration, the argument sorts are separated by the sign ' $\times$ '; the result sort is preceded by ' $\rightarrow$ ', which is input as ' $\rightarrow$ '.

The ' $\dots \text{ when } \dots \text{ else } \dots$ ' construct used above is itself an abbreviation, and:

*min*(*x*, *y*) = *x* when *x* < *y* else *y*

abbreviates:

$(x < y \Rightarrow min(x, y) = x) \wedge (\neg(x < y) \Rightarrow min(x, y) = y)$

In CASL specifications, visibility is *linear*, i.e., any symbol must be declared before being used. In the above example, *min* should be declared before being used to define *max*.

Linear visibility does not imply, however, that a fixed scheme is to be used when writing specifications: the specifier is free to present the required declarations and axioms in any order, as long as the linear visibility rule is respected. For instance, one may prefer to declare first all sorts and all operation

---

<sup>2</sup> See Chap. 6 for constructs allowing the explicit restriction of the set of symbols exported by a specification.

or predicate symbols needed, and then specify their properties by the relevant axioms. Or, in contrast, one may prefer to have each operation or predicate symbol declaration immediately followed by the axioms constraining its interpretations. Both styles are equally fine, and can even be mixed if desired. This flexibility is illustrated in the following variant of the `TOTAL_ORDER_WITH_MINMAX` specification, where for explanatory purposes we refrain from using the useful abbreviations explained above.

```
spec  VARIANT_OF_TOTAL_ORDER_WITH_MINMAX =
      TOTAL_ORDER
then  vars  x, y : Elem
      op  min : Elem × Elem → Elem
          • x < y ⇒ min(x, y) = x
          • ¬(x < y) ⇒ min(x, y) = y
      op  max : Elem × Elem → Elem
          • x < y ⇒ max(x, y) = y
          • ¬(x < y) ⇒ max(x, y) = x
end
```

Note that in order to avoid the tedious repetition of the declaration of the variables  $x$  and  $y$  for each list of axioms, we have used here a *global* variable declaration which introduces  $x$  and  $y$  for the rest of the specification. Variable declarations are of course not exported across specification extensions: the variables  $x$  and  $y$  declared in `VARIANT_OF_TOTAL_ORDER_WITH_MINMAX` are not visible in any of its extensions.

*Symbols may be conveniently displayed as usual mathematical symbols by means of **%display** annotations.*

```
%display _<=_ %LATEX _ ≤ _
```

```
spec  PARTIAL_ORDER =
      STRICT_PARTIAL_ORDER
then  pred  _ ≤ _ (x, y : Elem) ⇔ (x < y ∨ x = y)
end
```

The above example relies on a **%display** annotation: while, for obvious reasons, the specification text should be *input* using the ISO Latin-1 character set, it is often convenient to *display* some symbols differently, e.g., as mathematical symbols. This is the case here where the ‘<=’ predicate symbol is more conveniently displayed as ‘≤’. Display annotations, as any other CASL annotations, are auxiliary parts of a specification, for use by tools, and do not affect the semantics of the specification.<sup>3</sup>

<sup>3</sup> Display annotations should be provided at the beginning of a library, and are explained in more detail in Chap. 9.

In the above example, we have again used the facility of simultaneously declaring and defining a symbol (here, the predicate symbol ' $\leq$ ') in order to obtain a more concise specification.

*The **%implies** annotation is used to indicate that some axioms are supposedly redundant, being consequences of others.*

```
spec PARTIAL_ORDER_1 =
  PARTIAL_ORDER
then %implies
   $\forall x, y, z : Elem \bullet x \leq y \wedge y \leq z \Rightarrow x \leq z$       %(transitive)%
end
```

The **%implies** annotation above is used to emphasize that the transitivity of ' $\leq$ ' should follow from the other axioms, or, in other words, that the model class of `PARTIAL_ORDER_1` is exactly the same as the model class of `PARTIAL_ORDER`. The **%implies** annotation applies to the whole of the specification extension where it occurs (which happens here to introduce a single axiom).

Note however that an annotation does not affect the semantics of a specification, hence removing the **%implies** annotation does not change the class of models of the above specification. The sole aim of an **%implies** annotation is to stress the specifier's intentions, and it will also help readers confirm their understanding. Some tools may of course use such annotations to generate corresponding proof obligations. For instance, here, the proof obligation is:

$$PARTIAL\_ORDER \models \forall x, y, z : Elem \bullet x \leq y \wedge y \leq z \Rightarrow x \leq z$$

Discharging these proof obligations increases the trustworthiness of a specification.

To fully understand that an **%implies** annotation has no effect on the semantics, the best is to consider an example where the corresponding proof obligation cannot be discharged, as shown below.

```
spec IMPLIES_DOES_NOT_HOLD =
  PARTIAL_ORDER
then %implies
   $\forall x, y : Elem \bullet x < y \vee y < x \vee x = y$       %(total)%
end
```

Since the loose specification `PARTIAL_ORDER` has models where ' $<$ ' is interpreted by a partial ordering relation, the proof obligation corresponding to the above **%implies** annotation cannot be discharged. However, since annotations have no impact on the semantics, the specification `IMPLIES_DOES_NOT_HOLD` is well-formed and just constrains the interpretation of ' $<$ ' to be a total ordering relation. The fact that the proof obligation cannot be discharged merely points out a potential mistake in the specification.

*Attributes may be used to abbreviate axioms for associativity, commutativity, idempotence, and unit properties.*

```
spec MONOID =
  sort Monoid
  ops 1      : Monoid;
      _ * _ : Monoid × Monoid → Monoid, assoc, unit 1
end
```

The above example introduces a constant symbol  $1$  of sort *Monoid*, then a binary operation symbol  $*$ , which is asserted to be associative and to have  $1$  as unit element. (Note that there is no  $\rightarrow$  sign before the sort when declaring a constant.) The *assoc* attribute abbreviates, as expected, the following axiom:

$$\forall x, y, z : \text{Monoid} \bullet (x * y) * z = x * (y * z)$$

The ‘*unit 1*’ attribute abbreviates:

$$\forall x : \text{Monoid} \bullet (x * 1 = x) \wedge (1 * x = x)$$

Note that to make the use of ‘*unit 1*’ legal, it is necessary to have previously declared the constant  $1$ , to respect the linear visibility rule.

Other available attributes are *comm*, which abbreviates the obvious axiom stating that a binary operation is commutative, and *idem*, which can be used to assert the idempotence of a binary operation  $f$  (i.e., that  $f(x, x) = x$ ).

Asserting ‘ $*$ ’ to be associative using the attribute *assoc* makes the term  $x * y * z$  well-formed (assuming  $x, y, z$  of the right sort), while otherwise grouping parentheses would be required. Moreover, it is expected that some tools (e.g., systems based on rewriting) may make special use of the *assoc* attribute, so it is generally advisable to use this attribute instead of stating the same property by an axiom (the same applies to the other attributes).

*Genericity of specifications can be made explicit using parameters.*

```
spec GENERIC_MONOID [sort Elem] =
  sort Monoid
  ops inj      : Elem → Monoid;
      1        : Monoid;
      _ * _    : Monoid × Monoid → Monoid, assoc, unit 1
  ∀ x, y : Elem • inj(x) = inj(y) ⇒ x = y
end
```

The above example describes monoids built over arbitrary elements (of sort *Elem*). The intention here is to reuse the specification `GENERIC_MONOID` to derive from it specifications of monoids built over, say, characters, symbols, etc. In such cases, it is appropriate to emphasize the intended genericity of the specification by making explicit, in a distinguished *parameter part* (which is here `[sort Elem]`), the piece of specification that is intended to vary in the derived specifications. In these, it will then be possible to *instantiate* the parameter part as desired in order to specialize the specification as appropriate (to obtain, e.g., a specification of monoids built over characters). A named specification with one or more parameter(s) is called *generic*.

The body of the generic specification `GENERIC_MONOID` is an extension of what is specified in the parameter part. Hence an alternative to the above generic specification `GENERIC_MONOID` is the following, less elegant, non-generic specification (which cannot be specialized by instantiation):

```
spec NON_GENERIC_MONOID =
  sort Elem
then sort Monoid
  ops inj   : Elem → Monoid;
      1     : Monoid;
      _ * _ : Monoid × Monoid → Monoid, assoc, unit 1
  ∀x, y : Elem • inj(x) = inj(y) ⇒ x = y
end
```

A generic specification may have more than one parameter, and parameters can be arbitrary specifications, named or not. When reused by reference to its name, a generic specification *must* be instantiated. Generic specifications and how to instantiate them are discussed in detail later in Chap. 7. Using generic specifications when appropriate improves the reusability of specification definitions.

*References to generic specifications always instantiate the parameters.*

```
spec GENERIC_COMMUTATIVE_MONOID [sort Elem] =
  GENERIC_MONOID [sort Elem]
then ∀x, y : Monoid • x * y = y * x
end
```

The above (generic) specification `GENERIC_COMMUTATIVE_MONOID` is defined as an extension of `GENERIC_MONOID`, which should therefore be instantiated, as explained above. Instantiating a generic specification is done by providing an argument specification that ‘fits’ the parameter part of the generic specification to be instantiated.



It is however quite frequent that the instantiation is ‘trivial’, i.e., the argument specification is identical to the parameter one. This is the case for the above example, where the generic specification `GENERIC_MONOID` is instantiated by providing the same argument specification ‘`sort Elem`’ as the original parameter.

```
spec  GENERIC_COMMUTATIVE_MONOID_1 [sort Elem] =
      GENERIC_MONOID [sort Elem]
then  op __ * __ : Monoid × Monoid → Monoid, comm
end
```

`GENERIC_COMMUTATIVE_MONOID_1` is an alternative version of the former specification where, instead of requiring explicitly with an axiom the commutativity of the operation ‘`*`’, we require it using the attribute *comm*. As explained before, it is in general better to describe such requirements using attributes rather than explicit axioms, since it is expected that some tools will rely on these attributes for specialized algorithms (e.g., AC term rewriting).

This example illustrates also an important feature of CASL, the ‘*same name, same thing*’ principle. The operation symbol ‘`*`’ is indeed declared twice, with the same profile, first in `GENERIC_MONOID` and then again in `GENERIC_COMMUTATIVE_MONOID_1` (the second declaration being enriched by the attribute *comm*). This is perfectly fine and defines only one binary operation symbol ‘`*`’ with the corresponding profile, according to the ‘same name, same thing’ principle. This principle applies to sorts, as well as to operation and predicate symbols. It applies both to symbols defined locally and to symbols imported from an extended specification, as it is the case here for ‘`*`’. Of course, it does not apply between separate named specifications, i.e., the same symbol may be used in different named specifications with entirely different interpretations.

Note that for operation and predicate symbols, the ‘same name, same thing’ principle is a little more subtle than for sorts: the ‘name’ of an operation (or of a predicate) includes its profile of argument and result sorts, so two operations defined with the same symbol but with different profiles do not have the same ‘name’, the symbol is just *overloaded*. When an overloaded symbol is used, the intended profile is to be determined by the context (e.g., the sorts of the arguments to which the symbol is applied).<sup>4</sup> Explicit disambiguation can be used when needed, by specifying the profile (or result sort) in an application.<sup>5</sup> Note that overloaded constants are allowed in CASL (e.g., *empty* may be declared to be a constant of various sorts of collections).

---

<sup>4</sup> See also the discussion of overloading in presence of subsorts in Chap. 5, p. 61.

<sup>5</sup> For instance, depending on the context, the term  $t1 * t2$  can be disambiguated by writing  $(op * : Monoid \times Monoid \rightarrow Monoid)(t1, t2)$ , or just  $(t1 : Monoid) * (t2 : Monoid)$ , or even  $(t1 * t2) : Monoid$ .

*Datatype declarations may be used to abbreviate declarations of sorts and constructors.*

```
spec CONTAINER [sort Elem] =
  type Container ::= empty | insert(Elem; Container)
  pred __is_in__ : Elem × Container
  ∀e, e' : Elem; C : Container
    • ¬(e is_in empty)
    • e is_in insert(e', C) ⇔ (e = e' ∨ e is_in C)
end
```

Specifications of ‘datatypes’ with constructors are frequently needed. CASL provides special constructs for datatype declarations to abbreviate the corresponding rather tedious declarations. For instance, the above datatype declaration:

```
type Container ::= empty | insert(Elem; Container)
```

abbreviates:

```
sort Container
ops empty : Container;
    insert : Elem × Container → Container
```

A datatype declaration looks like a context-free grammar in a variant of BNF. It declares the symbols on the left of ‘::=’ as sorts, and for each alternative on the right it declares a constructor.

A datatype declaration as the one above is *loose* since it does not imply any constraint on the values of the declared sorts: there may be some values of sort *Container* that are not built by any of the declared constructors, and the same value may be built by different applications of the constructors to some arguments.

Datatype declarations may also be specified as generated (see Sect. 3.2) or free (see Sect. 3.3). Moreover, selectors, which are usually partial operations, may be specified for each component (see Chap. 4).

*Loose datatype declarations are appropriate when further constructors may be added in extensions.*

```
spec MARKING_CONTAINER [sort Elem] =
  CONTAINER [sort Elem]
then type Container ::= mark_insert(Elem; Container)
```

```

pred __is_marked_in__ : Elem × Container
  ∀e, e' : Elem; C : Container
    • e is_in mark_insert(e', C) ⇔ (e = e' ∨ e is_in C)
    • ¬(e is_marked_in empty)
    • e is_marked_in insert(e', C) ⇔ e is_marked_in C
    • e is_marked_in mark_insert(e', C) ⇔ (e = e' ∨ e is_marked_in C)
end

```

The above specification extends `CONTAINER` (trivially instantiated) by introducing another constructor `mark_insert` for the sort `Container` (hence, values added to a container may now be ‘marked’ or not). Note that we heavily rely on the ‘same name, same thing’ principle here, since it ensures that the sort `Container` introduced by the datatype declaration of `CONTAINER` and the sort `Container` introduced by the datatype declaration of `MARKING_CONTAINER` are the same sort, which implies that the combination of both datatype declarations is equivalent to:

```

type Container ::= empty | insert(Elem; Container)
                      | mark_insert(Elem; Container)

```

Note that since ‘new’ values may be constructed by `mark_insert`, it is necessary to extend the specification of the predicate symbol `is_in` by an extra axiom taking care of the newly introduced constructor.

## 3.2 Generated Specifications

*Sorts may be specified as generated by their constructors.*

```

spec GENERATED_CONTAINER [sort Elem] =
  generated type Container ::= empty | insert(Elem; Container)
  pred __is_in__ : Elem × Container
    ∀e, e' : Elem; C : Container
      • ¬(e is_in empty)
      • e is_in insert(e', C) ⇔ (e = e' ∨ e is_in C)
end

```

When a datatype is declared as *generated*, as in the above example, the corresponding sort is constrained to be generated by the declared constructors, which means that any value of this sort is built by application of constructors. This constraint is sometimes referred to as the ‘no junk’ principle. For instance, in the above example, having declared the datatype `Container` to be generated entails that in any model of `GENERATED_CONTAINER`, any value of

sort *Container* is denotable by a term built with *empty*, *insert*, and variables of sort *Elem* only.

As a consequence, properties of values of sort *Container* can be proved by induction on the declared constructors. A major benefit of generated datatypes is indeed that induction on the declared constructors is a sound proof principle.

The construct ‘**generated type ...**’ used above is just an abbreviation for ‘**generated { type ... }**’, and ‘**generated**’ can be used around arbitrary signature declarations, enclosed in ‘{’ and ‘}’.

*Generated specifications are in general loose.*

```
spec  GENERATED_CONTAINER_MERGE [sort Elem] =
      GENERATED_CONTAINER [sort Elem]
then  op __merge__ : Container × Container → Container
      ∀e : Elem; C, C' : Container
      • e is_in (C merge C') ⇔ (e is_in C ∨ e is_in C')
end
```

A generated specification is in general loose. For instance, `GENERATED_CONTAINER` is loose since, although all values of sort *Container* are specified to be generated by *empty* and *insert*, the behavior of the *insert* constructor is still loosely specified (nothing is said about the case where an element is inserted into a container which already contains this element). Hence `GENERATED_CONTAINER` admits several non-isomorphic models.

`GENERATED_CONTAINER_MERGE` is as loose as `GENERATED_CONTAINER` with respect to *insert*, and in addition, the newly introduced operation symbol *merge* is also loosely specified: nothing is said about what happens when merging two containers which share some elements.

It is important to understand that looseness of a specification is not a problem, but on the contrary avoids unnecessary overspecification. In particular, loose specifications are in general well-suited to capturing requirements.

The fact that *merge* is loosely specified does not mean that it can produce new values of the sort *Container*. On the contrary, since this sort has been specified as being generated by *empty* and *insert*, it follows that any value denotable by a term of the form *merge*(..., ...) can also be denoted by a term built with *empty* and *insert* (and no *merge*). Hence, for the specification `GENERATED_CONTAINER_MERGE`, proofs by induction on *Container* only need to consider *empty* and *insert*, and not *merge*, as was the case for `GENERATED_CONTAINER`.

*Generated specifications need not be loose.*

```

spec GENERATED_SET [sort Elem] =
  generated type Set ::= empty | insert(Elem; Set)
  pred __is_in__ : Elem × Set
  ops  {__}(e : Elem) : Set = insert(e, empty);
        __ ∪ __      : Set × Set → Set;
        remove      : Elem × Set → Set
  ∀e, e' : Elem; S, S' : Set
    • ¬(e is_in empty)
    • e is_in insert(e', S) ⇔ (e = e' ∨ e is_in S)
    • S = S' ⇔ (∀x : Elem • x is_in S ⇔ x is_in S')      %(equal_sets)%
    • e is_in (S ∪ S') ⇔ (e is_in S ∨ e is_in S')
    • e is_in remove(e', S) ⇔ (¬(e = e') ∧ e is_in S)
then %implies
  ∀e, e' : Elem; S : Set
    • insert(e, insert(e, S)) = insert(e, S)
    • insert(e, insert(e', S)) = insert(e', insert(e, S))
  generated type Set ::= empty | {__}(Elem) | __ ∪ __ (Set; Set)
  op __ ∪ __ : Set × Set → Set, assoc, comm, idem, unit empty
end

```

Although generated specifications are in general loose, they need not be so, as illustrated by the above `GENERATED_SET` specification, where the axiom `%(equal_sets)%`, combined with the axioms defining `is_in`, fully constrains (up to isomorphism) the interpretations of the sort `Set` and of the constructors `empty` and `insert`, once an interpretation for the sort `Elem` is chosen.

Note also that this example displays the power of the annotation `%implies`. Remember that this annotation applies to the whole of the specification extension where it occurs, so here it applies not only to the two explicit axioms about `insert`, but also to the properties corresponding to the attributes of ‘`∪`’ as well as to the generatedness constraint. Hence, the `%implies` annotation is used here not only to stress that the usual properties of `insert` are expected to follow from the preceding declarations and axioms, but also that an alternative induction scheme, based on `empty`, `{__}` and `__ ∪ __`, can be used for sets. Moreover, it asserts that `__ ∪ __` is expected to be associative, commutative, idempotent (i.e.,  $S \cup S = S$ ), and to have `empty` as unit. Note again that this `%implies` part heavily relies on the ‘same name, same meaning’ principle.

*Generated types may need to be declared together.*

The following specification fragment illustrates what may go wrong.

```
sort Node
generated type Tree ::= mktree(Node; Forest)
generated type Forest ::= empty | add(Tree; Forest)
```

The above is *incorrect*, due to the linear visibility rule. This can easily be fixed by replacing ‘**sort** Node’ by ‘**sorts** Node, Tree, Forest’. Even when corrected, the above is *wrong*, since the corresponding semantics is not what a naive reader may expect. One may expect that only models where the carrier sets of the sorts *Tree* and *Forest* are generated by *mktree*, *empty* and *add* are acceptable, but more models satisfy the above two *separate* sort generatedness constraints. For instance, a model with both a junk tree *jt* and a junk forest *jf* fulfills the above declarations (assuming that the interpretations of *mktree* and *add* on *jt* and *jf* in this model are such that  $jt = mktree(n, jf)$  for any node  $n$  and that  $jf = add(jt, jf)$ ). Hence, one must write instead:

```
sort Node
generated types Tree ::= mktree(Node; Forest);
               Forest ::= empty | add(Tree; Forest)
```

Here, the mutually recursive datatypes *Tree* and *Forest* are correctly defined simultaneously within the same **generated types** construct, and the resulting semantics is the expected one (without junk values for trees and forests). Note that therefore, the linear visibility rule is *not* applicable *within* a **generated types** construct (to allow such mutually recursive definitions), but that this is the only exception to the linear visibility principle. Only mutually recursive generated datatypes need to be declared together; in simpler cases, it makes no difference to have a sequence of successive generated datatype declarations or just one introducing all the desired datatypes.<sup>6</sup>

### 3.3 Free Specifications

*Free specifications provide initial semantics and avoid the need for explicit negation.*

```
spec NATURAL = free type Nat ::= 0 | suc(Nat)
```

<sup>6</sup> The same explanations apply to free datatypes, introduced in the next subsection.

A **free** datatype declaration corresponds to the so-called ‘*no junk, no confusion*’ principle: there are no other values of sort *Nat* than those denoted by the constructor terms built with *0* and *suc*, and all distinct constructor terms denote different values.

Hence, a **free** datatype declaration such as the one above has *exactly* the same effect as the corresponding **generated** datatype declaration, *together* with axioms stating that *suc* is injective, and that *0* cannot be the successor of a natural number. An alternative to the above ‘**free type** *Nat* ::= *0* | *suc*(*Nat*)’ is therefore:

$$\begin{aligned} &\textbf{generated type } Nat ::= 0 \mid suc(Nat) \\ &\forall x, y : Nat \bullet suc(x) = suc(y) \Rightarrow x = y \\ &\forall x : Nat \bullet \neg(0 = suc(x)) \end{aligned}$$

*Free datatype declarations are particularly convenient for defining enumerated datatypes.*

```
spec COLOR =
  free type RGB ::= Red | Green | Blue
  free type CMYK ::= Cyan | Magenta | Yellow | Black
end
```

Using ‘**free**’ instead of ‘**generated**’ for defining enumerated datatypes saves the writing of many explicit distinctness assertions (for instance, here,  $\neg(Red = Green)$ ,  $\neg(Red = Blue)$ , ...).

*Free specifications can also be used when the constructors are related by some axioms.*

```
spec INTEGER =
  free { type Int ::= 0 | suc(Int) | pre(Int)
        \forall x : Int \bullet suc(pre(x)) = x
        \bullet pre(suc(x)) = x }
end
```

When some relations are to be imposed between the constructors (as is the case here for *suc* and *pre* which are inverses of each other), a **free** datatype declaration cannot be used, since the contradiction between the ‘no confusion’ principle and the axioms imposed on the constructors would lead to an inconsistent specification. Instead, one should impose a ‘*freeness constraint*’ around the datatype declaration followed by the required axioms. A freeness

constraint, expressed by the keyword **free**, can be imposed around any specification.

In the case of the above **INTEGER** specification, the freeness constraint imposes that the semantics of the specification is the class of all algebras isomorphic to the quotient of the constructor terms by (the minimal congruence induced by) the given axioms. This is exactly the desired semantics. More generally, a freeness constraint around a specification indicates its initial model, which may not exist, of course. It is however well-known that initial models of basic specifications with axioms restricted to Horn clauses (of which equations as in **INTEGER** are a special case) always exist.<sup>7</sup> Remember also that equality holds minimally in initial models of equational specifications.

*Predicates hold minimally in models of free specifications.*

```
spec NATURAL_ORDER =
  NATURAL
then free { pred _ < _ : Nat × Nat
            ∀x, y : Nat
            • 0 < suc(x)
            • x < y ⇒ suc(x) < suc(y) }
end
```

A freeness constraint imposed around a predicate declaration followed by some defining axioms has the effect that the predicate only holds when this follows from the given axioms, and does not hold otherwise. For instance, in the above example, it is not necessary to explicitly state that ‘ $\neg(0 < 0)$ ’, since this will follow from the imposed freeness constraint. Hence, in such cases a freeness constraint has exactly the same effect as the so-called ‘*negation as failure*’ or ‘*closed world assumption*’ principles in logic programming. More generally, it is often convenient to define a predicate within a freeness constraint, since by doing so, one has to specify the ‘positive’ cases only.

*Operations and predicates may be safely defined by induction on the constructors of a free datatype declaration.*

```
spec NATURAL_ARITHMETIC =
  NATURAL_ORDER
```

<sup>7</sup> Strictly speaking, existence of initial models depends on a further requirement: namely the existence of a ground term for each sort. This ensures that the term-algebra has non-empty carriers and hence is a CASL model.



```

then ops 1      :  $Nat = suc(0)$ ;
          -- + -- :  $Nat \times Nat \rightarrow Nat$ , assoc, comm, unit 0;
          -- * -- :  $Nat \times Nat \rightarrow Nat$ , assoc, comm, unit 1
           $\forall x, y : Nat$ 
          •  $x + suc(y) = suc(x + y)$ 
          •  $x * 0 = 0$ 
          •  $x * suc(y) = (x * y) + x$ 
end

```

To define some operation on a free datatype, it is generally recommended to make a case distinction with respect to the various constructors defined. This is illustrated by the above definitions of ‘+’ and ‘\*’ (although for the ‘+’ operation, the case for the constructor  $0$  is already taken care of by the attribute ‘*unit 0*’).<sup>8</sup>

*More care may be needed when defining operations or predicates on free datatypes when there are axioms relating the constructors.*

```

spec INTEGER_ARITHMETIC =
  INTEGER
then ops 1      :  $Int = suc(0)$ ;
          -- + -- :  $Int \times Int \rightarrow Int$ , assoc, comm, unit 0;
          -- - -- :  $Int \times Int \rightarrow Int$ ;
          -- * -- :  $Int \times Int \rightarrow Int$ , assoc, comm, unit 1
           $\forall x, y : Int$ 
          •  $x + suc(y) = suc(x + y)$ 
          •  $x + pre(y) = pre(x + y)$ 
          •  $x - 0 = x$ 
          •  $x - suc(y) = pre(x - y)$ 
          •  $x - pre(y) = suc(x - y)$ 
          •  $x * 0 = 0$ 
          •  $x * suc(y) = (x * y) + x$ 
          •  $x * pre(y) = (x * y) - x$ 
end

```

While a case distinction with respect to the constructors of a free datatype is harmless, this may not be the case for a datatype defined within a freeness constraint, since, due to the axioms relating the constructors to each other, some cases may overlap. This does not mean, however, that one cannot use the case distinction, but just that more attention should be paid than for a free datatype – since one needs to ensure that the definitions lead to the same

<sup>8</sup> In specification libraries, ordinary decimal notation for natural numbers can be provided by use of so-called literal syntax annotations, see Chap. 9.

results for overlapping cases. For instance, in the above example no problem arises. But one should be more careful with the next one, since a negative integer can be of the form  $\text{suc}(x)$ , hence asserting, e.g.,  $0 \leq \text{suc}(x)$ , would of course be wrong.

```

spec INTEGER_ARITHMETIC_ORDER =
  INTEGER_ARITHMETIC
then preds  $-- \leq --, -- \geq --, -- < --, -- > -- : Int \times Int$ 
   $\forall x, y : Int$ 
  •  $0 \leq 0$ 
  •  $\neg(0 \leq \text{pre}(0))$ 
  •  $0 \leq x \Rightarrow 0 \leq \text{suc}(x)$ 
  •  $\neg(0 \leq x) \Rightarrow \neg(0 \leq \text{pre}(x))$ 
  •  $\text{suc}(x) \leq y \Leftrightarrow x \leq \text{pre}(y)$ 
  •  $\text{pre}(x) \leq y \Leftrightarrow x \leq \text{suc}(y)$ 
  •  $x \geq y \Leftrightarrow y \leq x$ 
  •  $x < y \Leftrightarrow (x \leq y \wedge \neg(x = y))$ 
  •  $x > y \Leftrightarrow y < x$ 
end

```

*Generic specifications often involve free extensions of (loose) parameters.*

```

spec LIST [sort Elem] = free type List ::= empty | cons(Elem; List)

```

The parameter of a generic specification should be loose to cope with the various expected instantiations. On the other hand, it is a frequent situation that the body of the generic specification should have a free, initial interpretation. This is illustrated by the above example, where we want to combine a loose interpretation for the sort *Elem* with a free interpretation for lists. The following example is similar in spirit.

```

spec SET [sort Elem] =
  free { type Set ::= empty | insert(Elem; Set)
    pred _is_in_ : Elem  $\times$  Set
     $\forall e, e' : Elem; S : Set$ 
    •  $\text{insert}(e, \text{insert}(e, S)) = \text{insert}(e, S)$ 
    •  $\text{insert}(e, \text{insert}(e', S)) = \text{insert}(e', \text{insert}(e, S))$ 
    •  $\neg(e \text{ is\_in } \text{empty})$ 
    •  $e \text{ is\_in } \text{insert}(e, S)$ 
    •  $e \text{ is\_in } \text{insert}(e', S) \text{ if } e \text{ is\_in } S$  }
end

```

As for the LIST example, we want to have a loose interpretation for the sort *Elem* and a free interpretation for sets. Since some axioms are required to hold for the *Set* constructors *empty* and *insert*, we cannot use a free datatype declaration, hence we use a freeness constraint.

Note that since, as already explained, predicates hold minimally in models of free specifications, it would have been enough, in the above example, to define the predicate *is\_in* by the sole axiom  $e \text{ is\_in } \text{insert}(e, S)$ .<sup>9</sup> However, doing so would have decreased the comprehensibility of the specification and this is the reason why we have preferred a more verbose axiomatization of the predicate *is\_in*.

Note also the use of the keyword ‘*if*’ to write an implication in the reverse order:

$$e \text{ is\_in } \text{insert}(e', S) \text{ if } e \text{ is\_in } S$$

is equivalent to:

$$e \text{ is\_in } S \Rightarrow e \text{ is\_in } \text{insert}(e', S)$$

The following example specifies the transitive closure of an arbitrary binary relation *R* on some sort *Elem* (both provided by the parameter).

```
spec TRANSITIVE_CLOSURE [sort Elem pred _R_ : Elem × Elem] =
  free { pred _R+_ : Elem × Elem
        ∀x, y, z : Elem
        • x R y ⇒ x R+ y
        • x R+ y ∧ y R+ z ⇒ x R+ z }
```

In the above example, it is crucial that predicates hold minimally in models of free specifications, since this property ensures that what we define as ‘*R+*’ is actually the *smallest* transitive relation including *R*. Without requiring the freeness constraint, one would allow arbitrary transitive relations containing *R* (and these undesired relations cannot be eliminated merely by specifying further first-order axioms).

*Loose extensions of free specifications can avoid overspecification.*

```
spec NATURAL_WITH_BOUND =
  NATURAL_ARITHMETIC
then op max_size : Nat
  • 0 < max_size
end
```

---

<sup>9</sup> If an element *e* belongs to a set *S'*, then this set *S'* can always be denoted by a constructor term of the form *insert*(*e*, *S*), due to the axioms constraining the constructor *insert*.

The above example shows another benefit of mixing loose and initial semantics. Assume that at this stage we want to introduce some bound, of sort *Nat*, without fixing its value yet (this value is likely to be fixed later in some refinement, and all that we need for now is the existence of some bound). This is provided by the above specification `NATURAL_WITH_BOUND`, where we mix an initial interpretation for the sort *Nat* (defined using a free datatype declaration in `NATURAL`) and a loose interpretation for the constant *max\_size*. Each model of `NATURAL_WITH_BOUND` will provide a fixed interpretation of the constant *max\_size*, and all these models are captured by `NATURAL_WITH_BOUND`, which is in this sense loose. Using such loose extensions is in general appropriate to avoid unnecessary overspecification.

```
spec SET_CHOOSE [sort Elem] =
  SET [sort Elem]
then op choose : Set → Elem
  ∀S : Set • ¬(S = empty) ⇒ choose(S) is_in S
end
```

This example shows again the benefit of mixing initial and loose semantics. Here, we want to extend sets, defined using a free constraint in `SET`, by a loosely specified operation *choose*.<sup>10</sup> At this stage, the only property required for *choose* is to provide some element belonging to the set to which it is applied, and we do not want to specify more precisely which specific element is to be chosen. Note that each model of `SET_CHOOSE` will provide a function implementing some specific choice strategy, and that since all these interpretations of *choose* have to be *functions*, they are necessarily ‘deterministic’ (e.g., applied twice to the same set argument, they return the same result).

*Datatypes with observer operations or predicates can be specified as generated instead of free.*

```
spec SET_GENERATED [sort Elem] =
  generated type Set ::= empty | insert(Elem; Set)
  pred __is_in__ : Elem × Set
  ∀e, e' : Elem; S, S' : Set
    • ¬(e is_in empty)
    • e is_in insert(e', S) ⇔ (e = e' ∨ e is_in S)
    • S = S' ⇔ (∀x : Elem • x is_in S ⇔ x is_in S')
end
```

<sup>10</sup> For the purpose of this example, we disregard the fact that *choose* should be undefined on the empty set, and we just leave this case unspecified. Partial functions are discussed in Chap. 4.

The above specification is an alternative to the specification SET (see p. 40). Both SET and SET\_GENERATED define exactly the same class of models. The former specification relies on a freeness constraint, while SET\_GENERATED relies on the observer *is\_in* to specify when two sets are equal. Indeed, the last axiom of SET\_GENERATED expresses directly that two sets having exactly the same elements are equal values. This axiom, together with the first two axioms defining *is\_in*, will entail as well the expected properties on the constructor *insert* (see GENERATED\_SET p. 35). Note also that since, in SET\_GENERATED, the predicate *is\_in* is not defined within a freeness constraint, we specify when it holds using ' $\Leftrightarrow$ ' rather than a one-way implication.

While a freeness constraint may be unavoidable to define a predicate, as illustrated by TRANSITIVE\_CLOSURE, the choice between relying on a freeness constraint to define a datatype such as *Set*, or using instead a generated datatype declaration together with some observers to unambiguously determine the values of interest, is largely a matter of convenience. One may argue that SET is more suitable for prototyping tools based on term rewriting, while SET\_GENERATED is more suitable for theorem-proving tools.

*The %def annotation is useful to indicate that some operations or predicates are uniquely defined.*

```
spec SET_UNION [sort Elem] =
  SET [sort Elem]
then %def
  ops  $--\cup--$  : Set  $\times$  Set  $\rightarrow$  Set, assoc, comm, idem, unit empty;
      remove : Elem  $\times$  Set  $\rightarrow$  Set
   $\forall e, e' : \text{Elem}; S, S' : \text{Set}$ 
    •  $S \cup \text{insert}(e', S') = \text{insert}(e', S \cup S')$ 
    •  $\text{remove}(e, \text{empty}) = \text{empty}$ 
    •  $\text{remove}(e, \text{insert}(e, S)) = \text{remove}(e, S)$ 
    •  $\text{remove}(e, \text{insert}(e', S)) = \text{insert}(e', \text{remove}(e, S))$  if  $\neg(e = e')$ 
end
```

The annotation **%def** expresses that SET\_UNION is a *definitional extension* of SET, i.e., that each model of SET can be *uniquely* extended to a model of SET\_UNION, which means that the operations introduced in SET\_UNION are uniquely defined. As with the **%implies** annotation, the **%def** annotation has no impact on the semantics, but a corresponding proof obligation can be generated, to be discharged by theorem proving tools. The **%def** annotation is especially useful to stress that the specifier's intention is to impose a unique interpretation of what is defined within the scope of this annotation (once an interpretation for the part which is extended has been chosen).

*Operations can be defined by axioms involving observer operations, instead of inductively on constructors.*

```

spec  SET_UNION_1 [sort Elem] =
      SET_GENERATED [sort Elem]
then %def
  ops  -- ∪ -- : Set × Set → Set, assoc, comm, idem, unit empty;
        remove : Elem × Set → Set
  ∀e, e' : Elem; S, S' : Set
    • e is_in (S ∪ S') ⇔ (e is_in S ∨ e is_in S')
    • e is_in remove(e', S) ⇔ (¬(e = e') ∧ e is_in S)
end

```

The specification `SET_UNION_1` is an alternative to `SET_UNION` and defines exactly the same model class. While an inductive definition style was chosen for the operations ‘`∪`’ and `remove` in `SET_UNION`, in `SET_UNION_1` these operations are defined ‘implicitly’ by characterizing their results through the observer `is_in`. Note that this ‘observer’ style does not prevent us providing a unique definition of both operations, as claimed by the `%def` annotation.

Similarly to the discussion on the respective merits of `SET` and of `SET_GENERATED`, the choice between an inductive definition style and an ‘observer’ definition style is partly a matter of taste. One may argue that the ‘observer’ definition style is more abstract in the sense that there is no hint to any algorithmic computation of the so-defined operations, while the inductive definition style mimics a recursive definition in a functional programming language. Again, the inductive definition style may be more suitable for prototyping tools based on term rewriting, while the ‘observer’ definition style may be more suitable for theorem-proving tools.

*Sorts declared in free specifications are not necessarily generated by their constructors.*

```

spec  UNNATURAL =
  free { type  UnNat ::= 0 | suc(UnNat)
        op    -- + -- : UnNat × UnNat → UnNat,
                assoc, comm, unit 0
        ∀x, y : UnNat • x + suc(y) = suc(x + y)
        ∀x : UnNat • ∃y : UnNat • x + y = 0 }
end

```

This rather peculiar example illustrates the fact that a sort defined within a freeness constraint need not be generated by its constructors. In `UNNATURAL`, the specification enclosed within the `free { ... }` construct specifies

Abelian groups with one generator  $\text{succ}(0)$ , and the integers are the free such Abelian group. Hence, the (unique up to isomorphism) model of `UNNATURAL` corresponds to integers, and not to natural numbers as one may expect – just consider the last axiom. This example points out why in general datatypes defined using freeness constraints can be more difficult to understand than datatypes defined using generatedness constraints. However, the reader should be aware that the specification `UNNATURAL` uses a proper first-order formula with an existential quantifier in the axioms. The specification `UNNATURAL` is provided here for explanatory purposes only, and clearly the writing of similar specifications should be discouraged. When only Horn clauses are used as axioms in a freeness constraint, then the datatype will indeed be generated by its constructors.

## Partial Functions

*Partial functions arise naturally.*

Partial functions arise in a number of situations. CASL provides means for the declaration of partial functions, the specification of their domains of definition, and more generally the specification of system properties involving partial functions. The aim of this chapter is to discuss and illustrate how to handle partial functions in CASL specifications.

### 4.1 Declaring Partial Functions

*Partial functions are declared differently from total functions.*

```
spec SET_PARTIAL_CHOOSE [sort Elem] =
    GENERATED_SET [sort Elem]
then op choose : Set →? Elem
end
```

The *choose* function on sets is naturally a partial function, expected to be undefined on the empty set. In CASL, a partial function is declared similarly to a total one, *but* for the question mark ‘?’ following the arrow in the profile. It is therefore quite easy to distinguish the functions declared to be total from the ones declared to be partial.

A function declared to be partial may happen to be total in some of the models of the specification. For instance, the above specification SET\_PARTIAL\_CHOOSE does not exclude models where the function symbol *choose* is interpreted by a total function, defined on all set values. Axioms can be



used to specify the domain of definition of a partial function, and how to do this is detailed later in this chapter.

*Terms containing partial functions may be undefined, i.e., they may fail to denote any value.*

For instance, the (value of the) term  $choose(empty)$  may be undefined.<sup>1</sup> This is more natural than insisting that  $choose(empty)$  has to denote some arbitrary but fixed element of  $Elem$ .

Note that variables range only over defined values, and therefore a variable always denotes a value, in contrast to terms containing partial functions.

*Functions, even total ones, propagate undefinedness.*

If the term  $choose(S)$  is undefined for some value of  $S$ , then the term  $insert(choose(S), S')$  is undefined as well for this value of  $S$ , although  $insert$  is a total function.

*Predicates do not hold on undefined arguments.*

CASL is based on classical two-valued logic. A predicate symbol is interpreted by a relation, and when the value of some argument term is undefined, the application of a predicate to this term does not hold. For instance, if the term  $choose(S)$  is undefined, then the atomic formula  $choose(S) \text{ is\_in } S$  does not hold.

*Equations hold when both terms are undefined.*

In CASL, equations are by default *strong*, which means that they hold not only when both sides denote equal values, but also when both sides are simultaneously undefined. For instance, let us consider the equation:

$$insert(choose(S), insert(choose(S), empty)) = insert(choose(S), empty)$$

---

<sup>1</sup> Note that the term  $choose(empty)$  is well-formed and therefore is a ‘correct term’. It is *its value* which may be undefined. To avoid unnecessary pedantry, in the following we will simply write that a term is undefined to mean that its value is so. Obviously, a term with variables may be defined for some values of the variables and undefined for other values.

Either  $choose(S)$  is defined and then both sides are defined and denote equal values due to the axioms on  $insert$ , or  $choose(S)$  is undefined and then both sides are undefined, and the strong equation ‘holds trivially’.

CASL provides also so-called *existential equations*, explained at the end of this chapter.

*Special care is needed in specifications involving partial functions.*

Partial functions are intrinsically more difficult to understand and specify than total ones. This is why special care is needed when writing the axioms of specifications involving partial functions. The point is that an axiom may imply the definedness of terms containing partial functions, and as a consequence that these functions are total, which may not be what the specifier intended. Here are three typical cases:

- Asserting  $choose(S) \text{ is\_in } S$  as an axiom implies that  $choose(S)$  is defined, for any  $S$ . The point here is that since predicates applied to an undefined term do not hold, in any model satisfying  $choose(S) \text{ is\_in } S$ , the function  $choose$  must be total (i.e., always defined).
- Asserting  $remove(choose(S), insert(choose(S), empty)) = empty$  as an axiom implies that  $choose(S)$  is defined for any  $S$ , since the term  $empty$  is always defined. To understand this, assume that  $choose$  is undefined for some set value of  $S$ ; then the above equation cannot hold for this value, since the undefinedness of  $choose(S)$  implies the undefinedness of  $remove(choose(S), insert(choose(S), empty))$ , giving a contradiction with the definedness of  $empty$ . Hence, an equation between a term involving a partial function  $PF$  and a term involving total functions only may imply that the partial function  $PF$  is always defined.
- Asserting  $insert(choose(S), S) = S$  as an axiom implies that  $choose(S)$  is defined for any  $S$ , since a variable always denotes a defined value. This case is indeed similar to the previous one, the only difference being that now the right-hand side of the equation is a variable (instead of a term involving total functions only).

Moreover, the ‘same name, same thing’ principle has a subtle side-effect regarding partial operations: if an operation is declared both as a total operation and as a partial operation with the same profile (i.e., the same argument sorts and the same result sort) then it is interpreted as a total operation in all models of the specification.

## 4.2 Specifying Domains of Definition

*The definedness of a term can be checked or asserted.*

```
spec  SET_PARTIAL_CHOOSE_1 [sort Elem] =
      SET_PARTIAL_CHOOSE [sort Elem]
then  •  $\neg \text{def choose}(\text{empty})$ 
       $\forall S : \text{Set} \bullet \text{def choose}(S) \Rightarrow \text{choose}(S) \text{ is\_in } S$ 
end
```

A definedness assertion, written ‘*def t*’, where *t* is a term, is a special kind of atomic formula: it holds if and only if the value of the term *t* is defined. For instance, in the above example,  $\neg \text{def choose}(\text{empty})$  explicitly asserts that *choose* is undefined when applied to *empty*. Note that this axiom does not say anything about the definedness of *choose* applied to values other than *empty*, which means that *choose* may well be undefined on those values too. The second axiom of the above example asserts *choose(S) is\_in S* under the condition *def choose(S)*, to avoid undesired definedness induced by axioms, as explained in the previous section.

Note that if the two axioms of the above example were to be replaced by:

$$\forall S : \text{Set} \bullet \neg(S = \text{empty}) \Rightarrow \text{choose}(S) \text{ is\_in } S$$

then we could conclude that *choose(S)* is defined when *S* is not equal to *empty*, but nothing about the undefinedness of *choose(empty)*.

*The domains of definition of partial functions can be specified exactly.*

```
spec  SET_PARTIAL_CHOOSE_2 [sort Elem] =
      SET_PARTIAL_CHOOSE [sort Elem]
then   $\forall S : \text{Set} \bullet \text{def choose}(S) \Leftrightarrow \neg(S = \text{empty})$ 
       $\forall S : \text{Set} \bullet \text{def choose}(S) \Rightarrow \text{choose}(S) \text{ is\_in } S$ 
end
```

In the above example, the domain of definition of the partial function *choose* is exactly specified by the axiom  $\text{def choose}(S) \Leftrightarrow \neg(S = \text{empty})$ .

*Loosely specified domains of definition may be useful.*

```
spec NATURAL_WITH_BOUND_AND_ADDITION =
  NATURAL_WITH_BOUND
then op  $_{++?} : Nat \times Nat \rightarrow? Nat$ 
   $\forall x, y : Nat$ 
  •  $def(x+?y)$  if  $x + y < max\_size$ 
  % {  $x + y < max\_size$  implies both
     $x < max\_size$  and  $y < max\_size$  } %
  •  $def(x+?y) \Rightarrow x+?y = x + y$ 
end
```

In some cases, it is useful to loosely specify the domain of definition of a partial function, as illustrated in the above example for ‘+?’, which is required to be defined for all arguments  $x$  and  $y$  such that  $x + y < max\_size$ , but may well be defined on larger natural numbers as well. The point in loose specifications of definition domains is to avoid unnecessary constraints on the models of the specification. For instance, the above example does not exclude a model where ‘+?’ is interpreted by a total function (which would then coincide with ‘+’).<sup>2</sup>

Indeed, in some cases, specifying exactly domains of definition can be considered as overspecification. In most specifications, however, one would expect an exact specification of domains of definition, even for otherwise loosely specified functions (see, e.g., *choose* in SET\_PARTIAL\_CHOOSE\_2).

*Domains of definition can be specified more or less explicitly.*

```
spec SET_PARTIAL_CHOOSE_3 [sort Elem] =
  SET_PARTIAL_CHOOSE [sort Elem]
then •  $\neg def choose(empty)$ 
   $\forall S : Set$  •  $\neg(S = empty) \Rightarrow choose(S) is\_in S$ 
end
```

SET\_PARTIAL\_CHOOSE\_3 specifies exactly the domain of definition of *choose*, but does this too implicitly, since some reasoning is needed to conclude that the above specification entails  $def choose(S) \Leftrightarrow \neg(S = empty)$ .

<sup>2</sup> In this example, it is essential to choose a new name ‘+?’ for our partial addition operation. Otherwise, since ‘+’ is (rightly) declared as a total operation in NATURAL\_WITH\_BOUND, the declaration **op**  $_{++} : Nat \times Nat \rightarrow? Nat$  would be useless: the same name, same thing principle would lead to models with just one, total, addition operation.

To improve the clarity of specifications, it is in general advisable to specify definition domains as explicitly as possible, and `SET_PARTIAL_CHOOSE_2` is somehow easier to understand than `SET_PARTIAL_CHOOSE_3` (both specifications define the same class of models).

```
spec NATURAL_PARTIAL_PRE =
  NATURAL_ARITHMETIC
then op pre : Nat →? Nat
  • ¬ def pre(0)
  ∀x : Nat • pre(suc(x)) = x
end
```

In the above example, one can consider that the domain of definition of *pre* is (exactly) specified in an explicit enough way, since the first axiom states exactly that *pre*(0) is undefined while the second one implies that *pre* is defined for all natural numbers of the form *suc*(*x*).

```
spec NATURAL_PARTIAL_SUBTRACTION_1 =
  NATURAL_PARTIAL_PRE
then op -- : Nat × Nat →? Nat
  ∀x, y : Nat
  • x - 0 = x
  • x - suc(y) = pre(x - y)
end
```

The above specification is perfect from a mathematical point of view, but is clearly not explicit enough, since there is no easy way to infer when *x - y* is defined. From a methodological point of view, the following alternative version is much better.

```
spec NATURAL_PARTIAL_SUBTRACTION =
  NATURAL_PARTIAL_PRE
then op -- : Nat × Nat →? Nat
  ∀x, y : Nat
  • def(x - y) ⇔ (y < x ∨ y = x)
  • x - 0 = x
  • x - suc(y) = pre(x - y)
end
```

The above examples clearly demonstrate why the explicit specification of definition domains is generally advisable from a methodological point of view. However, they also indicate that this recommendation should not be applied in too strict a way, and that deciding whether a specification is explicit enough or not is to some extent a matter of taste.

*Partial functions are minimally defined by default in free specifications.*

```

spec LIST_SELECTORS_1 [sort Elem] =
  LIST [sort Elem]
then free { ops head : List →? Elem;
            tail : List →? List
             $\forall e : \textit{Elem}; L : \textit{List}$ 
            • head(cons(e, L)) = e
            • tail(cons(e, L)) = L }
end

```

In the above example, the given axioms imply that *head* and *tail* are defined on lists of the form *cons*(*e*, *L*). The freeness constraint requires that these functions are minimally defined. Since the terms *head*(*empty*) and *tail*(*empty*) are not equated to any other term, the freeness constraint implies that these terms are undefined, and hence that the functions *head* and *tail* are undefined on *empty*. The situation here is similar to the fact that predicates hold minimally in models of free specifications (see Chap. 3, p. 38).

```

spec LIST_SELECTORS_2 [sort Elem] =
  LIST [sort Elem]
then ops head : List →? Elem;
        tail : List →? List
         $\forall e : \textit{Elem}; L : \textit{List}$ 
        •  $\neg \textit{def } \textit{head}(\textit{empty})$ 
        •  $\neg \textit{def } \textit{tail}(\textit{empty})$ 
        • head(cons(e, L)) = e
        • tail(cons(e, L)) = L
end

```

The above specification LIST\_SELECTORS\_2 is an alternative to LIST\_SELECTORS\_1; both specifications define exactly the same class of models. However, LIST\_SELECTORS\_2 is clearly easier to understand and can be considered as technically simpler, since it involves no freeness constraint.

Operations like *head* and *tail* are usually called *selectors*, and CASL provides abbreviations to specify selectors in a very concise way, as we see next.

### 4.3 Partial Selectors and Constructors

*Selectors can be specified concisely in datatype declarations, and are usually partial.*

```
spec LIST_SELECTORS [sort Elem] =
  free type List ::= empty | cons(head :? Elem; tail :? List)
```

The above free datatype declaration introduces, in addition to the constructors *empty* and *cons*, two partial selectors *head* and *tail* yielding the respective arguments of the constructor *cons*. Hence, this free datatype declaration with selectors has exactly the same effect as the ordinary free datatype declaration **free type** *List* ::= *empty* | *cons*(*Elem*; *List*), together with the operation declarations and axioms of LIST\_SELECTORS\_2 (i.e., LIST\_SELECTORS and LIST\_SELECTORS\_2 define exactly the same class of models). The following example is similar in spirit.

```
spec NATURAL_SUC_PRE = free type Nat ::= 0 | suc(pre :? Nat)
```

*Selectors are usually total when there is only one constructor.*

```
spec PAIR_1 [sorts Elem1, Elem2] =
  free type Pair ::= pair(first : Elem1; second : Elem2)
```

While selectors are usually partial operations when there is more than one alternative in the corresponding datatype declaration, they can be total, and this is generally the case when there is only one constructor, as in the above example. The free datatype declaration entails in particular axioms asserting that *first* and *second* yield the respective arguments of the constructor *pair* (i.e.,  $first(pair(e1, e2)) = e1$  and  $second(pair(e1, e2)) = e2$ ).

*Constructors may be partial.*

```
spec PART_CONTAINER [sort Elem] =
  generated type
    P_Container ::= empty | insert(Elem; P_Container)?
  pred addable : Elem × P_Container
  vars e, e' : Elem; C : P_Container
  • def insert(e, C) ⇔ addable(e, C)
```

```

pred is_in : Elem × P_Container
  •  $\neg(e \text{ is\_in } \text{empty})$ 
  •  $(e \text{ is\_in } \text{insert}(e', C) \Leftrightarrow (e = e' \vee e \text{ is\_in } C)) \text{ if } \text{addable}(e', C)$ 
end

```

The intention in the above example is to define a reusable specification of partial containers. The *insert* constructor is specified as a partial operation, defined if some condition on both the element  $e$  to be added and the container  $C$  to which the element is to be added holds. This condition is abstracted here in a predicate *addable*, so far left unspecified. Later on, instantiations of the PART\_CONTAINER specification can be adapted to specific purposes by extending them with axioms defining *addable*.

The above generated datatype declaration abbreviates as usual the declaration of a sort *P\_Container*, a constant constructor *empty*, and a partial constructor *insert* : *Elem* × *P\_Container*  $\rightarrow?$  *P\_Container*. It also entails the corresponding generatedness constraint.

## 4.4 Existential Equality

*Existential equality requires the definedness of both terms as well as their equality.*

```

spec NATURAL_PARTIAL_SUBTRACTION_2 =
      NATURAL_PARTIAL_SUBTRACTION_1
then  $\forall x, y, z : \text{Nat} \bullet y - x \stackrel{e}{=} z - x \Rightarrow y = z$ 
      % {  $y - x = z - x \Rightarrow y = z$  would be wrong,
         $\text{def}(y - x) \wedge \text{def}(z - x) \wedge y - x = z - x \Rightarrow y = z$ 
        is correct, but better abbreviated in the above axiom } %
end

```

An *existential* equation  $t1 \stackrel{e}{=} t2$  is equivalent to  $\text{def}(t1) \wedge \text{def}(t2) \wedge t1 = t2$ , so it holds if and only if both terms  $t1$  and  $t2$  are defined and denote the same value. Existential equality ' $\stackrel{e}{=}$ ' is input as ' $\text{=e=}$ '.

Note that while a trivial strong equation of the form  $t = t$  always holds, this is not the case for existential equations. For instance, the trivial existential equation  $x - y \stackrel{e}{=} x - y$  does not hold, since the term  $x - y$  may be undefined.

In general consequences of undefinedness are undesirable. Hence a conditional equation of the form  $t1 = t2 \Rightarrow t3 = t4$  is often wrong if  $t1$  and  $t2$  may be undefined, because the equality  $t3 = t4$  would be implied when both  $t1$  and  $t2$  are undefined (since then the strong equation  $t1 = t2$  would hold). The above specification provides a typical example of such a situation:  $y - x = z - x \Rightarrow y = z$  would be wrong, since it would entail that any two



arbitrary values  $y$  and  $z$  are equal (it is enough to choose an  $x$  greater than  $y$  and  $z$  to make  $y - x$  and  $z - x$  both undefined).

Therefore, to avoid such undesirable consequences of undefinedness, it is advisable to use existential equations instead of strong equations in the premises of conditional equations involving partial operations. An alternative is to add the relevant definedness assertions explicitly to the equations in the premises.

## Subsorting

*Subsorts and supersorts are often useful in CASL specifications.*

Many examples naturally involve subsorts and supersorts. CASL provides means for the declaration of a sort as a subsort of another one when the values of the subsort are regarded a special case of those in the other sort. The aim of this chapter is to discuss and illustrate how to handle subsorts and supersorts in CASL specifications.

### 5.1 Subsort Declarations and Definitions

*Subsort declarations directly express relationships between carrier sets.*

```
spec  GENERIC_MONOID_1 [sort Elem] =
      sorts Elem < Monoid
      ops  1      : Monoid;
          _ * _ : Monoid × Monoid → Monoid, assoc, unit 1
end
```

The above example declares the sort *Elem* to be a subsort of *Monoid*, or, symmetrically, the sort *Monoid* to be a supersort of *Elem*. Hence the specification `GENERIC_MONOID_1` is quite similar to the specification `GENERIC_MONOID` given in Chap. 3, p. 30, the only difference being the use of a subsorting relation between *Elem* and *Monoid* instead of an explicit *inj* operation to embed values of sort *Elem* into values of sort *Monoid*.

In contrast to most other algebraic specification languages providing subsorting facilities, subsorts in CASL are interpreted by arbitrary *embeddings* between the corresponding carrier sets. In the above example, the subsort declaration  $Elem < Monoid$  induces an implicit (unnamed) embedding from the carrier of the sort  $Elem$  into the carrier of the sort  $Monoid$ . Thus the main difference between `GENERIC_MONOID` and `GENERIC_MONOID_1` is that the embedding is explicit and named *inj* in `GENERIC_MONOID` while it is implicit in `GENERIC_MONOID_1`.

Note that interpreting subsorting relations by embeddings rather than inclusions does not exclude models where the (carrier of the) subsort happens to be a subset of (the carrier of) the supersort, and the embedding a proper inclusion. Embeddings are just slightly more general than inclusions, and technically not much more complex.

*Operations declared on a sort are automatically inherited by its subsorts.*

```
spec VEHICLE =
  NATURAL
then sorts Car, Bicycle < Vehicle
  ops   max_speed      : Vehicle → Nat;
        weight         : Vehicle → Nat;
        engine_capacity : Car → Nat
end
```

The above example introduces three sorts, *Car*, *Bicycle* and *Vehicle*, and declares both *Car* and *Bicycle* to be subsorts of *Vehicle*. A subsort declaration entails that any term of a subsort is also a term of the supersort, so here, any term of sort *Car* is also a term of sort *Vehicle*, and we can apply the operations *max\_speed* and *weight* to it (and similarly for a term of sort *Bicycle*).

In other words, with the single declaration  $max\_speed : Vehicle \rightarrow Nat$ , we get the effect of having declared also two other operations,  $max\_speed : Car \rightarrow Nat$  and  $max\_speed : Bicycle \rightarrow Nat$ .<sup>1</sup>

Obviously, operations that are only meaningful for some subsort should be defined at the appropriate level. This is the case here for the operation *engine\_capacity*, which is only relevant for cars, and therefore defined with the appropriate profile exploiting the subsort *Car*.

---

<sup>1</sup> Strictly speaking, there is just one *max\_speed* operation in the signature of `VEHICLE`. The difference between the kind of inheritance described here and operations actually declared on subsorts becomes important when writing symbol maps, see Chap. 7.

*Inheritance applies also for subsorts that are declared afterwards.*

```
spec MORE_VEHICLE = VEHICLE then sorts Boat < Vehicle
```

The order in which a subsort and an operation on the supersort are declared is irrelevant. In `MORE_VEHICLE`, we introduce a further subsort *Boat* of *Vehicle*, and as a consequence, we again get the effect of having both *max\_speed* and *weight* available for boats, as was already the case for cars and bikes.

*Subsort membership can be checked or asserted.*

```
spec SPEED_REGULATION =  
  VEHICLE  
then ops speed_limit : Vehicle → Nat;  
          car_speed_limit, bike_speed_limit : Nat  
  ∀v : Vehicle  
    • v ∈ Car ⇒ speed_limit(v) = car_speed_limit  
    • v ∈ Bicycle ⇒ speed_limit(v) = bike_speed_limit  
end
```

A subsort membership assertion, written ‘ $t \in s$ ’, where  $t$  is a term and  $s$  is a sort, is a special kind of atomic formula: it holds if and only if the value of the term  $t$  is the embedding of some value of sort  $s$ . For instance, in the above example,  $v \in \text{Car}$  holds if and only if  $v$  denotes a vehicle which is the embedding of a car value. Note that ‘ $\in$ ’ is input as ‘`in`’, but displayed as ‘ $\in$ ’.

*Datatype declarations can involve subsort declarations.*

The sequence of declarations:

```
sorts Car, Bicycle, Boat  
type Vehicle ::= sort Car | sort Bicycle | sort Boat
```

is equivalent to the declaration `sorts Car, Bicycle, Boat < Vehicle`. There may be some values of sort *Vehicle* which are not the embedding of any value of sort *Car*, *Bicycle*, or *Boat*. Intuitively, the above datatype declaration just means that *Vehicle* ‘contains’ the *union* (which may not be disjoint) of *Car*, *Bicycle* and *Boat*. Note that the subsorts used in the datatype declaration must already be declared beforehand.

The sequence of declarations:

```
sorts Car, Bicycle, Boat
generated type Vehicle ::= sort Car | sort Bicycle | sort Boat
```

is more restrictive, since the generatedness constraint implies that any value of the supersort *Vehicle* must be the embedding of some value of the declared subsorts *Car*, *Bicycle* and *Boat*. Intuitively, the above datatype declaration means that *Vehicle* ‘is exactly’ the union (which again may not be disjoint) of *Car*, *Bicycle* and *Boat*. In particular, this declaration prevents subsequent introduction of further subsorts (unless the values of the new subsorts are intended to correspond to some values of the already declared subsorts). For instance, if we were now to extend the above specification with **sorts** *Plane* < *Vehicle*, all values of sort *Plane* would have to correspond to *Car*, *Bicycle* or *Boat* values (which is presumably not what we were intending).

The sequence of declarations:

```
sorts Car, Bicycle, Boat
free type Vehicle ::= sort Car | sort Bicycle | sort Boat
```

entails the same generatedness constraint as in the previous example, and, moreover, the freeness constraint requires that there is no ‘common’ value in the subsorts of *Vehicle*. Intuitively, the above declaration means that *Vehicle* ‘is exactly’ the *disjoint* union of *Car*, *Bicycle* and *Boat*. This means in particular that the introduction of a further common subsort of both *Car* and *Boat* (say, **sorts** *Amphibious* < *Car, Boat*) is impossible.

*Subsorts may also arise as classifications of previously specified values, and their values can be explicitly defined.*

```
spec NATURAL_SUBSORTS =
  NATURAL_ARITHMETIC
then pred even : Nat
  • even(0)
  •  $\neg$  even(1)
   $\forall n : \text{Nat} \bullet \text{even}(\text{suc}(\text{suc}(n))) \Leftrightarrow \text{even}(n)$ 
sort Even = {x : Nat • even(x)}
sort Prime = {x : Nat •  $1 < x \wedge$ 
   $\forall y, z : \text{Nat} \bullet x = y * z \Rightarrow y = 1 \vee z = 1$ }
end
```

The subsort definition **sort** *Even* = {*x* : *Nat* • *even*(*x*)} is equivalent to the declaration of a subsort *Even* of *Nat* (i.e., **sorts** *Even* < *Nat*) together with the assertion  $\forall x : \text{Nat} \bullet x \in \text{Even} \Leftrightarrow \text{even}(x)$ .

The main advantage of defining the subsort *Even* in addition to the predicate *even* is that we may then use the subsort when declaring operations (e.g., **op** *times2* : *Nat* → *Even*) and variables.

The subsort definition for *Prime* above illustrates that it is not always necessary to introduce and define an explicit predicate characterizing the values of the subsort: the formula used in a subsort definition is not restricted to predicate applications. In fact whenever a (unary) predicate *p* on a sort *s* could be defined by **pred**  $p(x : s) \Leftrightarrow f$  for some formula *f*, we may instead define **sort**  $P = \{x : s \bullet f\}$ , and use sort membership assertions  $t \in P$  instead of predicate applications  $p(t)$ , avoiding the introduction of the predicate *p* altogether.

The following example is a further illustration of subsort definitions. We declare a subsort *Pos* of *Nat* and ensure that values of *Pos* correspond to non-zero values of *Nat*. (Several alternative ways of specifying the sort *Pos* will be considered later in this section.)

```
spec POSITIVE =
  NATURAL_PARTIAL_PRE
then sort Pos = {x : Nat • ¬(x = 0)}
```

## 5.2 Subsorts and Overloading

*It may be useful to redeclare previously defined operations, using the new subsorts introduced.*

```
spec POSITIVE_ARITHMETIC =
  POSITIVE
then ops 1      : Pos;
        suc     : Nat → Pos;
        -- + --, -- * -- : Pos × Pos → Pos;
        -- + -- : Pos × Nat → Pos;
        -- + -- : Nat × Pos → Pos
end
```

Since, in *POSITIVE*, we have defined *Pos* as a subsort of *Nat*, all operations defined on natural numbers, like *suc*, ‘+’ and ‘\*’ (see *NATURAL\_ARITHMETIC* in Chap. 3, p. 38, which is extended into *NATURAL\_PARTIAL\_PRE* in Chap. 4, p. 52), are automatically inherited by *Pos* and can be applied to terms of sort *Pos*. However, according to their declarations, these operations, when applied to terms of sort *Pos*, yield results of sort *Nat*. To indicate that these results always correspond to values in the subsort *Pos*, it is necessary to explicitly *overload* these operations by similar ones with the appropriate profiles. This is

the aim of the first three lines of operation declarations in the above example. The last two operation declarations further overload ‘+’ to specify that ‘+’ also yields a result of sort *Pos* as soon as one of its arguments is a term of sort *Pos*.

It is quite important to understand that the above overloading declarations are enough to achieve the desired effect, and that *no axioms are necessary*. The fundamental rule is that, in models of CASL specifications with subsorting, embedding and overloading have to be compatible: embeddings commute with overloaded operations. This can be rephrased into the following intuitive statement: *If terms look the same, then they have the same value in the same sort*. Thus, in our example, the value of ‘ $1 + 1$ ’ is the same in *Nat* whatever the combination of the overloaded constant ‘1’ and operation ‘+’ is chosen, and there is no need for any axiom to ensure this, since this is implicit in the semantics of subsorting.

### 5.3 Subsorts and Partiality

*A subsort may correspond to the definition domain of a partial function.*

```
spec POSITIVE_PRE =
  POSITIVE_ARITHMETIC
then op pre : Pos → Nat
```

Since we have introduced the subsort *Pos* of non-zero natural numbers, it makes sense to overload the partial *pre* operation on *Nat* by a *total* one on *Pos*, as illustrated above, to emphasize the fact that indeed *pre* is a total operation on its definition domain. Note again that no further axiom is necessary, and that the semantics of subsorting will ensure that both the partial and total *pre* operations will give the same value when applied to the same non-zero value.<sup>2</sup>

*Using subsorts may avoid the need for partial functions.*

```
spec NATURAL_POSITIVE_ARITHMETIC =
  free types Nat ::= 0 | sort Pos;
           Pos ::= suc(pre : Nat)
```

---

<sup>2</sup> This should not be confused with the same name, same meaning principle, which does not apply here: the total *pre* and the partial one have different profiles, and hence are just overloaded.

```

ops 1 : Pos = suc(0);
      -- + -- : Nat × Nat → Nat,  assoc,  comm,  unit 0;
      -- * -- : Nat × Nat → Nat,  assoc,  comm,  unit 1;
      -- + --, -- * -- : Pos × Pos → Pos;
      -- + -- : Pos × Nat → Pos;
      -- + -- : Nat × Pos → Pos
      ∀x, y : Nat
      • x + suc(y) = suc(x + y)
      • x * 0 = 0
      • x * suc(y) = x + (x * y)
end

```

It is indeed tempting to exploit subsorting to avoid the declaration of partial functions, as illustrated by the above NATURAL\_POSITIVE\_ARITHMETIC specification, which is an alternative to POSITIVE\_PRE and avoids the introduction of the partial predecessor operation. Note that in the above example, we have fully used the facilities for defining free datatypes with subsorts, and in particular non-linear visibility for the declared sorts, which allows us to refer to the subsort *Pos* in the first line before defining it in the second one.

Avoiding the introduction of the partial predecessor operation has some drawbacks, since some previously well-formed terms (with defined values) have now become ill-formed, e.g.,  $pre(pre(suc(1)))$ , where  $pre(suc(1))$  is a (well-formed) term of sort *Nat* but  $pre$  expects an argument of sort *Pos*. (The fact that  $pre(suc(1)) = 1$  is a consequence of the specified axioms, and that  $1$  is of sort *Pos*, does not of course entail that  $pre(suc(1))$  is of sort *Pos* too, since axioms are disregarded when checking for well-formedness.) See below for possible workarounds using explicit casts. Moreover, it is not always possible, or easy, to avoid the declaration of partial operations by using appropriate subsorts—just consider, e.g., subtraction on natural numbers.

As a last remark on this issue, the reader should be aware of the fact that, while overloading a partial operation on a supersort (say,  $pre$  on *Nat*) with a total one on a subsort ( $pre$  on *Pos*) is fine, overloading a total operation on a supersort with a partial one on a subsort forces the partial operation to be total, and hence, the latter would be better declared as total too.<sup>3</sup>

*Casting a term from a supersort to a subsort is explicit and the value of the cast may be undefined.*

<sup>3</sup> Overloading a total *cons* on *List* with a partial *cons* on the subsort *OrderedList* would either lead to a total *cons* operation on *OrderedList*, or to an inconsistent specification, depending on how the definition domain of the partial *cons* is specified.



In CASL, a term of a subsort can always be considered as a term of any supersort, and embeddings are implicit. On the contrary, casting a term from a supersort to a subsort is explicit, and since casting is essentially a partial operation, the resulting casted term may not denote any value. Casting a term  $t$  to a sort  $s$  is written  $t \text{ as } s$ . Consider the term  $\text{pre}(\text{pre}(\text{suc}(1)) \text{ as } \text{Pos})$  which is well-formed in the context of the `NATURAL_POSITIVE_ARITHMETIC` specification. This term does indeed denote a value, but the value is not a positive natural number, so the value of the term  $\text{pre}(\text{pre}(\text{suc}(1)) \text{ as } \text{Pos}) \text{ as } \text{Pos}$  is undefined.

Note that  $\text{def } (t \text{ as } s)$  is equivalent to  $t \in s$ , for a well-formed term  $t$  of a supersort of  $s$ .

*Supersorts may be useful when generalizing previously specified sorts.*

```

spec INTEGER_ARITHMETIC_1 =
  NATURAL_POSITIVE_ARITHMETIC
then free type Int ::= sort Nat | --(Pos)
  ops --+ -- : Int × Int → Int, assoc, comm, unit 0;
      -- - -- : Int × Int → Int;
      -- * -- : Int × Int → Int, assoc, comm, unit 1;
      abs    : Int → Nat
  ∀x : Int; n : Nat; p, q : Pos
  • suc(n) + (-1) = n
  • suc(n) + (-suc(q)) = n + (-q)
  • (-p) + (-q) = -(p + q)
  • x - 0 = x
  • x - p = x + (-p)
  • x - (-q) = x + q
  • 0 * (-q) = 0
  • p * (-q) = -(p * q)
  • (-p) * (-q) = p * q
  • abs(n) = n
  • abs(-p) = p
end

```

The specification `INTEGER_ARITHMETIC_1` extends `NATURAL_POSITIVE_ARITHMETIC` and defines the sort `Int` as a supersort of the sort `Nat`. As a consequence, terms which have two parses in sort `Int`, depending whether the embedding from `Nat` to `Int` is applied to the arguments or to the result of overloaded operations, are required by the semantics of subsorting to have the same value for both parses, and they can therefore be used without explicit disambiguation.

The situation would be quite different if one would be using a combination of `NATURAL_ARITHMETIC` and `INTEGER_ARITHMETIC` (see Chap. 3), say by extending both in a structured specification (see the next chapter for more details on structured specifications). In such a combination, a term such as  $\text{succ}(0)$  would have two parses, one in sort *Nat* and one in sort *Int*; in the absence of any subsort declaration relating *Nat* and *Int*, and hence of any implicit embedding, this term would then be ambiguous, and would require explicit disambiguation to become a well-formed term.

*Supersorts may also be used for extending the intended values by new values representing errors or exceptions.*

```
spec SET_ERROR_CHOOSE [sort Elem] =
  GENERATED_SET [sort Elem]
then sorts Elem < ElemError
  op choose : Set → ElemError
  pred _is_in_ : ElemError × Set
  ∀S : Set • ¬(S = empty) ⇒ choose(S) ∈ Elem ∧ choose(S) is_in S
end
```

The above specification `SET_ERROR_CHOOSE` is another variant of the various specifications of sets equipped with a partial choose function given in Chap. 4. This variant avoids the declaration of a partial function *choose* by using instead a supersort of *Elem*, namely *ElemError*, as the target sort of *choose*. The idea here is that values of *ElemError* which are not (embeddings of) values of *Elem* represent errors, e.g., the application of *choose* to the empty set. Note that to obtain the desired effect, it is necessary to explicitly state that  $\text{choose}(S) \in \text{Elem}$  when *S* is not the empty set; moreover, to make the term  $\text{choose}(S) \text{ is\_in } S$  well-formed, we have to explicitly overload the predicate  $\_is\_in\_ : \text{Elem} \times \text{Set}$  provided by `GENERATED_SET` by a predicate  $\_is\_in\_ : \text{ElemError} \times \text{Set}$  as shown above. This example demonstrates that avoiding partial functions by the use of ‘error supersorts’ is not as innocuous as it may seem, since in general one would need to enlarge the signatures considerably by adding all required overloads.

```
spec SET_ERROR_CHOOSE_1 [sort Elem] =
  GENERATED_SET [sort Elem]
then sorts Elem < ElemError
  op choose : Set → ElemError
  ∀S : Set • ¬(S = empty) ⇒ (choose(S) as Elem) is_in S
end
```

The specification `SET_ERROR_CHOOSE_1` above is a last attempt to avoid the use of partial functions; again, we introduce a supersort *ElemError* as

in `SET_ERROR_CHOOSE`, but to avoid the need for overloading the predicate *is\_in*, we explicitly *cast* the term *choose*(*S*) in *(choose*(*S*) *as Elem*) is\_in *S*. Note that when, for some value of *S*, *(choose*(*S*) *as Elem*) is\_in *S* holds, this implies that *choose*(*S*) *as Elem* is defined, and hence that *choose*(*S*)  $\in$  *Elem* holds as well. This last version may seem preferable to the previous one. However, one should be aware that here, despite our attempt to avoid the use of partial functions, we rely on explicit casts, hence on terms that may not denote values: partiality has not been eliminated, the partial functions have merely been factorized as compositions of total functions with casting.

## Structuring Specifications

*Large and complex specifications are easily built out of simpler ones by means of (a small number of) specification-building operations.*

In the previous chapters, we have focused attention on basic specifications and detailed how to use the various constructs of CASL to write meaningful, but relatively simple, specifications. The aim of this chapter is to discuss and illustrate how to assemble simple pieces of specifications into more complex, structured ones. In particular we explain how to extend specifications, make the union of several specifications, as well as how to rename or hide symbols when assembling specifications. Parametrization and instantiation of generic specifications are explained in the next chapter.

### 6.1 Union and Extension

*Union and extension can be used to structure specifications.*

```
spec LIST_SET [sort Elem] =
  LIST_SELECTORS [sort Elem]
and GENERATED_SET [sort Elem]
then op elements_of_ : List → Set
  ∀e : Elem; L : List
  • elements_of empty = empty
  • elements_of cons(e, L) = {e} ∪ elements_of L
end
```

The above example shows how to make the union (expressed by ‘**and**’) of two specifications `LIST_SELECTORS` (see Chap. 4, p. 54) and `GENERATED_SET` (see Chap. 3, p. 35), and then further extend this union by an operation and some axioms (using ‘**then**’). Union and extension are the most commonly used specification-building operations. In contrast with extension, whose purpose is to extend a given piece of specification by new symbols and axioms, union is generally used to combine two self-contained specifications. Union of specifications is obviously associative and commutative.

All symbols introduced by a specification are by default exported by it and visible in its extensions or in its unions with other specifications. (Variables are not considered as symbols, and never exported.) Remember also the ‘*same name, same thing*’ principle: in the above `LIST_SET` specification, it is therefore the same sort *Elem* which is used to construct both lists and sets.<sup>1</sup>

*Specifications may combine parts with loose, generated, and free interpretations.*

```

spec LIST_CHOOSE [ sort Elem ] =
  LIST_SELECTORS [ sort Elem ]
and SET_PARTIAL_CHOOSE_2 [ sort Elem ]
then ops elements_of __ : List → Set;
          choose : List → ? Elem
           $\forall e : \textit{Elem}; L : \textit{List}$ 
          • elements_of empty = empty
          • elements_of cons(e, L) = {e} ∪ elements_of L
          • def choose(L) ⇔ ¬(L = empty)
          • choose(L) = choose(elements_of L)
end

spec SET_TO_LIST [ sort Elem ] =
  LIST_SET [ sort Elem ]
then op list_of __ : Set → List
           $\forall S : \textit{Set}$  • elements_of(list_of S) = S
end

```

The specification `LIST_CHOOSE` is built as an extension of the union of `LIST_SELECTORS` and `SET_PARTIAL_CHOOSE_2` (see Chap. 4, p. 50). This extension introduces an operation *elements\_of* (as in `LIST_SET`) and a partial operation *choose*, which are defined by some axioms. In `LIST_SELECTORS`, lists are defined by a free datatype construct (with selectors), hence have a

<sup>1</sup> The constant *empty* is overloaded, since we have a constant *empty* : *List* for lists and a constant *empty* : *Set* for sets.

free interpretation. `SET_PARTIAL_CHOOSE_2` is itself an extension of (`SET_PARTIAL_CHOOSE` which is an extension of) `GENERATED_SET`, where sets are defined by a generated datatype construct. However, note that as discussed in Chap. 3, p. 35, the apparently loose specification `GENERATED_SET` is in fact not so. Moreover, the *choose* partial function on sets is loosely defined in `SET_PARTIAL_CHOOSE_2`, and so is therefore also the *choose* partial function on lists defined in `LIST_CHOOSE`. It is easy to see that the operation *elements\_of* is uniquely defined. The sort *Elem* has of course a loose interpretation.

Thus the specification `LIST_CHOOSE` combines parts with a free interpretation, parts with a generated interpretation, and parts with a loose interpretation. The situation is similar to that with `LIST_SET` (and `SET_TO_LIST`), where the operation *list\_of* is loosely defined with the help of the operation *elements\_of*.

## 6.2 Renaming

*Renaming may be used to avoid unintended name clashes, or to adjust names of sorts and change notations for operations and predicates.*

```
spec  STACK [sort Elem] =
      LIST_SELECTORS [sort Elem] with sort List ↦ Stack,
                                     ops  cons ↦ push_onto_,
                                          head ↦ top,
                                          tail  ↦ pop
end
```

While the ‘same name, same thing’ principle has proven to be appropriate in numerous examples given in the previous chapters and above, it may still happen that, when combining specifications, this principle leads to unintended name clashes. An unintended name clash arises for instance when one combines two specifications that both export the same symbol (with the same profile in case of an operation or a predicate), while this symbol is not intended to denote the same ‘thing’ in the combination. In such cases, it is necessary to explicitly *rename* some of the symbols exported by the specifications put together in order to avoid the unintended name clashes.

When reusing a named specification, it may be convenient to rename some of its symbols; moreover, in the case of operation or predicate symbols, one may also change the style of notation. This is illustrated in the specification `STACK` above which is obtained as a renaming of the specification `LIST_SELECTORS`. (A renaming is introduced by the keyword ‘**with**’.) First, the sort *List* is renamed into *Stack*, then the operation *cons* is renamed into a mixfix

operation *push\_onto\_*, and finally the selectors *head* and *tail* are renamed into *top* and *pop*, respectively. Note that ‘ $\mapsto$ ’ is input as ‘ $|->$ ’.

The user only needs to indicate how symbols provided by the renamed specification are mapped to new symbols. A *signature morphism* is automatically deduced from this ‘symbol map’. For instance, the signature morphism inferred from the symbol map specified in `STACK` maps the operation symbol *cons* :  $Elem \times List \rightarrow List$  to the operation symbol *push\_onto\_* :  $Elem \times Stack \rightarrow Stack$ : not only the operation name is changed, but also its profile according to the renaming of *List* into *Stack*.

In a symbol map, one can qualify the symbol to be renamed by its kind, using the keywords **sort**, **op**, and **pred** (or their plural forms), as appropriate; this is illustrated in `STACK` above. Qualification in symbol maps is generally recommended since it improves their readability.

While it is possible to change the syntax of an operation or predicate symbol, as illustrated above for *cons* mapped to *push\_onto\_*, it is not possible to change the order of the arguments of the renamed operation or predicate.

In general, one does not need to rename all the symbols provided by the specification to be renamed. In the symbol map describing the intended renaming, it is indeed enough to mention only the symbols that change. By default, any symbol not explicitly mentioned is left unchanged (although its profile may be updated according to the renaming specified for some sorts). This is illustrated here in `STACK` where there is no need to rename the constant *empty*, which will therefore have the same name for both lists and stacks. However, the induced signature morphism maps the constant symbol *empty* : *List* into the constant symbol *empty* : *Stack*.

One can also explicitly rename a symbol to itself, say by writing ‘*empty*  $\mapsto$  *empty*’, or just mention it without providing a new name, as in ‘**with** *empty*’, which is equivalent to ‘**with** *empty*  $\mapsto$  *empty*’.

By default, overloaded symbols are renamed simultaneously. For instance, in `INTEGER_ARITHMETIC_1` **with**  $__ + __ \mapsto plus$ , all the five overloaded infix ‘+’ operations exported by `INTEGER_ARITHMETIC_1` (see Chap. 5, p. 64) are renamed into five *plus* operations, with a functional syntax and the appropriate profiles.

In general, it is possible to specifically rename one of some overloaded symbols, by specifying its profile in the symbol map. For instance, in `LIST_SET` **with** *empty* : *List*  $\mapsto nil$ , only the constant *empty* of sort *List* is renamed into *nil*, while the constant *empty* of sort *Set* remains unchanged. However, care is needed in the presence of subsorts, since the signature morphism induced by the specified symbol map should preserve the overloading relations associated with subsorts. For instance, if we attempt to only rename in the specification `INTEGER_ARITHMETIC_1` the addition ‘+’ of two positive numbers into *plus* and write `INTEGER_ARITHMETIC_1` **with**  $__ + __ : Pos \times Pos \rightarrow Pos \mapsto plus$ , we merely obtain an ill-formed specification. Thus in the specification `INTEGER_ARITHMETIC_1`, all the five overloaded ‘+’ operations must be renamed simultaneously, again into five overloaded symbols.

*When combining specifications, origins of symbols can be indicated.*

```
spec LIST_SET_1 [sort Elem] =
  LIST_SELECTORS [sort Elem] with empty, cons
and GENERATED_SET [sort Elem] with empty, {--}, -- ∪ --
then op elements_of_-- : List → Set
  ∀e : Elem; L : List
  • elements_of empty = empty
  • elements_of cons(e, L) = {e} ∪ elements_of L
end
```

Since, as explained above, ‘**with empty, cons**’ means ‘**with empty**  $\mapsto$  *empty*, *cons*  $\mapsto$  *cons*’, identity renaming can be used just to emphasize the fact that a given specification exports some symbols. This is illustrated in the specification LIST\_SET\_1 above, which is quite similar to LIST\_SET, but for the fact that here we emphasize that LIST\_SELECTORS exports in particular the operations *empty* and *cons*, and that GENERATED\_SET exports in particular the operations *empty*, ‘{--}’, and ‘∪’.

## 6.3 Hiding

*Auxiliary symbols used in structured specifications can be hidden.*

```
spec NATURAL_PARTIAL_SUBTRACTION_3 =
  NATURAL_PARTIAL_SUBTRACTION_1 hide suc, pre
end

spec NATURAL_PARTIAL_SUBTRACTION_4 =
  NATURAL_PARTIAL_SUBTRACTION_1
  reveal Nat, 0, 1, -- + --, -- - --, -- * --, -- < --
end
```

When writing large specifications, it is quite frequent to rely on auxiliary operations (and predicates) to specify the operations (and predicates) of interest. Once these are defined, the auxiliary operations are no longer needed, and are better removed from the exported signature of the specification, which should include only the symbols that were required to be specified. This is the purpose of the **hide** construct.

Consider for instance the specification NATURAL\_PARTIAL\_SUBTRACTION\_1 given in Chap. 4, p. 52. Once addition and subtraction are defined, the



two basic operations *suc* and *pre* are no longer needed (since *suc*(*x*) is more conveniently written  $x + 1$ , and similarly *pre*(*x*) is expressed by  $x - 1$ ), and can therefore be hidden. This is illustrated by the specification NATURAL\_PARTIAL\_SUBTRACTION\_3 given above.

Depending on the relative proportion of symbols to be hidden or not, in some cases it may be more convenient to explicitly list the symbols to be exported by a specification rather than those to be hidden. The construct ‘**reveal**’ can be used for that purpose, and ‘**hide**’ and ‘**reveal**’ are just two symmetric constructs to achieve the same effect. The use of ‘**reveal**’ is illustrated in NATURAL\_PARTIAL\_SUBTRACTION\_4 above, and the reader can convince himself that both NATURAL\_PARTIAL\_SUBTRACTION\_3 and NATURAL\_PARTIAL\_SUBTRACTION\_4 export exactly the same symbols. However, in this case the first specification is clearly more concise. A more convincing example of the use of ‘**reveal**’ is provided by the following example.

**spec** PARTIAL\_ORDER\_2 = PARTIAL\_ORDER **reveal** **pred**  $\_ \leq \_$  --

Similar rules to the ones explained for renaming apply to the **hide** and **reveal** constructs. One can qualify a symbol to be hidden or revealed by its kind (**sort**, **op** or **pred**), and by default, overloaded symbols are hidden (or revealed) simultaneously.

Note that hiding a sort entails hiding all the operations or predicates that use it in their profile. Similarly, revealing an operation or a predicate entails revealing all the sorts involved in its profile. For instance, in the specification PARTIAL\_ORDER\_2 above, revealing the predicate ‘ $\leq$ ’ entails revealing also the sort *Elem*.

As a consequence, hiding sorts should be used with care in the presence of subsorts. For instance, hiding the sort *Nat* in the specification POSITIVE given in Chap. 5, p. 61, leads to a specification of positive natural numbers with a sort *Pos* which has the expected carrier set, but without any operation or predicate available on it. Hiding the sort *Nat* in the specification POSITIVE\_ARITHMETIC (see Chap. 5, p. 61) may seem more appropriate, but one should still note that the predicate ‘ $<$ ’ is no longer available in POSITIVE\_ARITHMETIC **hide sort** *Nat*.

As a last remark, note that when convenient, **reveal** can be combined with a renaming of (some of) the exported symbols. For instance, in the above PARTIAL\_ORDER\_2 specification, we could have written ‘**reveal pred**  $\_ \leq \_ \mapsto leq$ ’ if, in addition to a restriction of the signature of PARTIAL\_ORDER, we wanted to rename the infix predicate ‘ $\_ \leq \_$ ’ into a predicate *leq* with a functional notation.

## 6.4 Local Specifications

*Auxiliary symbols can be made local when they do not need to be exported.*

```

spec LIST_ORDER [TOTAL_ORDER with sort Elem, pred -- < --] =
  LIST_SELECTORS [sort Elem]
then local  op insert : Elem × List → List
            ∀e, e' : Elem; L : List
            • insert(e, empty) = cons(e, empty)
            • insert(e, cons(e', L)) = cons(e', insert(e, L)) when e' < e
                                          else cons(e, cons(e', L))

            within op order : List → List
            ∀e : Elem; L : List
            • order(empty) = empty
            • order(cons(e, L)) = insert(e, order(L))
end

```

In many cases, auxiliary symbols are introduced for immediate use, and they do not need to be exported by the specification where they are declared. Then the best is to limit the scope of the declarations of such auxiliary symbols by using the ‘**local** ... **within** ...’ construct. This is illustrated in the above specification LIST\_ORDER, where the *insert* operation is introduced only for the purpose of the axiomatization of *order*. The declaration of *insert* has its scope limited to the part that follows ‘**within**’, and *insert* is therefore not exported by the specification LIST\_ORDER.

It is generally advisable to ensure that auxiliary symbols are declared in local parts of specifications.

```

spec LIST_ORDER_SORTED
  [TOTAL_ORDER with sort Elem, pred -- < --] =
  LIST_SELECTORS [sort Elem]
then local  pred --is_sorted : List
            ∀e, e' : Elem; L : List
            • empty is_sorted
            • cons(e, empty) is_sorted
            • cons(e, cons(e', L)) is_sorted ⇔
              cons(e', L) is_sorted ∧ ¬(e' < e)

            within op order : List → List
            ∀L : List • order(L) is_sorted
end

```

The specification LIST\_ORDER\_SORTED above is a variant of the specification LIST\_ORDER illustrating again the use of the ‘**local** ... **within** ...’

construct – this time to declare an auxiliary predicate. (Actually, the two specifications are not equivalent, since `LIST_ORDER_SORTED` is much looser and only requires that  $order(L)$  is a sorted list, but perhaps not with the same elements as  $L$ .)

*Care is needed with local sort declarations.*

```
spec WRONG_LIST_ORDER_SORTED
  [TOTAL_ORDER with sort Elem, pred _ < _] =
  LIST_SELECTORS [sort Elem]
then local  pred __is_sorted : List
            sort SortedList = {L : List • L is_sorted}
            ∀e, e' : Elem; L : List
            • empty is_sorted
            • cons(e, empty) is_sorted
            • cons(e, cons(e', L)) is_sorted ⇔
              cons(e', L) is_sorted ∧ ¬(e' < e)
            within op order : List → SortedList
end
```

Note that the above specification `WRONG_LIST_ORDER_SORTED`, which may at first glance be considered as a slight variant of `LIST_ORDER_SORTED`, is *ill-formed*: `order` is exported by `WRONG_LIST_ORDER_SORTED`, and hence all sorts occurring in its profile should also be exported, which cannot be, since the sort `SortedList` is auxiliary. So, if the specifier really intends to insist that the result sort of `order` is `SortedList`, this subsort should be exported, as shown below.

```
spec LIST_ORDER_SORTED_2
  [TOTAL_ORDER with sort Elem, pred _ < _] =
  LIST_SELECTORS [sort Elem]
then local  pred __is_sorted : List
            sort SortedList = {L : List • L is_sorted}
            ∀e, e' : Elem; L : List
            • empty is_sorted
            • cons(e, empty) is_sorted
            • cons(e, cons(e', L)) is_sorted ⇔
              cons(e', L) is_sorted ∧ ¬(e' < e)
            within sort SortedList = {L : List • L is_sorted}
            op order : List → SortedList
end
```

In fact the ‘`local ... within ...`’ construct abbreviates a combination of extension and explicit hiding. The specification `LIST_ORDER_SORTED_2` above, for instance, abbreviates:

```

spec LIST_ORDER_SORTED_3
  [TOTAL_ORDER with sort Elem, pred _ < _] =
  LIST_SELECTORS [sort Elem]
then {
  pred _is_sorted : List
  ∀e, e' : Elem; L : List
  • empty is_sorted
  • cons(e, empty) is_sorted
  • cons(e, cons(e', L)) is_sorted ⇔
    cons(e', L) is_sorted ∧ ¬(e' < e)
  then sort SortedList = {L : List • L is_sorted}
  op order : List → SortedList
} hide _is_sorted
end

```

The main advantage of using the ‘**local** ... **within** ...’ construct is that hiding the symbols introduced in the **local** part is left implicit. The convenience of this generally outweighs the danger of overlooking a locally-declared sort that is needed for the profile of an exported symbol. In any case, CASL allows both styles, and users can simply choose the one they prefer.

## 6.5 Named Specifications

*Naming a specification allows its reuse.*

It is in general advisable to define as many named specifications as felt appropriate, since this improves the reusability of specifications: a named specification can easily be reused by referring to its name.

Not only do the names serve as abbreviations when writing specifications, they also make it easy for readers to notice reuse. Moreover, when the name of a specification is aptly chosen, e.g., `NATURAL_ARITHMETIC`, readers may well be able to guess its signature – and perhaps even the specified axioms – from the name itself. (In Chap. 9, we shall see how named specifications and other items can be collected in libraries, and particular versions of them made available for use over the Internet.)

References to named specifications are particularly convenient for specifications structured using unions and extensions, where verbatim insertion of unnamed specifications would tend to obscure the structure. When needed, the signature of a referenced specification can be adjusted through appropriate combinations of renaming and hiding (although this should not often be necessary, provided that auxiliary symbols are made local, as explained in the previous section).

---

## Generic Specifications

*Making a specification generic (when appropriate) improves its reusability.*

As mentioned in the previous chapter, naming specifications is a good idea. In many cases, however, datatypes are naturally *generic*, having sorts, operations, and/or predicates that are deliberately left loosely specified, to be determined when the datatype is used. For instance, datatypes of lists and sets are generic regarding the sort of elements. *Generic specifications* allow the genericity of a datatype to be made explicit by declaring *parameters* when the specification is named: in the case of lists and sets, there is a single parameter that simply declares the sort *Elem*.<sup>1</sup> A fitting argument specification has to be provided for each parameter of a generic specification whenever it is referenced; this is called *instantiation* of the generic specification.

The aim of this chapter is to discuss and illustrate how to define generic specifications and instantiate them. We have seen plenty of simple examples of generic specifications and instantiations in the previous chapters. In more complicated cases, however, explicit fitting symbol maps may be required to determine the exact relationship between parameters and arguments in instantiations, and so-called *imports* should be separated from the bodies of generic specifications.

---

<sup>1</sup> Generic specifications are also useful to ensure loose coupling between several named specifications, replacing an explicit extension by a parameter including only the necessary symbols and their required properties. This is illustrated in the Steam-Boiler Control System case study, see Chap. 13.

## 7.1 Parameters and Instantiation

*Parameters are arbitrary specifications.*

Any specification, named or not, can be used as the parameter of a generic specification. Commonly, the parameter is a rather trivial specification consisting merely of a single sort declaration, as in most of the examples given in the previous chapters, e.g.:

```
spec  GENERIC_MONOID [ sort Elem ] =  %{ See Chap. 3, p. 30 }%
```

```
spec  LIST_SELECTORS [ sort Elem ] =  %{ See Chap. 4, p. 54 }%
```

However, the parameter can also be a more complex, possibly structured, specification, as in:

```
spec  LIST_ORDER [ TOTAL_ORDER with sort Elem, pred -- < -- ] =  
      %{ See Chap. 6, p. 73 }%
```

Recall that ‘**with**’ requires the signature of the specification to include the listed symbols; here, in fact, the signature of `TOTAL_ORDER` does not contain any further symbols, so those are all the symbols that have to be supplied when instantiating `LIST_ORDER`.

*The argument specification of an instantiation must provide symbols corresponding to those required by the parameter.*

```
spec  LIST_ORDER_NAT = LIST_ORDER [ NATURAL_ORDER ]
```

The correspondence between the symbols provided by the argument specification and those required by the parameter can be described by a fitting symbol map or left implicit when it is not ambiguous, which is often the case.

In the above example, the argument specification `NATURAL_ORDER` provides the sort *Nat*, the operation symbols *0* and *suc*, and the binary predicate symbol ‘<’. Hence this specification indeed provides symbols corresponding to those required by the parameter specification `TOTAL_ORDER` and the correspondence can be left implicit because the argument `NATURAL_ORDER` has only single symbols of the right kind. (The coincidence of the predicate symbol in the parameter and argument is irrelevant here.)

How to describe explicit fitting symbol maps and when they can be omitted is detailed later in this section.

*The argument specification of an instantiation must ensure that the properties required by the parameter hold.*

```
spec NAT_WORD = GENERIC_MONOID [NATURAL]
```

A (*fitting*) *signature morphism* from the signature of the parameter part to the signature of the argument specification is automatically deduced, taking into account the explicitly specified fitting symbol map if any (the situation here is quite similar to a renaming, where a signature morphism is deduced from a symbol map). The instantiation is *defined* if all models of the argument specification, when reduced along the induced fitting signature morphism, provide models of the parameter part. In particular the symbols provided by the argument specification must have the properties, if any, specified in the parameter for their counterparts. When this is the case, we get not only a signature morphism, but also a (*fitting*) *specification morphism* from the argument specification to the parameter specification.<sup>2</sup>

In the above NAT\_WORD example, since the parameter of GENERIC\_MONOID is trivial, it is obvious that the instantiation is defined.

The effect of the instantiation is to make the union of the argument specification and of the (non generic equivalent of the) generic specification, renamed according to the induced fitting signature morphism. In particular, a side-effect of the instantiation is to rename the symbols of the generic specification according to the fitting signature morphism induced by the instantiation. In our NAT\_WORD example, the operation symbol  $inj : Elem \rightarrow Monoid$  is renamed into  $inj : Nat \rightarrow Monoid$ , while the operation symbols ‘1’ and ‘\*’ are left unchanged (as well as the sort *Monoid*). Thus, the specification NAT\_WORD abbreviates the following specification:

```
NATURAL and { NON_GENERIC_MONOID with  $Elem \mapsto Nat$  }.
```

When convenient, the instantiation can be completed by a renaming, as illustrated in the following variant of NAT\_WORD.

```
spec NAT_WORD_1 =
  GENERIC_MONOID [NATURAL]
  with  $Monoid \mapsto Nat\_Word$ 
end
```

In the case of the specification LIST\_ORDER\_NAT above, checking the definedness of the instantiation corresponds to a non-trivial proof obligation. The instantiation is defined since the predicate ‘<’ provided by NATURAL\_ORDER is indeed a total ordering relation, hence the properties required by

---

<sup>2</sup> Note that consistency is entirely orthogonal to definedness: a defined instantiation may be consistent or not.

TOTAL\_ORDER are fulfilled, even if there is no syntactic correspondence between the axioms given in TOTAL\_ORDER and those in NATURAL\_ORDER.

*There must be no shared symbols between the argument specification and the body of the instantiated generic specification.*

**spec** THIS\_IS\_WRONG = GENERIC\_MONOID [ MONOID ]

The intention in the above example may have been to specify monoids of monoids. However, the above instantiation is ill-formed since the sort *Monoid* and the operation symbols ‘1’ and ‘\*’ are shared between the body of the generic specification GENERIC\_MONOID and the argument specification MONOID.

Section 7.3 provides useful hints on how to structure generic specifications in order to avoid as far as possible undesirable clashes of symbols in instantiations. A correct specification of monoids of monoids is provided in Sect. 7.2, p. 86.

*In instantiations, the fitting of parameter symbols to identical argument symbols can be left implicit.*

**spec** GENERIC\_COMMUTATIVE\_MONOID [ **sort** Elem ] =  
       GENERIC\_MONOID [ **sort** Elem ]  
**then** ...

When the parameter and the argument have symbols in common, these parameter symbols are implicitly taken to fit directly to the corresponding argument symbols. Thus it is *never* necessary to make explicit that a symbol is mapped identically. In the above example, for instance, the parameter specification of GENERIC\_MONOID is exactly the same as the argument specification in its instantiation, so the fitting can be left implicit.

*The fitting of parameter sorts to unique argument sorts can also be left implicit.*

When the argument specification has only a single sort, the fitting of all parameter sorts to that sort is obvious, and can again be left implicit, as illustrated earlier by the NAT\_WORD specification. Of course, this does not apply the other way round: if the parameter has a single sort (which is often



the case in practice) but the argument specification has more than one sort, the parameter sort could be mapped to any of the argument sorts, so the fitting symbol map has to be made explicit – except when the parameter sort is identical to one of the argument sorts, as previously explained, or when the fitting of sorts can be implied from the fitting of other symbols, as explained below.

*Fitting of operation and predicate symbols can sometimes be left implicit too, and can imply fitting of sorts.*

**spec** LIST\_ORDER\_POSITIVE = LIST\_ORDER [ POSITIVE ]

Fitting of operation and predicate symbols can imply fitting of sorts. For instance, when a parameter predicate symbol is fitted to an argument predicate symbol whose profile involves different sorts, this implies that the parameter sorts involved have to be fitted to the corresponding sorts in the argument specification.

This is illustrated in the above LIST\_ORDER\_POSITIVE specification. In a first step, the fitting of the parameter sort *Elem* to one of the argument sorts *Nat* and *Pos* provided by the specification POSITIVE (see Chap. 5, p. 61) may seem ambiguous. However, no explicit fitting of symbols is necessary here, since the argument specification provides only one binary predicate symbol, and the fitting of the corresponding binary predicate symbol of the parameter specification to it entails the fitting of the sort *Elem* to the sort *Nat* (Again, the coincidence of the predicate symbol in the parameter and argument is irrelevant here.)

As may be clear by now, the exact rules for when the fitting between parameter and argument symbols can be left implicit are quite sophisticated. It seems best to make the intended fitting explicit whenever it is not completely obvious, using the notation for fitting arguments illustrated in the following examples.

*The intended fitting of the parameter symbols to the argument symbols may have to be specified explicitly.*

**spec** NAT\_WORD\_2 =  
GENERIC\_MONOID [ NATURAL\_SUBSORTS **fit** *Elem*  $\mapsto$  *Nat* ]

The correspondence between the symbols required by the parameter and those provided by the argument specification can be made explicit using so-called *fitting symbol maps*. For instance, the above NAT\_WORD\_2 specification, which differs from NAT\_WORD only regarding the presence of subsorts of

*Nat*, is obtained as an instantiation of `GENERIC_MONOID`, fitting the parameter part ‘`sort Elem`’ to the `NATURAL_SUBSORTS` specification. The mapping between the parameter sort *Elem* and the sort *Nat* provided by `NATURAL_SUBSORTS` is described by the fitting symbol map ‘`fit Elem ↦ Nat`’.

*A generic specification may have more than one parameter.*

```
spec PAIR [sort Elem1] [sort Elem2] =
  free type Pair ::= pair(first : Elem1; second : Elem2)
```

Using several parameters is merely a notational convenience, since they are equivalent to their union. For instance, the above `PAIR` specification is nothing but a variant of the specification `PAIR_1` with just one parameter ‘`sorts Elem1, Elem2`’ defined in Chap. 4, p. 54.

Note that writing:

```
spec HOMOGENEOUS_PAIR_1 [sort Elem] [sort Elem] =
  free type Pair ::= pair(first : Elem; second : Elem)
```

merely defines pairs of values of the same sort, and `HOMOGENEOUS_PAIR_1` is (equivalent to and) better defined as follows:

```
spec HOMOGENEOUS_PAIR [sort Elem] =
  free type Pair ::= pair(first : Elem; second : Elem)
```

since the two parameters in `HOMOGENEOUS_PAIR_1` are equivalent to just one ‘`sort Elem`’ parameter.

From a methodological point of view, it is generally advisable to use as many parameters as convenient: the part of the specification that is intended to be specialized at instantiation time is better split into logically coherent units, each one corresponding to a parameter. Consider for instance:

```
spec TABLE [sort Key] [sort Val] = ...
```

Here, using two parameters in `TABLE` emphasizes that *Key* and *Val* are logically distinct entities which can be specialized as desired independently of each other.

*Instantiation of generic specifications with several parameters is similar to the case of just one parameter.*

```
spec PAIR_NATURAL_COLOR =
  PAIR [NATURAL_ARITHMETIC] [COLOR fit Elem2 ↦ RGB]
```

In the above example, the first parameter ‘**sort** *Elem1*’ of PAIR is instantiated by NATURAL\_ARITHMETIC, which exports only one sort *Nat*, hence no explicit fitting symbol map is necessary. The second parameter ‘**sort** *Elem2*’ of PAIR is instantiated by COLOR: in this case a fitting symbol map must be provided, since COLOR exports two sorts, *RGB* and *CMYK*.

Using the specification PAIR\_1 would require us to write:

```
spec PAIR_NATURAL_COLOR_1 =
  PAIR_1 [NATURAL_ARITHMETIC and COLOR
        fit Elem1  $\mapsto$  Nat, Elem2  $\mapsto$  RGB]
```

which clearly demonstrates the benefit of using two parameters as in PAIR instead of just one as in PAIR\_1.

When parameters are trivial ones (i.e., just one sort), one can always avoid explicit fitting maps. Consider for instance the following alternative to PAIR\_NATURAL\_COLOR:

```
spec PAIR_NATURAL_COLOR_2 =
  PAIR [sort Nat] [sort RGB]
and NATURAL_ARITHMETIC and COLOR
```

This may be convenient when the argument specification exports several sorts. Compare for instance:

```
spec PAIR_POS =
  HOMOGENEOUS_PAIR [sort Pos] and INTEGER_ARITHMETIC_1
```

with:

```
spec PAIR_POS_1 =
  HOMOGENEOUS_PAIR [INTEGER_ARITHMETIC_1 fit Elem  $\mapsto$  Pos]
```

Note that the instantiation:

```
HOMOGENEOUS_PAIR_1 [NATURAL] [COLOR fit Elem  $\mapsto$  RGB]
```

is ill-formed, since it entails mapping the sort *Elem* to both *Nat* and *RGB*.

More generally, care is needed when the several parameters of a generic specification share some symbols, which in general is not advisable.

As a last remark, note that it is easy to specialize a generic specification with several parameters, using a ‘partial instantiation’, as in the following version of TABLE:

```
spec MY_TABLE [sort Val] =
  TABLE [NATURAL_ARITHMETIC] [sort Val]
```

where we still have a parameter for the values to be stored, but have decided that the keys are natural numbers.

*Composition of generic specifications is expressed using instantiation.*

```
spec SET_OF_LIST [ sort Elem ] =
  GENERATED_SET [ LIST_SELECTORS [ sort Elem ] fit Elem  $\mapsto$  List ]
```

The above generic specification `SET_OF_LIST` describes sets of lists of arbitrary elements, and is obtained by an instantiation of the generic specification `GENERATED_SET`, whose parameter ‘`sort Elem`’ is instantiated by the specification `LIST_SELECTORS`, itself trivially instantiated. Since the (trivially instantiated) specification `LIST_SELECTORS` exports two sorts *Elem* and *List*, it is of course necessary to specify, in the instantiation of `GENERATED_SET`, the fitting symbol map from the parameter sort *Elem* to the argument sort *List*.

Note especially that the following specification:

```
spec MISTAKE [ sort Elem ] =
  GENERATED_SET [ LIST_SELECTORS [ sort Elem ] ]
```

does *not* provide sets of lists of elements: The sort *Elem* in the parameter of `GENERATED_SET` is mapped by the identity fitting symbol map to the sort *Elem* provided by the instantiation of the generic specification `LIST_SELECTORS [sort Elem]`, rather than to the sort *List*.<sup>3</sup> Thus `MISTAKE` just provides sets of arbitrary elements *and* lists of arbitrary elements. If this was indeed the desired effect, then one should rather write instead:

```
spec SET_AND_LIST [ sort Elem ] =
  GENERATED_SET [ sort Elem ] and LIST_SELECTORS [ sort Elem ]
```

As illustrated by `SET_OF_LIST`, composition of generic specifications is fairly easy in CASL. Note however that this composition is achieved by means of appropriate instantiations (some possibly trivial), and that CASL does not provide higher-order genericity.

It may be worth mentioning that the following composition of generic specifications is ill-formed:

```
spec THIS_IS_STILL_WRONG =
  GENERIC_MONOID [ GENERIC_MONOID [ sort Elem ]
    fit Elem  $\mapsto$  Monoid ]
```

---

<sup>3</sup> However, the situation would be different if the parameter of `GENERATED_SET` had been, e.g., ‘`sort Val`’, since then the absence of an explicit fitting symbol map would have led to an ambiguity: in that case the specifier would have to specify whether the sort *Val* is to be mapped to *Elem* or to *List*.

The above instantiation is ill-formed since the sort *Monoid* and the operation symbols ‘1’ and ‘\*’ are shared between the body of the generic specification `GENERIC_MONOID` and the argument specification `GENERIC_MONOID [sort Elem]` (where this time the generic specification `GENERIC_MONOID` is trivially instantiated). The next section provides (p. 86) a correct specification of monoids of monoids.

## 7.2 Compound Symbols

*Compound sorts introduced by a generic specification get automatically renamed on instantiation, which avoids name clashes.*

```
spec LIST_REV [sort Elem] =
  free type List[Elem] ::= empty |
                                cons(head :? Elem; tail :? List[Elem])
  ops  ++ ++ -- : List[Elem] × List[Elem] → List[Elem],
        assoc, unit empty;
        reverse : List[Elem] → List[Elem]
  ∀e : Elem; L, L1, L2 : List[Elem]
    • cons(e, L1) ++ L2 = cons(e, L1 ++ L2)
    • reverse(empty) = empty
    • reverse(cons(e, L)) = reverse(L) ++ cons(e, empty)
end
```

```
spec LIST_REV_NAT = LIST_REV [NATURAL]
```

A compound sort is a sort of the form ‘*Name*[*Name*<sub>1</sub>, . . . , *Name*<sub>N</sub>]’. In the specification `LIST_REV`, we introduce a compound sort `List[Elem]` to denote lists (of arbitrary elements), instead of the simple sort `List` used in the previous examples. When the specification `LIST_REV` is instantiated as in `LIST_REV_NAT`, the translation induced by the (implicit) fitting symbol map is applied to the component *Elem* also where it occurs in `List[Elem]`, providing a sort `List[Nat]`. Thus, compound sorts can be seen as a convenient way of implicitly completing the instantiation by an appropriate renaming of the (compound) sorts introduced by the generic specification.

```
spec TWO_LISTS =
  LIST_REV [NATURAL] %% Provides the sort List[Nat]
and LIST_REV [COLOR fit Elem ↦ RGB] %% Provides the sort List[RGB]
```

Using a compound sort `List[Elem]` proves particularly useful in the above example `TWO_LISTS`, where we make the union of two distinct instantiations

of `LIST_REV`. If we had used an ordinary sort *List*, then an unintentional name clash would have arisen,<sup>4</sup> and we would have to complete each instantiation by an explicit renaming of the sort *List*.

Note that in the specification `TWO_LISTS`, we have two sorts *List*[*Nat*] and *List*[*RGB*], hence two overloaded constants *empty* (one of each sort), which may need disambiguation when used in terms. (How to disambiguate terms is explained in Chap. 3, p. 31.)

Similarly, we have overloaded operation symbols *cons*, *head*, *tail*, *++*, and *reverse*, but in general their context of use in terms will be enough to disambiguate which one is meant.

```
spec TWO_LISTS_1 =
  LIST_REV [INTEGER_ARITHMETIC_1 fit Elem ↦ Nat]
and LIST_REV [INTEGER_ARITHMETIC_1 fit Elem ↦ Int]
```

Since the specification `INTEGER_ARITHMETIC_1` provides three sorts *Nat*, *Pos*, and *Int*, an explicit fitting symbol map is needed in the above instantiations, which provide the sorts *List*[*Nat*] and *List*[*Int*]. Note that the subsorting relation *Nat* < *Int* does *not* entail *List*[*Nat*] < *List*[*Int*], but of course this can be added if desired in an extension by a subsorting declaration.

Using compound sorts, we can now easily specify monoids of monoids.

```
spec MONOID_C [sort Elem] =
  sort Monoid[Elem]
  ops inj : Elem → Monoid[Elem];
      1 : Monoid[Elem];
      _ * _ : Monoid[Elem] × Monoid[Elem] → Monoid[Elem],
          assoc, unit 1
  ∀x, y : Elem • inj(x) = inj(y) ⇒ x = y
end
```

```
spec MONOID_OF_MONOID [sort Elem] =
  MONOID_C [MONOID_C [sort Elem] fit Elem ↦ Monoid[Elem]]
```

The instantiation in `MONOID_OF_MONOID` is now correct, since the use of a compound sort *Monoid*[*Elem*] ensures there is no clash of symbols between the body of the instantiated generic specification and the argument specification.

---

<sup>4</sup> And the specification `TWO_LISTS` would have been inconsistent, due to the same name, same thing principle and the fact that *List* is defined by a free type construct.

*Compound symbols can also be used for operations and predicates.*

```

spec LIST_REV_ORDER [TOTAL_ORDER] =
  LIST_REV [sort Elem]
then local op insert : Elem × List[Elem] → List[Elem]
  ∀e, e' : Elem; L : List[Elem]
    • insert(e, empty) = cons(e, empty)
    • insert(e, cons(e', L)) = cons(e', insert(e, L)) when e' < e
      else cons(e, cons(e', L))
  within op order[<] : List[Elem] → List[Elem]
  ∀e : Elem; L : List[Elem]
    • order[<](empty) = empty
    • order[<](cons(e, L)) = insert(e, order[<](L))
end

spec LIST_REV_WITH_TWO_ORDERS =
  LIST_REV_ORDER
  [INTEGER_ARITHMETIC_ORDER fit Elem ↦ Int, < ↦ <, > ↦ >]
  %% Provides the sort List[Int] and the operation order[<]
and LIST_REV_ORDER
  [INTEGER_ARITHMETIC_ORDER fit Elem ↦ Int, < ↦ <, > ↦ >]
  %% Provides the sort List[Int] and the operation order[>]
then %implies
  ∀L : List[Int] • order[<](L) = reverse(order[>](L))
end

```

The above example illustrates the use of compound identifiers for operation symbols, and the same rules apply to predicate symbols. While in most cases using compound identifiers for sorts will be sufficient, in some situations it is also convenient to use them for operation or predicate symbols, as done here for  $order[<]$ . When `LIST_REV_ORDER` is instantiated, not only does the sort  $List[Elem]$  get renamed (here, to  $List[Int]$ ), but also the operation symbol  $order[<]$ , according to the fitting symbol map corresponding to the instantiation. If we had not used a compound identifier for the  $order$  operation, then an unintentional name clash would have arisen. Note that on the other hand we rely on the same name, same thing principle to ensure that the sorts  $List[Int]$  provided by each of the two instantiations are the same, which indeed is what we want for this example.

Of course we do not bother to use a compound identifier for the  $insert$  operation symbol. This operation being local, it is not exported by `LIST_REV_ORDER` and cannot be the source of unintentional name clashes in instantiations.

### 7.3 Generic Specifications with Imports

*Parameters should be distinguished from references to fixed specifications that are not intended to be instantiated.*

```

spec LIST_WEIGHTED_ELEM [sort Elem op weight : Elem  $\rightarrow$  Nat]
      given NATURAL_ARITHMETIC =
      LIST_REV [sort Elem]
then op weight : List[Elem]  $\rightarrow$  Nat
       $\forall e : \textit{Elem}; L : \textit{List}[\textit{Elem}]$ 
      • weight(empty) = 0
      • weight(cons(e, L)) = weight(e) + weight(L)
end

```

In the above example, we specialize lists of arbitrary elements to lists of elements equipped with a *weight* operation, which is then overloaded by a *weight* operation on lists. Therefore we specify that the generic specification LIST\_WEIGHTED\_ELEM has for parameter a specification extending the ‘**given**’ specification NATURAL\_ARITHMETIC by a sort *Elem* and an operation symbol *weight*. Thereby the intention is to emphasize the fact that only the sort *Elem* and the operation *weight* are intended to be specialized when the specification LIST\_WEIGHTED\_ELEM is instantiated, and not the ‘fixed part’ NATURAL\_ARITHMETIC. In CASL, the specifications listed after the ‘**given**’ keyword are called *imports*. One could have written instead:

```

spec LIST_WEIGHTED_ELEM
      [NATURAL_ARITHMETIC then sort Elem op weight : Elem  $\rightarrow$  Nat]
      = ...

```

but the latter, which is correct, misses the essential distinction between the part which is intended to be specialized and the part which is ‘fixed’ (since, by definition, the parameter is the part which has to be specialized).

Note also that omitting the ‘**given** NATURAL\_ARITHMETIC’ clause would make the declaration:

```

spec LIST_WEIGHTED_ELEM [sort Elem op weight : Elem  $\rightarrow$  Nat] = ...

```

ill-formed, since the sort *Nat* is not available.

To summarize, the ‘**given**’ construct is useful to distinguish the ‘true’ parameter from the part which is ‘fixed’. Both the parameter and the body of the generic specification extend what is provided by the imports (i.e., by the specifications listed after the ‘**given**’ keyword), whose exported symbols are therefore available.



*Argument specifications are always implicitly regarded as extension of the imports.*

```
spec LIST_WEIGHTED_PAIR_NATURAL_COLOR =
  LIST_WEIGHTED_ELEM [PAIR_NATURAL_COLOR fit  $Elem \mapsto Pair$ ,
                       $weight \mapsto first$ ]
```

The instantiation specified in LIST\_WEIGHTED\_PAIR\_NATURAL\_COLOR is correct since the fitting symbol map is the identity on all the symbols exported by the ‘fixed part’ NATURAL\_ARITHMETIC (which happens here to be included in the argument specification PAIR\_NATURAL\_COLOR). More generally, the argument specification is always regarded as an extension of the imports, and the fitting symbol map should be the identity on all symbols provided by these imports. This is illustrated in the next example:

```
spec LIST_WEIGHTED_INSTANTIATED =
  LIST_WEIGHTED_ELEM [sort  $Value$  op  $weight : Value \rightarrow Nat$ ]
```

Here we rely on a rather trivial instantiation (whose purpose is merely to illustrate our point) where the fitting symbol map can be omitted since no ambiguity arises and where the argument specification ‘**sort**  $Value$  **op**  $weight : Value \rightarrow Nat$ ’ is well-formed because it is regarded as an extension of the imports of LIST\_WEIGHTED\_ELEM (i.e., as an extension of NATURAL\_ARITHMETIC), which implies that the sort  $Nat$  is available.

*Imports are also useful to prevent ill-formed instantiations.*

```
spec LIST_LENGTH [sort  $Elem$ ] given NATURAL_ARITHMETIC =
  LIST_REV [sort  $Elem$ ]
then op  $length : List[Elem] \rightarrow Nat$ 
   $\forall e : Elem; L : List[Elem]$ 
    •  $length(empty) = 0$ 
    •  $length(cons(e, L)) = length(L) + 1$ 
then %implies
   $\forall L : List[Elem]$  •  $length(reverse(L)) = length(L)$ 
end
```

The specification LIST\_LENGTH needs the sort  $Nat$  and the usual arithmetic operations provided by NATURAL\_ARITHMETIC to specify the *length* operation. In this case it is clear that the imports have nothing to do with the (trivial) parameter of LIST\_LENGTH. The reason to specify NATURAL\_ARITHMETIC as an import is that this will make instantiations of LIST\_LENGTH similar to the following one well-formed.

```
spec LIST_LENGTH_NATURAL =  
      LIST_LENGTH [NATURAL_ARITHMETIC]
```

To understand this point, consider the following variant of LIST\_LENGTH:

```
spec WRONG_LIST_LENGTH [sort Elem] =  
      NATURAL_ARITHMETIC and LIST_REV [sort Elem]  
then ...  
end
```

The specification WRONG\_LIST\_LENGTH is fine as long as one does not need to instantiate it with NATURAL\_ARITHMETIC as argument specification. The instantiation WRONG\_LIST\_LENGTH [NATURAL\_ARITHMETIC] is ill-formed since some symbols of the argument specification are shared with some symbols of the body (and not already occurring in the parameter) of the instantiated generic specification, which is wrong, as already explained p. 80. Of course the same problem will occur with any argument specification which provides, e.g., the sort *Nat*.

*In generic specifications, auxiliary required specifications should be imported rather than extended.*

As illustrated by the above examples, one should remember the following essential point. Since an instantiation is ill-formed as soon as there are some shared symbols between the argument specification and the body of the generic specification, when designing a generic specification, it is generally advisable to turn auxiliary required specifications (such as NATURAL\_ARITHMETIC for LIST\_LENGTH) into imports, and generic specifications of the form ' $F[X] = SP$  **then** ...' are better written ' $F[X]$  **given**  $SP = \dots$ ' to allow the instantiation ' $F[SP]$ '.

## 7.4 Views

*Views are named fitting maps, and can be defined along with specifications.*

```
view INTEGER_AS_TOTAL_ORDER :  
      TOTAL_ORDER to INTEGER_ARITHMETIC_ORDER =  
       $Elem \mapsto Int, \_ < \_ \mapsto \_ < \_$ 
```

```

view INTEGER_AS_REVERSE_TOTAL_ORDER :
  TOTAL_ORDER to INTEGER_ARITHMETIC_ORDER =
    Elem  $\mapsto$  Int,  $-- < -- \mapsto -- > --$ 

spec LIST_REV_WITH_TWO_ORDERS_1 =
  LIST_REV_ORDER [ view INTEGER_AS_TOTAL_ORDER ]
and LIST_REV_ORDER [ view INTEGER_AS_REVERSE_TOTAL_ORDER ]
then %implies
   $\forall L : \text{List}[\text{Int}] \bullet \text{order}[-- < --](L) = \text{reverse}(\text{order}[-- > --](L))$ 
end

```

A view is nothing but a convenient way to name a specification morphism (induced by a symbol map) from a (parameter) specification to an (argument) specification. The rules regarding the omission of ‘evident’ symbol maps in explicit fittings apply to views too. A view proves particularly useful when the same instantiation (with the same fitting symbol map) is intended to be used several times: naming a specification morphism once and for all makes its reuse easier. Once a view is defined, as e.g. `INTEGER_AS_TOTAL_ORDER` above, it can be referenced in instantiations as in `LIST_REV_ORDER [ view INTEGER_AS_TOTAL_ORDER ]`, where the keyword ‘**view**’ makes it clear that the argument is not merely a named specification with an implicit fitting map, which would be written differently.

Since a view is defined only when the given symbol map induces a specification morphism (i.e., all models of the target specification, when reduced along the signature morphism induced by the given symbol map, provide models of the source specification), it may be convenient to use views just to explicitly document the existence of some specification morphisms, even when these are not intended to be used in any instantiation. For instance, the view `INTEGER_AS_TOTAL_ORDER` can be seen as the assertion that `INTEGER_ARITHMETIC_ORDER` indeed specifies ‘`<`’ to be a total ordering relation, and would therefore make sense even without being used later on in instantiations.

*Views can also be generic.*

```

view LIST_AS_MONOID [ sort Elem ] :
  MONOID to LIST_REV [ sort Elem ] =
    Monoid  $\mapsto$  List[Elem],  $1 \mapsto \text{empty}, -- * -- \mapsto -- + --$ 

```

A view can be generic, being then defined with some parameters (as illustrated above in the `LIST_AS_MONOID` view) and possibly some imports. The reader should be aware that, in a generic view, the target specification (here, the trivially instantiated specification `LIST_REV`) is not interpreted as such, but as the body of a generic specification with the same parameters and

imports as the view. (The source specification is on the contrary interpreted exactly as provided.)

The above example illustrates again the use of a view as a ‘proof obligation’, asserting that lists (equipped with the ‘++’ operation) form a monoid.

---

## Specifying the Architecture of Implementations

*Architectural specifications impose structure on implementations, whereas specification-building operations only structure the text of specifications.*

As explained in the previous chapters, the specification of a complex system may be fairly large and should be structured into coherent, easy to grasp, pieces. CASL provides a number of specification-building operations to achieve this, as detailed in Chap. 6. Moreover, generic specifications, described in Chap. 7, provide pieces of specification that are easy to reuse in different contexts, where they can be adapted as desired by instantiating them.

Specification-building operations and generic specifications are useful *to structure the text of the specification* of the system under consideration. However, the models of a structured specification have no more structure than do those of a flat, unstructured, specification. Indeed, most examples given in the previous chapters could have been structured differently, with the same meaning (i.e., with the same models). Structured specifications are usually adequate at the requirements stage, where the focus is on the expected overall properties of the system under consideration.

In contrast, the aim of architectural specifications is *to prescribe the intended architecture of the implementation* of the system. Architectural specifications provide the means for specifying the various *components* from which the system will be built, and describing how these components are to be assembled to provide an implementation of the system of interest. At the same time, they allow the task of implementing a system to be split into independent, clearly-specified sub-tasks. Thus, architectural specifications are essential at the design stage, where the focus is on how to factor the implementation of the system into components.

The aim of this chapter is to discuss and illustrate both the role of architectural specifications and how to express them in CASL.

The idea underlying architectural specifications is that eventually in the process of systematic development of modular software from specifications, components are implemented as software modules in some chosen programming language. However, this step is beyond the scope of specification formalisms, so in CASL and in this chapter we identify components with models (and with functions from models to models, in the case of generic components). The modular structure of the software under development, as described by an architectural specification, is therefore captured here simply as an explicit, structural way to build CASL models.

*The examples in this chapter are artificially simple.*

Architectural specifications, and more generally component-oriented approaches, are intended for relatively large systems. In this chapter, however, we have to rely on simple small examples to illustrate and explain CASL architectural specification concepts and constructs. After reading this chapter, the reader is encouraged to study Chap. 13, which provides realistic examples of the use of architectural specifications. A more detailed account of the rationale behind architectural specifications in the context of formal software development by stepwise refinement can be found in [11].

The following structured specifications will be referred to later in this chapter when illustrating CASL architectural specifications:

```

spec  COLOR =  %{ As defined in Chap. 3, p. 37 }%
spec  NATURAL_ORDER =  %{ As defined in Chap. 3, p. 38 }%
spec  NATURAL_ARITHMETIC =  %{ As defined in Chap. 3, p. 38 }%

spec  ELEM = sort Elem

spec  CONT [ELEM] =
  generated type Cont[Elem] ::= empty | insert(Elem; Cont[Elem])
  preds  __is_empty : Cont[Elem];
          __is_in__ : Elem  $\times$  Cont[Elem]
  ops    choose : Cont[Elem]  $\rightarrow$ ? Elem;
          delete : Elem  $\times$  Cont[Elem]  $\rightarrow$  Cont[Elem]
   $\forall e, e' : \text{Elem}; C : \text{Cont}[\text{Elem}]$ 
    • empty is_empty
    •  $\neg \text{insert}(e, C) \text{ is\_empty}$ 
    •  $\neg e \text{ is\_in } \text{empty}$ 
    •  $e \text{ is\_in } \text{insert}(e', C) \Leftrightarrow (e = e' \vee e \text{ is\_in } C)$ 
    •  $\text{def choose}(C) \Leftrightarrow \neg C \text{ is\_empty}$ 
    •  $\text{def choose}(C) \Rightarrow \text{choose}(C) \text{ is\_in } C$ 
    •  $e \text{ is\_in } \text{delete}(e', C) \Leftrightarrow (e \text{ is\_in } C \wedge \neg(e = e'))$ 
end

```

```

spec  CONT_DIFF [ELEM] =
  CONT [ELEM]
then op diff : Cont[Elem] × Cont[Elem] → Cont[Elem]
  ∀e : Elem; C, C' : Cont[Elem]
    • e is_in diff(C, C') ⇔ (e is_in C ∧ ¬(e is_in C'))
end

spec  REQ = CONT_DIFF [NATURAL_ORDER]

spec  FLAT_REQ =
  free type Nat ::= 0 | suc(Nat)
  pred   __ < __ : Nat × Nat
  generated type Cont[Nat] ::= empty | insert(Nat; Cont[Nat])
  preds  __is_empty : Cont[Nat];
         __is_in__ : Nat × Cont[Nat]
  ops    choose : Cont[Nat] →? Nat;
         delete : Nat × Cont[Nat] → Cont[Nat];
         diff : Cont[Nat] × Cont[Nat] → Cont[Nat]
  ∀e, e' : Nat; C, C' : Cont[Nat]
    • 0 < suc(e)
    • ¬(e < 0)
    • suc(e) < suc(e') ⇔ e < e'
    • empty is_empty
    • ¬ insert(e, C) is_empty
    • ¬ e is_in empty
    • e is_in insert(e', C) ⇔ (e = e' ∨ e is_in C)
    • def choose(C) ⇔ ¬ C is_empty
    • def choose(C) ⇒ choose(C) is_in C
    • e is_in delete(e', C) ⇔ (e is_in C ∧ ¬(e = e'))
    • e is_in diff(C, C') ⇔ (e is_in C ∧ ¬(e is_in C'))
end

```

## 8.1 Architectural Specifications

Let's assume in the following that REQ describes our requirements about the system to be implemented. First, note that both REQ and FLAT\_REQ have the same models, which illustrates our point about the fact that the CASL specification-building operations are merely facilities to structure the text of specifications into coherent units.

*An architectural specification consists of a list of unit declarations, specifying the required components, and a result part, indicating how they are to be combined.*

```

arch spec SYSTEM =
units  N : NATURAL_ORDER;
        C : CONT [NATURAL_ORDER] given N;
        D : CONT_DIFF [NATURAL_ORDER] given C
result D

```

The SYSTEM architectural specification is intended to prescribe a specific architecture for implementing the system specified by REQ.

The first part, introduced by the keyword **units**, indicates that we require the implementation of our system to be made of three components *N*, *C*, and *D*. The second part, introduced by the keyword **result**, indicates that the component *D* provides the desired implementation.

Each component is provided with its specification. The line:

```
N : NATURAL_ORDER
```

declares a component *N* specified by NATURAL\_ORDER, which means simply that *N* should be a model of NATURAL\_ORDER.

The line:

```
C : CONT [NATURAL_ORDER] given N
```

declares a component *C* which, given the previously declared component *N*, provides a model of CONT [NATURAL\_ORDER]. It is essential to understand that the component *C* must *expand* the assumed component *N* into a model of CONT [NATURAL\_ORDER], which means that *C* reduced to the signature of NATURAL\_ORDER must be equal to *N*. This property reflects the fact that a software module is supposed to use what it is given exactly as supplied, without altering it.

Similarly, the line:

```
D : CONT_DIFF [NATURAL_ORDER] given C
```

declares a component *D* which, given the component *C*, expands it into a model of CONT\_DIFF [NATURAL\_ORDER].

The final result is therefore simply *D*. (More complex examples of result expressions will be illustrated in examples below.)

As in the rest of CASL, visibility is linear in architectural specifications, meaning that any component must be declared before being used (e.g., the component *N* should be declared before being referred to by '**given** *N*' in the declaration of the component *C* in the architectural specification SYSTEM).



Component names (such as  $N$ ,  $C$ , and  $D$  in `SYSTEM`) are *local* to the architectural specification where they are declared, and are not visible outside it.

*There can be several distinct architectural choices for the same requirements specification.*

```
arch spec SYSTEM_1 =
units  $N$  : NATURAL_ORDER;
       $CD$  : CONT_DIFF [NATURAL_ORDER] given  $N$ 
result  $CD$ 
```

The architectural specifications `SYSTEM` and `SYSTEM_1` both provide models of `REQ`. However, the former insists on an implementation made of three components, while the latter insists on an implementation made of two components. Thus the architectural specification `SYSTEM_1` corresponds to a different architectural choice for implementing our `REQ` specification. Of course, further design for implementing the component  $CD$  of `SYSTEM_1` may lead to splitting this implementation task exactly as in `SYSTEM` above. However, there are also other possibilities, including for instance an architectural design where we would split the task of implementing  $CD$  into two different tasks, one for implementing containers with all their operations (including *diff*) except *delete*, the other for implementing *delete* by means of *diff* and other operations.

*Each unit declaration listed in an architectural specification corresponds to a separate implementation task.*

For instance, in the architectural specification `SYSTEM`, the task of providing a component  $D$  expanding  $C$  and implementing `CONT_DIFF [NATURAL_ORDER]` is independent from the tasks of providing implementations  $N$  of `NATURAL_ORDER` and  $C$  of `CONT [NATURAL_ORDER]` **given**  $N$ . Hence, when providing the component  $D$ , one cannot make any further assumption on how the component  $C$  is (or will be) implemented, besides what is expressly ensured by its specification.

To understand this, let us consider again the requirements specification `REQ` (or its variant `FLAT_REQ`). Among its models, there is one where containers are implemented by sorted lists (in increasing order, without repetitions), and in this model we can choose to implement *diff* by the following algorithm:

$$\begin{aligned}
diff(L, L') &= nil \text{ when } L = nil \\
&\text{else } L \text{ when } L' = nil \\
&\text{else } insert(head(L), diff(tail(L), L')) \text{ when } head(L) < head(L') \\
&\text{else } diff(tail(L), tail(L')) \text{ when } head(L) = head(L') \\
&\text{else } diff(L, tail(L'))
\end{aligned}$$

In this model, however, we rely on knowledge about the implementation of containers to decide how to implement *diff* – which is fine, since both are simultaneously implemented in the same component. In contrast, in the architectural specification **SYSTEM**, we request that containers are to be implemented in the component *C* while *diff* is to be provided by a separate component *D*. Imposing that the component *D* can be developed independently of the component *C* means that for *D* it is no longer possible to implement *diff* as sketched above, since this specific implementation choice may not be compatible with an independently chosen realization for *C* (where containers may be implemented by bags, for instance). Hence an implementation of *diff* in the component *D* can only rely on the operations provided by *C* (e.g., *choose* and *delete*); this may turn out to be less efficient for some particular realization of *C*, but should be compatible with any independently chosen realization for *C* (bags, for instance). In the case of the architectural specification **SYSTEM\_1**, since both containers and the *diff* operation are implemented in the same component *CD*, we can of course decide to implement containers by ordered lists without repetitions and *diff* as sketched above.

Thus the component *D* should expand *any* given implementation *C* of **CONT** [**NATURAL\_ORDER**] and provide an implementation of **CONT\_DIFF** [**NATURAL\_ORDER**], which is tantamount to providing a *generic* implementation *G* of **CONT\_DIFF** [**NATURAL\_ORDER**] which takes the particular implementation of **CONT** [**NATURAL\_ORDER**] as a parameter to be expanded. Then we obtain *D* by simply applying *G* to *C*.

Genericity here arises from the independence of the developments of *C* and *D*, rather than from the desire to build multiple implementations of **CONT\_DIFF** [**NATURAL\_ORDER**] using different implementations of **CONT** [**NATURAL\_ORDER**]. This is reflected by the fact that *G* is left implicit in the architectural specification **SYSTEM**.

*A unit can be implemented only if its specification is a conservative extension of the specifications of its given units.*

For instance, the component *D* can exist only if the specification **CONT\_DIFF** [**NATURAL\_ORDER**] is a conservative extension of **CONT** [**NATURAL\_ORDER**], i.e., if any model of the latter specification can be expanded into a model of the former one, which is indeed the case here. Similarly, the component *C* can exist since **CONT** [**NATURAL\_ORDER**] is a conservative extension of **NATURAL\_ORDER**.

Consider now the following variant of `CONT_DIFF [NATURAL_ORDER]` and the associated variant of the architectural specification `SYSTEM`.

```

spec  CONT_DIFF_1 =
      CONT [NATURAL_ORDER]
then op diff : Cont[Nat] × Cont[Nat] → Cont[Nat]
      ∀x, y : Nat; C, C' : Cont[Nat]
      • diff (C, empty) = C
      • diff (empty, C') = empty
      • diff (insert(x, C), insert(y, C')) =
          insert(x, diff (C, insert(y, C'))) when x < y
          else diff (C, C') when x = y
          else diff (insert(x, C), C')
      • x is_in diff (C, C') ⇔ (x is_in C ∧ ¬(x is_in C'))
end

arch spec INCONSISTENT =
units  N : NATURAL_ORDER;
        C : CONT [NATURAL_ORDER] given N;
        D : CONT_DIFF_1 given C
result D

```

The specification `CONT_DIFF_1` is consistent (has some models, for instance sorted lists, in increasing order, without repetitions), but is not a conservative extension of `CONT [NATURAL_ORDER]` (since, for instance, a model of `CONT [NATURAL_ORDER]` where containers are realized by arbitrary lists, possibly with repetitions, cannot be expanded into a model of `CONT_DIFF_1` – in that case, the last two axioms are contradictory). As a consequence, in the architectural specification `INCONSISTENT`, the specification of the component *D* is inconsistent, since no component can expand *all* implementations *C* of `CONT [NATURAL_ORDER]` into models of `CONT_DIFF_1`. The architectural specification `INCONSISTENT` is therefore itself inconsistent.

To summarize, architectural specifications not only prescribe the intended architecture of the implementation of the system, but they also ensure that the specified components can be developed independently of each other (which imposes a certain degree of genericity for these components).

## 8.2 Generic Components

*Genericity of components can be made explicit in architectural specifications.*

```

arch spec SYSTEM_G =
units  N : NATURAL_ORDER;
        F : NATURAL_ORDER → CONT [NATURAL_ORDER];
        G : CONT [NATURAL_ORDER] → CONT_DIFF [NATURAL_ORDER]
result G [F [N]]

```

The architectural specification `SYSTEM_G` is a variant of `SYSTEM`; here we choose to specify the second and third components as explicit *generic components*.

The line:

$$F : \text{NATURAL\_ORDER} \rightarrow \text{CONT} [\text{NATURAL\_ORDER}]$$

declares a generic component  $F$ . Given any component implementing (i.e., model of) `NATURAL_ORDER`,  $F$  should expand it into an implementation of `CONT [NATURAL_ORDER]`. The models of the generic-component specification  $\text{NATURAL\_ORDER} \rightarrow \text{CONT} [\text{NATURAL\_ORDER}]$  are functions that map any model of `NATURAL_ORDER` to a model of `CONT [NATURAL_ORDER]`. These functions are required to be persistent, meaning that the result model expands the argument model.

The third component  $G$  is also specified as a generic component: given any implementation of `CONT [NATURAL_ORDER]`,  $G$  should expand it into an implementation of `CONT_DIFF [NATURAL_ORDER]`.

Hence the whole system is obtained by the composition of applications  $G [F [N]]$ , as described in the result part. In CASL, such combinations of components are called *unit terms*. (More complex examples of unit terms will be illustrated in examples below.)

The component  $C$  of `SYSTEM` corresponds to the application  $F [N]$  in `SYSTEM_G`, and similarly the component  $D$  in `SYSTEM` corresponds to  $G [C]$ , i.e., to  $G [F [N]]$  in `SYSTEM_G`.

The models of a specification of the form  $SP1 \rightarrow SP2$  are generic components  $GC$  that should always expand their argument into a model of the target specification. This only makes sense as long as the signature of the target specification contains the signature of  $SP1$ . This is why in CASL,  $SP2$  is always considered as an implicit extension of  $SP1$ , and  $SP1 \rightarrow SP2$  abbreviates  $SP1 \rightarrow \{ SP1 \text{ then } SP2 \}$ .<sup>1</sup> Moreover, since the generic component  $GC$

<sup>1</sup> When  $SP2$  is already defined as an extension of  $SP1$ , as it is the case for instance here for `CONT_DIFF [NATURAL_ORDER]`,  $SP2$  is equivalent to  $SP1 \text{ then } SP2$ .

should expand *any* model of  $SP1$ , the specification  $SP1 \rightarrow SP2$  is *consistent* (i.e., has some models) if and only if the specification  $SP1$  **then**  $SP2$  is a conservative extension of  $SP1$ . Forgetting this fact is a potential source of inconsistent specifications of generic components in architectural specifications. For instance, the specification  $\text{CONT} [\text{NATURAL\_ORDER}] \rightarrow \text{CONT\_DIFF\_1}$  is inconsistent, for the reasons explained at the end of the previous section.

*A generic component may be applied to an argument richer than required by its specification.*

```

arch spec SYSTEM_A =
units  NA : NATURAL_ARITHMETIC;
        F  : NATURAL_ORDER  $\rightarrow$  CONT [NATURAL_ORDER];
        G  : CONT [NATURAL_ORDER]  $\rightarrow$  CONT_DIFF [NATURAL_ORDER]
result G [F [NA]]

```

The above architectural specification  $\text{SYSTEM\_A}$  is a variant of  $\text{SYSTEM\_G}$ . Here we require a component  $NA$  implementing the specification  $\text{NATURAL\_ARITHMETIC}$ , instead of a component  $N$  implementing  $\text{NATURAL\_ORDER}$  as in  $\text{SYSTEM\_G}$  (perhaps because we know that such a component is already available in some collection of previously-implemented components.)

The generic component  $F$  requires a component fulfilling the specification  $\text{NATURAL\_ORDER}$ , but can of course be applied to a richer argument, as in  $F [NA]$ . A similar reasoning applies to  $G$ .

More generally, a generic component can be applied to any component (or to any unit term) that can be reduced along some morphism to an argument of the required ‘type’ (i.e., to a model of the required specification). When necessary, a fitting symbol map can be used to describe the correspondence between the symbols provided by the argument and those expected by the generic component. We do not detail here the technicalities related to these fitting symbol maps, since they are quite similar to those used in instantiations of generic specifications and the notations are the same.

As a last remark, note that, similarly to what happens when instantiating a generic specification by an argument specification, when a generic component is applied to an argument richer than required, the extra symbols are kept in the result. Hence the result of the architectural specification  $\text{SYSTEM\_A}$  above contains also the interpretations of the arithmetic and ordering operations on natural numbers, as they are provided by the component  $NA$ . This means in particular that the implementations described by  $\text{SYSTEM\_A}$  have a larger signature than the ones described by  $\text{SYSTEM\_G}$ .

*Specifications of components can be named for further reuse.*

**unit spec** CONT\_COMP = ELEM  $\rightarrow$  CONT [ELEM]

**unit spec** DIFF\_COMP = CONT [ELEM]  $\rightarrow$  CONT\_DIFF [ELEM]

**arch spec** SYSTEM\_G1 =

**units**  $N$  : NATURAL\_ORDER;

$F$  : CONT\_COMP;

$G$  : DIFF\_COMP

**result**  $G[F[N]]$

In the above example, we give the name CONT\_COMP to the specification (of generic components) ELEM  $\rightarrow$  CONT [ELEM]. Similarly, we give the name DIFF\_COMP to the specification CONT [ELEM]  $\rightarrow$  CONT\_DIFF [ELEM]. Then both named specifications can be reused in the architectural specification SYSTEM\_G1 which is similar to the architectural specification SYSTEM\_G.

In the architectural specification SYSTEM\_G1, we use again the fact that the generic component  $F$  can be applied to richer arguments than models of ELEM (and similarly for  $G$ ). Since ELEM is more general (has more models) than NATURAL\_ORDER, there are potentially fewer possibilities for implementing the generic component specified by CONT\_COMP (which should be compatible with any model of ELEM) than there are for implementing the generic component specified by NATURAL\_ORDER  $\rightarrow$  CONT [NATURAL\_ORDER] (which only needs to be compatible with models of NATURAL\_ORDER; a similar argument holds for DIFF\_COMP). As a consequence, the architectural specifications SYSTEM\_G and SYSTEM\_G1 do not describe the same implementations of the requirements specification REQ.

*Both named and unnamed specifications can be used to specify components.*

**unit spec** DIFF\_COMP\_1 =

CONT [ELEM]  $\rightarrow$  { **op**  $diff : Cont[Elem] \times Cont[Elem] \rightarrow Cont[Elem]$   
 $\forall e : Elem; C, C' : Cont[Elem]$   
 $\bullet e \text{ is\_in } diff(C, C') \Leftrightarrow$   
 $(e \text{ is\_in } C \wedge \neg(e \text{ is\_in } C'))$  }

So far we have always used named (structured) specifications to specify components. unnamed specifications can be used as well, as illustrated by the above variant DIFF\_COMP\_1 of DIFF\_COMP. Here, for the sake of the example, we directly specify the *diff* operation instead of referring to the

named specification `CONT_DIFF`. Remember that in a specification of a generic component of the form  $SP1 \rightarrow SP2$ ,  $SP2$  is always considered as an implicit extension of  $SP1$ , which explains why the above example is well-formed.

*Specifications of generic components should not be confused with generic specifications.*

Generic specifications naturally give rise to specifications of generic components, which can be named for later reuse, as illustrated above by `CONT_COMP`. However, the reader should not confuse a generic specification (which is nothing other than a piece of specification that can easily be adapted by instantiation) with the corresponding specification of a generic component: the latter cannot be instantiated, it is the specified *generic component* which gets *applied* to suitable components.

Conservative extensions of the form ‘**spec**  $SP2 = SP1$  **then**  $SP$ ’ also naturally give rise to specifications of generic components of the form  $SP1 \rightarrow SP2$ , as illustrated by `DIFF_COMP` above.

*A generic component may be applied more than once in the same architectural specification.*

```
arch spec OTHER_SYSTEM =
units  N : NATURAL_ORDER;
       C : COLOR;
       F : CONT_COMP
result F [N] and F [C fit Elem ↦ RGB]
```

The above architectural specification requires a component  $N$  specified by `NATURAL_ORDER`, a component  $C$  specified by `COLOR`, and a generic component  $F$  specified by `CONT_COMP`. Then, as described by the result part, the desired system is obtained by applying  $F$  to  $N$  and applying  $F$  to  $C$  (in this case, an explicit fitting symbol map is necessary, since `COLOR` exports two sorts *RGB* and *CMYK*). Finally both application results are combined, which is expressed by ‘**and**’.

Apart from ‘**free**’, all specification-building operations for structured specifications have natural counterparts at the level of components, which are expressed using the same keywords.<sup>2</sup> The reader should remember that specification-building operations work with specifications defining classes of

<sup>2</sup> The situation is however a bit different with specification extensions, which lead to specifications of generic components, as explained above, or to specifications of components expanding a given component, as illustrated in the previous section.

models (e.g., union of specifications, denoted by ‘**and**’), while in architectural specifications we work with individual models (corresponding to components, as is the case here in `OTHER_SYSTEM` where ‘**and**’ is used to combine the two components  $F[N]$  and  $F[C \text{ fit } Elem \mapsto RGB]$ ).

Hence renaming and hiding also have natural counterparts at the level of components. For instance, remember that the implementations described by `SYSTEM_A` have a larger signature than the implementations described by `SYSTEM_G`. It is however easy to modify the result part of `SYSTEM_A` if what we really want are implementations with the same signature as the implementations described by `SYSTEM_G`: one has just to hide the extra symbols resulting from the component `NA` as follows:

**result**  $G[F[NA]]$  **hide**  $1, \_ + \_, \_ * \_$

or:

**result**  $G[F[NA \text{ hide } 1, \_ + \_, \_ * \_]]$

Symbol maps used in renaming and hiding at the level of components follow the same rules as symbol maps used in renaming and hiding at the level of structured specifications (see Chap. 6).

*Several applications of the same generic component is different from applications of several generic components with similar specifications.*

```
arch spec OTHER_SYSTEM_1 =
units  N  : NATURAL_ORDER;
        C  : COLOR;
        FN : NATURAL_ORDER → CONT [NATURAL_ORDER];
        FC : COLOR → CONT [COLOR fit Elem ↦ RGB]
result FN[N] and FC[C]
```

The above architectural specification `OTHER_SYSTEM_1` is a variant of `OTHER_SYSTEM`. However, in `OTHER_SYSTEM`, we insist on choosing one implementation for containers in the generic component  $F$ , and then we apply it twice, first to a component  $N$  implementing `NATURAL_ORDER`, and then to a component  $C$  implementing `COLOR`. In contrast, in `OTHER_SYSTEM_1`, we may choose two different implementations for containers, one for containers of natural numbers in the component  $FN$  and another one for containers of colors in the component  $FC$ .

The architectural specifications `OTHER_SYSTEM` and `OTHER_SYSTEM_1` are therefore similar but clearly different. Neither is better than the other: each corresponds to a different architectural decision, and selecting one rather than the other is a matter of architectural design. Components that are more widely reusable tend to have less efficient implementations, in general. (Here



the fact that *RGB* has only three values might be exploited in *FC* to give a more space-efficient representation of containers than is possible for *FN*.)

*Generic components may have more than one argument.*

```

unit spec SET_COMP = ELEM → GENERATED_SET [ELEM]

spec  CONT2SET [ELEM] =
      CONT [ELEM] and GENERATED_SET [ELEM]
then op elements_of __ : Cont[Elem] → Set
      ∀ e : Elem; C : Cont[Elem]
      • elements_of empty = empty
      • elements_of insert(e, C) = {e} ∪ elements_of C
end

arch spec ARCH_CONT2SET_NAT =
units  N : NATURAL_ORDER;
        C : CONT_COMP;
        S : SET_COMP;
        F : CONT [ELEM] × GENERATED_SET [ELEM]
            → CONT2SET [ELEM]
result F [C [N]] [S [N]]

```

The architectural specification *ARCH\_CONT2SET\_NAT* requires a component *N* implementing *NATURAL\_ORDER*, a generic component *C* implementing *CONT\_COMP*, i.e., containers, and a generic component *S* implementing *SET\_COMP*, i.e., sets. Then it further requires a generic component *F* that, given *any pair of compatible* models *X* of *CONT [ELEM]* and *Y* of *GENERATED\_SET [ELEM]*, expands them into a model of *CONT2SET [ELEM]*.

Models *X* and *Y* are said to be *compatible* if they share a common interpretation for all symbols they have in common. Here the only symbol they have in common is the sort *Elem*, so the compatibility condition means that *X* and *Y* have the same carrier set for *Elem*. Compatibility is a natural condition, since it is obviously necessary that *X* and *Y* have a common interpretation of their common symbols, otherwise they cannot be both expanded to the same more complex component.

The result is then produced by applying *F* to the pair obtained by applying *C* to *N* and *S* to *N*. Here the pair of arguments *C [N]* and *S [N]* are obviously compatible, since their common symbols (the sort *Nat* equipped with the operations *0* and *suc*) all come from the *same* component *N* which provides their interpretation, which is expanded (hence cannot be modified) in *C [N]* and in *S [N]*, thus compatibility is guaranteed.

*Open systems can be described by architectural specifications using generic unit expressions in the result part.*

```

arch spec ARCH_CONT2SET =
units  C : CONT_COMP;
        S : SET_COMP;
        F : CONT [ELEM] × GENERATED_SET [ELEM]
            → CONT2SET [ELEM]
result λ X : ELEM • F [C [X]] [S [X]]

```

```

arch spec ARCH_CONT2SET_USED =
units  N : NATURAL_ORDER;
        CSF : arch spec ARCH_CONT2SET
result CSF [N]

```

So far our example architectural specifications have described ‘closed’, stand-alone systems where all components necessary to build the desired system were declared in the architectural specification of interest. In CASL, it is however possible to describe ‘open’ systems, i.e., systems made of some components that would require further components to provide a ‘closed’ system. This is illustrated by the architectural specification `ARCH_CONT2SET` which describes a system with a generic component *C* implementing containers, a generic component *S* implementing sets, and a generic component *F* that expands them to provide an implementation of the operation *elements\_of*. The result part is therefore a generic structured component, i.e., an ‘open’ system, which, given any component *X* implementing `ELEM`, provides a system built by applying *F* to the pair made of the applications of *C* to *X* and of *S* to *X*. In CASL, ‘λ’ is input as ‘`lambda`’.

As illustrated by `ARCH_CONT2SET_USED`, we can then describe a ‘closed’ system made of a component *N* implementing `NATURAL_ORDER`, and of an ‘open’ system *CSF* specified by `ARCH_CONT2SET`, which is then applied to *N* in the result part.

### 8.3 Writing Meaningful Architectural Specifications

In the previous sections we have already pointed out potential sources of inconsistent specifications of components. Another issue which deserves some attention when designing an architectural specification is *compatibility* between components (or, more generally, unit terms) that are to be combined together, either by ‘**and**’, or by fitting them to a generic component with multiple arguments.

*When components are to be combined, it is best to check that any shared symbol originates from the same non-generic component.*

```

arch spec ARCH_CONT2SET_NAT_1 =
units  N : NATURAL_ORDER;
        C : CONT_COMP;
        S : SET_COMP;
        G : { CONT [ELEM] and GENERATED_SET [ELEM] }
            → CONT2SET [ELEM]
result G [C [N] and S [N] fit Cont[Elem] ↦ Cont[Nat]]

```

The architectural specification ARCH\_CONT2SET\_NAT\_1 is a variant of ARCH\_CONT2SET\_NAT where, instead of declaring a generic component  $F$  with two arguments, we now declare a generic component  $G$  with a single argument, which must be a model of the specification  $\{ \text{CONT [ELEM]} \text{ **and** GENERATED\_SET [ELEM]} \}$ , obtained as the union of the two (trivially instantiated) specifications of containers and sets.

As a consequence, to obtain the desired system, in the result part we apply the generic component  $G$  to the combination (denoted by ‘**and**’) of  $C$  applied to  $N$  and of  $S$  applied to  $N$ .<sup>3</sup> This combination makes sense only if both  $C [N]$  and  $S [N]$  share the same interpretation of their common symbols. Here their common symbols (the sort  $Nat$  equipped with the operations  $\emptyset$  and  $suc$ ) all come from the *same* component  $N$  which provides their interpretation, which is expanded (hence cannot be modified) in  $C [N]$  and in  $S [N]$ , thus compatibility is guaranteed.

There is a clear analogy here between the application of the generic component  $F$  with multiple arguments in ARCH\_CONT2SET\_NAT and the combination of  $C [N]$  and  $S [N]$  in ARCH\_CONT2SET\_NAT\_1: in both cases the result is meaningful because we can trace shared symbols like the sort  $Nat$  and the operations  $\emptyset$  and  $suc$  to a single component  $N$  introducing them.

Let us emphasize again that compatibility is a natural requirement: since each unit declaration corresponds to a separate implementation task (and hence each unit subterm to an independently developed subsystem), obviously the combination of components or unit terms makes sense only when some compatibility conditions are fulfilled.

Let us now consider an example where the compatibility condition is violated.

---

<sup>3</sup> In the application of the generic component  $G$  we need an explicit fitting symbol map since otherwise the sort  $\text{Cont}[Elem]$  can ambiguously be mapped to either  $\text{Cont}[Nat]$  or  $\text{Set}$ .

```

arch spec WRONG_ARCH_SPEC =
units  CN : CONT [NATURAL_ORDER];
        SN : GENERATED_SET [NATURAL_ORDER];
        F  : CONT [ELEM] × GENERATED_SET [ELEM]
              → CONT2SET [ELEM]
result F [CN] [SN]

```

The architectural specification `WRONG_ARCH_SPEC` is a variant of `ARCH_CONT2SET_NAT` where, instead of requiring a component  $N$  implementing `NATURAL_ORDER` and two generic components implementing containers and sets respectively, we just require a component  $CN$  implementing containers of natural numbers and a component  $SN$  implementing sets of natural numbers. However, then the application  $F [CN] [SN]$  makes no sense since there is no way to ensure that the common symbols of  $CN$  and  $SN$  have the same interpretation. It may indeed be the case that natural numbers are interpreted in some way in  $CN$  and in a different way in  $SN$ , which makes the application of  $F$  impossible. (Hence a similar problem would arise if one would use the combination of components ‘ $CN$  and  $SN$ ’.)

Let us now consider a more complex example.

```

arch spec BADLY_STRUCTURED_ARCH_SPEC =
units  N : NATURAL_ORDER;
        A : NATURAL_ORDER → NATURAL_ARITHMETIC;
        C : CONT_COMP;
        S : SET_COMP;
        F : CONT [ELEM] × GENERATED_SET [ELEM]
              → CONT2SET [ELEM]
result F [C [A [N]]] [S [A [N]]]

```

The architectural specification `BADLY_STRUCTURED_ARCH_SPEC` is a variant of `ARCH_CONT2SET_NAT` where, in addition to the component  $N$  implementing `NATURAL_ORDER`, we require a generic component  $A$  which is used to expand  $N$  into an implementation of `NATURAL_ARITHMETIC`. In the architectural specification `ARCH_CONT2SET_NAT`, the compatibility condition in the application  $F [C [N]] [S [N]]$  was easy to discharge. Here, in the result unit term  $F [C [A [N]]] [S [A [N]]]$  of `BADLY_STRUCTURED_ARCH_SPEC`, we apply  $F$  to the pair made of  $C [A [N]]$  and  $S [A [N]]$ . In this case only a semantic analysis can ensure that these two arguments are compatible, since the common symbols cannot be traced to the same non-generic component, but only to two applications of the same generic component  $A$  to similar arguments. (Actually the arguments are just the same here, but in general checking this would require non-trivial semantic reasoning.)

It is advisable to use unit terms where compatibility can be checked by a simple static analysis. CASL provides additional constructs which make it easy to follow this recommendation, as explained below.

*Auxiliary unit definitions or local unit definitions may be used to avoid repetition of generic unit applications.*

```

arch spec WELL_STRUCTURED_ARCH_SPEC =
units  N  : NATURAL_ORDER;
        A  : NATURAL_ORDER → NATURAL_ARITHMETIC;
        AN = A [N];
        C  : CONT_COMP;
        S  : SET_COMP;
        F  : CONT [ELEM] × GENERATED_SET [ELEM]
              → CONT2SET [ELEM]
result F [C [AN]] [S [AN]]

arch spec ANOTHER_WELL_STRUCTURED_ARCH_SPEC =
units  N  : NATURAL_ORDER;
        A  : NATURAL_ORDER → NATURAL_ARITHMETIC;
        C  : CONT_COMP;
        S  : SET_COMP;
        F  : CONT [ELEM] × GENERATED_SET [ELEM]
              → CONT2SET [ELEM]
result local AN = A [N] within F [C [AN]] [S [AN]]

```

The problem illustrated in BADLY\_STRUCTURED\_ARCH\_SPEC can be fixed easily. An auxiliary unit definition may be used to avoid the repetition of generic unit applications, such as ‘ $AN = A[N]$ ’ in WELL\_STRUCTURED\_ARCH\_SPEC. An alternative is to make the definition of  $AN$  local to the result unit term, as illustrated in ANOTHER\_WELL\_STRUCTURED\_ARCH\_SPEC. In both cases common symbols can be traced to a non-generic unit, and compatibility can be checked by an easy static analysis.

---

## Libraries

*Libraries are named collections of named specifications.*

In the foregoing chapters, we have seen many examples of named specifications, and of references to them in later specifications. This chapter explains how a collection of named specifications can itself be named, as a *library*. The creation of libraries facilitates the *reuse* of specifications. For practical applications, it is important to be able to reuse (at least) existing specifications of basic datatypes, such as those described in Chap. 12.

*Local libraries are self-contained.*

A library is called *local* when it is self-contained, i.e., for each reference to a specification name in the library, the library includes a specification with that name. Local libraries might appear at first sight to be all that we need, but actually they provide poor support for reuse of specifications. The problem is that when a specification from one local library is reused in another, it has to be repeated *verbatim*. There is no formal link between the original specification and the copy, despite them having the same name: the names used in a library can be chosen freely, and different libraries could use the same name for completely different specifications.

*Distributed libraries support reuse.*

*Distributed libraries* allow duplication of specifications to be avoided altogether. Instead of making an explicit copy of a named specification from

one library for use in another, the second library merely indicates that the specification concerned can be *downloaded* from the first one.

*Different versions of the same library are distinguished by hierarchical version numbers.*

In practice, specifications *evolve*, e.g., to provide further operations or predicates on the specified sorts, or to define new subsorts. The libraries containing the specifications can evolve too, by adding or removing named specifications. Without some form of version control, even a trifling change in one library might cause specifications in other libraries to become ill-formed, or affect their meanings. CASL allows different versions of the same library to coexist (distinguishing them by hierarchical version numbers), and allows downloadings in a library to indicate that a particular version of another library is required.

Creation of new libraries is essential in connection with larger specification projects, and projects of any scale can benefit from reuse of specifications from existing libraries. The rest of this chapter illustrates the constructs used to specify local libraries, distributed libraries, and versions, and gives some advice on the organization of libraries.

## 9.1 Local Libraries

*Local libraries are self-contained collections of specifications.*

```
library USERMANUAL/EXAMPLES
```

```
...
```

```
spec NATURAL = ...
```

```
...
```

```
spec NATURAL_ORDER = NATURAL then ...
```

```
...
```

The collection of all the illustrative examples given in the foregoing chapters is self-contained, so it could be made into a local library and named USERMANUAL/EXAMPLES, as outlined above. To provide a separate local library for each chapter would however involve a considerable amount of duplication, since many of the specifications that are defined in the earlier chapters are also referenced in later chapters (e.g., SET\_PARTIAL\_CHOOSE in Chap. 4 in-

stantiates `GENERATED_SET`, which is defined in Chap. 3). Using *distributed* libraries, as explained in Sect. 9.2, this duplication can be avoided.<sup>1</sup>

The ‘same name, same thing’ principle of CASL applies only within specifications, and it is possible for a library to include alternative specifications for the same symbols (e.g., using different sets of axioms). However, when such alternative specifications are both extended (perhaps indirectly) in the same specification, the principle does apply, and unintended name clashes might then arise. Thus in general, it is advisable for the developers of a library to respect the ‘same name, same thing’ principle when choosing symbols throughout the library. In any case, this is obviously helpful to those who might later browse the library. Alternative specifications for the same symbols should therefore be given in separate libraries.<sup>2</sup>

*Specifications can refer to previous items in the same library.*

**library** USERMANUAL/EXAMPLES

...

**spec** STRICT\_PARTIAL\_ORDER = ...

...

**spec** TOTAL\_ORDER = STRICT\_PARTIAL\_ORDER **then** ...

...

**spec** PARTIAL\_ORDER = STRICT\_PARTIAL\_ORDER **then** ...

...

Although we may often regard libraries as *sets* of named specifications, they are actually *sequences*, and the order in which the specifications occur is significant.

Specification names have linear visibility: each specification can refer only to the names of the specifications that precede it. Thus a series of extensions has to be presented in a bottom-up fashion, starting with a specification that is entirely self-contained, containing no references to other specifications at all. Each specification name in a library has a unique defining occurrence, so overriding cannot arise. Extensions that do not refer to each other may be given in any order (e.g., `PARTIAL_ORDER` above could just as well be given before `TOTAL_ORDER`).

Linear visibility of specification names means that mutual recursion between specifications is prohibited. When two specifications each make use of symbols declared in the other, the declarations of those symbols have to be duplicated, or moved to a preceding specification that can then be referenced by them both.

<sup>1</sup> A distributed library for each chapter of Part II is available via the CoFI web pages; copies are provided on the CD-ROM accompanying this book.

<sup>2</sup> If we intended our comprehensive USERMANUAL/EXAMPLES library for general use, we would remove all the illustrative alternative specifications.



*All kinds of named specifications can be included in libraries.*

```
library USERMANUAL/EXAMPLES
...
spec STRICT_PARTIAL_ORDER = ...
...
spec GENERIC_MONOID [sort Elem] = ...
...
view INTEGER_AS_TOTAL_ORDER : ...
...
view LIST_AS_MONOID [sort Elem] : ...
...
arch spec SYSTEM = ...
...
unit spec CONT_COMP = ...
...
```

Items in libraries can be any kind of named specification, as illustrated above: simple named specifications, generic specifications, named view definitions, generic view definitions, and architectural and unit specifications. We shall henceforth refer to them generally as *library items*.

Libraries themselves never include anonymous specifications, such as declarations of sorts and operations. Moreover, the symbols declared by a library item are *not* automatically available for use in subsequent items: an explicit reference to the name of the library item is required to ‘import’ the item.

Technically, each library item is said to be *closed*, being interpreted without any pre-declared symbols at all. This facilitates downloading items from distributed libraries, see Sect. 9.2.

*Display, parsing, and literal syntax annotations apply to entire libraries.*

```
library USERMANUAL/EXAMPLES
...
%display __<=__      %LATEX _ ≤ _
%display __>=__      %LATEX _ ≥ _
%display __union__   %LATEX _ ∪ _
%prec {__+__, __-__} < {__*__, __/__}
%left_assoc __+__, __*__
...
spec STRICT_PARTIAL_ORDER = ...
...
```

```

spec PARTIAL_ORDER = STRICT_PARTIAL_ORDER then ...  $\leq$  ...
...
spec GENERATED_SET [sort Elem] = ...  $\cup$  ...
...
spec INTEGER_ARITHMETIC_ORDER = ...  $\leq$  ...  $\geq$  ...
...
```

Annotations affecting the way terms are written or displayed apply to an entire library, and have to be collected at the beginning of the library. These annotations include display and precedence annotations, illustrated above.

Recall that various reserved words and symbols in CASL specifications are input in ASCII, but displayed as mathematical signs (e.g., universal quantification is input as ‘forall’, and displayed as ‘ $\forall$ ’ when this sign is available in the current display format). *Display annotations* provide analogous flexibility for declared symbols. For example, the display annotations illustrated above determine how infix symbols input as ‘<=’, ‘>=’, and ‘union’ are displayed when using L<sup>A</sup>T<sub>E</sub>X to format the specification. Note that a display annotation applies to all occurrences of the input symbol in the library, regardless of overloading.

Display annotations can give alternative displays for different formats: apart from L<sup>A</sup>T<sub>E</sub>X, both RTF and HTML are presently supported. The display of the annotation itself shows only the input syntax of the symbol and the result produced by the current formatter. The input form of one of the above annotations might be as follows:

```
%display __union__ %HTML __<sup>U</sup> %LATEX __\cup__
```

When no display annotation is given for a particular format, the input format itself is displayed. Thus the symbol displayed as ‘ $\cup$ ’ in the present L<sup>A</sup>T<sub>E</sub>X version of this User Manual would be displayed as ‘union’ in an RTF version, unless the above annotation were to be extended with an RTF part.

*Parsing annotations* allow omission of grouping parentheses when terms are input. A single annotation can indicate the relative precedence or the associativity (left or right) of a group of operation symbols. The precedence annotation for infix arithmetic operations given above, namely:

```
%prec {__+__, __-__} < {__*__, __/__}
```

allows a term such as  $a + (b * c)$  to be input (and hence also displayed) as  $a + b * c$ . The left-associativity annotation for + and \*:

```
%left_assoc __+__, __*__
```

allows  $(a + b) + c$  to be input as  $a + b + c$ , and similarly for \*; but the parentheses cannot be omitted in  $(a + b) - c$  (not even if ‘\_\_-\_\_’ were to be included in the same left-associativity annotation).

When an operation symbol is declared with the associativity attribute *assoc*, an associativity *annotation* for that symbol is provided automatically.<sup>3</sup> Thus in practice, explicit associativity annotations are needed only for non-associative operations such as subtraction and division.

*Libraries and library items can have author and date annotations.*

```
library USERMANUAL/EXAMPLES
%authors( Michel Bidoit <bidoit@lsv.ens-cachan.fr>,
          Peter D. Mosses <pdmosses@brics.dk>      )%
%dates 15 Oct 2003, 1 Apr 2000
...
spec STRICT_PARTIAL_ORDER = ...
...
%authors Michel Bidoit <bidoit@lsv.ens-cachan.fr>
%dates 10 July 2003
spec INTEGER_ARITHMETIC_ORDER =
...
```

An *author annotation* at the beginning of a library indicates the collective authorship of the entire library; one preceding an individual library item indicates its specific authorship.

A *date annotation* at the beginning of a library should indicate the release date of the current version of the library. It may also give the release dates of some previous major versions, possibly including that of the original version. A date annotation on an individual library item should indicate when that item was last changed, and (optionally) the dates of previous changes.

## 9.2 Distributed Libraries

*Libraries can be installed on the Internet for remote access.*

```
library BASIC/NUMBERS
...
%left_assoc _@@_
%number _@@_
%floating _::_, _E_
%prec { _E_ } < { _:::_ }
```

<sup>3</sup> This implicit parsing annotation is local to the enclosing specification and to specifications that reference it, in contrast to ordinary parsing annotations.

```

...
spec NAT =
  free type Nat ::= 0 | suc(Nat)
  ...
  ops   1 : Nat = suc(0); ...; 9 : Nat = suc(8);
        --@@--(m, n : Nat) : Nat = (m * suc(9)) + n
  ...
spec INT = NAT then ...
spec RAT = INT then ...
spec DECIMALFRACTION = RAT then
  ...
  ops   ___::___ : Nat × Nat → Rat;
        --E_ : Rat × Int → Rat
  ...

```

The above example is an extract from one of the CASL libraries of basic datatypes, described in Chap. 12 and available on the Internet. It illustrates the overall structure of a library intended for general use, as well as some helpful annotations concerning literal syntax for numbers, which are explained below.

*Validated libraries can be registered for public access.*

CoFI will maintain a register of useful libraries. Registered CASL libraries are identified by hierarchical path names. For instance, all the CASL libraries of basic datatypes have names starting with ‘BASIC/’, and path names starting with ‘CASL/’ are reserved for libraries connected with the CASL language itself (e.g., the specification of the abstract syntax of CASL in CASL).

Registered libraries will be mirrored at several sites, to ensure their continuous accessibility. The URLs of a library can be obtained from the library name using a table provided on the CoFI web pages.

Libraries have to be *validated* before registration. The validation of a library ensures not only that it is well-formed, but also that semantic annotations expressing consistency of specifications (or conservativity over the parameters, in case of generic or unit specifications) have been added, and that all proof obligations (corresponding both to well-formedness conditions and to semantic annotations in the library) have been verified.

It is likely that new versions of existing libraries will be produced, e.g., providing further operations whose usefulness was not realized beforehand. Although the assignment and use of library version numbers allows users to protect their specifications from changes due to new versions (see Sect. 9.3), at least the names used in a registered library should not change much between versions.

*Libraries should include appropriate annotations.*

In particular, parsing and display annotations can be provided, as explained in Sect. 9.1. The above example illustrates a further kind of annotation, used to provide *literal syntax* for numbers in CASL. The effect of the illustrated annotations is that, after downloading the appropriate specifications from the library BASIC/NUMBERS, conventional decimal notation can be used for integers and decimal fractions, e.g.,  $42$ ,  $2.718$ ,  $10E-12$ . The digits  $42$  are interpreted as the term  $4@@2$ , and  $2.718$  is interpreted as the term  $2::718$  (where  $718$  is subsequently interpreted as  $(7@@1)@@8$ ). The definition of the operation  $@@$  is shown above; those of  $---$  and  $..E..$  are a bit more involved, and omitted here. Notice that the library BASIC/NUMBERS is *not* hard-wired into CASL, and users could provide annotations to interpret the literal syntax for integers and decimal fractions as terms involving different operations, e.g., on different sorts.

*Libraries can include items downloaded from other libraries.*

**library** BASIC/STRUCTUREDDATATYPES

...

**from** BASIC/NUMBERS **get** NAT, INT

...

**spec** LIST [**sort** *Elem*] **given** NAT = ...

...

**spec** ARRAY ... **given** INT = ...

...

Individual specifications and other items can be *downloaded* from other libraries. For example, the library BASIC/STRUCTUREDDATATYPES, outlined above, does not itself provide the specifications of natural numbers and integers that are needed in the specifications LIST and ARRAY, but instead downloads NAT and INT from the BASIC/NUMBERS library.

The names of the items to be downloaded from a library have to be listed explicitly: one cannot request the downloading of all the items that happen to be provided by a library. However, although the name of each item provided by a downloading is always explicit, no indication is given of its kind (i.e., whether it is an ordinary, generic, or architectural specification, or a view) nor of what symbols it declares. Thus the well-formedness of a library depends on what items are actually downloaded from other libraries.

The construct '**from** ... **get** ...' above has the effect of downloading the specifications that are named NAT and INT in BASIC/NUMBERS, preserving their names. It is also possible to give downloaded specifications different

names, e.g., to avoid clashes with specification names that are already in use locally:

**from** BASIC/NUMBERS **get** NAT  $\mapsto$  NATURAL, INT  $\mapsto$  INTEGER

Items that are referenced by downloaded items are *not* themselves automatically downloaded, e.g., downloading INT does not entail the downloading of NAT. This is because downloading involves the *semantics* of the named items, not their *text*. The semantics of INT consists of a signature and a class of models, and is a self-contained entity – recall from Chap. 6 that the models of a structured specification have no more structure than do those of a flat, unstructured specification. Thus downloading INT gives exactly the same result as if the reference to NAT in its text had already been replaced by the text of NAT *before* downloading. For the same reason, the presence of another item with the name NAT in the current library makes no difference to the result of downloading INT. In terms of software packages, downloading specifications from CASL libraries is analogous to installing packages from binaries, rather than from sources.

Downloading any item from another library B in a library A causes all the parsing and display annotations of B to be inserted at the beginning of A. (Conflicting annotations from different libraries are ignored altogether, and local annotations override conflicting downloaded annotations.) The copied annotations allow terms to be written and displayed in A in the same way as in B.

*Substantial libraries of basic datatypes are already available.*

The organization of the following libraries of basic datatypes is explained in Chap. 12:

BASIC/NUMBERS: natural numbers, integers, and rationals.

BASIC/RELATIONSANDORDERS: reflexive, symmetric, and transitive relations, equivalence relations, partial and total orders, boolean algebras.

BASIC/ALGEBRA\_I: monoids, groups, rings, integral domains, and fields.

BASIC/SIMPLEDATATYPES: booleans, characters.

BASIC/STRUCTURED DATATYPES: sets, lists, strings, maps, bags, arrays, trees.

BASIC/GRAPHS: directed graphs, paths, reachability, connectedness, colorability, and planarity.

BASIC/ALGEBRA\_II: monoid and group actions on a space, euclidean and factorial rings, polynomials, free monoids, and free commutative monoids.

BASIC/LINEARALGEBRA\_I: vector spaces, bases, and matrices.

BASIC/LINEARALGEBRA\_II: algebras over a field.

BASIC/MACHINENUMBERS: bounded subtypes of naturals and integers.

These libraries form a coherent collection of highly-polished specifications. The libraries themselves are provided in full in the *CASL Reference Manual* [20], and are available on the Internet.

*Libraries need not be registered for public access.*

```
library http://www.cofi.info/CASL/Test/Security
...
from http://casl:password@www.cofi.info/CASL/RSA get KEY
...
spec DECRYPT = KEY then ...
...
```

Libraries under development, and libraries provided for restricted groups of users, are named and accessed by their URLs (instead of the simple path names used for registered libraries). This allows the CASL library constructs to be fully exploited for libraries that are not yet – and perhaps never will be – registered for public access. Moreover, validation of libraries can be a demanding and time-consuming process, and getting a library approved and registered is appropriate only when it provides specifications that are likely to be found useful by persons not directly involved in its development.

The primary Internet access protocols HTTP and FTP both support password protection of libraries and the insertion of usernames and passwords in URLs allows downloading between protected libraries (With HTTP, the username and password can be unrelated to those used for the host file system.)

### 9.3 Version Control

*Subsequent versions of a library are distinguished by explicit version numbers.*

```
library BASIC/NUMBERS version 1.0
...
spec NAT = ...
...
spec INT = NAT then ...
...
spec RAT = INT then ...
...
```

As illustrated above, a library can be assigned an explicit version number, allowing it to be distinguished from previous and future versions of the same library. CASL allows conventional hierarchical version numbers, familiar from version numbers of software packages: the initial digits indicate a major version, digits after a dot indicate sub-versions, and digits after a further dot indicate patches to correct bugs. (Distinctions between alpha, beta, and other pre-release versions are not supported.)

The smallest version number is written simply ‘0’, and can be omitted when specifying the initial version of a library; this is the case with the version of BASIC/NUMBERS shown in Sect. 9.2, it implicitly has version number ‘0’, but in general the first-installed version of a library could have any version number at all. The numbers of successively installed versions do not have to be contiguous, nor even increasing: e.g., a patched version 0.99.1 could be installed after version 1.0.

Individual library items do not have separate version numbers. Date annotations can be used to indicate which items have changed between two versions of a library.

*Libraries can refer to specific versions of other libraries.*

**library** BASIC/STRUCTURED DATATYPES **version** 1.0

...

**from** BASIC/NUMBERS **version** 1.0 **get** NAT, INT

...

**spec** LIST [**sort** *Elem*] **given** NAT = ...

...

**spec** ARRAY ... **given** INT = ...

...

Downloading items from particular versions of libraries is necessary if one wants to ensure coherence between libraries. For example, as illustrated above, version 1.0 of BASIC/STRUCTURED DATATYPES downloads NAT and INT from version 1.0 of BASIC/NUMBERS. Omitting the version number when downloading gives implicitly the *current version* of the library, which may of course change. By the way, the current version of a library is *not* necessarily the one most recently installed: it is the one with the *largest version number*. As previously mentioned, a patched version 0.99.1 could be installed after version 1.0, but a downloading without an explicit version number would still refer to version 1.0.

Even though the developers of libraries may try to ensure backwards compatibility between versions, it could happen that symbols introduced in a new version of a downloaded specification clash with symbols already in use in the library that specified the downloading, causing ill-formedness or inconsistency. So for safety, it is advisable to give explicit version numbers when



downloading (also when downloading from version ‘0’ of another library). If one subsequently wants to use symbols that are introduced only in some later version of another library, all that is needed is to change the version number in the downloading(s).

An alternative strategy is to ensure consistency with the *current* versions of all libraries from which specifications are downloaded, by observing the changes in the new versions and adapting the downloading library accordingly. For instance, one might download INT from the current version of BASIC/NUMBERS, instead of from version 1.0 of that library. This may involve extra work when a new version of BASIC/NUMBERS appears, but it has several advantages over the more cautious approach. CASL leaves the choice to the user, although registered libraries will generally be required to use explicit version numbers when downloading from other libraries.

*All downloadings should be collected at the beginning of a library.*

Although CASL allows downloadings to be interleaved with specification definitions, it is advisable to collect the downloadings at the beginning of libraries (together with any parsing and display annotations). This makes it easy to see dependencies between libraries, and to ensure that different downloadings from the same library all refer to the same version of it.

## Part III

---

### Carrying On

---

## Foundations

Donald Sannella and Andrzej Tarlecki

*A complete presentation of CASL is in the Reference Manual.*

This *User Manual* has introduced the potential user to the features of CASL mainly by means of illustrative examples. It has presented and discussed the typical ways in which the language concepts and constructs are expected to be used in the course of building system specifications. Thus, the presentation in Part II focused on *what* the constructs and concepts of CASL are *for*, and *how* they should (and should not) be used. We tried to make these points as clear as possible by referring to simple examples, and by discussing both the general ideas and some details of CASL specifications. We hope that this has given the reader a sufficient feel of the formalism, and enough understanding, to look through the presentation of a basic library of CASL specifications in Chap. 12, to follow the case study in Chap. 13, and to start experimenting with the use of CASL for writing specifications – perhaps employing the support tools presented in Chap. 11.

By no means, however, should this book be regarded as a complete presentation of the CASL specification formalism – this is given in the accompanying volume, the *CASL Reference Manual* [20].

*CASL has a definitive summary.*

The CASL Reference Manual begins with a definitive *summary* of the entire CASL language: all the language constructs are listed there systematically, together with the syntax used to write them down and a detailed explanation of their intended meaning. However, although it tries to be precise and complete, the CASL Summary still relies on natural language to present CASL.

This inherently leaves some room for interpretation and ambiguity in various corners of the language, for example where details of different constructs interact.

*CASL has a complete formal definition.*

A key aim of the entire CoFI initiative is to avoid any potential ambiguities by providing a complete formal definition for CASL, and sound mathematical foundations for the advocated methodology of its use in software specification and development.

*Abstract and concrete syntax of CASL are defined formally.*

The next part of the Reference Manual is a formal definition of the syntax of CASL. *Abstract syntax* is given, where each phrase is written in a way that directly indicates its components, thus making evident its internal structure. In essence, the use of each construct of the language is explicitly labeled here. This is convenient for formal manipulation and analysis, but is not so readable. Therefore, the so-called *concrete syntax* of CASL (as used for instance in the examples throughout this book) is given as well, retaining a direct correspondence with the abstract syntax. This offers to the user of CASL a convenient and readable way of writing down CASL specifications, in a way that makes clear the formal structure of the phrases and constructs used to build them. As usual, the syntax is given as a context-free grammar, using a variant of the BNF notation, relying on well-established theory to give its formal meaning, and on a variety of tools and techniques available for syntactic analysis of languages presented in such a style.

*CASL has a complete formal semantics.*

The ultimate definition of the meaning of CASL specifications is provided by the *semantics* of CASL in the Reference Manual. The semantics first defines mathematical entities that formally model the intended meaning of various concepts underlying CASL, as already hinted at in Chap. 2, and further introduced and discussed throughout the summary.

The key concepts here are that of CASL signature, model and formula, together with the satisfaction relation between models and formulas over a common signature. In fact, these are variants of the standard algebraic and logical notions, thus linking work on CASL to well-established mathematical theories and ideas.

In a more or less standard way, we use these concepts to build the *semantic domains* which the meanings of phrases in various syntactic categories of CASL inhabit. We have chosen to give the semantics in the form of so-called *natural semantics*, with formal deduction rules to derive judgments concerning the meaning of each CASL phrase from the meanings of its constituent parts. Not only is this a mathematically well-established formalism with an unambiguous interpretation, but we also hope that this makes the semantics itself more readable, with details easier to follow than in some other approaches.

The overall semantics of CASL consists of two parts. The *static semantics* captures a form of static analysis of CASL specifications, in which they are checked for well-formedness – for example, that axioms are well-typed, and references to named entities are in scope. Then the *model semantics* takes a well-formed CASL specification and assigns a model-theoretic meaning to it.

*CASL specifications denote classes of models.*

In CASL, well-formed specifications denote signatures (static semantics) and classes of models (model semantics). Basic specifications, which in essence present a signature and a set of axioms over this signature, denote the class of models that satisfy all the axioms. The semantics of basic specifications is split into two parts: first many-sorted basic specifications are described and then features for defining and using subsorts are added.

*The semantics is largely institution-independent.*

A few more concepts are needed to explain the semantics of structured specifications. Key here is the notion of signature morphism, and the model reducts and translation of sentences that signature morphisms induce. Having introduced those, we obtain the *institution* [23] of CASL, that is, the underlying logical system of CASL equipped with extra structure to capture ways of moving between signatures linked by signature morphisms. One important point of the semantics is that all the layers of the semantics “above” basic specifications are *institution-independent*, i.e., well-defined for any institution chosen to build basic specifications (as long as the institution comes with a bit of extra structure concerned with forming unions of signatures and defining signature morphisms).

Next, we have the semantics of architectural specifications, which relies on a formal counterpart of the concept of a unit (module) of a system to be developed: self-contained units are viewed simply as models of the underlying institution, and parametrized units as functions mapping such parameter models to result models. Architectural specifications provide a way of specifying the component units of a system and indicating how the overall system

is built by putting these units together. This intuition is captured by the semantics of architectural specifications, which denote a class of permitted unit bindings and a function that maps each such environment to a result unit.

Finally, the libraries layer of CASL is given a rather standard description with libraries modeled as environments giving names to the entities introduced (specifications, architectural specifications, etc.).

*The semantics is the ultimate reference for the meanings of all CASL constructs.*

Overall, the formal mathematical semantics is crucial in the understanding of CASL specifications. It provides their unambiguous meaning, and thus gives the ultimate reference point for all questions concerning the interpretation of any CASL phrase in any context.

We have already experienced how important such a formal semantics may be in the design of CASL itself. In many cases, the intended semantics was prominent in internal discussions on the details of the constructs under consideration, and provided guidelines for many choices in the design of CASL. Indeed, the concrete syntax of CASL was designed only after the semantics was settled.

*Proof systems for various layers of CASL are provided.*

The semantics is also a necessary prerequisite for the development of mechanisms for formal reasoning about CASL specifications. This is the role of *proof calculi* that support reasoning about the various layers of CASL. The starting point is a formal system of deduction rules which determines a proof-theoretic counterpart of the consequence relation between sets of formulas, thus providing a way for deriving consequences of sets of axioms in CASL specifications. This is then extended to systems of rules for deriving consequences of structured specifications and for proving inclusions between classes of models of such specifications. These systems are also used in rules for formal verification of the internal correctness of system designs as captured by architectural specifications. For all these systems, their soundness is proved and completeness discussed by reference to the formal semantics of CASL.

One point of interest is that, again, the extension of the basic proof system for consequences between sets of formulas to structured and architectural specifications does not really rely on the specifics of the underlying institution, but just reflects the way in which the structuring and architectural constructs are defined for an arbitrary institution.

*The foundations of our CASL are rock-solid!*

All this work on the mathematical underpinnings of CASL and related specification and development methodology, as documented in the Reference Manual, should make it exceptionally trustworthy – at least in the sense that it provides a formal point of reference against which all the claims may (and should) be checked.

## Tools

Till Mossakowski

This chapter gives an overview of the CASL tools. Analysis tools for CASL like parsers and static checkers, as well as formatters, are stable now and cover the whole of CASL. Proof tools are available but are less mature.

CASL has been designed with the goal of subsuming many previous specification languages. Most of these languages come with specific tools, and of course, these tools should be reusable in the context of CASL. Hence, a central issue is to build bridges to existing tools (rather than building new tools from scratch). Using an interchange format generated by the analysis tools, CASL has been interfaced in this way to rewriting engines and theorem provers, usually working for a subset of CASL.

Naturally, due to the ongoing development of these tools, detailed descriptions would become outdated sooner or later. Therefore, we give here just an appetizer, intended to encourage the reader to install the tools and experiment with them (and to convince her/him that this is rather easy). More detailed descriptions of the tools, as well as their latest versions and other tools that may be developed in the future, are available by following the links on the CoFI tools home page [21]: <http://www.cofi.info/Tools>.

The analysis tools for CASL have been used to check all the examples contained in this book, as well as the CASL Basic Libraries [20]. Moreover, some proofs from a case study in refinement have been carried out with the proof tools.

*CASL specifications can be checked for well-formedness using a form-based web page.*

The easiest way to check a CASL specification for well-formedness is to visit the web interface. Using a web form, you can submit your specification (without the need to install anything on your machine), and get immediate



feedback about the well-formedness of the specification. Follow the “web interface” link on the CoFI tools home page.

## 11.1 The Heterogeneous Tool Set (HETS)

*The Heterogeneous Tool Set (HETS) is the main analysis tool for CASL.*

HETS is a tool set for the analysis of specifications written in CASL, its extensions and sublanguages – hence the name “heterogeneous”. HETS consists of logic-specific tools for the parsing and static analysis of the different CASL extensions and sublanguages, as well as a logic-independent parsing and static analysis tool for structured and architectural specifications and libraries (which of course calls the logic-specific tools whenever a basic specification is encountered). In order to be able to tackle proof obligations occurring in (statically well-formed) specifications, HETS is interfaced with various logic-specific theorem proving, rewriting and consistency checking tools. On top of this, there is a logic-independent proof engine called MAYA, which manages the proof obligations. MAYA uses so-called *development graphs*, a graphical representation of CASL structured specifications.

The architecture of HETS is shown in Fig. 11.1. The latest version can be obtained from the CoFI tools home page [21]. Installation is easy; just follow the instructions.

Consider the first example in this book:

```
spec STRICT_PARTIAL_ORDER =
  sort Elem
  pred -- < -- : Elem × Elem
  ∀x, y, z : Elem
    • ¬(x < x)                                %(strict)%
    • x < y ⇒ ¬(y < x)                        %(asymmetric)%
    • x < y ∧ y < z ⇒ x < z                  %(transitive)%
  %{ Note that there may exist x, y such that
    neither x < y nor y < x. }%
end
```

*HETS can be used for parsing and checking static well-formedness of specifications.*

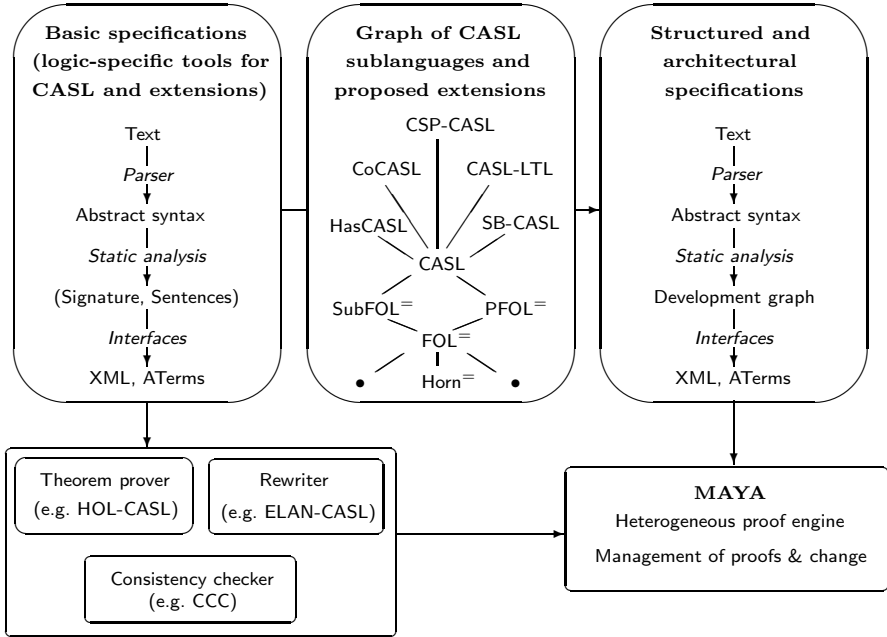


Fig. 11.1. Architecture of the heterogeneous tool set.

Let us assume that the example is in a file named `Order.casl` (actually, this file is provided on the web and on the CD-ROM coming with this volume). Then you can check the well-formedness of the specification by typing (into some shell):

```
hets Order.casl
```

HETS checks both the correctness of this specification with respect to the CASL syntax, as well as its correctness with respect to the static semantics (e.g. whether all identifiers have been declared before they are used, whether operators are applied to arguments of the correct sorts, whether the use of overloaded symbols is unambiguous, and so on).

It is also possible to generate a pretty printed  $\text{\LaTeX}$  version of `Order.casl` by typing:

```
hets -o pp.tex Order.casl
```

One use of `Order.casl` might be to express the fact that the natural numbers form a strict partial order as a view, as follows:

```
spec NATURAL = free type Nat ::= 0 | suc(Nat) end
```

```

spec NATURAL_ORDER_2 =
  NATURAL
then pred  $-- < -- : Nat \times Nat$ 
   $\forall x, y : Nat$ 
  •  $0 < suc(x)$ 
  •  $\neg x < 0$ 
  •  $suc(x) < suc(y) \Leftrightarrow x < y$ 
end

view v1 : STRICT_PARTIAL_ORDER to NATURAL_ORDER_2 =
   $Elem \mapsto Nat$ 
end

```


Again, these specifications can be checked with HETS. However, this only checks syntactic and static semantic well-formedness – it is *not* checked whether the predicate ‘ $-- < --$ ’ introduced in NATURAL\_ORDER\_2 actually is constrained to be interpreted by a strict partial ordering relation. Checking this requires theorem proving, which will be discussed below.

*HETS also displays and manages proof obligations, using development graphs.*

However, before coming to theorem proving, let us first inspect the proof obligations arising from a specification. This can be done with:

```
hets -g Order.casl
```

(assuming that the above two specifications and the view have been added to the file `Order.casl`). HETS now displays a so-called development graph (which is just an overview graph showing the overall structure of the specifications in the library), see Fig. 11.2.



**Fig. 11.2.** Sample development graph.

*Nodes in a development graph correspond to CASL specifications.  
Arrows show how specifications are related by the structuring constructs.*

The solid arrow denotes an ordinary import of specifications (generated by the extension), while the dashed<sup>1</sup> arrow denotes a proof obligation (corresponding to the view). This proof obligation needs to be discharged in order to show that the view is well-formed.

As a more complex example, consider the following loose specification of a sorting function, taken from Chap. 6:

```
spec LIST_ORDER_SORTED
  [TOTAL_ORDER with sort Elem, pred _ < _] =
  LIST_SELECTORS [sort Elem]
then local pred __is_sorted : List
  ∀e, e' : Elem; L : List
  • empty is_sorted
  • cons(e, empty) is_sorted
  • cons(e, cons(e', L)) is_sorted ⇔
    (cons(e', L) is_sorted ∧ ¬(e' < e))
  within op order : List → List
    ∀L : List • order(L) is_sorted
end
```

The following specification of sorting by insertion also is taken from Chap. 6:

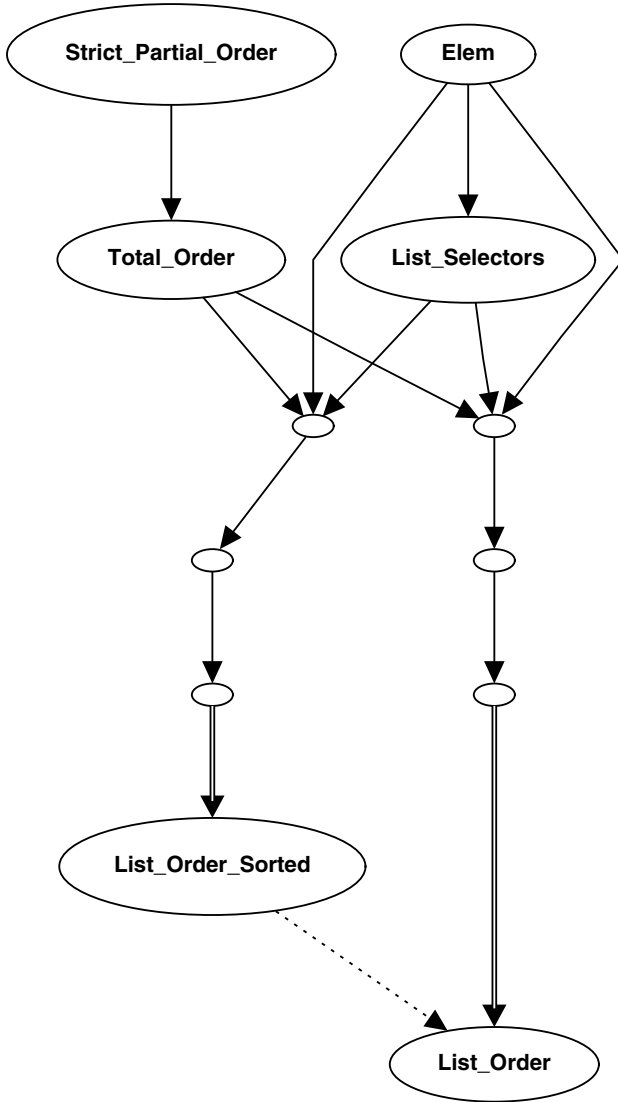
```
spec LIST_ORDER [TOTAL_ORDER with sort Elem, pred _ < _] =
  LIST_SELECTORS [sort Elem]
then local op insert : Elem × List → List
  ∀e, e' : Elem; L : List
  • insert(e, empty) = cons(e, empty)
  • insert(e, cons(e', L)) = cons(e', insert(e, L)) when e' < e
    else cons(e, cons(e', L))
  within op order : List → List
    ∀e : Elem; L : List
    • order(empty) = empty
    • order(cons(e, L)) = insert(e, order(L))
end
```

Both specifications are related. To see this, we first inspect their signatures. This is possible with:

```
hets -g Sorting.casl
```

<sup>1</sup> Actually, the dashed arrow will be displayed as solid and in red by HETS; we do not have colors available here.

assuming that `Sorting.casl` contains the above specifications. HETS now displays a more complex development graph, see Fig. 11.3.



**Fig. 11.3.** Development graph for the two sorting specifications.

*Internal nodes in a development graph correspond to unnamed parts of a structured specification.*

In the above-mentioned development graph, we have two types of nodes. The named ones correspond to named specifications, but there are also unnamed nodes corresponding to anonymous basic specifications like the declaration of the *insert* operation in `LIST_ORDER` above. Basically, there is an internal node for each structured specification that is not named.

Again, the simple solid arrows denote an ordinary import of specifications (corresponding to the extensions and unions in the specifications), while the double arrows denote hiding (corresponding to the local specification).

By clicking on the nodes, one can inspect their signatures. In this way, we can see that both `LIST_ORDER_SORTED` and `LIST_ORDER` have the same signature. Hence, it is legal to add a view:

```
view v2[TOTAL_ORDER] : LIST_ORDER_SORTED[TOTAL_ORDER] to LIST_ORDER[TOTAL_ORDER]
end
```

We have already added this view to `Sorting.casl`. The corresponding proof obligation between `LIST_ORDER_SORTED` and `LIST_ORDER` is displayed in Fig. 11.3 as a dotted arrow.

*Proof obligations can be discharged in various ways.*

Trivial proof obligations can be discharged by HETS alone using the “Proofs” menu. The proof obligation in Fig. 11.3, indicated by the lower dotted arrow between `LIST_ORDER_SORTED` and `LIST_ORDER`, states that insertion sort, as defined by the operation *order* in `LIST_ORDER`, actually has the properties of a sorting algorithm. Here, one has to choose a theorem prover that is to be used to discharge the proof obligation, which is then done by using commands specific to the theorem prover (cf. e.g. Section 11.2). Alternatively, one can state that one just conjectures the obligation to be true.

HETS is written in Haskell. Its parser uses combinator parsing. The user-defined (also known as “mixfix”) syntax of CASL calls for a two-pass approach. In the first pass, the skeleton of a CASL abstract syntax tree is derived, in order to extract user-defined syntax rules. In a second pass, which is performed during static analysis, these syntax rules are used to parse any expressions that use mixfix notation. The output is stored in the so-called ATerm format [12], which is used as interchange format for interfacing with other tools.

HETS provides an abstract interface for institutions, so that new logics can be integrated smoothly. In order to do so, a parser, a static checker and a prover for basic specifications in the logic have to be provided.

HETS has been built based on experiences with its precursors, CATS and MAYA. The CASL Tool Set (CATS) [28, 30] comes with roughly the same analysis tools as HETS. The management of development graphs is not integrated in CATS, but is provided with a stand-alone version of the tool MAYA [5, 4]. CATS and MAYA can be obtained from the CoFI tools home page [21].

## 11.2 HOL-CASL

*HOL-CASL is an interactive theorem prover for CASL, based on the tactical theorem prover ISABELLE.*

The HOL-CASL system [28] provides an interface between CASL and the theorem proving system ISABELLE/HOL [31]. We have chosen ISABELLE because it has a very small core guaranteeing correctness, and its provers, like the simplifier or the tableaux prover, are built on top of this core. Furthermore, there is over ten years of experience with it, and several mathematical textbooks have been partially verified with ISABELLE.

*CASL is linked to ISABELLE/HOL by an encoding.*

Since subsorting and partiality are present in CASL but not in ISABELLE/HOL, we have to encode these features, as explained in [29]. This means that theorem proving is not done in the CASL logic directly, but in the logic HOL (for higher-order logic) of ISABELLE. HOL-CASL tries to make the user's life easy by:

- choosing a shallow embedding of CASL into HOL, which means that e.g. CASL's logical implication  $\Rightarrow$  is mapped directly to ISABELLE/HOL's logical implication  $-->$  (and the same holds for other logical connectives and quantifiers); and
- adapting ISABELLE/HOL's syntax to conform with the CASL syntax, e.g. ISABELLE/HOL's  $-->$  is displayed as  $\Rightarrow$ , as in CASL.

However, it is essential to be aware of the fact that the ISABELLE/HOL logic is different from the CASL logic. Therefore, the formulas appearing in subgoals of proofs with HOL-CASL will not fully conform to the CASL syntax: they may use features of ISABELLE/HOL such as higher-order functions that are not present in CASL. HOL-CASL can be obtained from the CoFI tools home page [21].

To start the HOL-CASL system, follow the installation instructions, and then type:

HOL-CASL

You can load the above specification file `Order.casl` by typing:

```
use_casl "Order";
```

Let us try to prove part of the view `v1` above. To prepare for conducting a proof in the target specification of the view, `NATURAL_ORDER_2`, type in:

```
CASL_context Natural_Order_2.casl;
```

```
AddsimpAll();
```

The first command just selects the specification as the current proof context; the second one adds all the axioms of the specification to ISABELLE's simplifier (a rewriting engine). Note that the latter is advisable only if the axioms are terminating, when considered as a set of rewrite rules.

To prove the first property expressed by the view, we first have to type in the goal. Then we chose to perform induction over the variable  $x$ , and the rest can be done with automatic simplification. Finally, we name the theorem for later reference:

```
Goal "forall x:Nat . not x<x";
by (induct_tac "x" 1);
by Auto_tac;
qed "Nat_irreflexive";
```

Both the stand-alone MAYA as well as the MAYA part of HETS also provide an interface to HOL-CASL, so that it can be used to discharge proof obligations arising in development graphs.

## 11.3 ASF+SDF Parser and Syntax-Directed Editor

*ASF+SDF was used to prototype the CASL syntax.*

The algebraic specification formalism ASF+SDF [22] and the ASF+SDF Meta-Environment [13] have been deployed to prototype CASL's concrete syntax, and to develop a mapping for the concrete syntax to an abstract syntax representation using so-called ATerms [12]. Currently, only the first pass of parsing (i.e. without mixfix analysis) is realized in SDF. Parsing is performed based on the underlying Scannerless Generalized LR parsing technology. A prototype of the mapping from the concrete to abstract representation is written in ASF rewrite rules.



*The ASF+SDF Meta-Environment provides syntax-directed editing of CASL specifications.*

Given the concrete syntax definition of CASL in SDF, syntax-directed editors within the ASF+SDF Meta-Environment come for free. Recent developments in the Meta-Environment even allow for the development of a CASL specification environment.

The ASF+SDF Meta-Environment provides a built-in library mechanism which contains a collection of grammars, among others CASL. Via the library the user of the Meta-Environment has access to the CASL syntax in SDF and a collection of ASF equations to map CASL specifications into an interchange format named CasFix. The `asfc` tool compiles the CASL grammar into a stand-alone C program.

A link to the ASF+SDF Meta-Environment with further information and a download possibility can be found at the CoFI tools home page [21]. The built-in library of the Meta-Environment (Version 1.5 and higher) will provide the full CASL grammar in SDF and the mapping to CasFix as an ASF specification.

## 11.4 Other Tools

The following tools are at a prototype stage at the time of appearance of this volume. Please refer to at the CoFI tools home page [21], where the latest versions can be downloaded.

### *Translation to OCAML*

A translation from CASL into OCAML has been developed at Paris and Poitiers. The translation works for a “functional programming” sublanguage of CASL that includes free datatypes and recursive definitions of operations over these types.

### *Translation to Haskell*

A translation from CASL into Haskell has been developed in Bremen. Actually, the translation works on an executable subset of HasCASL (a higher-order extension of CASL). Using an embedding of CASL into HasCASL, this translation can also be used for CASL.

### *ELAN-CASL*

ELAN-CASL is a rewriting engine for the  $\text{HORN}^=$  sublanguage of CASL. It is based on a translation of that sublanguage to the input language of the rewriting engine ELAN.

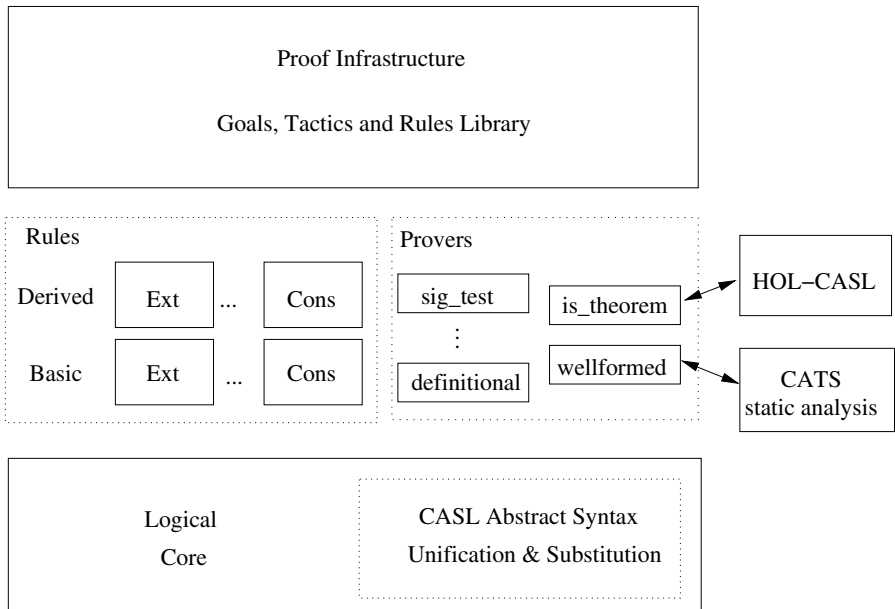
Given a CASL specification and a query term, ELAN-CASL computes the normal forms of the term. Note that because the set of rules is not required to be terminating nor confluent, a query term may have several normal forms, or may not terminate.

ELAN can be called either as an interpreter or as a compiler.

### *CASL consistency checker*

The CASL consistency checker (CCC) has been developed in Bremen and Swansea. With CCC, one can interactively check whether a CASL specification is consistent. It is a faithful implementation of a consistency calculus for full CASL [32].

Besides using certain syntactical criteria, the consistency calculus relies heavily on the CASL structuring mechanisms and their semantic annotations. Consequently, consistency proofs follow the structure of the given specification. In this way, the calculus highlights the (usually few) ‘hot spots’ of a specification (e.g. views requiring theorem proving), while the (lengthy) ‘trivial’ parts of the consistency argument are discharged automatically. As the consistency calculus works along the structure of the specification, the need to construct (and prove correct) actual models of specifications is avoided as far as possible.



**Fig. 11.4.** CCC System Architecture

The CASL Consistency Checker consists of four parts (see Fig. 11.4): first, a *logical core* implementing a representation of the propositions to be proved, and the basic proof rules, along with a generic unification package for the CASL abstract syntax; secondly, representations of the *calculi*: the base rules which are stated axiomatically, and derived rules; thirdly, *proof procedures* (automatic or semi-automatic decision procedures), which serve to discharge specific proof obligations; and finally, *proof infrastructure* such as a package which helps users to conduct goal-directed proofs, tactics that support writing advanced proof procedures, and a simple database which allows storage and retrieval of proved theorems. The system is encapsulated by a single interface which does not allow the end-user (i.e. the working specifier) to add any more axiomatic rules or provers; thus, the typing is used to both increase confidence in the correctness of the implementation in the style of LCF, and to protect the integrity of the system from user intervention.

## Basic Libraries

Till Mossakowski

*The CASL Basic Libraries contain the standard datatypes.*

The CASL Basic Libraries consist of specifications of often-needed datatypes and views between them, freeing the specifier from re-inventing well-known things. This can be compared to standard libraries in programming languages. While this book often discusses several styles of specification with CASL, the basic datatypes consistently follow a specific style described in [20].

*Here we show a cut-down version without axioms.*

Here, we describe two of the libraries (see the overview in Fig. 12.1): the libraries of numbers and of structured datatypes. We also provide stripped-down versions of the libraries themselves, with some of the specifications and all axioms and annotations removed. These stripped-down versions can serve for getting a first overview of the signatures of the specified datatypes.

The full CASL Basic Libraries with complete specifications is presented in the CASL Reference Manual [20], and is also included in the CD-ROM coming with this volume. The latest version is available at:

<http://www.cofi.info/Libraries>

*HETS can be used to get an overview of the Basic Libraries.*

The HETS tool described in Chap. 11 allows the structure of the specifications in the libraries to be displayed as a graph and their signatures to be inspected. This is recommended as a way of obtaining a better overview, and also for answering specific questions that arise when using the basic datatypes.

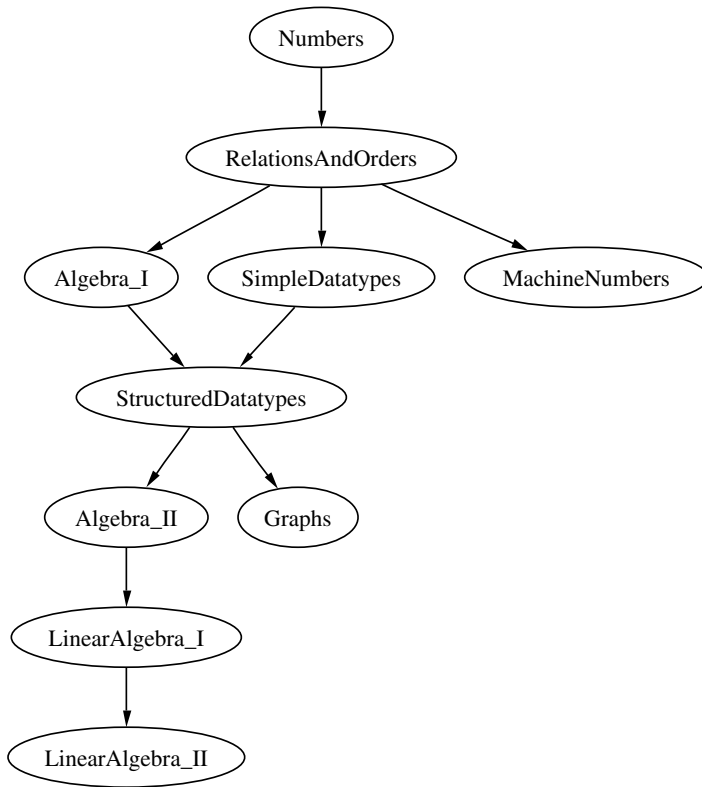


Fig. 12.1. Dependency graph of the libraries of basic datatypes.

## 12.1 Library BASIC/NUMBERS

This library provides specifications of natural numbers, integers, and rational numbers.

*The natural numbers are specified as a free datatype.*

In the specification NAT, the natural numbers are specified as a free datatype, together with a collection of predicates and operations over the sort *Nat* of natural numbers.

Note that the names for the *partial operations* subtraction  $\_ - ? \_$  and division  $\_ / ? \_$  include a question mark. This is to avoid overloading with the *total operations*  $\_ - \_$  on integers and  $\_ / \_$  on rationals, which would lead to inconsistencies as both these specifications import the specification NAT. The *total operation* for subtraction differs from subtraction on the integers as

well. It is written  $--!-$ , and it is 0 whenever the partial subtraction  $-?$  is undefined, while it otherwise coincides with the latter.

The digits are introduced as constants, together with an operation  $--@@--$  for concatenation of digits. Together with an annotation (see Chap. 9):

**%number**  $--@@--$

this allows one to write the usual literals (like e.g. 8364) for natural numbers.

The introduction of the subsort *Pos*, consisting of the positive naturals, gives rise to certain new operations, e.g.:

$-- \times -- : Pos \times Pos \rightarrow Pos$ ,

whose semantics is completely determined by overloading.

**library** BASIC/NUMBERS

**spec** NAT =

**free type** *Nat* ::= 0 | *suc*(*pre*:?*Nat*)

**preds**  $--\leq--$ ,  $--\geq--$ ,  $--<--$ ,  $-->-- : Nat \times Nat$ ;  
*even*, *odd* : *Nat*

**ops**  $--! : Nat \rightarrow Nat$ ;

$--+--$ ,  $--\times--$ ,  $--^--$ , *min*, *max*,  $--!-- : Nat \times Nat \rightarrow Nat$ ;  
 $--?--$ ,  $--/?--$ ,  $--div--$ ,  $--mod--$ , *gcd* : *Nat*  $\times$  *Nat*  $\rightarrow ?$  *Nat*

%% Operations to represent natural numbers with digits:

**ops** 1: *Nat* = *suc*(0); % (1\_def\_Nat)%  
2: *Nat* = *suc*(1); % (2\_def\_Nat)%  
3: *Nat* = *suc*(2); % (3\_def\_Nat)%  
4: *Nat* = *suc*(3); % (4\_def\_Nat)%  
5: *Nat* = *suc*(4); % (5\_def\_Nat)%  
6: *Nat* = *suc*(5); % (6\_def\_Nat)%  
7: *Nat* = *suc*(6); % (7\_def\_Nat)%  
8: *Nat* = *suc*(7); % (8\_def\_Nat)%  
9: *Nat* = *suc*(8); % (9\_def\_Nat)%  
 $--@@--(m: Nat; n: Nat): Nat = (m \times suc(9)) + n$   
% (decimal\_def)%

**sort** *Pos* = {*p*: *Nat* • *p* > 0}

**ops** 1: *Pos* = *suc*(0); % (1\_as\_Pos\_def)%  
 $--\times-- : Pos \times Pos \rightarrow Pos$ ;  
 $--+-- : Pos \times Nat \rightarrow Pos$ ;  
 $--+-- : Nat \times Pos \rightarrow Pos$ ;  
*suc* : *Nat*  $\rightarrow$  *Pos*

**end**

*The integers are specified as difference pairs of naturals.*

The specification `INT` of integers is built on top of the specification `NAT`: integers are defined as equivalence classes of pairs of naturals written as differences — the axioms (which are omitted in the specification below) specify that two pairs are equivalent if their differences are equal:

$$\begin{aligned} &\forall a, b, c, d: \text{Nat} \\ &\bullet a - b = c - d \Leftrightarrow a + d = c + b \end{aligned} \quad \text{\% (equality\_Int) \%}$$

The sort *Nat* is then declared to be a subsort of *Int*. Besides the division operator `--/?--`, the specification `INT` also provides the function pairs *div/mod* and *quot/rem*, respectively, as constructs for division — behaving differently on negative numbers, see [20] for a discussion. The operation *sign* gives the sign of an integer (which is either -1, 0, or 1).

```
spec INT =
  NAT
then generated type Int ::= --(Nat; Nat)
  sort   Nat < Int
         %% a system of representatives for sort Int is
         %% a - 0 and 0 - p, where a: Nat and p: Pos
  preds  --≤--, --≥--, --<--, -->-- : Int × Int;
         even, odd : Int
  ops    --, sign : Int → Int;
         abs : Int → Nat;
         --+--, --×--, --÷--, min, max : Int × Int → Int;
         --^-- : Int × Nat → Int;
         --/?--, --div--, --quot--, --rem-- : Int × Int →? Int;
         --mod-- : Int × Int →? Nat
end
```

*The rationals are specified as fractions of integers.*

The specification `RAT` of rational numbers follows the same scheme as the specification of integers discussed above. This time, the specification `INT` is imported. The rationals are then defined as equivalence classes of pairs consisting of an integer and a positive number written as fractions, using the axiom:

$$\begin{aligned} &\forall i, j: \text{Int}; p, q: \text{Pos} \\ &\bullet i / p = j / q \Leftrightarrow i \times q = j \times p \end{aligned} \quad \text{\% (equality\_Rat) \%}$$

The sort *Int* is then declared to be a subsort of *Rat*. Note that thanks to the behavior of subsorted overloading in CASL, the declaration of the operation:

$--/_ : Rat \times Rat \rightarrow? Rat;$

allows rationals to be written also as  $x/y$ , for arbitrary integers  $x$  and  $y \neq 0$ .

```
spec RAT =
  INT
then generated type Rat ::= --/_(Int; Pos)
  sort   Int < Rat
  preds  --<--, --<--, -->--, -->_ : Rat × Rat
  ops    --, abs : Rat → Rat;
         --+--, --−, --×--, min, max : Rat × Rat → Rat;
         --/_ : Rat × Rat →? Rat;
         --^_ : Rat × Int → Rat
end
```

## 12.2 Library BASIC/STRUCTURED DATATYPES

This library provides specifications of the familiar structured datatypes as used e.g. for the design of algorithms or within programming languages. Its main focus is data structures like (finite) sets, lists, strings, (finite) maps, (finite) bags, arrays, and various kinds of trees. Common to all these concepts is that they are generic. Consequently, all of the specifications of this library are generic.

*Finite sets, maps and bags are specified as generated datatypes, with equality determined by means of observers.*

Finite sets, finite maps and finite bags are specified using a generated sort. An observer operation or predicate is then introduced in order to define equality on this sort. Concerning finite sets, equality on the sort  $Set[Elem]$  is characterized using the predicate  $--eps--$  (displayed as  $\epsilon$ ) in the specification GENERATESET. This leads to the extensionality axiom:

$$\bullet M = N \Leftrightarrow \forall x : Elem \bullet x \in M \Leftrightarrow x \in N \quad \% (equality\_Set) \%$$

**library** BASIC/STRUCTURED DATATYPES



```

spec GENERATESSET [sort Elem] =
  generated type Set[Elem] ::= {} | _+_(Set[Elem]; Elem)
  pred _ε_ : Elem × Set[Elem]
    %% a system of representatives for sort Set[Elem] is
    %%
    %% {} and {} + x_1 + x_2 + ... + x_n
    %%
    %% where x_1 < x_2 < ... < x_n, n >= 1, x_i of type Elem,
    %% for an arbitrary strict total order _ε_ on Elem.
end

```

```

spec SET [sort Elem] given NAT =
  GENERATESSET [sort Elem]
then %def
  preds isNonEmpty : Set[Elem];
    _⊆_ : Set[Elem] × Set[Elem]
  ops   {_} : Elem → Set[Elem];
    #_ : Set[Elem] → Nat;
    _+_ : Elem × Set[Elem] → Set[Elem];
    _-_- : Set[Elem] × Elem → Set[Elem];
    _∩_, _∪_, _-_- , _symDiff_ :
      Set[Elem] × Set[Elem] → Set[Elem]
end

```

Finite maps, i.e. elements of the sort *Map*[*S*, *T*], are considered to be identical if looking up any value in *S* yields the same result in both cases:

- $M = N \Leftrightarrow \forall s : S \bullet \text{lookup}(s, M) = \text{lookup}(s, N)$       %(*equality\_Map*)%

On top of this, the specification MAP adds e.g. predicates for elementhood ( $\epsilon$ ) as well as for determining the profile of a map ( $f :: x \rightarrow y$  means that  $f$  is a map from  $x$  to  $y$ ).

The specification TOTALMAP restricts maps to everywhere-defined maps (which are isomorphic to tuples). Since maps are finite, totality is only possible for maps over finite argument sorts. The latter is specified in the specification FINITE, using a partial surjection from the natural numbers. Since this specification is rather unusual, we make an exception and also show its axioms.

```

spec GENERATEMAP [sort S] [sort T] =
  generated type Map[S, T] ::= empty | _[_/_]_(Map[S, T]; T; S)
  op   lookup : S × Map[S, T] →? T
end

```

```

spec MAP [sort S] [sort T] given NAT =
  GENERATEMAP [sort S] [sort T]
and SET [sort S]
and SET [sort T]
then %def
  free type Entry[S,T] ::= [_/_](target:T; source:S)
  preds isEmpty : Map[S,T];
    --ε-- : Entry[S,T] × Map[S,T];
    --::--→-- : Map[S,T] × Set[S] × Set[T]
  ops
    --+-- , -- - : Map[S,T] × Entry[S,T] → Map[S,T];
    -- - : Map[S,T] × S → Map[S,T];
    -- - : Map[S,T] × T → Map[S,T];
    dom : Map[S,T] → Set[S];
    range : Map[S,T] → Set[T];
    --∪-- : Map[S,T] × Map[S,T] →? Map[S,T]
end
spec FINITE [sort Elem] =
  {
    NAT
    then op f : Nat →? Elem
      • ∀ x : Elem • ∃ n : Nat • f(n) = x          %(f_surjective)%
      • ∃ n : Nat • ∀ m : Nat • def f(m) ⇒ m < n    %(f_bounded)%
    }
  reveal Elem
end

spec TOTALMAP [FINITE [sort S]] [sort T] =
  {
    MAP [sort S] [sort T]
    then sort TotalMap[S,T] =
      {M : Map[S,T] • ∀ x : S • def lookup(x, M)}
    ops
      --[_/_] : TotalMap[S,T] × T × S → TotalMap[S,T];
      lookup : S × TotalMap[S,T] → T;
      --+-- : TotalMap[S,T] × Entry[S,T] → TotalMap[S,T];
      range : TotalMap[S,T] → Set[T];
      --∪-- : TotalMap[S,T] × TotalMap[S,T]
        →? TotalMap[S,T]
    pred --ε-- : Entry[S,T] × TotalMap[S,T]
  }
  hide Map[S,T]
end

```

In the specification GENERATEBAG, those elements of sort *Bag*[Elem] are identified that show the same number of occurrences (observed by the operation *freq*) for all entries:

$$\bullet M = N \Leftrightarrow \forall x : Elem \bullet freq(M, x) = freq(N, x) \quad \% (equality\_Bag) \%$$

```

spec GENERATEBAG [sort Elem] given NAT =
  generated type Bag[Elem] ::= {} | _+_ (Bag[Elem]; Elem)
  op   freq : Bag[Elem] × Elem → Nat
end

spec BAG [sort Elem] given NAT =
  GENERATEBAG [sort Elem]
then preds isEmpty : Bag[Elem];
             --ε-- : Elem × Bag[Elem];
             --⊆-- : Bag[Elem] × Bag[Elem]
  ops   --+-- : Elem × Bag[Elem] → Bag[Elem];
         --_-- : Bag[Elem] × Elem → Bag[Elem];
         --_--, _∪_, _∩_ : Bag[Elem] × Bag[Elem] → Bag[Elem]
end

```

*Lists are specified as a free datatype.*

In the specification GENERATELIST, lists are built up from the empty list by adding elements in front. The usual list operations are provided: *first* and *last* select the first or last element of a list, while *rest* or *front* select the remaining list; #\_ counts the number of elements in a list, while *freq* counts the number of occurrences of a given element; *take* takes the first *n* elements of a list, while *drop* drops them.

```

spec GENERATELIST [sort Elem] =
  free type List[Elem] ::= [] | _::_(first:?Elem; rest:?List[Elem])
end

spec LIST [sort Elem] given NAT =
  GENERATELIST [sort Elem]
then preds isEmpty : List[Elem];
             --ε-- : Elem × List[Elem]
  ops   --+-- : List[Elem] × Elem → List[Elem];
         first, last : List[Elem] →? Elem;
         front, rest : List[Elem] →? List[Elem];
         #_ : List[Elem] → Nat;
         --++_ : List[Elem] × List[Elem] → List[Elem];
         reverse : List[Elem] → List[Elem];
         _!_ : List[Elem] × Nat →? Elem;
         take, drop : Nat × List[Elem] →? List[Elem];
         freq : List[Elem] × Elem → Nat
end

```

*Arrays are specified as certain finite maps.*

The specification `ARRAY` includes the condition  $\min \leq \max$  as an axiom in its first parameter. This ensures a non-empty index set. Arrays are defined as finite maps from the sort *Index* to the sort *Elem*, where the typical array operations `lookup` (`__!__`) and `assignment` (`__!__ := __`) are introduced in terms of finite map operations. Finally, revealing the essential signature elements yields the desired datatype.

```
spec ARRAY [ops min, max : Int • min ≤ max %(Cond_nonEmptyIndex)%]
  [sort Elem]
  given INT =
  sort Index = {i: Int • min ≤ i ∧ i ≤ max}
then {
  MAP [sort Index] [sort Elem]
  with sort Map[Index,Elem] ↦ Array[Elem],
  op empty ↦ init
  then ops
    __!__ := __ : Array[Elem] × Index × Elem → Array[Elem];
    __!__ : Array[Elem] × Index →? Elem
  }
  reveal sort Array[Elem], ops init, __!__, __!__ := __
end
```

*Several kinds of tree are available, differing in the branching and in the positions of elements.*

The library concludes with several specifications concerning trees. There are specifications of binary trees (`BINTREE`, `BINTREE2`), *k*-branching trees (`KTREE`), and trees with possibly different branching at each node (`NTREE`). Each of these branching structures can be equipped with data in different ways: Either all nodes of a tree carry data (as is the case in `BINTREE`, `KTREE`, and `NTREE`), or just the leaves of a tree have a data entry (as in `BINTREE2`).

*Binary trees admit two children for each internal node.*

```

spec GENERATEBINTREE [sort Elem] =
  free type
    BinTree[Elem] ::= nil
                      | binTree(entry:?Elem; left:?BinTree[Elem];
                                right:?BinTree[Elem])
  end

spec BINTREE [sort Elem] given NAT =
  GENERATEBINTREE [sort Elem] and SET [sort Elem]
  then preds isEmpty, isLeaf : BinTree[Elem];
           isCompoundTree : BinTree[Elem];
           _ε_ : Elem × BinTree[Elem]
  ops    height : BinTree[Elem] → Nat;
           leaves : BinTree[Elem] → Set[Elem]
  end

spec GENERATEBINTREE2 [sort Elem] =
  free type NonEmptyBinTree2[Elem] ::=
    leaf(entry:?Elem)
    | binTree(left:?NonEmptyBinTree2[Elem];
              right:?NonEmptyBinTree2[Elem])
  free type BinTree2[Elem] ::= nil | sort NonEmptyBinTree2[Elem]
  end

spec BINTREE2 [sort Elem] given NAT =
  GENERATEBINTREE2 [sort Elem] and SET [sort Elem]
  then %def
    preds isEmpty, isLeaf : BinTree2[Elem];
           isCompoundTree : BinTree2[Elem];
           _ε_ : Elem × BinTree2[Elem]
    ops    height : BinTree2[Elem] → Nat;
           leaves : BinTree2[Elem] → Set[Elem]
  end

```

*k*-trees admit *k* children for each internal node (with *k* fixed).

We now come to *k*-branching trees. The branching is specified by using arrays of trees of size *k*, which are used to contain the children of a node in the tree.

```

spec GENERATEKTREE [op k : Int • k ≥ 1 %(Cond_nonEmptyBranching)%]
  [sort Elem] given INT =
  ARRAY [ops 1 : Int; k : Int
    fit ops min : Int ↦ 1, max : Int ↦ k]
    [sort KTree[k,Elem]]
then free type
  KTree[k,Elem] ::= nil
    | kTree(entry: ?Elem;
      branches: ?Array[KTree[k,Elem]])
end

spec KTREE [op k : Int • k ≥ 1          %(Cond_nonEmptyBranching)%]
  [sort Elem]
  given INT =
  GENERATEKTREE [op k : Int] [sort Elem]
and SET [sort Elem]
then %def
  preds isEmpty, isLeaf : KTree[k,Elem];
    isCompoundTree : KTree[k,Elem];
    _ε_ : Elem × KTree[k,Elem]
  ops height : KTree[k,Elem] → Nat;
    maxHeight : Index × Array[KTree[k,Elem]] → Nat;
    leaves : KTree[k,Elem] → Set[Elem];
    allLeaves : Index × Array[KTree[k,Elem]] → Set[Elem]
end

```

*n-trees admit arbitrary branching.*

Finally,  $n$ -trees are trees with possibly different branching at each node. This is specified by equipping each node in a tree with a list of child trees.

```

spec GENERATENTREE [sort Elem] =
  LIST [sort NTree[Elem]]
then free type
    NTree[Elem] ::= nil
                  | nTree(entry:?Elem;
                           branches:?List[NTree[Elem]])
end

spec NTREE [sort Elem] given NAT =
  GENERATENTREE [sort Elem] and SET [sort Elem]
then preds isEmpty, isLeaf : NTree[Elem];
             isCompoundTree : NTree[Elem];
             _ε_ : Elem × NTree[Elem]
ops height : NTree[Elem] → Nat;
     maxHeight : List[NTree[Elem]] → Nat;
     leaves : NTree[Elem] → Set[Elem];
     allLeaves : List[NTree[Elem]] → Set[Elem]
end

```

## Case Study: The Steam-Boiler Control System

In this chapter we illustrate the use of CASL on a fairly large and complex case study, the *steam-boiler control system*. This case study is particularly interesting since it has been used several times as a competition problem, and many other specification frameworks have been illustrated with it, see [1]. Here we describe how to derive a CASL specification of the steam-boiler control system, starting from the informal requirements provided to the participants of the Dagstuhl meeting *Methods for Semantics and Specification*, organized jointly by Jean-Raymond Abrial, Egon Börger and Hans Langmaack in June 1995. The aim of this formalization process is to analyze the informal requirements, to detect inconsistencies and loose ends, and to translate the requirements into a CASL specification. During this process we have to provide interpretations for the unclear or missing parts. We explain how we can keep track of these additional interpretations by localizing very precisely in the formal specification where they lead to specific axioms, thereby taking care of the traceability issues. We also explain how the CASL specification is obtained in a stepwise way by successive analysis of various parts of the problem description. Finally we discuss the validation of the CASL requirements specification resulting from the formalization process, and in a last step we refine the requirements specification in a sequence of architectural specifications that describe the intended architecture of the steam-boiler control system.<sup>1</sup>

*The reader not already familiar with the steam-boiler control system case study may want to start by reading App. C, where the original description of the problem is reproduced.*

<sup>1</sup> This chapter partially relies on an earlier work published in [8] where the PLUSS specification language [7, 9] was used together with the LARCH prover [25]. However, the specification methodology illustrated in this chapter is significantly improved, and moreover CASL provides several features that lead to a much more concise and perspicuous specification, as illustrated later in this chapter.



### 13.1 Introduction

The aim of this chapter is to illustrate how one can solve the steam-boiler control specification problem with CASL. For this we have to provide a CASL specification of the software system that controls the level of water in the steam-boiler. Our work plan can be described as follows:

1. The main task is to derive, starting from the informal requirements, a requirements specification, written in CASL, of the steam-boiler control system. In particular, this task involves the following activities:
  - a) We must perform an in-depth analysis of the informal requirements. Obviously, this is necessary to gain a sufficient understanding of the problem to be specified, and this preliminary task may not seem worth mentioning. Let us stress, however, that the kind of preliminary analysis required for writing a formal specification proves especially useful to detect discrepancies in the informal requirements that would otherwise be very difficult to detect. Indeed, from our practical experience, this step is usually very fruitful from an engineering point of view, and one could argue that the benefits to be expected here are enough in themselves to justify the use of formal methods, even if for lack of time (or other resources) no full formal development of the system is performed.
  - b) Once we have a sufficient understanding of the problem to be specified, we must translate the informal requirements into a formal specification. This step will require us to provide interpretations for the unclear or missing parts of the informal requirements. Moreover, this formalization process will also be helpful to further detect inconsistencies and loose ends in the informal requirements. Here, a very important issue is to keep track of the interpretations made during the formalization process, in order to be able, later on, to take into account further modifications and changes of the informal requirements.
  - c) When we have written the formal requirements specification, we must carefully check its adequacy with respect to the informal requirements: this part is called the validation of the formal specification.

In principle there should be some interaction between the specification team and the team who has designed the informal requirements, in particular to check whether the suggested interpretations of the detected loose ends are adequate. In the framework of this case study, however, such interactions were not possible, and we can only use our intuition to assess the soundness of the interpretations made during the writing of the formal specification.

2. Once a validated requirements specification is obtained, we can proceed toward a program by a sequence of refinements. Here a crucial step is the choice of an architecture of the desired implementation, expressed by an architectural specification as explained in Chap. 8. Each refinement step

leads to proof obligations which allow the correctness of the performed refinement to be assessed. In a last step, a program is derived from the final design specification.

Before starting to explain how to write a CASL requirements specification of the steam-boiler control system, let us make a few comments on this case study. First, note that, although in principle a hybrid system, the steam-boiler control system turns out to be merely a reactive system, not even a ‘hard real-time’ system (see e.g. the assumptions made in App. C.3). Moreover, even if the whole system, i.e., the control program and its physical environment is distributed, this is not the case, at least at the requirements level, for the steam-boiler control system. CASL turns out to be especially well-suited to capture the features of systems like the steam-boiler control system, where data and control are equally important (in particular, here data play a prominent role in failure detection). The various constructs provided by CASL allow the specifications to be formulated straightforwardly and perspicuously – and significantly more concisely than in other algebraic specification languages.

As a last remark we must make clear that for the sake of simplicity the *initialization* phase of the steam-boiler control system (see App. C.4.1) is not specified. However, it should be clear that it would be straightforward to extend our specification so as to take the initialization phase into account, following exactly the same methodology as for the rest of the case study.

This chapter is organized as follows. In Sect. 13.2 we start by providing some elementary specifications that will be useful for the rest of the case study. In Sect. 13.3 we explain how we will proceed to derive the CASL requirements specification in a stepwise way. Then in Sect. 13.4 we detail the specification of the mode of operation of the steam-boiler control system. In Sect. 13.5 we specify the detection of the various equipment failures, and in Sect. 13.6 we explain how we can compute, at each cycle, some predicted values for the messages to be received at the next cycle. In Sect. 13.9 we explain how our CASL requirements specification can be validated, and in Sect. 13.10 we refine the CASL requirements specification in a sequence of architectural specifications that describe the intended architecture of the implementation of the steam-boiler control system.

## 13.2 Getting Started

As explained in App. C.3, in each cycle the steam-boiler control system collects the messages received, performs some analysis of the information contained in them, and then sends messages to the physical units. We will therefore start with the specification of some elementary datatypes, such as “messages received” and “messages sent”. To specify the messages sent and received, we follow App. C.5 and C.6. Note that some messages have parameters (e.g. pump number, pump state, pump controller state, mode of operation), and we must

therefore specify the corresponding datatypes as well. For the sake of clarity, we group together all similar messages (e.g. all “repaired” messages, all “failure acknowledgement” messages) by introducing a suitable parameter “physical unit”. A physical unit is either a pump, a pump controller, the water level measuring device or the steam output measuring device. Remember that we do not specify the physical units as such, since we do not specify the physical environment of the steam-boiler (we do not specify the steam-boiler either, we only specify the steam-boiler control system). Hence the datatype “physical unit” is just an elementary datatype that says that we have some pumps, some pump controllers, and the two measuring devices.

Some messages have a value  $v$  as parameter. From the informal requirements we can infer that these values are (approximations of) real numbers, but it is not necessary at this level to make any decision about the exact specification of these values. In our case study, we will therefore rely on a very abstract (loose) specification `VALUE`, introducing a sort *Value* together with some operations and predicates, which are left unspecified (we expect of course that these operations and predicates will have the intuitive interpretation suggested by their names). This means that we consider `VALUE` as being a general parameter of our specification.<sup>2</sup> This point is discussed again in Sect. 13.10. Note also that we will abstract from measuring units (such as liter, liter/sec), since ensuring that these units are consistently used is a very minor aspect of this particular case study.<sup>3</sup>

This first analysis leads to the following specifications: `VALUE`, `BASICS`, `MESSAGES_SENT`, and `MESSAGES_RECEIVED`.

**from** `BASIC/NUMBERS` **get** `NAT`

**%display** `_half` **%LATEX** `_{1/2}`

**%display** `_square` **%LATEX** `_{^2}`

**spec** `VALUE =`

**%%** At this level we don’t care about the exact specification of values.

`NAT`

**then** **sorts** `Nat < Value`

**ops** `-- + _ : Value × Value → Value, assoc, comm, unit 0;`

`-- − _ : Value × Value → Value;`

`-- × _ : Value × Value → Value, assoc, comm, unit 1;`

---

<sup>2</sup> We leave `VALUE` as an implicit parameter of our specifications, rather than using generic specifications taking `VALUE` as a parameter, since our specifications are not to be instantiated by argument specifications describing several kinds of values, but on the contrary should all refer to the same abstract datatype of values.

<sup>3</sup> It is of course possible to take measuring units into account, following for instance the method described in [18]. Appropriate CASL libraries supporting measuring units are currently being developed.

```

--/2, --2 : Value → Value;
min, max : Value × Value → Value
preds -- < --, -- ≤ -- : Value × Value
end

spec BASICS =
  free type PumpNumber ::= Pump1 | Pump2 | Pump3 | Pump4;
  free type PumpState ::= Open | Closed;
  free type PumpControllerState ::= Flow | NoFlow;
  free type PhysicalUnit ::= Pump(PumpNumber)
                        | PumpController(PumpNumber)
                        | SteamOutput | WaterLevel;
  free type Mode ::= Initialization | Normal | Degraded
                  | Rescue | EmergencyStop;
end

spec MESSAGES_SENT =
  BASICS
then free type
  S_Message ::= MODE(Mode) | PROGRAM_READY | VALVE
              | OPEN_PUMP(PumpNumber)
              | CLOSE_PUMP(PumpNumber)
              | FAILURE_DETECTION(PhysicalUnit)
              | REPAIRED_ACKNOWLEDGEMENT(PhysicalUnit);
end

spec MESSAGES_RECEIVED =
  BASICS and VALUE
then free type
  R_Message ::= STOP | STEAM_BOILER_WAITING
              | PHYSICAL_UNITS_READY
              | PUMP_STATE(PumpNumber; PumpState)
              | PUMP_CONTROLLER_STATE(PumpNumber;
                                      PumpControllerState)
              | LEVEL(Value) | STEAM(Value)
              | REPAIRED(PhysicalUnit)
              | FAILURE_ACKNOWLEDGEMENT(PhysicalUnit)
              | junk;
end

```

In addition to the “messages received” specified in App. C.6, we add an extra constant message *junk*. This message will represent any message received which does not belong to the class of recognized messages. We do not add a similar message to the messages sent, since we may assume that the steam-boiler control system will only send proper messages. Obviously, receiving a

*junk* message will lead to the detection of a failure of the message transmission system.

In the SBCS\_CONSTANTS specification we describe the various constants that characterize the steam-boiler (these constants are explained in App. C.2.6).

```
spec SBCS_CONSTANTS =
  VALUE
then ops C, M1, M2, N1, N2, W, U1, U2, P : Value;
      dt : Value %% Time duration between two cycles (5 sec.)
      %% These constants must verify some obvious properties:
      • 0 < M1 • M1 < N1 • N1 < N2 • N2 < M2 • M2 < C
      • 0 < W • 0 < U1 • 0 < U2 • 0 < P
end
```

We will also specify the datatypes “set of messages received” and “set of messages sent” since, as suggested at the end of App. C.3, all messages are supposed to be received (or emitted) simultaneously at each cycle. The two latter specifications are obtained by instantiation of a generic specification SET of “sets of elements”, which is imported from the library BASIC/STRUCTURED DATATYPES.

```
from BASIC/STRUCTURED DATATYPES get SET
```

```
spec PRELIMINARY =
  SET [MESSAGES_RECEIVED fit Elem ↦ R_Message]
and SET [MESSAGES_SENT fit Elem ↦ S_Message]
and SBCS_CONSTANTS
end
```

♡ As illustrated by the above specifications, it is particularly convenient to *structure* our formal specification into coherent, easy to grasp, *named specifications* that will be easily reused later on by referring to their names (and this of course will prove even more important in the sequel). Moreover, *free datatypes* are especially useful here to obtain concise specifications. On the other hand, *loose specifications* are useful to avoid overspecification of values in VALUE and of the steam-boiler constants in SBCS\_CONSTANTS. Declaring that *Nat* is a *subsort* of *Value* ensures that natural numbers can be used as arguments of operations on values. Reusing standard specifications of usual datatypes from the BASIC *libraries* avoids the need to specify them again, and of course it is essential that these specifications are *generic* in order to easily adapt them as desired when they are reused. Finally, *display annotations* are useful to conveniently display some symbols as usual mathematical symbols.<sup>4</sup> ♡

---

<sup>4</sup> In this chapter, metacomments about the adequacy of CASL features will be highlighted like this.

## 13.3 Carrying On

As emphasized in App. C.3, the steam-boiler control system is a typical example of a control-command system. The specification of such systems always follows the same pattern:

- A preliminary set of specifications group all the basic information about the system to be controlled, such as the specification of the various messages to be exchanged between the system and its environment, and the specification of the various constants related to the system of interest. This is indeed the aim of the specification `PRELIMINARY` introduced in the previous section.
- Then, the various states of the control system should be described. At this stage, however, it would be much too early to determine which state variables are needed. Thus states will be represented by values of a (loosely specified) sort *State*, equipped with some observers (corresponding to access to state variables). During the requirements analysis and formalization phase we may need further observers, to be introduced on a by-need basis.
- Then a (group of) specification(s) will take care of the analysis of the messages received – here, of failure detection in particular. On the basis of this analysis, some actions should be taken, corresponding to the messages to be sent to the environment. State variables are also updated according to the result of the analysis of the messages received and to the messages to be sent.
- Finally a specification describes the overall control-command system as a labeled transition system.

A very rough preliminary sketch of the steam-boiler control system specification looks therefore as follows:

```

library SBCS
from BASIC/NUMBERS get NAT
from BASIC/STRUCTURED DATATYPES get SET
%display _half %LATEX _/2
%display _square %LATEX _^2
:
spec  PRELIMINARY =  %{ See previous section. }%

spec  SBCS_STATE =
      PRELIMINARY
then  sort State
      ops  %% Needed state observers are introduced here.
           %% E.g., we need an observer for the mode of operation:
           mode : State → Mode;
           ...
end
```

```

spec SBCS_ANALYSIS =
  SBCS_STATE
then %% Analysis of messages received and in particular failure detection.
  %% Computation of the messages to be sent.
  op messages_to_send : State  $\times$  Set[R_Message]  $\rightarrow$  Set[S_Message];
  %% Computation of the updates of the state variables.
  %%{ For each observer obs defined in SBCS_STATE,
    we introduce an operation next_obs that computes the
    corresponding update according to the analysis of the messages
    received in this round. For instance, we specify here an operation
    next_mode corresponding to the update of the observer mode. }%
  ops next_mode : State  $\times$  Set[R_Message]  $\rightarrow$  Mode;
  ...
end

spec STEAM_BOILER_CONTROL_SYSTEM =
  SBCS_ANALYSIS
then op init : State
  pred is_step : State  $\times$  Set[R_Message]  $\times$  Set[S_Message]  $\times$  State
  %% Specification of the initial state init by means of the observers, e.g:
  • mode(init) = ...
  • ...
  %%{ Specification of is_step by means of the observers
    and of the updating operations, e.g.: }%
   $\forall s, s' : \textit{State}; \textit{msgs} : \textit{Set}[\textit{R\_Message}]; \textit{Smsg} : \textit{Set}[\textit{S\_Message}]$ 
  • is_step(s, msgs, Smsg, s')  $\Leftrightarrow$ 
     $\textit{mode}(s') = \textit{next\_mode}(s, \textit{msgs}) \wedge \dots \wedge$ 
     $\textit{Smsg} = \textit{messages\_to\_send}(s, \textit{msgs})$ 
then %% Specification of the reachable states:
  free { pred reach : State
     $\forall s, s' : \textit{State}; \textit{msgs} : \textit{Set}[\textit{R\_Message}]; \textit{Smsg} : \textit{Set}[\textit{S\_Message}]$ 
    • reach(init)
    • reach(s)  $\wedge$  is_step(s, msgs, Smsg, s')  $\Rightarrow$  reach(s') }
end

```

Of course the specification SBCS\_ANALYSIS is likely to be further structured into smaller pieces of specifications. Indeed, since the informal requirements are too complex to be handled as a whole, we will therefore successively concentrate on various parts of them. The study and formalization of each chunk of requirements will lead to specifications that will later on be put together to obtain the SBCS\_ANALYSIS specification. As already pointed out, it is likely that when analyzing a chunk of requirements we will discover the need for new observers on states (i.e., new state variables). This means that the specification SBCS\_STATE will be subject to iterated extensions where we introduce the new observers that are needed.

For instance, in App. C.6 it is explained that when the *STOP* message has been received three times in a row, the program must go into the *EmergencyStop* mode. We need therefore an observer (i.e., a state variable) to record the number of times we have successively received the *STOP* message. So in the sequel we will start from the following specification of states:

```
spec  SBCS_STATE_1 =
      PRELIMINARY
then  sort  State
      ops   mode : State → Mode;
           numSTOP : State → Nat
end
```

Introducing the new observer *numSTOP* means that we will have to specify a corresponding *next\_numSTOP* operation in the *SBCS\_ANALYSIS* specification.

♡ Let us insist again on the importance of *structuring* our formal specification into coherent, easy to grasp, *named specifications* that will be easy to reuse later on by referring to their names. As explained above it is moreover essential to rely on a *loose specification* of states so that we can introduce later on as many observers as needed. Using a *predicate* (such as *is\_step*) to describe a labeled transition system is quite convenient here, and provides us with an elegant way of handling both input and output for each transition. Finally, it is essential to use a *free constraint* to specify the reachable states, and thus we need to *combine parts with a loose interpretation and parts with an initial interpretation in the same specification*. ♡

## 13.4 Specifying the Mode of Operation

Our next step is the specification of the various operating modes in which the steam-boiler control system operates. (As explained in Sect. 13.1 we do not take into account the *Initialization* mode in this specification.) According to App. C.4, the operating mode of the steam-boiler control system depends on which failures have been detected (see e.g. “all physical units [are] operating correctly”, “a failure of the water level measuring unit”, “detection of an erroneous transmission”). It depends also on the expected evolution of the water level (see “If the water level is risking to reach...”).

We will therefore assume that the specification *SBCS\_ANALYSIS* will provide the following predicates which, given a known state and newly received messages, should reflect the failures detected by the steam-boiler control system.<sup>5</sup>

<sup>5</sup> It is important to make a subtle distinction between the actual failures, about which we basically know nothing, and the failures detected by the steam-boiler



- *Transmission\_OK* :  $State \times Set[R\_Message]$   
should hold iff we rely on the message transmission system,
- *PU\_OK* :  $State \times Set[R\_Message] \times PhysicalUnit$   
should hold iff we rely on the corresponding physical unit,
- *DangerousWaterLevel* :  $State \times Set[R\_Message]$   
should hold iff we estimate that the water level risks reaching the min (*M1*) or max (*M2*) limits.

However, at this stage our understanding of the steam-boiler control system is still quite preliminary, and it is therefore too early to attempt to specify these predicates. Therefore, our specification *MODE\_EVOLUTION*, where we specify the new operating mode according to the previous one and the newly received messages (i.e., the operation *next\_mode*), will be made *generic* w.r.t. these predicates. Let us emphasize that here genericity is used to ensure a *loose coupling* between the current specification of interest, *MODE\_EVOLUTION*, and other specifications expected to provide the needed predicates.

Let us now explain how to specify the new mode of operation. At first glance the informal requirements (see App. C.4) look quite complicated, mainly because they explain, for each operating mode, under which conditions the steam-boiler control system should stay in the same operating mode or switch to another one. However, things get simpler if we analyze under which conditions the next mode is one of the specified operating modes. In particular, a careful analysis of the requirements shows that, except for switching to the *EmergencyStop* mode, we can determine the new operating mode (after receiving some messages) without even taking into account the previous one.

To improve the legibility of our specification it is better to introduce some auxiliary predicates (*Everything\_OK*, *AskedToStop*, *SystemStillControllable*, and *Emergency*) that will facilitate the characterization of the conditions under which the system switches from one mode to another:

- The aim of the predicate *Everything\_OK* is to express that we believe that all physical units are operating correctly, including the message transmission system.
- The aim of the predicate *AskedToStop* is to determine if we have received the *STOP* message three times in a row.
- The aim of the predicate *SystemStillControllable* is to characterize the conditions under which the steam-boiler control system will operate in *Rescue* mode. Let us point out that the corresponding part of the informal requirements (see App. C.4.4) is not totally clear, in particular the exact meaning of the sentence “if one can rely upon the information which comes from the units for controlling the pumps”. There is a double ambiguity here: on the one hand it is unclear whether “the pumps” means “all pumps” or “at least

---

control system. The behavior of the steam-boiler control system is induced by the failures detected, whatever the actual failures are.

one pump”; on the other hand there are two ways of “controlling” each pump (the information sent by the pump and the information sent by the pump controller), and it is unclear whether “controlling” refers to both of them or only to the pump controller. Our interpretation will be as follows: we consider it is enough that at least one pump is “correctly working”, and for us correctly working will mean we rely on both the pump and the associated pump controller. As with all interpretations made during the formalization process, we should in principle interact with the designers of the informal requirements in order to clarify what was the exact intended meaning and to check that our interpretation is adequate. The important point is that our interpretation is entirely localized in the axiomatization of *SystemStillControllable*, and it will therefore be fairly easy to change our specification in case of misinterpretation.

- The aim of the predicate *Emergency* is to characterize when we should switch to the *EmergencyStop* mode. In App. C.4.2, it is said that the steam-boiler control system should switch from *Normal* mode to *Rescue* mode as soon as a failure of the water level measuring unit is detected. However, in App. C.4.4, it is explained that the steam-boiler control system can only operate in *Rescue* mode if some additional conditions hold (represented by our predicate *SystemStillControllable*). We decide therefore that when in *Normal* mode, if a failure of the water level measuring unit is detected, the steam-boiler control system will switch to *Rescue* mode only if *SystemStillControllable* holds, otherwise it will switch (directly) to *EmergencyStop* mode.<sup>6</sup>

The axiomatization of the next mode of operation is now both simple and clear, as illustrated by the `MODE_EVOLUTION` specification.<sup>7</sup>

```
spec  MODE_EVOLUTION
      [preds Transmission_OK : State × Set[R_Message];
          PU_OK : State × Set[R_Message] × PhysicalUnit;
          DangerousWaterLevel : State × Set[R_Message] ]
      given SBCS_STATE_1 =
local  %% Auxiliary predicates to structure the specification of next_mode.
```

<sup>6</sup> If our interpretation is incorrect, then in some cases we may have replaced a sequence *Normal* → *Rescue* → *EmergencyStop* by a sequence *Normal* → *EmergencyStop*. Note that a sequence *Normal* → *Rescue* → *Normal* or *Degraded* is not possible since several cycles are necessary between a failure detection and the decision that the corresponding unit is again fully operational, see Sect. 13.5, i.e., we must have a sequence of the form *Normal* → *Rescue* → ... → *Rescue* → *Normal* or *Degraded* in such cases.

<sup>7</sup> Note that once in the *EmergencyStop* mode, we specify that we stay in this mode forever, rather than specifying that the steam-boiler control system actually stops. Note also that we realize that the operation *next\_numSTOP* is better specified in this `MODE_EVOLUTION` specification.

```

preds Everything_OK, AskedToStop, SystemStillControllable,
        Emergency : State × Set[R_Message]
∀ s : State; msgs : Set[R_Message]
    • Everything_OK(s, msgs) ⇔
        (Transmission_OK(s, msgs) ∧
         (∀ pu : PhysicalUnit • PU_OK(s, msgs, pu)))
    • AskedToStop(s, msgs) ⇔ numSTOP(s) = 2 ∧ STOP ∈ msgs
    • SystemStillControllable(s, msgs) ⇔
        (PU_OK(s, msgs, SteamOutput) ∧
         (∃ pn : PumpNumber • PU_OK(s, msgs, Pump(pn))
          ∧ PU_OK(s, msgs, PumpController(pn))))
    • Emergency(s, msgs) ⇔
        ( mode(s) = EmergencyStop ∨
          AskedToStop(s, msgs) ∨
          ¬ Transmission_OK(s, msgs) ∨
          DangerousWaterLevel(s, msgs) ∨
          (¬ PU_OK(s, msgs, WaterLevel) ∧
           ¬ SystemStillControllable(s, msgs)) )

within ops next_mode : State × Set[R_Message] → Mode;
            next_numSTOP : State × Set[R_Message] → Nat
∀ s : State; msgs : Set[R_Message]
%% Emergency stop mode:
    • Emergency(s, msgs) ⇒ next_mode(s, msgs) = EmergencyStop
%% Normal mode:
    • ¬ Emergency(s, msgs) ∧
      Everything_OK(s, msgs) ⇒ next_mode(s, msgs) = Normal
%% Degraded mode:
    • ¬ Emergency(s, msgs) ∧
      ¬ Everything_OK(s, msgs) ∧
      PU_OK(s, msgs, WaterLevel) ∧
      Transmission_OK(s, msgs) ⇒ next_mode(s, msgs) = Degraded
%% Rescue mode:
    • ¬ Emergency(s, msgs) ∧
      ¬ PU_OK(s, msgs, WaterLevel) ∧
      SystemStillControllable(s, msgs) ∧
      Transmission_OK(s, msgs) ⇒ next_mode(s, msgs) = Rescue
%% next_numSTOP:
    • next_numSTOP(s, msgs) = numSTOP(s) + 1 when STOP ∈ msgs
      else 0
end

```

In the next step of our formalization process, we will specify the predicates assumed by *MODE\_EVOLUTION*, which amounts to specifying the detection of equipment failures. This will be the topic of the next section.

♡ Two essential features of CASL have been used in the specification `MODE_EVOLUTION`. On the one hand, the use of a *generic specification* (*with imports*) ensures loose coupling of the current specification of interest with the rest of the steam-boiler control system specification. On the other hand, *auxiliary predicates* improve the legibility of the specification, and declaring them in the *local part* of the specification ensures they are hidden and therefore not exported. ♡

## 13.5 Specifying the Detection of Equipment Failures

The detection of equipment failures is described in App. C.7. It is quite clear that this detection is the most difficult part to formalize, mainly because both our intuition and the requirements (see e.g. “knows from elsewhere”, “incompatible with the dynamics”) suggest that we should take into account some inter-dependencies when detecting the various possible failures.

For instance, if we ask a pump to stop, and if in the next cycle the pump state still indicates that the pump is open, we may in principle infer either a failure of the message transmission system (e.g. the stop order was not properly sent or was not received, or the message indicating the pump state has been corrupted) or a failure of the pump (which was not able to execute the stop order or which sends incorrect state messages). Our understanding of the requirements is that in such a case we must conclude there has been a failure of the pump, not of the message transmission system. Let us stress again that it is important to distinguish between the actual failures of the various pieces of equipment, and the diagnosis we will make. Only the latter is relevant in our specification.

### 13.5.1 Understanding the Detection of Equipment Failures

Before starting to specify the detection of equipment failures, we must proceed to a careful analysis of App. C.7, in order to clarify the inter-dependencies mentioned above. Only then will we be able to understand how to structure our specification of this crucial part of the problem.

A first rough analysis of the part of App. C.7 devoted to the description of potential failures of the physical units (i.e. of the pumps, the pump controllers and the two measuring devices) shows that these failures are detected on the basis of the information contained in the messages received: we must check that the received values are in accordance with some *expected* values (according to the history of the system, i.e. according to the “dynamics of the system” and to the messages previously sent by the steam-boiler control system). In particular, the detection of failures of the physical units relies on the fact that we have effectively received the necessary messages. If we have not received these messages, then we should conclude there has been a failure of the message transmission system (see below), and in these cases (see the

MODE\_EVOLUTION specification), the steam-boiler control system switches to the *EmergencyStop* mode. The further detection of failures of the physical units (in addition to the already detected failure of the message transmission system) is therefore irrelevant in such cases.

Let us now consider the message transmission system. The description of potential failures of the message transmission system in App. C.7 is quite short. Basically, it tells us that we should check that the steam-boiler control system has received all the messages it was expecting, and that none of the messages received is aberrant. However, it is important to note that the involved analysis of the messages received combines two aspects: on the one hand, there is some ‘static’ analysis of the messages received in order to check that all messages that must be present in each transmission are effectively present (see App. C.6). These messages are exactly the messages required to proceed to the detection of the failures of the physical units. On the other hand, the steam-boiler control system expects to receive (or, on the contrary, not to receive) some specific messages according to the history of the system (for instance, the steam-boiler control system expects to receive a “failure acknowledgement” from a physical unit once it has detected a corresponding failure and sent a “failure” message to this unit, but not before), and here some ‘dynamic’ analysis is required. Obviously, the static analysis of the messages can be made on the basis of the messages received only, while the dynamic analysis must take into account, in addition to the messages received, the history of the system, and more precisely the history of the failures detected so far and of the “failure acknowledgement” and “repaired” messages received so far.

From this first analysis we draw the following conclusions on how to specify the detection of equipment failures:

1. In a first step we should keep track of the failure status of the physical units. This will lead to a new observer *status* on states, and to a specification STATUS\_EVOLUTION of how this status evolves, i.e., of a *next\_status* operation.
2. Then we specify the detection of the message transmission system failures (hence *Transmission\_OK*) in the specification MESSAGE\_TRANSMISSION\_SYSTEM\_FAILURE. As explained above, in a first step we take care of the static analysis of the messages received, and then in a second step we take care of the dynamic analysis of the messages received, using how we have kept track of the “status” of the physical units, i.e., using the observer *status*.
3. Then, for each physical unit, we specify the detection of its failures by comparing the message received with the *expected* one. For this comparison we can freely assume that the static analysis of the messages received has been successful, i.e., that the message sent by the physical unit has been received.

The corresponding specifications are described in the next subsections.

### 13.5.2 Keeping Track of the Status of the Physical Units

Remember that to perform the dynamic analysis of the messages received, as explained above, we must check that we receive “failure acknowledgement” and “repaired” messages when appropriate. In order to do this, we must keep track of the failures detected and of the “failure acknowledgement” and “repaired” messages received. Since the same reasoning applies for all physical units, we can do the analysis in a generic way. For each physical unit, we will keep track of its status, which can be either *OK*, *FailureWithoutAck* or *FailureWithAck*. The status of a physical unit will then be updated accordingly to the detection of failures, and receipt of “failure acknowledgement” and “repaired” messages.

Thus, in a first step we should extend the specification `SBCS_STATE_1` to add an observer related to the failure status of physical units:

```
spec SBCS_STATE_2 =
  SBCS_STATE_1
then free type Status ::= OK | FailureWithoutAck | FailureWithAck
  op status : State × PhysicalUnit → Status;
end
```

Now the specification of how the status of a physical unit evolves, i.e., of the operation *next\_status* in `STATUS_EVOLUTION`, is quite straightforward. We rely again on the predicate *PU\_OK*.<sup>8</sup>

```
spec STATUS_EVOLUTION
[pred PU_OK : State × Set[R_Message] × PhysicalUnit]
given SBCS_STATE_2 =
op next_status : State × Set[R_Message] × PhysicalUnit → Status
∀ s : State; msgs : Set[R_Message]; pu : PhysicalUnit
• status(s, pu) = OK ∧ PU_OK(s, msgs, pu)
  ⇒ next_status(s, msgs, pu) = OK
• status(s, pu) = OK ∧ ¬ PU_OK(s, msgs, pu)
  ⇒ next_status(s, msgs, pu) = FailureWithoutAck
• status(s, pu) = FailureWithoutAck ∧
  FAILURE_ACKNOWLEDGEMENT(pu) ∈ msgs
  ⇒ next_status(s, msgs, pu) = FailureWithAck
• status(s, pu) = FailureWithoutAck ∧
  ¬ (FAILURE_ACKNOWLEDGEMENT(pu) ∈ msgs)
  ⇒ next_status(s, msgs, pu) = FailureWithoutAck
```

---

<sup>8</sup> The reader may detect that the specification `STATUS_EVOLUTION` is not completely correct. However, we prefer to give here the text of the specification as it was originally written, and we will explain in Sect. 13.9 how we detect, when validating the specification of the steam-boiler control system, that something is not correct, and how the problem can be fixed.

- $status(s, pu) = FailureWithAck \wedge REPAIRED(pu) \in msgs$   
 $\Rightarrow next\_status(s, msgs, pu) = OK$
- $status(s, pu) = FailureWithAck \wedge \neg (REPAIRED(pu) \in msgs)$   
 $\Rightarrow next\_status(s, msgs, pu) = FailureWithAck$

end

♡ Here again we rely on a *generic specification with imports* to ensure loose coupling. As claimed earlier, the *loose specification* of states makes it easy to introduce further observers (hence further state variables). ♡

### 13.5.3 Detection of the Message Transmission System Failures

As explained above, we first specify the static analysis of the messages received, and then we specify the dynamic analysis of these messages.

To specify the static analysis of messages, it is necessary to check that all “indispensable” messages are present. In addition, a set of messages received is “acceptable” if there are no “duplicated” messages in this set. Since we have specified the collection of messages received as a set, we cannot have several occurrences of exactly the same message in this set. (Note that this means that our choice of using “sets” instead of “bags”, for instance, is therefore not totally innocent: either we assume that receiving several occurrences of exactly the same message will never happen, and this is an assumption about the environment, or we assume that this case should not lead to the detection of a failure of the message transmission system, and this is an assumption about the requirements.) However, specifying the collection of messages received as a set does not imply that a set of messages received cannot contain several  $LEVEL(v)$  messages, with distinct values (for instance). Hence we must check this explicitly.

Remember that receiving “unknown” messages (i.e., messages that do not belong to the list of messages as specified in App. C.6) is taken into account via the extra constant *junk* message (see the specification `MESSAGES_RECEIVED`). Another erroneous situation is when we simultaneously receive a failure acknowledgement and a repaired message for the same physical unit, i.e., that at least one cycle is needed between acknowledging the failure and repairing the unit. We will check this as well.<sup>9</sup>

We focus now on the dynamic analysis of the messages received. As explained above, to perform this dynamic analysis, we check that we receive “failure acknowledgement” and “repaired” messages when appropriate, according to the current status of each physical unit. We understand that for

---

<sup>9</sup> We must confess that this belief is induced by our intuition about the behavior of the system. Indeed nothing in the requirements allows us to make either this interpretation or the opposite one. Although not essential, this assumption will simplify the axiomatization.

each failure signaled by the steam-boiler control system, the corresponding physical unit will send just one failure acknowledgement. Moreover, we will specify the steam-boiler control system in such a way that when it receives a “repaired” message, the steam-boiler control system acknowledges it immediately. Hence, if there is no problem with the message transmission system, and due to the fact that transmission time can be neglected, the steam-boiler control system must in principle receive only one repaired message for a given failure. Note that this does not contradict the “until...” part of the sentences describing the “repaired” messages in the informal requirements (see App. C.6). To summarize, we consider that we have received an unexpected message when:

- the program receives initialization messages but is no longer in initialization mode; or
- the program receives for some physical unit a “failure acknowledgement” without having previously sent the corresponding failure detection message, or receives redundant failure acknowledgements; or
- the program receives for some physical unit a “repaired message”, but the unit is OK or its failure is not yet acknowledged.

We now have all the ingredients required to specify the *Transmission\_OK* predicate, taking into account both static and dynamic aspects, which leads to the following MESSAGE\_TRANSMISSION\_SYSTEM\_FAILURE specification.

```

spec MESSAGE_TRANSMISSION_SYSTEM_FAILURE =
  SBCS_STATE_2
then local %% Static analysis:
  pred __is_static_OK : Set[R_Message]
   $\forall msgs : Set[R\_Message]$ 
  • msgs is_static_OK  $\Leftrightarrow$ 
    (  $\neg(junk \in msgs) \wedge$ 
      ( $\exists!v : Value \bullet LEVEL(v) \in msgs$ )  $\wedge$ 
      ( $\exists!v : Value \bullet STEAM(v) \in msgs$ )  $\wedge$ 
      ( $\forall pn : PumpNumber \bullet \exists!ps : PumpState \bullet$ 
        PUMP_STATE(pn, ps)  $\in msgs$ )  $\wedge$ 
      ( $\forall pn : PumpNumber \bullet \exists!pcs : PumpControllerState \bullet$ 
        PUMP_CONTROLLER_STATE(pn, pcs)  $\in msgs$ )  $\wedge$ 
      ( $\forall pu : PhysicalUnit \bullet$ 
         $\neg(FAILURE\_ACKNOWLEDGEMENT(pu) \in msgs$ 
           $\wedge REPAIRED(pu) \in msgs))$  )
  %% Dynamic analysis:
  pred __is_NOT_dynamic_OK_for__ : Set[R_Message]  $\times$  State

```



```

forall  $s : \text{State}; \text{msgs} : \text{Set}[\text{R\_Message}]$ 
  •  $\text{msgs is\_NOT\_dynamic\_OK\_for } s \Leftrightarrow$ 
    ( (  $\neg(\text{mode}(s) = \text{Initialization}) \wedge$ 
      (  $\text{STEAM\_BOILER\_WAITING} \in \text{msgs} \vee$ 
         $\text{PHYSICAL\_UNITS\_READY} \in \text{msgs}$  ) )
     $\vee$  (  $\exists pu : \text{PhysicalUnit}$  •
       $\text{FAILURE\_ACKNOWLEDGEMENT}(pu) \in \text{msgs} \wedge$ 
      (  $\text{status}(s, pu) = \text{OK} \vee \text{status}(s, pu) = \text{FailureWithAck}$  ) )
     $\vee$  (  $\exists pu : \text{PhysicalUnit}$  •
       $\text{REPAIRED}(pu) \in \text{msgs} \wedge$ 
      (  $\text{status}(s, pu) = \text{OK} \vee \text{status}(s, pu) = \text{FailureWithoutAck}$  ) ) )

within
pred  $\text{Transmission\_OK} : \text{State} \times \text{Set}[\text{R\_Message}]$ 
forall  $s : \text{State}; \text{msgs} : \text{Set}[\text{R\_Message}]$ 
  •  $\text{Transmission\_OK}(s, \text{msgs}) \Leftrightarrow$ 
    (  $\text{msgs is\_static\_OK} \wedge \neg(\text{msgs is\_NOT\_dynamic\_OK\_for } s)$  )
end

```

♡ Here again *auxiliary predicates* declared in the *local part* of the specification are quite useful to improve the legibility of the specification. Note also the use of nested *quantifiers* in axioms (‘ $\forall$ ’, ‘ $\exists$ ’ as well as ‘ $\exists!$ ’) – without them the axioms would be much more intricate, or further auxiliary operations would be needed. ♡

### 13.5.4 Detection of the Pump and Pump Controller Failures

We start by considering the detection of the failures of the pumps.

As explained in Sec 13.5.1, we rely on the predicted pump state message. Thus, in a first step we should extend the specification `SBCS_STATE_2` to add an observer related to the prediction of pump state messages. The prediction (*Open* or *Closed*) can however only be made when the status of the corresponding pump is *OK*. This is why we extend the sort *PumpState* to introduce a constant *Unknown\_PS*:

```

spec SBCS_STATE_3 =
  SBCS_STATE_2
then free type ExtendedPumpState ::= sort PumpState | Unknown_PS
  op  $\text{PS\_predicted} : \text{State} \times \text{PumpNumber} \rightarrow \text{ExtendedPumpState};$ 
  % {  $\text{status}(s, \text{Pump}(pn)) = \text{OK} \Leftrightarrow$ 
     $\neg(\text{PS\_predicted}(s, pn) = \text{Unknown\_PS})$  } %
end

```

The specification of the detection of pump failures is now straightforward and is given in the `PUMP_FAILURE` specification. Remember that the meaning of *Pump\_OK* is only relevant when *Transmission\_OK* holds, which in particular implies that for each pump, there is only one *PUMP\_STATE* message for

it in *msgs*. Moreover, we check the received value only if the predicted value is not *Unknown\_PS*.

```

spec PUMP_FAILURE =
  SBCS_STATE_3
then pred Pump_OK : State  $\times$  Set[R_Message]  $\times$  PumpNumber
   $\forall s : \text{State}; \text{msgs} : \text{Set}[\text{R\_Message}]; \text{pn} : \text{PumpNumber}$ 
    • Pump_OK(s, msgs, pn)  $\Leftrightarrow$ 
      PS_predicted(s, pn) = Unknown_PS  $\vee$ 
      PUMP_STATE(pn, PS_predicted(s, pn) as PumpState)  $\in$  msgs
end

```

Let us now consider the detection of the failures of the pump controllers. Again we rely on the predicted pump state controller message. Here, we must be a bit careful in order to reflect the fact that stopping a pump has an instantaneous effect, while starting it takes five seconds (see App. C.2.3). Since five seconds is, unfortunately, exactly the elapsed time between two cycles, when we decide to activate a pump we may have to wait two cycles to receive a corresponding *Flow* pump controller state. This is why, in addition to the constant *Unknown\_PCS*, used for the cases where no prediction can be made since the pump controller is not working correctly, we also introduce a constant *SoonFlow* to be used for the prediction related to a just activated pump.

```

spec SBCS_STATE_4 =
  SBCS_STATE_3
then free type
  ExtendedPumpControllerState ::= sort PumpControllerState
    | SoonFlow | Unknown_PCS
op PCS_predicted : State  $\times$  PumpNumber
   $\rightarrow$  ExtendedPumpControllerState;
  % { status(s, PumpController(pn)) = OK  $\Rightarrow$ 
     $\neg$  (PCS_predicted(s, pn) = Unknown_PCS) } %
end

```

The specification of the detection of pump controller failures is now straightforward and is given in the PUMP\_CONTROLLER\_FAILURE specification. Remember that the meaning of *Pump\_Controller\_OK* is only relevant when *Transmission\_OK* holds, which in particular implies that for each pump, there is only one *PUMP\_CONTROLLER\_STATE* message for it in *msgs*. Moreover, we check the received value only if the predicted value is either *Flow* or *NoFlow*, since if it is *SoonFlow* or *Unknown\_PCS* we cannot conclude.

```

spec PUMP_CONTROLLER_FAILURE =
  SBCS_STATE_4
then pred Pump_Controller_OK : State  $\times$  Set[R_Message]  $\times$  PumpNumber
   $\forall s : \text{State}; \text{msgs} : \text{Set}[\text{R\_Message}]; \text{pn} : \text{PumpNumber}$ 
    • Pump_Controller_OK(s, msgs, pn)  $\Leftrightarrow$ 
      PCS_predicted(s, pn) = Unknown_PCS
       $\vee$  PCS_predicted(s, pn) = SoonFlow
       $\vee$  PUMP_CONTROLLER_STATE(pn,
        PCS_predicted(s, pn) as PumpControllerState)  $\in$  msgs
end

```

♡ In the above specifications, using *supersorts* to extend previously defined datatypes is particularly convenient, and avoids the need to explicitly relate values of *PumpState* and values of *ExtendedPumpState* (and similarly for *PumpControllerState*). Note the use of explicit *castings* in the axioms – in particular, the fact that *predicates do not hold on undefined arguments* resulting from castings is used in the above specifications. ♡

### 13.5.5 Detection of the Steam and Water Level Measurement Device Failures

To specify the failures of the steam and water level measurement devices, we must again rely on some predicted values. Here however we cannot predict an exact value, but only an interval in which the received value should be contained. This leads to the following extension of *SBCS\_STATE\_4*:

```

spec SBCS_STATE_5 =
  SBCS_STATE_4
then free type Valpair ::= pair(low : Value; high : Value)
ops steam_predicted, level_predicted : State  $\rightarrow$  Valpair;
  % { low(steam_predicted(s)) is the minimal steam output predicted,
    high(steam_predicted(s)) is the maximal steam output predicted,
    and similarly for level_predicted. } %
end

```

The specification of the failures of the measurement devices is again straightforward and is given in the *STEAM\_FAILURE* and *LEVEL\_FAILURE* specifications. Remember that the meaning of *Steam\_OK* (*Level\_OK* resp.) is only relevant when *Transmission\_OK* holds, which in particular implies that there is only one *STEAM*(*v*) (*LEVEL*(*v*) resp.) message in *msgs* (hence only one possible *v* in the quantifications  $\forall v : \text{Value} \dots$  below). Note also that here we assume that the predicted values will take care of the static limits (*0* and *W* for the steam, *0* and *C* for the water level), thus we do not need to check these static limits explicitly here.

```

spec STEAM_FAILURE =
  SBCS_STATE_5
then pred Steam_OK : State  $\times$  Set[R_Message]
   $\forall s : \text{State}; \text{msgs} : \text{Set}[\text{R\_Message}]$ 
    • Steam_OK(s, msgs)  $\Leftrightarrow$ 
      ( $\forall v : \text{Value}$  • STEAM(v)  $\in$  msgs  $\Rightarrow$ 
        (low(steam_predicted(s))  $\leq$  v)  $\wedge$ 
        (v  $\leq$  high(steam_predicted(s))) )
end

```

```

spec LEVEL_FAILURE =
  SBCS_STATE_5
then pred Level_OK : State  $\times$  Set[R_Message]
   $\forall s : \text{State}; \text{msgs} : \text{Set}[\text{R\_Message}]$ 
    • Level_OK(s, msgs)  $\Leftrightarrow$ 
      ( $\forall v : \text{Value}$  • LEVEL(v)  $\in$  msgs  $\Rightarrow$ 
        (low(level_predicted(s))  $\leq$  v)  $\wedge$ 
        (v  $\leq$  high(level_predicted(s))) )
end

```

### 13.5.6 Summing Up

We now have all the ingredients necessary for the specification of the predicate *PU\_OK*. This is done in the FAILURE\_DETECTION specification, which integrates together all the specifications related to failure detection.

```

spec FAILURE_DETECTION =
{
  MESSAGE_TRANSMISSION_SYSTEM_FAILURE
  and PUMP_FAILURE and PUMP_CONTROLLER_FAILURE
  and STEAM_FAILURE and LEVEL_FAILURE
  then pred PU_OK : State  $\times$  Set[R_Message]  $\times$  PhysicalUnit
     $\forall s : \text{State}; \text{msgs} : \text{Set}[\text{R\_Message}]; \text{pn} : \text{PumpNumber}$ 
      • PU_OK(s, msgs, Pump(pn))  $\Leftrightarrow$  Pump_OK(s, msgs, pn)
      • PU_OK(s, msgs, PumpController(pn))  $\Leftrightarrow$ 
        Pump_Controller_OK(s, msgs, pn)
      • PU_OK(s, msgs, SteamOutput)  $\Leftrightarrow$  Steam_OK(s, msgs)
      • PU_OK(s, msgs, WaterLevel)  $\Leftrightarrow$  Level_OK(s, msgs)
} hide ops Pump_OK, Pump_Controller_OK, Steam_OK, Level_OK
end

```

♡ In the above specification, we rely on explicit *hiding* of operations that are no longer needed. Moreover, the ‘*same name, same thing*’ principle is essential here: each of the five specifications extended in FAILURE\_DETECTION is itself an extension of some specification SBCS\_STATE\_*i* of states, but with the ‘*same name, same thing*’ principle we get the effect that each of them extends SBCS\_STATE\_5. ♡

## 13.6 Predicting the Behavior of the Steam-Boiler

In the previous section we have explained that failure detection was to a large extent based on a comparison between the messages received and the expected ones. For this purpose we have extended the specification `SBCS_STATE` by several observers, which means we have assumed that at each cycle, we record in some state variables the information needed to compute the expected messages at the next cycle. According to our explanations in Sect. 13.3, we must now specify, for each observer *obs* introduced, a corresponding *next\_obs* operation. This is the aim of this section.

We have already defined the operation *next\_mode* in the generic specification `MODE_EVOLUTION` (see Sect. 13.4) and the operation *next\_status* in the generic specification `STATUS_EVOLUTION` (see Sect. 13.5.2). Thus what is left is the specification of the operations *next\_PS\_predicted*, *next\_PCS\_predicted*, *next\_steam\_predicted* and *next\_level\_predicted*.

As explained in Sect. 13.5, the informal requirements suggest that we should take into account some inter-dependencies when predicting values to be received at the next cycle. For instance, the water level in the steam-boiler depends on how much steam is produced, but also on how much water is poured into the steam boiler by the pumps which are open. The information provided by the water level prediction is obviously crucial to decide whether we should activate or stop some pumps. On the other hand, to predict the pump state and pump controller state messages to be received at the next cycle, we must know which pumps have been ordered to be activated or to be stopped.

From this first analysis we draw the following conclusions on how to specify the needed predictions:

1. In a first step we should predict the interval in which the steam output is expected to stay during the next cycle: this prediction relies only on the just received value *STEAM(v)* (if we trust it) or on the previously predicted values for the steam production. This is because the production of steam is expected to vary according to its maximum gradients of increase and decrease, and nothing else.
2. In the next step we should decide whether some pumps have to be ordered to activate or to stop. This decision, plus the knowledge about the current state of the pumps (as much as we trust it), and the predicted evolution of the steam production, should allow us to predict the evolution of the water level.
3. Then, on the basis of the current states of the pumps and pump controllers, together with the choice of pumps to be activated or stopped, we can predict the states of the pumps and of the pump controllers at the next cycle.

Of course all these predictions are only meaningful as long as no failure of the message transmission system has been detected (but if this is not the case the

steam-boiler control system switches to the *EmergencyStop* mode and stops, so no predictions are needed anyway). The corresponding specifications are described in the next subsections.

### 13.6.1 Predicting the Steam Output and the Water Level

To predict the intervals in which the steam output and the water level are expected to stay during the next cycle, we will proceed as follows (taking into account the “Additional Information” provided in [1, pp. 507–509]):

1. Following the analysis sketched above, when we are in the state  $s$  and have received the messages  $msgs$ , to predict the interval in which the steam output is expected to stay during the next cycle, we first should compute the *adjusted\_steam* interval: this interval is either the (interval reduced to the) *received\_steam* value if we can rely on it (i.e., if  $PU\_OK(s, msgs, SteamOutput)$  holds), or the *steam\_predicted* interval (stored in the state  $s$  at the previous cycle).
2. Then, we use the maximum gradients of increase and decrease (i.e.,  $U1$  and  $U2$ ), to predict the interval in which the steam output is expected to stay during the next cycle.
3. We proceed similarly for the water level: first we compute the *adjusted\_level* interval, which is either the (interval reduced to the) *received\_level* value if we can rely on it (i.e., if  $PU\_OK(s, msgs, WaterLevel)$  holds), or the *level\_predicted* interval (stored in the state  $s$  at the previous cycle).
4. Then we should consider *broken\_pumps* (the pumps  $pn$  for which either  $PU\_OK(s, msgs, Pump(pn))$  does not hold or  $PU\_OK(s, msgs, PumpController(pn))$  does not hold – or both) and the *reliable\_pumps*, which are not broken and are therefore known to be either *Open* or *Closed*.
5. At this point we must decide which pumps are ordered to activate or to stop.

*However, the specific control strategy for deciding which pumps should be activated or stopped need not to be detailed in this requirements specification: this can be left to a further refinement towards an implementation of the steam-boiler control system.* (Obviously the strategy should compare the *adjusted\_level* with the recommended interval ( $N1, N2$ ) and decide accordingly.)

We will therefore rely on a loosely specified *chosen\_pumps* operation, for which we just impose some soundness conditions (e.g., a pump ordered to activate should be currently considered as “reliable” and *Closed*, a pump ordered to stop should be currently considered as “reliable” and *Open*).

6. Now we can compute the minimal and maximal amounts of water that will be poured into the steam-boiler during the next cycle. To compute *minimal\_pumped\_water*, we consider that only the pumps which are “reliable” and already *Open* will pour some water in; the *broken\_pumps*, the pumps which are just ordered to activate, and the pumps which are ordered to stop are all considered not to be pouring water in. Similarly, to

compute *maximal\_pumped\_water*, we consider that the pumps which are “reliable” and already *Open*, the pumps which are just ordered to activate, as well as all the *broken\_pumps*, may pour some water in; only the “reliable” pumps just ordered to stop or already stopped are known not to be pouring any water in.

7. Finally, we can predict the interval in which the water level is expected to stay during the next cycle.
8. This prediction is the basis for deciding whether the water level risks to reach a *DangerousWaterLevel* (i.e., below *M1* or above *M2*).

Note that the intervals in which the steam output and the water level are expected to stay during the next cycle are predicted without considering the *next\_status* of these devices. This is indeed necessary for the *Degraded* and *Rescue* operating modes. This leads to the following STEAM\_AND\_LEVEL\_PREDICTION specification.

```
spec STEAM_AND_LEVEL_PREDICTION =
  FAILURE_DETECTION and SET [sort PumpNumber]
then local
  ops received_steam : State × Set[R_Message] → Value;
     adjusted_steam : State × Set[R_Message] → Valpair;
     received_level : State × Set[R_Message] → Value;
     adjusted_level : State × Set[R_Message] → Valpair;
     broken_pumps : State × Set[R_Message] → Set[PumpNumber];
     reliable_pumps :
       State × Set[R_Message] × PumpState → Set[PumpNumber]
  ∇s : State; msgs : Set[R_Message]; pn : PumpNumber; ps : PumpState
  %% Axioms for STEAM:
  • Transmission_OK(s, msgs) ⇒
    STEAM(received_steam(s, msgs)) ∈ msgs
  • adjusted_steam(s, msgs) =
    pair(received_steam(s, msgs), received_steam(s, msgs))
    when (Transmission_OK(s, msgs) ∧ PU_OK(s, msgs, SteamOutput))
    else steam_predicted(s)
  %% Axioms for LEVEL:
  • Transmission_OK(s, msgs) ⇒
    LEVEL(received_level(s, msgs)) ∈ msgs
  • adjusted_level(s, msgs) =
    pair(received_level(s, msgs), received_level(s, msgs))
    when (Transmission_OK(s, msgs) ∧ PU_OK(s, msgs, WaterLevel))
    else level_predicted(s)
  %% Axioms for auxiliary pumps operations:
  • pn ∈ broken_pumps(s, msgs) ⇔
    ¬ ( PU_OK(s, msgs, Pump(pn)) ∧
        PU_OK(s, msgs, PumpController(pn)) )
```

- $pn \in \text{reliable\_pumps}(s, \text{msgs}, ps) \Leftrightarrow$   
 $\neg (pn \in \text{broken\_pumps}(s, \text{msgs})) \wedge$   
 $\text{PUMP\_STATE}(pn, ps) \in \text{msgs}$

**within**

**ops**  $\text{next\_steam\_predicted} : \text{State} \times \text{Set}[\text{R\_Message}] \rightarrow \text{Valpair};$   
 $\text{chosen\_pumps} :$   
 $\text{State} \times \text{Set}[\text{R\_Message}] \times \text{PumpState} \rightarrow \text{Set}[\text{PumpNumber}];$   
 $\text{minimal\_pumped\_water}, \text{maximal\_pumped\_water} :$   
 $\text{State} \times \text{Set}[\text{R\_Message}] \rightarrow \text{Value};$   
 $\text{next\_level\_predicted} : \text{State} \times \text{Set}[\text{R\_Message}] \rightarrow \text{Valpair}$   
**pred**  $\text{DangerousWaterLevel} : \text{State} \times \text{Set}[\text{R\_Message}]$

%% Axioms for STEAM:

$\forall s : \text{State}; \text{msgs} : \text{Set}[\text{R\_Message}]; pn : \text{PumpNumber}$

- $\text{low}(\text{next\_steam\_predicted}(s, \text{msgs})) =$   
 $\max(0, \text{low}(\text{adjusted\_steam}(s, \text{msgs})) - (U2 \times dt))$
- $\text{high}(\text{next\_steam\_predicted}(s, \text{msgs})) =$   
 $\min(W, \text{high}(\text{adjusted\_steam}(s, \text{msgs})) + (U1 \times dt))$

%% Axioms for PUMPS:

- $pn \in \text{chosen\_pumps}(s, \text{msgs}, \text{Open}) \Rightarrow$   
 $pn \in \text{reliable\_pumps}(s, \text{msgs}, \text{Closed})$
- $pn \in \text{chosen\_pumps}(s, \text{msgs}, \text{Closed}) \Rightarrow$   
 $pn \in \text{reliable\_pumps}(s, \text{msgs}, \text{Open})$
- $\text{minimal\_pumped\_water}(s, \text{msgs}) =$   
 $dt \times P \times \#(\text{reliable\_pumps}(s, \text{msgs}, \text{Open})$   
 $\quad - \text{chosen\_pumps}(s, \text{msgs}, \text{Closed}))$
- $\text{maximal\_pumped\_water}(s, \text{msgs}) =$   
 $dt \times P \times \#((\text{reliable\_pumps}(s, \text{msgs}, \text{Open})$   
 $\quad \cup \text{chosen\_pumps}(s, \text{msgs}, \text{Open})$   
 $\quad \cup \text{broken\_pumps}(s, \text{msgs}))$   
 $\quad - \text{chosen\_pumps}(s, \text{msgs}, \text{Closed}))$

%% Axioms for LEVEL:

- $\text{low}(\text{next\_level\_predicted}(s, \text{msgs})) =$   
 $\max(0, (\text{low}(\text{adjusted\_level}(s, \text{msgs}))$   
 $\quad + \text{minimal\_pumped\_water}(s, \text{msgs}))$   
 $\quad - ( (dt^2 \times U1/2)$   
 $\quad \quad + (dt \times \text{high}(\text{adjusted\_steam}(s, \text{msgs}))) ) )$
- $\text{high}(\text{next\_level\_predicted}(s, \text{msgs})) =$   
 $\min(C, (\text{high}(\text{adjusted\_level}(s, \text{msgs}))$   
 $\quad + \text{maximal\_pumped\_water}(s, \text{msgs}))$   
 $\quad - ( (dt^2 \times U2/2)$   
 $\quad \quad + (dt \times \text{low}(\text{adjusted\_steam}(s, \text{msgs}))) ) )$
- $\text{DangerousWaterLevel}(s, \text{msgs}) \Leftrightarrow$   
 $(\text{low}(\text{next\_level\_predicted}(s, \text{msgs})) \leq M1) \vee$   
 $(M2 \leq \text{high}(\text{next\_level\_predicted}(s, \text{msgs})))$



```

hide ops minimal_pumped_water, maximal_pumped_water
end

```

♡ Note the combination of *implicit hiding of auxiliary operations* declared in the *local part* and of *explicit hiding*: the operations *minimal\_pumped\_water* and *maximal\_pumped\_water* cannot be made local since their specification relies on *chosen\_pumps* which must be exported. ♡

### 13.6.2 Predicting the Pump and Pump Controller States

Specifying the predicted state of each pump at the next cycle is almost trivial. The next pump state is *Unknown\_PS* if the *next\_status* of the pump is not *OK*, otherwise it should be *Open* if:

- it is *Open* now and the pump is not ordered to stop, or
- the pump is ordered to activate;

otherwise, it should be *Closed* since:

- it is *Closed* now and the pump is not ordered to activate, or
- it is ordered to stop.

This leads to the following PUMP\_STATE\_PREDICTION specification. This specification extends STEAM\_AND\_LEVEL\_PREDICTION (since we rely on *chosen\_pumps* for our predictions), and STATUS\_EVOLUTION (which provides *next\_status*) instantiated by FAILURE\_DETECTION (which provides the predicate *PU\_OK* parameter of STATUS\_EVOLUTION).

```

spec PUMP_STATE_PREDICTION =
  STATUS_EVOLUTION [FAILURE_DETECTION]
  and STEAM_AND_LEVEL_PREDICTION
then op next_PS_predicted :
   $State \times Set[R\_Message] \times PumpNumber \rightarrow ExtendedPumpState$ 
   $\forall s : State; msgs : Set[R\_Message]; pn : PumpNumber$ 
  •  $next\_PS\_predicted(s, msgs, pn) =$ 
     $Unknown\_PS \text{ when } \neg(next\_status(s, msgs, Pump(pn)) = OK)$ 
     $\text{else } Open \text{ when } ( PUMP\_STATE(pn, Open) \in msgs \wedge$ 
       $\neg(pn \in chosen\_pumps(s, msgs, Closed)) )$ 
       $\vee pn \in chosen\_pumps(s, msgs, Open)$ 
     $\text{else } Closed$ 
end

```

The reasoning to predict the pump controller state is similar, but we must take into account that two cycles may be needed before a just activated pump leads to a *Flow* state (provided the pump is not stopped meanwhile). Thus, the next pump controller state is *Unknown\_PCS* if the *next\_status* of the pump controller is not *OK*, or if the *next\_status* of the corresponding pump is not *OK*, otherwise the predicted pump controller state value is:

- *Flow* when the pump is not ordered to stop and it is currently *Flow*, or it is currently *NoFlow* but *PCS\_predicted SoonFlow*;
- *NoFlow* if the pump is ordered to stop, or if it is currently *NoFlow* and is not *PCS\_predicted SoonFlow* and the pump is not ordered to activate;
- *SoonFlow* otherwise.

This leads to the following PUMP\_CONTROLLER\_STATE\_PREDICTION specification.

```

spec PUMP_CONTROLLER_STATE_PREDICTION =
  STATUS_EVOLUTION [ FAILURE_DETECTION ]
  and STEAM_AND_LEVEL_PREDICTION
then op next_PCS_predicted :
  State  $\times$  Set[R_Message]  $\times$  PumpNumber
     $\rightarrow$  ExtendedPumpControllerState
   $\forall s : \text{State}; \text{msgs} : \text{Set}[R\_Message]; \text{pn} : \text{PumpNumber}$ 
  • next_PCS_predicted( $s, \text{msgs}, \text{pn}$ ) =
    Unknown_PCS when
       $\neg ( \text{next\_status}(s, \text{msgs}, \text{PumpController}(\text{pn})) = \text{OK} \wedge$ 
         $\text{next\_status}(s, \text{msgs}, \text{Pump}(\text{pn})) = \text{OK} )$ 
    else Flow when
      ( PUMP_CONTROLLER_STATE( $\text{pn}, \text{Flow}$ )  $\in \text{msgs} \vee$ 
        ( PUMP_CONTROLLER_STATE( $\text{pn}, \text{NoFlow}$ )  $\in \text{msgs} \wedge$ 
          PCS_predicted( $s, \text{pn}$ ) = SoonFlow ) )
       $\wedge \neg (\text{pn} \in \text{chosen\_pumps}(s, \text{msgs}, \text{Closed}))$ 
    else NoFlow when
      ( $\text{pn} \in \text{chosen\_pumps}(s, \text{msgs}, \text{Closed})$ )
       $\vee ( \text{PUMP\_CONTROLLER\_STATE}(\text{pn}, \text{NoFlow}) \in \text{msgs} \wedge$ 
         $\neg (\text{PCS\_predicted}(s, \text{pn}) = \text{SoonFlow}) \wedge$ 
         $\neg (\text{pn} \in \text{chosen\_pumps}(s, \text{msgs}, \text{Open})) )$ 
    else SoonFlow
end

```

All our predictions are summarized in the following PU\_PREDICTION specification.

```

spec PU_PREDICTION =
  PUMP_STATE_PREDICTION
  and PUMP_CONTROLLER_STATE_PREDICTION
  % { Both specifications extend STATUS_EVOLUTION
    (instantiated by FAILURE_DETECTION)
    and STEAM_AND_LEVEL_PREDICTION } %
end

```

♡ Since the specification FAILURE\_DETECTION provides the predicate *PU\_OK* required by STATUS\_EVOLUTION, we can now put pieces

together as illustrated by PU\_PREDICTION. Again the ‘*same name, same thing*’ principle is essential here. ♡

### 13.7 Specifying the Messages to Send

At this stage we are left with the specification of the messages to send at each cycle. This is easily specified, following App. C.5, and leads to the following SBCS\_ANALYSIS specification.

The specification SBCS\_ANALYSIS is obtained by instantiating the MODE\_EVOLUTION specification by PU\_PREDICTION, and extending the result by the specification of the operation *messages\_to\_send*.

```

spec SBCS_ANALYSIS =
  MODE_EVOLUTION [ PU_PREDICTION ]
then local
  ops PumpMessages, FailureDetectionMessages :
    State  $\times$  Set[R_Message]  $\rightarrow$  Set[S_Message];
    RepairedAcknowledgementMessages :
    Set[R_Message]  $\rightarrow$  Set[S_Message]
   $\forall s : \text{State}; \text{msgs} : \text{Set}[R\_Message]; \text{Smsg} : S\_Message$ 
  •  $\text{Smsg} \in \text{PumpMessages}(s, \text{msgs}) \Leftrightarrow$ 
    (  $\exists pn : \text{PumpNumber}$  •
      (  $pn \in \text{chosen\_pumps}(s, \text{msgs}, \text{Open})$ 
         $\wedge \text{Smsg} = \text{OPEN\_PUMP}(pn)$  )
       $\vee$  (  $pn \in \text{chosen\_pumps}(s, \text{msgs}, \text{Closed})$ 
         $\wedge \text{Smsg} = \text{CLOSE\_PUMP}(pn)$  ) )
  •  $\text{Smsg} \in \text{FailureDetectionMessages}(s, \text{msgs}) \Leftrightarrow$ 
    (  $\exists pu : \text{PhysicalUnit}$  •
       $\text{Smsg} = \text{FAILURE\_DETECTION}(pu) \wedge$ 
       $\text{next\_status}(s, \text{msgs}, pu) = \text{FailureWithoutAck}$  )
  •  $\text{Smsg} \in \text{RepairedAcknowledgementMessages}(\text{msgs}) \Leftrightarrow$ 
    (  $\exists pu : \text{PhysicalUnit}$  •
       $\text{Smsg} = \text{REPAIRED\_ACKNOWLEDGEMENT}(pu) \wedge$ 
       $\text{next\_status}(s, \text{msgs}, pu) = \text{FailureWithAck}$  )
  within
  op messages_to_send : State  $\times$  Set[R_Message]  $\rightarrow$  Set[S_Message]
   $\forall s : \text{State}; \text{msgs} : \text{Set}[R\_Message]$ 
  •  $\text{messages\_to\_send}(s, \text{msgs}) =$ 
    (  $\text{PumpMessages}(s, \text{msgs}) \cup$ 
       $\text{FailureDetectionMessages}(s, \text{msgs}) \cup$ 
       $\text{RepairedAcknowledgementMessages}(\text{msgs})$  )
    +  $\text{MODE}(\text{next\_mode}(s, \text{msgs}))$ 
end

```

♡ We rely again here on *auxiliary operations* declared in the *local part*, and their axiomatization is fairly easy using *existential quantifiers*. ♡

## 13.8 The Steam-Boiler Control System Specification

According to our work plan detailed in Sect. 13.3, we have already specified the main parts of our case study. First, let us display a basic (flat) specification equivalent to SBCS\_STATE\_5 and where all the state observers are listed together.

```

spec SBCS_STATE =
  PRELIMINARY
then sort State
  free type Status ::= OK | FailureWithoutAck | FailureWithAck
  free type ExtendedPumpState ::= sort PumpState | Unknown_PCS
  free type
    ExtendedPumpControllerState ::= sort PumpControllerState
    | SoonFlow | Unknown_PCS
  free type Valpair ::= pair(low : Value; high : Value)
  ops mode : State → Mode;
    numSTOP : State → Nat;
    status : State × PhysicalUnit → Status;
    PS_predicted : State × PumpNumber
      → ExtendedPumpState;
    PCS_predicted : State × PumpNumber
      → ExtendedPumpControllerState;
    steam_predicted, level_predicted : State → Valpair
end

```

We are now ready to provide the specification of the steam-boiler control system, considered as a labeled transition system. We leave partly unspecified the initial state *init*, since in our specification this state represents the state immediately following the receipt of the *PHYSICAL\_UNITS\_READY* message. Hence intuitively the omitted axioms should take into account the messages sent and received during the initialization phase (at least at the end of it). It is therefore better to leave open for now the value of most observers on *init*, and to note that this would have to be taken care of when specifying the initialization phase. The value of *mode(init)* is specified according to the end of App. C.4.1.

```

spec STEAM_BOILER_CONTROL_SYSTEM =
  SBCS_ANALYSIS
then op init : State
  pred is_step : State × Set[R_Message] × Set[S_Message] × State

```

```

%% Specification of the initial state init:
•  $mode(init) = Normal \vee mode(init) = Degraded$ 
%% Specification of is_step:
 $\forall s, s' : State; msgs : Set[R\_Message]; Smsg : Set[S\_Message]$ 
•  $is\_step(s, msgs, Smsg, s') \Leftrightarrow$ 
   $mode(s') = next\_mode(s, msgs) \wedge$ 
   $numSTOP(s') = next\_numSTOP(s, msgs) \wedge$ 
   $(\forall pu : PhysicalUnit \bullet$ 
     $status(s', pu) = next\_status(s, msgs, pu)) \wedge$ 
   $(\forall pn : PumpNumber \bullet$ 
     $PS\_predicted(s', pn) = next\_PS\_predicted(s, msgs, pn) \wedge$ 
     $PCS\_predicted(s', pn) = next\_PCS\_predicted(s, msgs, pn) ) \wedge$ 
   $steam\_predicted(s') = next\_steam\_predicted(s, msgs) \wedge$ 
   $level\_predicted(s') = next\_level\_predicted(s, msgs) \wedge$ 
   $Smsg = messages\_to\_send(s, msgs)$ 
then %% Specification of the reachable states:
  free { pred reach : State
     $\forall s, s' : State; msgs : Set[R\_Message]; Smsg : Set[S\_Message]$ 
    •  $reach(init)$ 
    •  $reach(s) \wedge is\_step(s, msgs, Smsg, s') \Rightarrow reach(s') \quad \}$ 
end

```

### 13.9 Validation of the CASL Requirements Specification

Once the formalization of the informal requirements is completed, we must now face the following question: is our formal specification adequate? This is a difficult question to answer since there is no formal way to establish the adequacy of a formal specification w.r.t. informal requirements, i.e., we cannot *prove* this adequacy. However, we can try to *test* it, by performing various ‘experiments’. When these experiments are successful, our confidence in the formal specification is increased. If some experiment fails, then we can inspect the specification and try to understand the causes of the failure, possibly detecting some flaw in the specification.

We will base our validation process on theorem proving, i.e., we will check that some formulas are logical consequences of our requirements specification STEAM\_BOILER\_CONTROL\_SYSTEM. For this purpose we use the tools described in Chap. 11. During this validation process we can consider two kinds of proof obligations:

1. We can inspect the text of the specification and derive from this inspection some formulas that are expected to be logical consequences of our specification. This can be considered as a kind of internal validation of the formal specification.

2. We can check that some expected properties inferred from the informal requirements are logical consequences of our specification (external validation). To do this, we must first reanalyze the informal specification, state some expected properties, translate them into formulas, and then attempt to prove that these formulas are logical consequences of our specification. This task is not easy, since in general one has the feeling that all expected properties were already detected and included in the axioms during the formalization process.

The application of these principles to the requirements specification of the steam-boiler control system leads to various proofs. Below we give only a few illustrative examples.

For instance, let us consider the specification of *next\_mode* in *MODE\_EVOLUTION*: it is advisable to prove that all the cases considered in the specification of *next\_mode* are mutually exclusive, and that their disjunction is equivalent to true. This is a typical example of internal validation of the specification, since we just consider the text of the specification to decide which proof attempt will be performed, without considering the informal requirements again. We do not spell out the corresponding proofs here, but the reader can easily check that indeed the operation *next\_mode* is well-defined (i.e., all cases are mutually exclusive and their disjunction is equivalent to true). In the same spirit we can prove that the same pump cannot simultaneously be ordered to activate and to stop, that we never resignal a failure which has already been signaled, that as long as the operating mode is not set to *EmergencyStop* the water level is safe, etc.

Let us now consider an example of external validation. According to our understanding of failure detection (see Sect. 13.5), if we have detected a failure of some physical unit *pu* (so *PU\_OK* does not hold for *pu*), then the status of this physical unit should not be set to *OK*. The corresponding proof obligation reads as follows:

$$\begin{aligned} \text{STEAM\_BOILER\_CONTROL\_SYSTEM} & \models \\ \forall s : \text{State}; \text{msgs} : \text{Set}[\text{R\_Message}]; \text{pu} : \text{PhysicalUnit} & \\ \bullet \text{Transmission\_OK}(s, \text{msgs}) \wedge \neg \text{PU\_OK}(s, \text{msgs}, \text{pu}) & \\ \wedge \text{reach}(s) \Rightarrow \neg (\text{next\_status}(s, \text{msgs}, \text{pu}) = \text{OK}) & \end{aligned}$$

However here we are unable to discharge this proof obligation. A careful analysis of the proof attempt shows that the proof fails since it could be the case that, simultaneously with the receipt of a repaired message for the physical unit *pu*, we nevertheless detect again a failure of the same unit. From this analysis we conclude that the following axiom in *STATUS\_EVOLUTION* is not adequate:

$$\begin{aligned} \bullet \text{status}(s, \text{pu}) = \text{FailureWithAck} \wedge \text{REPAIRED}(\text{pu}) \text{ is\_in } \text{msgs} & \\ \Rightarrow \text{next\_status}(s, \text{msgs}, \text{pu}) = \text{OK} & \end{aligned}$$

This means we must fix the *STATUS\_EVOLUTION* specification and replace the above axiom by:

- $status(s, pu) = FailureWithAck \wedge REPAIRED(pu) \text{ is\_in } msgs$   
 $\Rightarrow next\_status(s, msgs, pu) = OK \text{ when } PU\_OK(s, msgs, pu)$   
 $\text{else } FailureWithoutAck$

Once the specification STATUS\_EVOLUTION is modified as explained above, we can prove that the expected property holds.

To conclude, the reader should keep in mind that the validation of the specification is a very important task that deserves some serious attention. In this section we have only briefly illustrated some typical proof attempts that would naturally arise when validating the STEAM\_BOILER\_CONTROL\_SYSTEM specification, and obviously many other proof attempts are required to reach a stage where we can trust our requirements specification of the steam-boiler control system.

### 13.10 Designing the Architecture

We now have a validated requirements specification STEAM\_BOILER\_CONTROL\_SYSTEM of the steam-boiler control system. The next step is to refine it into an architectural specification, thereby prescribing the intended architecture of the implementation of the steam-boiler control system. Indeed, the explanations given in Sect. 13.3 suggest the following rather obvious architecture for the steam-boiler control system:

```

arch spec ARCH_SBCS =
units  P : VALUE → PRELIMINARY;
        S : PRELIMINARY → SBCS_STATE;
        A : SBCS_STATE → SBCS_ANALYSIS;
        C : SBCS_ANALYSIS → STEAM_BOILER_CONTROL_SYSTEM
result λ V : VALUE • C [A [S [P [V]]]]
end
```

Note that we decide to describe the implementation of the steam-boiler control system as an *open system*, relying on an external component  $V$  implementing VALUE. This is consistent with our explanations in Sect. 13.2: choosing a specific implementation of VALUE is obviously orthogonal to designing the implementation of the steam-boiler control system. This means in particular that the component  $V$  implementing VALUE will encapsulate the chosen representation of natural numbers and values, together with operations and predicates operating on them.

♡ As illustrated by ARCH\_SBCS, the intended architecture of the steam-boiler control system is easily described by an *architectural specification*. Then we can proceed with four separate implementation tasks, which are independent of each other. ♡

In a next step, we can refine the specification  $\text{VALUE} \rightarrow \text{PRELIMINARY}$  of the component  $P$  into the following architectural specification.

```

arch spec ARCH_PRELIMINARY =
units SET : { sort Elem }  $\times$  NAT  $\rightarrow$  SET [ sort Elem ];
      B   : BASICS;
      MS  : MESSAGES_SENT given B;
      MR  : VALUE  $\rightarrow$  MESSAGES_RECEIVED given B;
      CST : VALUE  $\rightarrow$  SBCS_CONSTANTS
result  $\lambda V : \text{VALUE}$  • SET [MS fit Elem  $\mapsto$  S_Message] [V]
      and SET [MR[V] fit Elem  $\mapsto$  R_Message] [V]
      and CST [V]
end

```

Here we decide to implement (generic) sets in a component  $SET$ , reused both for sets of messages received and sets of messages sent. Since the implementation of natural numbers is provided by the (external) component  $V$ , we use  $V$  for the second argument of the generic component  $SET$  in the result unit term.

♡ Note how the *generic specification with imports* SET is transposed into a *specification of a generic component* SET. Note also, for the component  $MR$ , the use of a *specification of a generic component* extending a *given unit*. ♡

The specification of the components  $C$  and  $S$  of ARCH\_SBCS are simple enough that they do not need to be further architecturally refined. The specification of the component  $S$  (which implements the states of the steam-boiler control system) can be refined into the following specification UNIT\_SBCS\_STATE, which provides a concrete implementation of states as a record of all the observable values.

```

from BASIC/STRUCTURED DATATYPES get TOTALMAP

spec SBCS_STATE_IMPL =
  PRELIMINARY
then free type Status ::= OK | FailureWithoutAck | FailureWithAck
      free type ExtendedPumpState ::= sort PumpState | Unknown_PS
      free type ExtendedPumpControllerState ::=
        sort PumpControllerState | SoonFlow | Unknown_PCS
      free type Valpair ::= pair(low : Value; high : Value)
then TOTALMAP [BASICS fit S  $\mapsto$  PhysicalUnit] [ sort Status ]
and TOTALMAP [BASICS fit S  $\mapsto$  PumpNumber] [ sort ExtendedPumpState ]
and TOTALMAP [BASICS fit S  $\mapsto$  PumpNumber]
      [ sort ExtendedPumpControllerState ]

```



```

then free type State ::= mk_state(
  mode : Mode;
  numSTOP : Nat;
  status : TotalMap[PhysicalUnit, Status];
  PS_predicted :
    TotalMap[PumpNumber, ExtendedPumpState];
  PCS_predicted :
    TotalMap[PumpNumber, ExtendedPumpControllerState];
  steam_predicted, level_predicted : Valpair )
ops status(s : State; pu : PhysicalUnit) : Status
    = lookup(pu, status(s));
  PS_predicted(s : State; pn : PumpNumber) : ExtendedPumpState
    = lookup(pn, PS_predicted(s));
  PCS_predicted(s : State; pn : PumpNumber)
    : ExtendedPumpControllerState
    = lookup(pn, PCS_predicted(s))
end

unit spec UNIT_SBCS_STATE =
  PRELIMINARY → SBCS_STATE_IMPL

```

♡ During the formalization process it was convenient to rely on a *loose specification* of states. At the design stage, this loose specification is refined into a specification where state variables are now explicit. ♡

The specification  $SBCS\_STATE \rightarrow SBCS\_ANALYSIS$  of the component  $A$  of ARCH\_SBCS can be refined into the following architectural specification:

```

arch spec ARCH_ANALYSIS =
units FD   : SBCS_STATE → FAILURE_DETECTION;
      PR   : FAILURE_DETECTION → PU_PREDICTION;
      ME   : PU_PREDICTION → MODE_EVOLUTION [PU_PREDICTION];
      MTS  : MODE_EVOLUTION [PU_PREDICTION] → SBCS_ANALYSIS
result λ S : SBCS_STATE • MTS [ME [PR [FD [S]]]]
end

```

In the above architectural specification ARCH\_ANALYSIS, the component  $FD$  provides an implementation of failure detection, the component  $PR$  an implementation of the predicted state variables for the next cycle, the component  $ME$  provides an implementation of *next\_mode* (and of *next\_numSTOP*), and the component  $MTS$  provides an implementation of *messages\_to\_send*.

The specifications of the components  $ME$  and  $MTS$  are simple enough to be directly implemented. The specifications of the components  $FD$  and  $PR$  can be refined as follows.

```

arch spec ARCH_FAILURE_DETECTION =
units MTSF : SBCS_STATE
      → MESSAGE_TRANSMISSION_SYSTEM_FAILURE;
    PF : SBCS_STATE → PUMP_FAILURE;
    PCF : SBCS_STATE → PUMP_CONTROLLER_FAILURE;
    SF : SBCS_STATE → STEAM_FAILURE;
    LF : SBCS_STATE → LEVEL_FAILURE;
    PU : MESSAGE_TRANSMISSION_SYSTEM_FAILURE
      × PUMP_FAILURE × PUMP_CONTROLLER_FAILURE
      × STEAM_FAILURE × LEVEL_FAILURE
      → FAILURE_DETECTION
result  $\lambda S : \text{SBCS\_STATE} \bullet$ 
      PU [MTSF[S]] [PF[S]] [PCF[S]] [SF[S]] [LF[S]]
      hide Pump_OK, Pump_Controller_OK, Steam_OK, Level_OK
end

```

The above architectural specification ARCH\_FAILURE\_DETECTION refines the specification SBCS\_STATE → FAILURE\_DETECTION of the component *FD* in ARCH\_ANALYSIS and introduces a component for each kind of failure detection. Then the component *PU* implements *PU\_OK*, and in the result unit expression we hide the auxiliary predicates provided by the components *PF*, *PCF*, *SF*, and *LF*.<sup>10</sup>

We refine the specification FAILURE\_DETECTION → PU\_PREDICTION of the component *PR* of the architectural specification ARCH\_ANALYSIS as follows:

```

arch spec ARCH_PREDICTION =
units SE : FAILURE_DETECTION →
      STATUS_EVOLUTION [FAILURE_DETECTION];
    SLP : FAILURE_DETECTION → STEAM_AND_LEVEL_PREDICTION;
    PP : STATUS_EVOLUTION [FAILURE_DETECTION]
      × STEAM_AND_LEVEL_PREDICTION
      → PUMP_STATE_PREDICTION;
    PCP : STATUS_EVOLUTION [FAILURE_DETECTION]
      × STEAM_AND_LEVEL_PREDICTION
      → PUMP_CONTROLLER_STATE_PREDICTION
result  $\lambda FD : \text{FAILURE\_DETECTION} \bullet$ 
      local SEFD = SE [FD]; SLPFD = SLP [FD] within
      PP [SEFD] [SLPFD] and PCP [SEFD] [SLPFD]
end

```

<sup>10</sup> These auxiliary predicates are already hidden in the specification FAILURE\_DETECTION. However, remember that in the specification of a generic component, the target specification is always an implicit extension of the argument specifications. This is why it is necessary to hide the auxiliary predicates at the level of the result unit expression.

In the above architectural specification, the component *SE* provides an implementation of *next\_status*. The component *SLP* provides an implementation of *next\_steam\_predicted*, *next\_level\_predicted*, *chosen\_pumps*, and *Dangerous-WaterLevel*. The component *PP* provides an implementation of *next\_PS\_predicted*, and the component *PCP* provides an implementation of *next\_PCS\_predicted*.

We are now left with specifications of components that are simple enough to be directly implemented, and this concludes our case study.

## Appendices

# A

---

## CASL Quick Reference

This appendix provides an overview of the (concrete) syntax of each part of CASL.

### *Basic specifications*

- declarations, definitions:
  - sorts, subsorts
  - functions: total, partial
  - constants: total, partial
  - predicates
  - datatypes
  - sort generation constraints
- variables, axioms
  - formulas
  - terms
- symbols
- comments
- annotations

### *Architectural specifications*

- named architectures, units
- architectural specifications
- unit specifications
- unit declarations, definitions
- unit expressions, terms

### *Structured specifications*

- specification structure
  - translation
  - hiding, revealing
  - union, extension
  - free extension, initiality
  - hiding local symbols
  - reference
  - instantiation
- named, generic specifications
  - fitting arguments
- named, generic views
  - fitting views
- symbol lists, maps

### *Libraries*

- named libraries
- downloadings
- library names, versions

## A.1 Basic Specifications

$\dots; \dots$	list of items (‘;’ optional)
<b>sorts</b> $\dots$	sort declarations and definitions
<b>ops</b> $\dots$	operation declarations and definitions
<b>preds</b> $\dots$	predicate declarations and definitions
<b>types</b> $\dots$	datatype declarations and definitions
<b>generated</b> $\{ \dots \}$	sort generation constraint
<b>vars</b> $\dots$	global variable declarations
$\forall \dots \bullet F_1 \dots \bullet F_n$	universally-quantified list of axioms
$\bullet F_1 \dots \bullet F_n$	unquantified list of axioms

### A.1.1 Declarations and Definitions

#### Sort Declarations and Definitions

<b>sort</b> $s$	sort declaration
<b>sorts</b> $s_1, \dots, s_n$	sorts declaration
<b>sorts</b> $s < s'$	subsort declaration
<b>sorts</b> $s_1, \dots, s_n < s'$	subsorts and supersort declaration
<b>sorts</b> $s < s_1; \dots; s < s_n$	subsort and supersorts declaration
<b>sorts</b> $s_1 = \dots = s_n$	isomorphic sorts declaration
<b>sort</b> $s = \{v : s' \bullet F\}$	subsort definition

#### Function Declarations and Definitions

<b>op</b> $f : s_1 \times \dots \times s_n \rightarrow s$	total function declaration
<b>op</b> $f : s_1 \times \dots \times s_n \rightarrow ? s$	partial function declaration
<b>op</b> $f : s \times s \rightarrow s, \textit{assoc}$	associative binary function
<b>op</b> $f : s \times s \rightarrow s', \textit{comm}$	commutative binary function
<b>op</b> $f : s \times s \rightarrow s, \textit{idem}$	idempotent binary function
<b>op</b> $f : s \times s \rightarrow s, \textit{unit } T$	unit term for binary function
<b>op</b> $f : s \times s \rightarrow s, \dots, \dots$	multiple function attributes
<b>ops</b> $f_1, \dots, f_n : \dots$	functions declaration
<b>op</b> $f(v_1 : s_1; \dots; v_n : s_n) : s = T$	total function definition
<b>op</b> $f(v_1 : s_1; \dots; v_n : s_n) : ? s = T$	partial function definition
<b>op</b> $f(\dots v_{i_1}, \dots, v_{i_m} : s_i \dots) \dots$	abbreviated arguments
<b>ops</b> $\dots; \dots$	multiple declarations/definitions

#### Constant Declarations and Definitions

<b>op</b> $c : s$	constant declaration
<b>op</b> $c : ? s$	partial constant declaration
<b>ops</b> $c_1, \dots, c_n : s$	constants declaration
<b>op</b> $c : s = T$	constant definition
<b>op</b> $c : ? s = T$	partial constant definition
<b>ops</b> $\dots; \dots$	multiple declarations/definitions

## Predicate Declarations and Definitions

<b>pred</b> $p : s_1 \times \dots \times s_n$	predicate declaration
<b>pred</b> $p : ()$	constant predicate declaration
<b>preds</b> $p_1, \dots, p_n : \dots$	predicates declaration
<b>pred</b> $p(v_1 : s_1; \dots; v_n : s_n) \Leftrightarrow F$	predicate definition
<b>pred</b> $p \Leftrightarrow F$	constant predicate definition
<b>pred</b> $p(\dots v_{i_1}, \dots, v_{i_m} : s_i \dots) \dots$	abbreviated arguments
<b>preds</b> $\dots; \dots$	multiple declarations/definitions

## Datatype Declarations

<b>type</b> $s ::= A$	datatype declaration with alternatives
<b>types</b> $s_1 ::= A_1;$ $\dots;$ $s_n ::= A_n$	multi-sorted datatype declaration
<b>generated types</b> $\dots$	generated datatype declaration
<b>free types</b> $\dots$	free datatype declaration

## Alternatives ( $A$ )

$f(s'_1; \dots; s'_k)$	total constructor function
$f(s'_1; \dots; s'_k)?$	partial constructor function
$f(\dots f_i : s_i \dots)$	total constructor and selector functions
$f(\dots f_i : ?s_i \dots)$	total constructor, partial selector functions
$f(\dots f_{i_1}, \dots, f_{i_m} : s_i \dots)$	abbreviated selectors
$c$	constant constructor value
$sort\ s$	subsort
$sorts\ s'_1, \dots, s'_k$	subsorts
$A_1 \mid \dots \mid A_m$	multiple alternatives

## Sort Generation Constraints

<b>generated</b> { <b>sorts</b> $\dots$	generated sorts
<b>ops</b> $\dots$	generating operations
<b>preds</b> $\dots$	
<b>types</b> $\dots$	generated sorts
<b>}</b>	and generating constructors

### A.1.2 Variables and Axioms

<b>var</b> $v : s$	global variable declaration
<b>vars</b> $v_1 : s_1; \dots; v_n : s_n$	global variables declaration
<b>vars</b> $\dots v_1, \dots, v_n : s_n \dots$	abbreviated variables declaration
<b>vars</b> $\dots; \dots$	multiple global variable declarations
$\forall v : s \bullet F_1 \dots \bullet F_n$	universally-quantified list of axioms
$\forall v_1, \dots, v_n : s \bullet \dots$	abbreviated quantifications
$\forall \dots; \dots \bullet \dots$	multiple quantifications
$\bullet F_1 \dots \bullet F_n$	unquantified list of axioms

### Formulas ( $F$ )

$\forall \dots \bullet F$	universal quantification on formula
$\exists \dots \bullet F$	existential quantification
$\exists! \dots \bullet F$	unique-existential quantification
$F_1 \wedge \dots \wedge F_n$	conjunction
$F_1 \vee \dots \vee F_n$	disjunction
$F \Rightarrow F'$	implication
$F' \text{ if } F$	reverse implication
$F \Leftrightarrow F'$	equivalence
$\neg F$	negation
<i>true</i>	truth
<i>false</i>	falsity
$p(T_1, \dots, T_n)$	predicate application
$t_0 \ T_1 \ t_1 \dots T_n \ t_n$	mixfix predicate application
$q$	constant predicate
$T = T'$	ordinary (strong) equality
$T \stackrel{e}{=} T'$	existential equality
<i>def</i> $T$	definedness
$T \in s$	subsort membership

### Terms ( $T$ )

$f(T_1, \dots, T_n)$	application
$t_0 \ T_1 \ t_1 \dots T_n \ t_n$	mixfix application
$t_0 \ T_1, \dots, T_n \ t_1$	literal syntax
$c$	constant
$v$	variable
$T : s$	sorted term
$T \text{ as } s$	projection to subsort
$T \text{ when } F \text{ else } T'$	conditional choice



### A.1.3 Symbols

Character set: ASCII (with optional use of ISO Latin-1).

#### Key Words and Signs

Reserved key words (always lowercase):

*and arch as axiom axioms closed def else end exists false fit  
forall free from generated get given hide if in lambda library local  
not op ops pred preds result reveal sort sorts spec then to true  
type types unit units var vars version view when with within*

Reserved key signs:

: :? ::= = => <=> ¬ . · | |-> \ /

Unreserved key signs:

< \* × -> ? ! [ ] { }

Key words and signs representing mathematical symbols:

*forall exists exists! not in lambda* =e= -> => <=> . · |-> /\ \/  
 $\forall \quad \exists \quad \exists! \quad \neg \in \quad \lambda \quad \overset{e}{=} \rightarrow \Rightarrow \Leftrightarrow \bullet \bullet \mapsto \wedge \vee$

#### Identifiers

Identifiers for sorts and variables are simple words (other than reserved words) possibly containing digits, primes, and *single* underscores:

*Elem Y\_1 Z2' A\_Rather\_Long\_Identifier*

Sort identifiers can also be compound:

*List[Int] Map[Index, Elem]*

Identifier for operations and predicates can moreover be sequences of (unreserved) signs, with any brackets [ ] { } balanced:

+ - \* / \ & = < > [ ] { } ! ? : . \$ @ # ^ ~ ¡ ¢ × ÷ £ © ± ¶ § <sup>1 2 3</sup> · ¸ ° ¬ μ |

or single decimal digits *1 2 3 4 5 6 7 8 9 0*, or single quoted characters *'c'*.

The signs ( ) ; , ‘ ’ % are *not* allowed in identifiers, nor are the ISO Latin-1 signs for general currency, yen, broken vertical bar, registered trade mark, masculine and feminine ordinals, left and right angle quotes, fractions, soft hyphen, acute accent, cedilla, macron, and umlaut.

Operation and predicate identifiers can also be compound:

*order[-- < --]*

Function and predicate identifiers can also be infixes, prefixes, postfixes, and general mixfixes, formed from words and/or sequences of signs separated by *double* underscores (indicating the positions of the arguments), with any brackets [ ] { } balanced:

```
--++--  ||--||  {[--]}  push_onto__  ___select1
```

Invisible mixfix identifiers (such as `----`) with two or more arguments are allowed. (Subsort embeddings give the effect of invisible unary functions.)

An operation, or predicate identifier can be compound, with a list of identifiers appended to its final token.

## Literal Strings and Numbers

```
"this is a string"  42  3.14159  1E-9  27.3e6
```

## Library Identifiers

Names of libraries are either paths, e.g.:

```
BASIC/NUMBERS  BASIC/ALGEBRA_II
```

or URLs formed from A...ZA...Z0...9\$\_@.&+!\*"'(),~ and hexadecimal codes %XX, and prefixed by HTTP://, FTP://, or FILE:///.

Version numbers of libraries are hierarchical: 0, 0.999, 1, 1.0, 1.0.2.

### A.1.4 Comments

```
%% This is a comment at the end of a line...
```

```
...%{ This is an in-line comment }% ...
```

```
...%{ This a comment that might take
      several lines }%
```

```
%[ This is for commenting-out text
```

```
%{ including other kinds of comment }% ]%
```

### A.1.5 Annotations

A label is of the form `%(text)%`.

An end-of-line annotation is of the general form `%word ...` with a *space* following the word.

A possibly multi-line annotation is of the general form `%word( ... )%` with *no space* preceding the ‘(’.

## A.2 Structured Specifications

### A.2.1 Specifications ( $SP$ )

$SP$ with $SM$	symbol translation
$SP$ hide $SL$	hiding listed symbols
$SP$ reveal $SM$	revealing/translating listed symbols
$SP_1$ and ... and $SP_n$	union
$SP_1$ then ... then $SP_n$	extension
free $SP$	free or initial
local $SP$ within $SP'$	hiding of local symbols
closed $SP$	self-contained
$SN$	reference to named specification
$SN[FA_1] \dots [FA_n]$	instantiation of generic specification

### A.2.2 Named and Generic Specifications

spec $SN = SP$ end	named specification ( <b>end</b> optional)
spec $SN[SP_1] \dots [SP_n] = SP$ end	generic specification ( <b>end</b> optional)
spec $SN[SP_1] \dots [SP_n]$	generic specification
given $SP''_1, \dots, SP''_m = SP$ end	with imports ( <b>end</b> optional)

### Fitting Arguments ( $FA$ )

$SP$ fit $SM$	fitting by symbol map
$SP$	implicit fitting
$FV$	fitting view

### A.2.3 Named and Generic Views

view $VN: SP$ to $SP' = SM$ end	named view ( <b>end</b> optional)
view $VN[SP_1] \dots [SP_n]$	generic view
: $SP$ to $SP' = SM$ end	( <b>end</b> optional)
view $VN[SP_1] \dots [SP_n]$	generic view
given $SP''_1, \dots, SP''_m$	with imports
: $SP$ to $SP' = SM$ end	( <b>end</b> optional)

### Fitting Views ( $FV$ )

view $VN$	reference to named view
view $VN[FA_1] \dots [FA_n]$	instantiation of generic view

### A.2.4 Symbol Lists ( $SL$ ) and Maps ( $SM$ )

$SY_1, \dots, SY_n$	lists (maybe with <b>sorts</b> , <b>ops</b> , <b>preds</b> )
$SY_1 \mapsto SY'_1, \dots, SY_n \mapsto SY'_n$	maps (maybe with <b>sorts</b> , <b>ops</b> , <b>preds</b> )
$\dots, SY_i, \dots$	in a map, abbreviates $\dots, SY_i \mapsto SY_i, \dots$

## A.3 Architectural Specifications

### A.3.1 Named Architectures and Units

**arch spec**  $ASN = ASP$  **end**                      named arch. spec. (**end** optional)  
**unit spec**  $SN = USP$  **end**                      named unit spec. (**end** optional)

### A.3.2 Architectural Specifications ( $ASP$ )

$ASN$     arch. spec. name  
**units**  $UD_1; \dots; UD_n$  **result**  $UE$                       basic arch. spec.

### A.3.3 Unit Specifications ( $USP$ )

$SP$     unit specification  
 $SP_1 \times \dots \times SP_n \rightarrow SP$                       generic-unit specification  
**closed**  $USP$                                       self-contained  
**arch spec**  $ASP$                                       models of arch. spec.

### A.3.4 Unit Declarations and Definitions ( $UD$ )

$UN : USP$                                       unit declaration  
 $UN : USP$  **given**  $UT_1, \dots, UT_n$                       importing units  
 $UN = UE$                                       unit definition

### A.3.5 Unit Expressions ( $UE$ )

$UT$     unit term  
 $\lambda UN_1 : SP_1; \dots; UN_n : SP_n \bullet UT$                       unit composition

### A.3.6 Unit Terms ( $UT$ )

$UT$  **with**  $SM$                                       symbol translation  
 $UT$  **hide**  $SL$                                       hiding listed symbols  
 $UT$  **reveal**  $SM$                                       revealing/translating listed symbols  
 $UT_1$  **and**  $\dots$  **and**  $UT_n$                       amalgamation  
**local**  $UD_1; \dots; UD_n$  **within**  $UT$                       local units  
 $UN$     unit name  
 $UN[UT_1] \dots [UT_n]$                       generic-unit application  
 $UN[UT_1 \text{ fit } SM_1] \dots [UT_n \text{ fit } SM_n]$                       with fitting by symbol maps

## A.4 Libraries

**library**  $LN$  ...                      named library of downloadings,  
specifications, views

### A.4.1 Downloadings

**from**  $LN$  **get**  $IN_1, \dots, IN_n$  **end**      downloads listed items  
**from**  $LN$  **get** ...  $IN \mapsto IN'$  ... **end**      renames downloaded items

### A.4.2 Library Names ( $LN$ )

BASIC/NUMBERS                      greatest version registered  
BASIC/ALGEBRA\_II **version**  $0.999$       specified version registered  
HTTP://...                      greatest version unregistered  
HTTP://... **version**  $1.0.2$               specified version unregistered

# B

---

## Points to Bear in Mind

### B.1 Introduction

- CoFI aims at establishing a wide consensus. . . . . 4
- The focus of CoFI is on *algebraic* techniques. . . . . 5
- CoFI has already achieved its main aims. . . . . 5
- CoFI is an open, voluntary initiative. . . . . 6
- CoFI has received funding as an ESPRIT Working Group, and is sponsored by IFIP WG 1.3. . . . . 6
- New participants are welcome! . . . . . 7
- CASL has been designed as a general-purpose algebraic specification language, subsuming many existing languages. . . . . 7
- CASL is at the center of a *family* of languages. . . . . 8
- CASL itself has several major *parts*. . . . . 9

### B.2 Underlying Concepts

- CASL is based on standard concepts of algebraic specification. . . . . 11
- A basic specification declares symbols, and gives axioms and constraints. . . . . 11
- The semantics of a basic specification is a signature and a class of models. . . . . 12
- CASL specifications may declare sorts, subsorts, operations, and predicates. . . . . 12
- Sorts are interpreted as carrier sets. . . . . 12
- Subsorts declarations are interpreted as embeddings. . . . . 13
- Operations may be declared as total or partial. . . . . 13
- Predicates are different from boolean-valued operations. . . . . 13
- Operation symbols and predicate symbols may be overloaded. . . . . 14
- Axioms are formulas of first-order logic. . . . . 14

• Sort generation constraints eliminate ‘junk’ from specific carrier sets.	15
• The semantics of a structured specification is simply a signature and a class of models. ....	16
• A translation merely renames symbols. ....	17
• Hiding symbols removes parts of models. ....	17
• Union of specifications identifies common symbols. ....	17
• Extension of specifications identifies common symbols too. ....	18
• Free specifications restrict models to being free, with initiality as a special case. ....	18
• Generic specifications have parameters, and have to be instantiated when referenced. ....	18
• The semantics of an architectural specification reflects its modular structure. ....	19
• Architectural specifications involve the notions of persistent function and conservative extension. ....	19
• The semantics of a library of specifications is a mapping from the names of the specifications to their semantics. ....	20

## B.3 Getting Started

• Simple specifications may be written in CASL essentially as in many other algebraic specification languages. ....	23
• CASL provides also useful abbreviations. ....	23
• CASL allows loose, generated and free specifications. ....	24
• CASL syntax for declarations and axioms involves familiar notation, and is mostly self-explanatory. ....	24
• Specifications can easily be extended by new declarations and axioms.	25
• In simple cases, an operation (or a predicate) symbol may be declared and its intended interpretation defined at the same time. ...	26
• Symbols may be conveniently displayed as usual mathematical symbols by means of <b>%display</b> annotations. ....	27
• The <b>%implies</b> annotation is used to indicate that some axioms are supposedly redundant, being consequences of others. ....	28
• Attributes may be used to abbreviate axioms for associativity, commutativity, idempotence, and unit properties. ....	29
• Genericity of specifications can be made explicit using parameters. ...	29
• References to generic specifications always instantiate the parameters.	30
• Datatype declarations may be used to abbreviate declarations of sorts and constructors. ....	32
• Loose datatype declarations are appropriate when further constructors may be added in extensions. ....	32
• Sorts may be specified as generated by their constructors. ....	33
• Generated specifications are in general loose. ....	34
• Generated specifications need not be loose. ....	35

- Generated types may need to be declared together. . . . . 36
- Free specifications provide initial semantics and avoid the need for explicit negation. . . . . 36
- Free datatype declarations are particularly convenient for defining enumerated datatypes. . . . . 37
- Free specifications can also be used when the constructors are related by some axioms. . . . . 37
- Predicates hold minimally in models of free specifications. . . . . 38
- Operations and predicates may be safely defined by induction on the constructors of a free datatype declaration. . . . . 38
- More care may be needed when defining operations or predicates on free datatypes when there are axioms relating the constructors. . . 39
- Generic specifications often involve free extensions of (loose) parameters. . . . . 40
- Loose extensions of free specifications can avoid overspecification. . . 41
- Datatypes with observer operations or predicates can be specified as generated instead of free. . . . . 42
- The `%def` annotation is useful to indicate that some operations or predicates are uniquely defined. . . . . 43
- Operations can be defined by axioms involving observer operations, instead of inductively on constructors. . . . . 44
- Sorts declared in free specifications are not necessarily generated by their constructors. . . . . 44

## B.4 Partial Functions

- Partial functions arise naturally. . . . . 47
- Partial functions are declared differently from total functions. . . . . 47
- Terms containing partial functions may be undefined, i.e., they may fail to denote any value. . . . . 48
- Functions, even total ones, propagate undefinedness. . . . . 48
- Predicates do not hold on undefined arguments. . . . . 48
- Equations hold when both terms are undefined. . . . . 48
- Special care is needed in specifications involving partial functions. . . 49
- The definedness of a term can be checked or asserted. . . . . 50
- The domains of definition of partial functions can be specified exactly. 50
- Loosely specified domains of definition may be useful. . . . . 51
- Domains of definition can be specified more or less explicitly. . . . . 51
- Partial functions are minimally defined by default in free specifications. . . . . 53
- Selectors can be specified concisely in datatype declarations, and are usually partial. . . . . 54
- Selectors are usually total when there is only one constructor. . . . . 54
- Constructors may be partial. . . . . 54



- Existential equality requires the definedness of both terms as well as their equality. . . . . 55

**B.5 Subsorting**

- Subsorts and supersorts are often useful in CASL specifications. . . . . 57
- Subsort declarations directly express relationships between carrier sets. . . . . 57
- Operations declared on a sort are automatically inherited by its subsorts. . . . . 58
- Inheritance applies also for subsorts that are declared afterwards. . . . 59
- Subsort membership can be checked or asserted. . . . . 59
- Datatype declarations can involve subsort declarations. . . . . 59
- Subsorts may also arise as classifications of previously specified values, and their values can be explicitly defined. . . . . 60
- It may be useful to redeclare previously defined operations, using the new subsorts introduced. . . . . 61
- A subsort may correspond to the definition domain of a partial function. . . . . 62
- Using subsorts may avoid the need for partial functions. . . . . 62
- Casting a term from a supersort to a subsort is explicit and the value of the cast may be undefined. . . . . 63
- Supersorts may be useful when generalizing previously specified sorts. 64
- Supersorts may also be used for extending the intended values by new values representing errors or exceptions. . . . . 65

**B.6 Structuring Specifications**

- Large and complex specifications are easily built out of simpler ones by means of (a small number of) specification-building operations. . . 67
- Union and extension can be used to structure specifications. . . . . 67
- Specifications may combine parts with loose, generated, and free interpretations. . . . . 68
- Renaming may be used to avoid unintended name clashes, or to adjust names of sorts and change notations for operations and predicates. . . . . 69
- When combining specifications, origins of symbols can be indicated. . 71
- Auxiliary symbols used in structured specifications can be hidden. . . 71
- Auxiliary symbols can be made local when they do not need to be exported. . . . . 73
- Care is needed with local sort declarations. . . . . 74
- Naming a specification allows its reuse. . . . . 75

## B.7 Generic Specifications

- Making a specification generic (when appropriate) improves its reusability. . . . . 77
- Parameters are arbitrary specifications. . . . . 78
- The argument specification of an instantiation must provide symbols corresponding to those required by the parameter. . . . . 78
- The argument specification of an instantiation must ensure that the properties required by the parameter hold. . . . . 79
- There must be no shared symbols between the argument specification and the body of the instantiated generic specification. . . 80
- In instantiations, the fitting of parameter symbols to identical argument symbols can be left implicit. . . . . 80
- The fitting of parameter sorts to unique argument sorts can also be left implicit. . . . . 80
- Fitting of operation and predicate symbols can sometimes be left implicit too, and can imply fitting of sorts. . . . . 81
- The intended fitting of the parameter symbols to the argument symbols may have to be specified explicitly. . . . . 81
- A generic specification may have more than one parameter. . . . . 82
- Instantiation of generic specifications with several parameters is similar to the case of just one parameter. . . . . 82
- Composition of generic specifications is expressed using instantiation. 84
- Compound sorts introduced by a generic specification get automatically renamed on instantiation, which avoids name clashes. . 85
- Compound symbols can also be used for operations and predicates. . . 87
- Parameters should be distinguished from references to fixed specifications that are not intended to be instantiated. . . . . 88
- Argument specifications are always implicitly regarded as extension of the imports. . . . . 89
- Imports are also useful to prevent ill-formed instantiations. . . . . 89
- In generic specifications, auxiliary required specifications should be imported rather than extended. . . . . 90
- Views are named fitting maps, and can be defined along with specifications. . . . . 90
- Views can also be generic. . . . . 91

## B.8 Specifying the Architecture of Implementations

- Architectural specifications impose structure on implementations, whereas specification-building operations only structure the text of specifications. . . . . 93

- An architectural specification consists of a list of unit declarations, specifying the required components, and a result part, indicating how they are to be combined. . . . . 96
- There can be several distinct architectural choices for the same requirements specification. . . . . 97
- Each unit declaration listed in an architectural specification corresponds to a separate implementation task. . . . . 97
- A unit can be implemented only if its specification is a conservative extension of the specifications of its given units. . . . . 98
- Genericity of components can be made explicit in architectural specifications. . . . . 100
- A generic component may be applied to an argument richer than required by its specification. . . . . 101
- Specifications of components can be named for further reuse. . . . . 102
- Both named and unnamed specifications can be used to specify components. . . . . 102
- Specifications of generic components should not be confused with generic specifications. . . . . 103
- A generic component may be applied more than once in the same architectural specification. . . . . 103
- Several applications of the same generic component is different from applications of several generic components with similar specifications. 104
- Generic components may have more than one argument. . . . . 105
- Open systems can be described by architectural specifications using generic unit expressions in the result part. . . . . 106
- When components are to be combined, it is best to check that any shared symbol originates from the same non-generic component. . . . 107
- Auxiliary unit definitions or local unit definitions may be used to avoid repetition of generic unit applications. . . . . 109

## B.9 Libraries

- Libraries are named collections of named specifications. . . . . 111
- Local libraries are self-contained. . . . . 111
- Distributed libraries support reuse. . . . . 111
- Different versions of the same library are distinguished by hierarchical version numbers. . . . . 112
- Local libraries are self-contained collections of specifications. . . . . 112
- Specifications can refer to previous items in the same library. . . . . 113
- All kinds of named specifications can be included in libraries. . . . . 114
- Display, parsing, and literal syntax annotations apply to entire libraries. . . . . 114
- Libraries and library items can have author and date annotations. . . 116
- Libraries can be installed on the Internet for remote access. . . . . 116

- Validated libraries can be registered for public access. . . . . 117
- Libraries should include appropriate annotations. . . . . 118
- Libraries can include items downloaded from other libraries. . . . . 118
- Substantial libraries of basic datatypes are already available. . . . . 119
- Libraries need not be registered for public access. . . . . 120
- Subsequent versions of a library are distinguished by explicit  
version numbers. . . . . 120
- Libraries can refer to specific versions of other libraries. . . . . 121
- All downloadings should be collected at the beginning of a library. . . 122

## B.10 Foundations

- A complete presentation of CASL is in the *Reference Manual*. . . . . 125
- CASL has a definitive summary. . . . . 125
- CASL has a complete formal definition. . . . . 126
- Abstract and concrete syntax of CASL are defined formally. . . . . 126
- CASL has a complete formal semantics. . . . . 126
- CASL specifications denote classes of models. . . . . 127
- The semantics is largely institution-independent. . . . . 127
- The semantics is the ultimate reference for the meanings of all CASL  
constructs. . . . . 128
- Proof systems for various layers of CASL are provided. . . . . 128
- The foundations of our CASL are rock-solid! . . . . . 129

## B.11 Tools

- CASL specifications can be checked for well-formedness using a  
form-based web page. . . . . 131
- The Heterogeneous Tool Set (HETS) is the main analysis tool for  
CASL. . . . . 132
- HETS can be used for parsing and checking static well-formedness  
of specifications. . . . . 133
- HETS also displays and manages proof obligations, using  
development graphs. . . . . 134
- Nodes in a development graph correspond to CASL specifications.  
Arrows show how specifications are related by the structuring  
constructs. . . . . 135
- Internal nodes in a development graph correspond to unnamed  
parts of a structured specification. . . . . 137
- HOL-CASL is an interactive theorem prover for CASL, based on the  
tactical theorem prover ISABELLE. . . . . 138
- CASL is linked to ISABELLE/HOL by an encoding. . . . . 138
- ASF+SDF was used to prototype the CASL syntax. . . . . 139

- The ASF+SDF Meta-Environment provides syntax-directed editing of CASL specifications. . . . . 140

## B.12 Basic Libraries

- The CASL Basic Libraries contain the standard datatypes. . . . . 143
- HETS can be used to get an overview of the Basic Libraries. . . . . 143

## The Steam-Boiler Control Specification Problem

For completeness, the text describing the steam-boiler control system case study, as originally provided by Jean-Raymond Abrial, is reproduced here (except for the “Additional Information” section, see [1, pp. 507–509]).

### C.1 Introduction

This text constitutes an informal specification of a program which serves to control the level of water in a steam-boiler. It is important that the program works correctly because the quantity of water present when the steam-boiler is working has to be neither too low nor too high; otherwise the steam-boiler or the turbine sitting in front of it might be seriously affected.

The proposed specification is derived from an original text that has been written by LtCol. J.C. Bauer for the Institute for Risk Research of the University of Waterloo, Ontario, Canada. The original text has been submitted as a competition problem to be solved by the participants to the International Software Safety Symposium organized by the Institute for Risk Research. It has been given to us by the Institut de Protection et de Sûreté Nucléaire, Fontenay-aux-Roses, France. We would like to thank the author, the Institute for Risk Research, and the Institut de Protection et de Sûreté Nucléaire for their kind permission to use their text.

The text to follow is severely biased to a particular implementation. This is very often the case with industrial specifications that are rarely independent from a certain implementation people have in mind. In that sense, this specification is realistic. Your first formal specification steps could be *much more abstract* if that seems important to you (in particular if your formalism allows you to do so). In other words, you are encouraged to *structure* your specification in a way that is not necessarily the same as the one proposed in what follows. But in any case, you are asked to demonstrate that your specification can be refined to an implementation that is close enough to the functional requirements of the “specification” proposed below.

You might also judge that the specification contain some loose ends and inconsistencies. Do not hesitate to point them out and to take yourself some appropriate decisions. The idea, however, is that such inconsistencies should be solely within the *organization* of the system and *not within its physical properties*.

We are aware of the fact that the text to follow does not propose any precise model of the physical evolution of the system, only elementary suggestions. As a consequence, you may have to take some simple, even simplistic, abstract decisions concerning such a physical model.

## C.2 Physical Environment

The system comprises the following units:

- the steam-boiler,
- a device to measure the quantity of water in the steam-boiler,
- four pumps to provide the steam-boiler with water,
- four devices to supervise the pumps (one controller for each pump),
- a device to measure the quantity of steam which comes out of the steam-boiler,
- an operator desk,
- a message transmission system.

### C.2.1 The Steam-Boiler

The steam-boiler is characterized by the following elements:

- A valve for evacuation of water. It serves only to empty the steam-boiler in its initial phase.
- Its total capacity  $C$  (indicated in liters).
- The minimal limit quantity  $M_1$  of water (in liters). Below  $M_1$  the steam-boiler would be in danger after five seconds, if the steam continued to come out at its maximum quantity without supply of water from the pumps.
- The maximal limit quantity  $M_2$  of waters (in liters). Above  $M_2$  the steam-boiler would be in danger after five seconds, if the pumps continued to supply the steam-boiler with water without possibility to evacuate the steam.
- The minimal normal quantity  $N_1$  of water in liters to be maintained in the steam-boiler during regular operation ( $M_1 < N_1$ ).
- The maximal normal quantity  $N_2$  of water (in liters) to be maintained in the steam-boiler during regular operation ( $N_2 < M_2$ ).
- The maximum quantity  $W$  of steam (in liters/sec) at the exit of the steam-boiler.
- The maximum gradient  $U_1$  of increase of the quantity of steam (in liters/sec/sec).

- The maximum gradient  $U_2$  of decrease of the quantity of steam (in liters/sec/sec).

### C.2.2 The Water Level Measurement Device

The device to measure the level of water in the steam-boiler provides the following information:

- the quantity  $q$  (in liters) of water in the steam-boiler.

### C.2.3 The Pumps

Each pump is characterized by the following elements:

- Its capacity  $P$  (in liters/sec).
- Its functioning mode: on or off.
- It's being started: after having been switched on, the pump needs five seconds to start pouring water into the boiler (this is due to the fact that the pump does not balance instantaneously the pressure of the steam-boiler).
- It's being stopped: with instantaneous effect.

### C.2.4 The Pump Control Devices

Each pump controller provides the following information:

- the water circulates from the pump to the steam-boiler or, on the contrary, it does not circulate.

### C.2.5 The Steam Measurement Device

The device to measure the quantity of steam which comes out of the steam-boiler provides the following information:

- a quantity of steam  $v$  (in liters/sec).

### C.2.6 Summary of Constants and Variables

The following table summarizes the various constants or physical variables of the system:



	Unit	Comment
<b>Quantity of water in the steam-boiler</b>		
$C$	liter	Maximal capacity
$M_1$	liter	Minimal limit
$M_2$	liter	Maximal limit
$N_1$	liter	Minimal normal
$N_2$	liter	Maximal normal
<b>Outcome of steam at the exit of the steam-boiler</b>		
$W$	liter/sec	Maximal quantity
$U_1$	liter/sec/sec	Maximum gradient of increase
$U_2$	liter/sec/sec	Maximum gradient of decrease
<b>Capacity of each pump</b>		
$P$	liter/sec	Nominal capacity
<b>Current measures</b>		
$q$	liter	Quantity of water in the steam-boiler
$p$	liter/sec	Throughput of the pumps
$v$	liter/sec	Quantity of steam exiting the steam-boiler

### C.3 The Overall Operation of the Program

The program communicates with the physical units through messages which are transmitted over a number of dedicated lines connecting each physical unit with the control unit. In first approximation, the time for transmission can be neglected.

The program follows a cycle and a priori does not terminate. This cycle takes place each five seconds and consists of the following actions:

- Reception of messages coming from the physical units.
- Analysis of informations which have been received.
- Transmission of messages to the physical units.

To simplify matters, and in first approximation, all messages coming from (or going to) the physical units are supposed to be received (emitted) *simultaneously* by the program at each cycle.

### C.4 Operation Modes of the Program

The program operates in different modes, namely *initialization*, *normal*, *degraded*, *rescue*, *emergency stop*.

### C.4.1 Initialization Mode

The *initialization* mode is the mode to start with. The program enters a state in which it waits for the message STEAM-BOILER\_WAITING to come from the physical units. As soon as this message has been received the program checks whether the quantity of steam coming out of the steam-boiler is really zero. If the unit for detection of the level of steam is defective (that is, when  $v$  is not equal to zero), the program enters the *emergency stop* mode. If the quantity of water in the steam-boiler is above  $N_2$  the program activates the valve of the steam-boiler in order to empty it. If the quantity of water in the steam-boiler is below  $N_1$  then the program activates a pump to fill the steam-boiler. If the program realizes a failure of the water level detection unit it enters the *emergency stop* mode. As soon as a level of water between  $N_1$  and  $N_2$  has been reached the program send continuously the signal PROGRAM\_READY to the physical units until it receives the signal PHYSICAL\_UNITS\_READY which must necessarily be emitted by the physical units. As soon as this signal has been received, the program enters either the mode *normal* if all the physical units operate correctly or the mode *degraded* if any physical unit is defective. A transmission failure puts the program into the mode *emergency stop*.

### C.4.2 Normal Mode

The normal mode is the standard operating mode in which the program tries to maintain the water level in the steam-boiler between  $N_1$  and  $N_2$  with all physical units operating correctly. As soon as the water level is below  $N_1$  or above  $N_2$  the level can be adjusted by the program by switching the pumps on or off. The corresponding decision is taken on the basis of the information which has been received from the physical units. As soon as the program recognizes a failure of the water level measuring unit it goes into *rescue* mode. Failure of any other physical unit puts the program into *degraded* mode. If the water level is risking to reach one of the limit values  $M_1$  or  $M_2$  the program enters the mode *emergency stop*. This risk is evaluated on the basis of a maximal behavior of the physical units. A transmission failure puts the program into *emergency stop* mode.

### C.4.3 Degraded Mode

The *degraded* mode is the mode in which the program tries to maintain a satisfactory water level despite of the presence of failure of some physical unit. It is assumed however that the water level measuring unit in the steam-boiler is working correctly. The functionality is the same as in the preceding case. Once all the units which were defective have been repaired, the program comes back to *normal* mode. As soon as the program sees that the water level measuring unit has a failure, the program goes into mode *rescue*. If the

water level is risking to reach one of the limit values  $M_1$  or  $M_2$  the program enters the mode *emergency stop*. A transmission failure puts the program into *emergency stop* mode.

#### C.4.4 *Rescue Mode*

The *rescue* mode is the mode in which the program tries to maintain a satisfactory water level despite of the failure of the water level measuring unit. The water level is then estimated by a computation which is done taking into account the maximum dynamics of the quantity of steam coming out of the steam-boiler. For the sake of simplicity, this calculation can suppose that exactly  $n$  liters of water, supplied by the pumps, do account for exactly the same amount of boiler contents (no thermal expansion). This calculation can however be done only if the unit which measures the quantity of steam is itself working and if one can rely upon the information which comes from the units for controlling the pumps. As soon as the water measuring unit is repaired, the program returns into mode *degraded* or into mode *normal*. The program goes into *emergency stop* mode if it realizes that one of the following cases holds: the unit which measures the outcome of steam has a failure, or the units which control the pumps have a failure, or the water level risks to reach one of the two limit values. A transmission failure puts the program into *emergency stop* mode.

#### C.4.5 *Emergency Stop Mode*

The *emergency stop* mode is the mode into which the program has to go, as we have seen already, when either the vital units have a failure or when the water level risks to reach one of its two limit values. This mode can also be reached after detection of an erroneous transmission between the program and the physical units. This mode can also be set directly from outside. Once the program has reached the *Emergency stop* mode, the physical environment is then responsible to take appropriate actions, and the program stops.

### C.5 Messages Sent by the Program

The following messages can be sent by the program:

- $\text{MODE}(m)$ : The program sends, at each cycle, its current mode of operation to the physical units.
- $\text{PROGRAM\_READY}$ : In *initialization* mode, as soon as the program assumes to be ready, this message is continuously sent until the message  $\text{PHYSICAL\_UNITS\_READY}$  coming from the physical units has been received.

- VALVE: In *initialization* mode this message is sent to the physical units to request opening and then closure of the valve for evacuation of water from the steam-boiler.
- OPEN\_PUMP( $n$ ): This message is sent to the physical units to activate a pump.
- CLOSE\_PUMP( $n$ ): This message is sent to the physical units to stop a pump.
- PUMP\_FAILURE\_DETECTION( $n$ ): This message is sent (until receipt of the corresponding acknowledgement) to indicate to the physical units that the program has detected a pump failure.
- PUMP\_CONTROL\_FAILURE\_DETECTION( $n$ ): This message is sent (until receipt of the corresponding acknowledgement) to indicate to the physical units that the program has detected a failure of the physical unit which controls a pump.
- LEVEL\_FAILURE\_DETECTION: This message is sent (until receipt of the corresponding acknowledgement) to indicate to the physical units that the program has detected a failure of the water level measuring unit.
- STEAM\_FAILURE\_DETECTION: This message is sent (until receipt of the corresponding acknowledgement) to indicate to the physical units that the program has detected a failure of the physical unit which measures the outcome of steam.
- PUMP\_REPAIRED\_ACKNOWLEDGEMENT( $n$ ): This message is sent by the program to acknowledge a message coming from the physical units and indicating that the corresponding pump has been repaired.
- PUMP\_CONTROL\_REPAIRED\_ACKNOWLEDGEMENT( $n$ ): This message is sent by the program to acknowledge a message coming from the physical units and indicating that the corresponding physical control unit has been repaired.
- LEVEL\_REPAIRED\_ACKNOWLEDGEMENT: This message is sent by the program to acknowledge a message coming from the physical units and indicating that the water level measuring unit has been repaired.
- STEAM\_REPAIRED\_ACKNOWLEDGEMENT: This message is sent by the program to acknowledge a message coming from the physical units and indicating that the unit which measures the outcome of steam has been repaired.

## C.6 Messages Received by the Program

The following messages can be received by the program:

- STOP: When the message has been received three times in a row by the program, the program must go into *emergency stop*.
- STEAM\_BOILER\_WAITING: When this message is received in *initialization* mode it triggers the effective start of the program.

- **PHYSICAL\_UNITS\_READY**: This message when received in *initialization* mode acknowledges the message **PROGRAM\_READY** which has been sent previously by the program.
- **PUMP\_STATE( $n, b$ )**: This message indicates the state of pump  $n$  (open or closed). This message must be present during each transmission.
- **PUMP\_CONTROL\_STATE( $n, b$ )**: This message gives the information which comes from the control unit of pump  $n$  (there is flow of water or there is no flow of water). This message must be present during each transmission.
- **LEVEL( $v$ )**: This message contains the information which comes from the water level measuring unit. This message must be present during each transmission.
- **STEAM( $v$ )**: This message contains the information which comes from the unit which measures the outcome of steam. This message must be present during each transmission.
- **PUMP\_REPAIRED( $n$ )**: This message indicates that the corresponding pump has been repaired. It is sent by the physical units until a corresponding acknowledgement message has been sent by the program and received by the physical units.
- **PUMP\_CONTROL\_REPAIRED( $n$ )**: This message indicates that the corresponding control unit has been repaired. It is sent by the physical units until a corresponding acknowledgement message has been sent by the program and received by the physical units.
- **LEVEL\_REPAIRED**: This message indicates that the water level measuring unit has been repaired. It is sent by the physical units until a corresponding acknowledgement message has been sent by the program and received by the physical units.
- **STEAM\_REPAIRED**: This message indicates that the unit which measures the outcome of steam has been repaired. It is sent by the physical units until a corresponding acknowledgement message has been sent by the program and received by the physical units.
- **PUMP\_FAILURE\_ACKNOWLEDGEMENT( $n$ )**: By this message the physical units acknowledge the receipt of the corresponding failure detection message which has been emitted previously by the program.
- **PUMP\_CONTROL\_FAILURE\_ACKNOWLEDGEMENT( $n$ )**: By this message the physical units acknowledge the receipt of the corresponding failure detection message which has been emitted previously by the program.
- **LEVEL\_FAILURE\_ACKNOWLEDGEMENT**: By this message the physical units acknowledge the receipt of the corresponding failure detection message which has been emitted previously by the program.
- **STEAM\_FAILURE\_ACKNOWLEDGEMENT**: By this message the physical units acknowledge the receipt of the corresponding failure detection message which has been emitted previously by the program.

## C.7 Detection of Equipment Failures

The following erroneous kinds of behavior are distinguished to decide whether certain physical units have a failure:

- **PUMP:** (1) Assume that the program has sent a start or stop message to a pump. The program detects that during the following transmission that pump does not indicate its having effectively been started or stopped. (2) The program detects that the pump changes its state spontaneously.
- **PUMP\_CONTROLLER:** (1) Assume that the program has sent a start or stop message to a pump. The program detects that during the second transmission after the start or stop message the pump does not indicate that the water is flowing or is not flowing; this despite of the fact that the program knows from elsewhere that the pump is working correctly. (2) The program detects that the unit changes its state spontaneously.
- **WATER\_LEVEL\_MEASURING\_UNIT:** (1) The program detects that the unit indicates a value which is out of the valid static limits (that is, between  $\theta$  and  $C$ ). (2) The program detects that the unit indicates a value which is incompatible with the dynamics of the system.
- **STEAM\_LEVEL\_MEASURING\_UNIT:** (1) The program detects that the unit indicates a value which is out of the valid static limits (that is, between  $\theta$  and  $W$ ). (2) The program detects that the unit indicates a value which is incompatible with the dynamics of the system.
- **TRANSMISSION:** (1) The program receives a message whose presence is aberrant. (2) The program does not receive a message whose presence is indispensable.

---

## References

1. J.-R. Abrial, E. Börger, and H. Langmaack, editors. *Formal Methods for Industrial Applications: Specifying and Programming the Steam Boiler Control*, LNCS Vol. 1165. Springer, 1996.
2. E. Astesiano, M. Bidoit, B. Krieg-Brückner, P. D. Mosses, D. Sannella, and A. Tarlecki. CASL: The Common Algebraic Specification Language. *Theoretical Comput. Sci.*, 286(2):153–196, 2002.
3. E. Astesiano, H.-J. Kreowski, and B. Krieg-Brückner, editors. *Algebraic Foundations of Systems Specification*. IFIP State-of-the-Art Reports. Springer, 1999.
4. S. Autexier, D. Hutter, T. Mossakowski, and A. Schairer. The development graph manager MAYA (system description). In H. Kirchner and C. Ringeissen, editors, *Algebraic Methods and Software Technology, 9th International Conference, AMAST 2002, Saint-Gilles-les-Bains, Reunion Island, France, Proceedings*, LNCS Vol. 2422, pages 495–502. Springer, 2002.
5. S. Autexier and T. Mossakowski. Integrating HOL-CASL into the development graph manager MAYA. In A. Armando, editor, *Frontiers of Combining Systems, 4th International Workshop, FroCoS 2002, Santa Margherita Ligure, Italy, Proceedings*, LNCS Vol. 2309, pages 2–17. Springer, 2002.
6. J. A. Bergstra, J. Heering, and P. Klint. The algebraic specification formalism ASF. In J. A. Bergstra, J. Heering, and P. Klint, editors, *Algebraic Specification*, ACM Press Frontier Series. Addison-Wesley, 1989.
7. M. Bidoit. Development of modular specifications by stepwise refinements using the PLUSS specification language. In C. Rattray and R. G. Clark, editors, *Unified Computation Laboratory: Modelling, Specifications, and Tools*, pages 171–192. Oxford Univ. Press, 1992.
8. M. Bidoit, C. Chevenier, C. Pellen, and J. Ryckbosch. An algebraic specification of the steam-boiler control system. In Abrial et al. [1], pages 79–108.
9. M. Bidoit, M.-C. Gaudel, and A. Mauboussin. How to make algebraic specifications more understandable? An experiment with the PLUSS specification language. *Science of Computer Programming*, 12(1):1–38, 1989.
10. M. Bidoit, H.-J. Kreowski, P. Lescanne, F. Orejas, and D. Sannella, editors. *Algebraic System Specification and Development*. LNCS Vol. 501. Springer, 1991.
11. M. Bidoit, D. Sannella, and A. Tarlecki. Architectural specifications in CASL. *Formal Aspects of Computing*, 13:252–273, 2002.

12. M. G. J. van den Brand, H. A. de Jong, P. Klint, and P. Olivier. Efficient annotated terms. *Software, Practice & Experience*, 30:259–291, 2000.
13. M. G. J. van den Brand, A. van Deursen, J. Heering, H. A. de Jong, M. de Jonge, T. Kuipers, P. Klint, L. Moonen, P. A. Olivier, J. Scheerder, J. J. Vinju, E. Visser, and J. Visser. The ASF+SDF Meta-Environment: A component-based language development environment. In R. Wilhelm, editor, *Compiler Construction, 10th International Conference, CC 2001, Genova, Italy, Proceedings*, LNCS Vol. 2027, pages 365–370. Springer, 2001.
14. M. Broy, C. Facchi, R. Grosu, R. Hettler, H. Hußmann, D. Nazareth, F. Regensburger, O. Slotosch, and K. Stølen. The requirement and design specification language SPECTRUM: An informal introduction (v 1.0). Technical Report TUM-I9311, TUM-I9312, Institut für Informatik, Technische Universität München, 1993.
15. R. M. Burstall and J. A. Goguen. The semantics of CLEAR, a specification language. In D. Bjørner, editor, *Abstract Software Specifications, 1979 Copenhagen Winter School, Proceedings*, LNCS Vol. 86, pages 292–332. Springer, 1980.
16. M. Cerioli, M. Gogolla, H. Kirchner, B. Krieg-Brückner, Z. Qian, and M. Wolf, editors. *Algebraic System Specification and Development: Survey and Annotated Bibliography*. BISS Monographs. Shaker, 2nd edition, 1997.
17. M. Cerioli and G. Reggio, editors. *Recent Trends in Algebraic Development Techniques, 15th International Workshop, WADT 2001, Joint with the CoFI WG Meeting, Genova, Italy, 2001, Selected Papers*, LNCS Vol. 2267. Springer, 2001.
18. F. Chen, G. Rosu, and R. P. Venkatesan. Rule-based analysis of dimensional safety. In R. Nieuwenhuis, editor, *Rewriting Techniques and Applications, 14th International Conference, RTA 2003, Valencia, Spain, Proceedings*, LNCS Vol. 2706, pages 197–207. Springer, 2003.
19. I. Claßen, H. Ehrig, and D. Wolz. *Algebraic Specification Techniques and Tools for Software Development*. AMAST Series in Computing Vol. 1. World Scientific, 1993.
20. CoFI (The Common Framework Initiative). *CASL Reference Manual*. LNCS, IFIP Series. Springer, 2004. To appear.
21. CoFI (The Common Framework Initiative) Tools Group. Home page. <http://www.cofi.info/Tools>.
22. A. van Deursen, J. Heering, and P. Klint, editors. *Language Prototyping: An Algebraic Specification Approach*. AMAST Series in Computing Vol. 5. World Scientific, 1996.
23. J. A. Goguen and R. M. Burstall. Institutions: Abstract model theory for specification and programming. *J. ACM*, 39(1):95–146, 1992.
24. J. A. Goguen, T. Winkler, J. Meseguer, K. Futatsugi, and J.-P. Jouannaud. Introducing OBJ3. In J. Goguen and G. Malcolm, editors, *Software Engineering with OBJ: Algebraic Specification in Action*. Kluwer, 1992.
25. J. V. Guttag and J. J. Horning. *Larch: Languages and Tools for Formal Specification*. Springer, 1993.
26. S. Kahrs, D. Sannella, and A. Tarlecki. The definition of Extended ML: A gentle introduction. *Theoretical Comput. Sci.*, 173:445–484, 1997.
27. J. Loeckx, H.-D. Ehrich, and M. Wolf. *Specification of Abstract Data Types*. Wiley/Teubner, 1996.
28. T. Mossakowski. CASL: From semantics to tools. In S. Graf and M. Schwartzbach, editors, *Tools and Algorithms for the Construction and Analysis of Sys-*



- tems, *6th International Conference, TACAS 2000, Berlin, Germany, Proceedings*, LNCS Vol. 1785, pages 93–108. Springer, 2000.
29. T. Mossakowski. Relating CASL with other specification languages: The institution level. *Theoretical Comput. Sci.*, 286:367–475, 2002.
  30. T. Mossakowski, Kolyang, and B. Krieg-Brückner. Static semantic analysis and theorem proving for CASL. In F. Parisi-Presicce, editor, *Recent Trends in Algebraic Development Techniques, 12th International Workshop, WADT'97, Tarquinia, Italy, 1997, Selected Papers*, LNCS Vol. 1376, pages 333–348. Springer, 1998.
  31. L. C. Paulson. *Isabelle: A Generic Theorem Prover*. LNCS Vol. 828. Springer, 1994.
  32. M. Roggenbach and L. Schröder. Towards trustworthy specifications I: Consistency checks. In Cerioli and Reggio [17], pages 305–327.
  33. D. Sannella. The Common Framework Initiative for algebraic specification and development of software: Recent progress. In Cerioli and Reggio [17], pages 328–343.
  34. D. Sannella and A. Tarlecki. *Foundations of Algebraic Specification and Formal Program Development*. To appear.
  35. D. Sannella and A. Tarlecki. Essential concepts of algebraic specification and program development. *Formal Aspects of Computing*, 9:229–269, 1997.
  36. M. Wirsing. Structured algebraic specifications: A kernel language. *Theoretical Comput. Sci.*, 42:123–249, 1986.
  37. M. Wirsing. Algebraic specification. In J. van Leeuwen, editor, *Handbook of Theoretical Computer Science*, volume B, chapter 13. Elsevier Science Publishers, Amsterdam; and MIT Press, 1990.

---

## List of Named Specifications

STRICT_ PARTIAL_ ORDER .....	24
TOTAL_ ORDER .....	25
TOTAL_ ORDER_ WITH_ MINMAX .....	26
VARIANT_ OF_ TOTAL_ ORDER_ WITH_ MINMAX .....	27
PARTIAL_ ORDER .....	27
PARTIAL_ ORDER_ 1 .....	28
IMPLIES_ DOES_ NOT_ HOLD .....	28
MONOID .....	29
GENERIC_ MONOID .....	29
NON_ GENERIC_ MONOID .....	30
GENERIC_ COMMUTATIVE_ MONOID .....	30
GENERIC_ COMMUTATIVE_ MONOID_ 1 .....	31
CONTAINER .....	32
MARKING_ CONTAINER .....	32
GENERATED_ CONTAINER .....	33
GENERATED_ CONTAINER_ MERGE .....	34
GENERATED_ SET .....	35
NATURAL .....	36
COLOR .....	37
INTEGER .....	37
NATURAL_ ORDER .....	38
NATURAL_ ARITHMETIC .....	38
INTEGER_ ARITHMETIC .....	39
INTEGER_ ARITHMETIC_ ORDER .....	40
LIST .....	40
SET .....	40
TRANSITIVE_ CLOSURE .....	41
NATURAL_ WITH_ BOUND .....	41
SET_ CHOOSE .....	42
SET_ GENERATED .....	42
SET_ UNION .....	43

SET_ UNION_ 1 .....	44
UNNATURAL .....	44
SET_ PARTIAL_ CHOOSE .....	47
SET_ PARTIAL_ CHOOSE_ 1 .....	50
SET_ PARTIAL_ CHOOSE_ 2 .....	50
NATURAL_ WITH_ BOUND_ AND_ ADDITION .....	51
SET_ PARTIAL_ CHOOSE_ 3 .....	51
NATURAL_ PARTIAL_ PRE .....	52
NATURAL_ PARTIAL_ SUBTRACTION_ 1 .....	52
NATURAL_ PARTIAL_ SUBTRACTION .....	52
LIST_ SELECTORS_ 1 .....	53
LIST_ SELECTORS_ 2 .....	53
LIST_ SELECTORS .....	54
NATURAL_ SUC_ PRE .....	54
PAIR_ 1 .....	54
PART_ CONTAINER .....	54
NATURAL_ PARTIAL_ SUBTRACTION_ 2 .....	55
GENERIC_ MONOID_ 1 .....	57
VEHICLE .....	58
MORE_ VEHICLE .....	59
SPEED_ REGULATION .....	59
NATURAL_ SUBSORTS .....	60
POSITIVE .....	61
POSITIVE_ ARITHMETIC .....	61
POSITIVE_ PRE .....	62
NATURAL_ POSITIVE_ ARITHMETIC .....	62
INTEGER_ ARITHMETIC_ 1 .....	64
SET_ ERROR_ CHOOSE .....	65
SET_ ERROR_ CHOOSE_ 1 .....	65
LIST_ SET .....	67
LIST_ CHOOSE .....	68
SET_ TO_ LIST .....	68
STACK .....	69
LIST_ SET_ 1 .....	71
NATURAL_ PARTIAL_ SUBTRACTION_ 3 .....	71
NATURAL_ PARTIAL_ SUBTRACTION_ 4 .....	71
PARTIAL_ ORDER_ 2 .....	72
LIST_ ORDER .....	73
LIST_ ORDER_ SORTED .....	73
WRONG_ LIST_ ORDER_ SORTED .....	74
LIST_ ORDER_ SORTED_ 2 .....	74
LIST_ ORDER_ SORTED_ 3 .....	75
LIST_ ORDER_ NAT .....	78
NAT_ WORD .....	79
NAT_ WORD_ 1 .....	79

THIS_ IS_ WRONG .....	80
LIST_ ORDER_ POSITIVE .....	81
NAT_ WORD_ 2 .....	81
PAIR .....	82
HOMOGENEOUS_ PAIR_ 1 .....	82
HOMOGENEOUS_ PAIR .....	82
TABLE .....	82
PAIR_ NATURAL_ COLOR .....	82
PAIR_ NATURAL_ COLOR_ 1 .....	83
PAIR_ NATURAL_ COLOR_ 2 .....	83
PAIR_ POS .....	83
PAIR_ POS_ 1 .....	83
MY_ TABLE .....	83
SET_ OF_ LIST .....	84
MISTAKE .....	84
SET_ AND_ LIST .....	84
THIS_ IS_ STILL_ WRONG .....	84
LIST_ REV .....	85
LIST_ REV_ NAT .....	85
TWO_ LISTS .....	85
TWO_ LISTS_ 1 .....	86
MONOID_ C .....	86
MONOID_ OF_ MONOID .....	86
LIST_ REV_ ORDER .....	87
LIST_ REV_ WITH_ TWO_ ORDERS .....	87
LIST_ WEIGHTED_ ELEM .....	88
LIST_ WEIGHTED_ PAIR_ NATURAL_ COLOR .....	89
LIST_ WEIGHTED_ INSTANTIATED .....	89
LIST_ LENGTH .....	89
LIST_ LENGTH_ NATURAL .....	90
INTEGER_ AS_ TOTAL_ ORDER .....	90
INTEGER_ AS_ REVERSE_ TOTAL_ ORDER .....	91
LIST_ REV_ WITH_ TWO_ ORDERS_ 1 .....	91
LIST_ AS_ MONOID .....	91
ELEM .....	94
CONT .....	94
CONT_ DIFF .....	95
REQ .....	95
FLAT_ REQ .....	95
SYSTEM .....	96
SYSTEM_ 1 .....	97
CONT_ DIFF_ 1 .....	99
INCONSISTENT .....	99
SYSTEM_ G .....	100
SYSTEM_ A .....	101

CONT_ COMP .....	102
DIFF_ COMP .....	102
SYSTEM_ G1 .....	102
DIFF_ COMP_ 1 .....	102
OTHER_ SYSTEM .....	103
OTHER_ SYSTEM_ 1 .....	104
SET_ COMP .....	105
CONT2SET .....	105
ARCH_ CONT2SET_ NAT .....	105
ARCH_ CONT2SET .....	106
ARCH_ CONT2SET_ USED .....	106
ARCH_ CONT2SET_ NAT_ 1 .....	107
WRONG_ ARCH_ SPEC .....	108
BADLY_ STRUCTURED_ ARCH_ SPEC .....	108
WELL_ STRUCTURED_ ARCH_ SPEC .....	109
ANOTHER_ WELL_ STRUCTURED_ ARCH_ SPEC .....	109
NATURAL_ ORDER_ 2 .....	134
v1 .....	134
v2 .....	137
NUMBERS .....	144
BASIC/NUMBERS .....	145
NAT .....	145
INT .....	146
RAT .....	147
STRUCTURED DATATYPES .....	147
BASIC/STRUCTURED DATATYPES .....	147
GENERATESET .....	148
SET .....	148
GENERATEMAP .....	148
MAP .....	149
FINITE .....	149
TOTALMAP .....	149
GENERATEBAG .....	150
BAG .....	150
GENERATELIST .....	150
LIST .....	150
ARRAY .....	151
GENERATEBINTREE .....	152
BINTREE .....	152
GENERATEBINTREE2 .....	152
BINTREE2 .....	152
GENERATEKTREE .....	153
KTREE .....	153
GENERATENTREE .....	154
NTREE .....	154

VALUE .....	158
BASICS .....	159
MESSAGES_ SENT .....	159
MESSAGES_ RECEIVED .....	159
SBCS_ CONSTANTS .....	160
PRELIMINARY .....	160
SBCS .....	161
SBCS_ STATE_ 1 .....	163
MODE_ EVOLUTION .....	165
SBCS_ STATE_ 2 .....	169
STATUS_ EVOLUTION .....	169
MESSAGE_ TRANSMISSION_ SYSTEM_ FAILURE .....	171
SBCS_ STATE_ 3 .....	172
PUMP_ FAILURE .....	173
SBCS_ STATE_ 4 .....	173
PUMP_ CONTROLLER_ FAILURE .....	174
SBCS_ STATE_ 5 .....	174
STEAM_ FAILURE .....	175
LEVEL_ FAILURE .....	175
FAILURE_ DETECTION .....	175
STEAM_ AND_ LEVEL_ PREDICTION .....	178
PUMP_ STATE_ PREDICTION .....	180
PUMP_ CONTROLLER_ STATE_ PREDICTION .....	181
PU_ PREDICTION .....	181
SBCS_ ANALYSIS .....	182
SBCS_ STATE .....	183
STEAM_ BOILER_ CONTROL_ SYSTEM .....	183
ARCH_ SBCS .....	186
ARCH_ PRELIMINARY .....	187
SBCS_ STATE_ IMPL .....	187
UNIT_ SBCS_ STATE .....	188
ARCH_ ANALYSIS .....	188
ARCH_ FAILURE_ DETECTION .....	189
ARCH_ PREDICTION .....	189

---

## Index of Library and Specification Names

ANOTHER\_WELL\_STRUCTURED\_ARCH\_ SPEC 109  
ARCH\_ANALYSIS 188  
ARCH\_CONT2SET 106  
ARCH\_CONT2SET\_NAT 105  
ARCH\_CONT2SET\_NAT\_1 107  
ARCH\_CONT2SET\_USED 106  
ARCH\_FAILURE\_DETECTION 189  
ARCH\_PREDICTION 189  
ARCH\_PRELIMINARY 187  
ARCH\_SBCS 186  
ARRAY 151  
  
BADLY\_STRUCTURED\_ARCH\_SPEC 108  
BAG 150  
BASIC/NUMBERS 145  
BASIC/STRUCTURED DATATYPES 147  
BASICS 159  
BINTREE 152  
BINTREE2 152  
  
COLOR 37  
CONT 94  
CONT2SET 105  
CONT\_COMP 102  
CONT\_DIFF 95  
CONT\_DIFF\_1 99  
CONTAINER 32  
  
DIFF\_COMP 102  
DIFF\_COMP\_1 102  
  
ELEM 94  
  
FAILURE\_DETECTION 175  
FINITE 149  
FLAT\_REQ 95  
  
GENERATEBAG 150  
GENERATEBINTREE 152  
GENERATEBINTREE2 152  
GENERATED\_CONTAINER 33  
GENERATED\_CONTAINER\_MERGE 34  
GENERATED\_SET 35  
GENERATEKTREE 153  
GENERATELIST 150  
GENERATEMAP 148  
GENERATENTREE 154  
GENERATESET 148  
GENERIC\_COMMUTATIVE\_MONOID 30  
GENERIC\_COMMUTATIVE\_MONOID\_1 31  
GENERIC\_MONOID 29  
GENERIC\_MONOID\_1 57  
  
HOMOGENEOUS\_PAIR 82  
HOMOGENEOUS\_PAIR\_1 82  
  
IMPLIES\_DOES\_NOT\_HOLD 28  
INCONSISTENT 99  
INT 146  
INTEGER 37  
INTEGER\_ARITHMETIC 39  
INTEGER\_ARITHMETIC\_1 64  
INTEGER\_ARITHMETIC\_ORDER 40  
INTEGER\_AS\_REVERSE\_TOTAL\_ORDER 91  
INTEGER\_AS\_TOTAL\_ORDER 90

- KTREE 153
- LEVEL\_FAILURE 175
- LIST 40, 150
- LIST\_AS\_MONOID 91
- LIST\_CHOOSE 68
- LIST\_LENGTH 89
- LIST\_LENGTH\_NATURAL 90
- LIST\_ORDER 73
- LIST\_ORDER\_NAT 78
- LIST\_ORDER\_POSITIVE 81
- LIST\_ORDER\_SORTED 73
- LIST\_ORDER\_SORTED\_2 74
- LIST\_ORDER\_SORTED\_3 75
- LIST\_REV 85
- LIST\_REV\_NAT 85
- LIST\_REV\_ORDER 87
- LIST\_REV\_WITH\_TWO\_ORDERS 87
- LIST\_REV\_WITH\_TWO\_ORDERS\_1 91
- LIST\_SELECTORS 54
- LIST\_SELECTORS\_1 53
- LIST\_SELECTORS\_2 53
- LIST\_SET 67
- LIST\_SET\_1 71
- LIST\_WEIGHTED\_ELEM 88
- LIST\_WEIGHTED\_INSTANTIATED 89
- LIST\_WEIGHTED\_PAIR\_NATURAL\_
  - COLOR 89
- MAP 149
- MARKING\_CONTAINER 32
- MESSAGE\_TRANSMISSION\_SYSTEM\_
  - FAILURE 171
- MESSAGES\_RECEIVED 159
- MESSAGES\_SENT 159
- MISTAKE 84
- MODE\_EVOLUTION 165
- MONOID 29
- MONOID\_C 86
- MONOID\_OF\_MONOID 86
- MORE\_VEHICLE 59
- MY\_TABLE 83
- NAT 145
- NAT\_WORD 79
- NAT\_WORD\_1 79
- NAT\_WORD\_2 81
- NATURAL 36
- NATURAL\_ARITHMETIC 38
- NATURAL\_ORDER 38
- NATURAL\_ORDER\_2 134
- NATURAL\_PARTIAL\_PRE 52
- NATURAL\_PARTIAL\_SUBTRACTION 52
- NATURAL\_PARTIAL\_SUBTRACTION\_1
  - 52
- NATURAL\_PARTIAL\_SUBTRACTION\_2
  - 55
- NATURAL\_PARTIAL\_SUBTRACTION\_3
  - 71
- NATURAL\_PARTIAL\_SUBTRACTION\_4
  - 71
- NATURAL\_POSITIVE\_ARITHMETIC 62
- NATURAL\_SUBSORTS 60
- NATURAL\_SUC\_PRE 54
- NATURAL\_WITH\_BOUND 41
- NATURAL\_WITH\_BOUND\_AND\_
  - ADDITION 51
- NON\_GENERIC\_MONOID 30
- NTREE 154
- NUMBERS 144
- OTHER\_SYSTEM 103
- OTHER\_SYSTEM\_1 104
- PAIR 82
- PAIR\_1 54
- PAIR\_NATURAL\_COLOR 82
- PAIR\_NATURAL\_COLOR\_1 83
- PAIR\_NATURAL\_COLOR\_2 83
- PAIR\_POS 83
- PAIR\_POS\_1 83
- PART\_CONTAINER 54
- PARTIAL\_ORDER 27
- PARTIAL\_ORDER\_1 28
- PARTIAL\_ORDER\_2 72
- POSITIVE 61
- POSITIVE\_ARITHMETIC 61
- POSITIVE\_PRE 62
- PRELIMINARY 160
- PU\_PREDICTION 181
- PUMP\_CONTROLLER\_FAILURE 174
- PUMP\_CONTROLLER\_STATE\_
  - PREDICTION 181
- PUMP\_FAILURE 173
- PUMP\_STATE\_PREDICTION 180
- RAT 147
- REQ 95



SBCS	161	STEAM_FAILURE	175
SBCS_ANALYSIS	182	STRICT_PARTIAL_ORDER	24
SBCS_CONSTANTS	160	STRUCTURED DATATYPES	147
SBCS_STATE	183	SYSTEM	96
SBCS_STATE_1	163	SYSTEM_1	97
SBCS_STATE_2	169	SYSTEM_A	101
SBCS_STATE_3	172	SYSTEM_G	100
SBCS_STATE_4	173	SYSTEM_G1	102
SBCS_STATE_5	174		
SBCS_STATE_IMPL	187	TABLE	82
SET	40, 148	THIS_IS_STILL_WRONG	84
SET_AND_LIST	84	THIS_IS_WRONG	80
SET_CHOOSE	42	TOTAL_ORDER	25
SET_COMP	105	TOTAL_ORDER_WITH_MINMAX	26
SET_ERROR_CHOOSE	65	TOTALMAP	149
SET_ERROR_CHOOSE_1	65	TRANSITIVE_CLOSURE	41
SET_GENERATED	42	TWO_LISTS	85
SET_OF_LIST	84	TWO_LISTS_1	86
SET_PARTIAL_CHOOSE	47		
SET_PARTIAL_CHOOSE_1	50	UNIT_SBCS_STATE	188
SET_PARTIAL_CHOOSE_2	50	UNNATURAL	44
SET_PARTIAL_CHOOSE_3	51		
SET_TO_LIST	68	v1	134
SET_UNION	43	v2	137
SET_UNION_1	44	VALUE	158
SPEED_REGULATION	59	VARIANT_OF_TOTAL_ORDER_WITH_	
STACK	69	MINMAX	27
STATUS_EVOLUTION	169	VEHICLE	58
STEAM_AND_LEVEL_PREDICTION	178		
STEAM_BOILER_CONTROL_SYSTEM	183	WELL_STRUCTURED_ARCH_SPEC	109
		WRONG_ARCH_SPEC	108
		WRONG_LIST_ORDER_SORTED	74

---

## Concept Index

- abbreviation 23
- abstract syntax 126
- algebraic signature 12
- analysis, static 127, 133
- annotation 115
  - associativity 115
  - author 116
  - date 116, 121
  - definitional extension 43
  - display 27, 115, 119
  - implies 28
  - label 25
  - literal syntax 118
  - parsing 115, 119
  - precedence 115
  - relative precedence 115
- architectural specification 93, 186
- argument
  - specification 19, 30
  - fitting 77
- ASF+SDF 139
- assertion
  - definedness 50
  - subsort membership 59
- associativity
  - annotation 115
  - attribute 29, 116
- ATerms 137, 139
- attribute 29
  - associativity 29, 116
  - commutativity 29
  - idempotence 29
  - unit 29
- author annotation 116
- auxiliary
  - operation 71
  - predicate 71
  - symbol 73
- axiom 11, 12, 14, 25
- basic
  - datatypes, libraries of 117, 119, 143
  - specification 11, 194
- body 30
  - specification 18
- carrier set 12
- casting 64
- CATS 138
- class of models 12
- closed
  - specification 114
  - system 106
  - world assumption 38
- comment 25
- commutativity attribute 29
- compatibility 105, 106
  - overloading with
    - embedding 62
    - renaming 70
    - subsorting 14
- completeness 128
- component 93
  - declaration 96
  - generic 100
    - application of 101
  - specification 96

- named 102
- composition 19
  - generic specification 84
- compound
  - sort 12, 85
  - symbol 85
    - operation 87
    - predicate 87
- concrete syntax 126
- consequence relation 128
- conservative extension 19, 98
- consistency checker 141
- consistent 12
- constant
  - operation 13
  - overloaded 31
  - symbol 29
- constraint 11, 12
  - sort generation 15, 33
- constructor 32
  - partial 55
- current version 121
- datatype
  - declaration 32, 59
  - enumerated 37
  - free 37
  - generated 33
  - structured 147
- date annotation 116, 121
- decimal notation 118
- declaration 11
  - before use *see* linear visibility
  - component 96
  - datatype 32, 59
  - function
    - partial 47
    - total 26
  - subsort 57, 59
  - symbol 11
  - unit 96
  - variable 25
    - global 27
- decomposition 19
- deduction rule 128
- definedness assertion 50
- definition
  - domain of 13, 50
  - operation 26

- predicate 28
  - style 44
  - subsort 60
  - unit 109
- definitional extension 43
- development graph 134
- disambiguation 31
- display
  - annotation 27, 115, 119
  - format 27
- distributed library 111, 116
- domain
  - of definition 13, 50
  - semantic 127
- downloaded specification 112, 118
- editor, syntax-directed 139
- embedding 13, 64
  - compatibility with overloading 62
- end-of-line comment 25
- enumerated datatype 37
- equation
  - existential 15, 55
  - strong 15, 48, 55
- error supersort 65
- evolution of specification 112
- existential
  - equation 15, 55
  - quantification 14, 25
- expansion 16
- explicit fitting 81
- exported symbol 16, 25
- expression, unit 106
- extension 25
  - conservative 19, 98
  - definitional 43
  - specification 18, 68
- fitting
  - argument specification 77
  - explicit 81
  - implicit 80
  - morphism 19
  - symbol map 81
- fixed part 88
- formal
  - specification 156
  - verification 128, 138
- formalization process 156

- formula 25
- free
  - datatype 37
  - model 18
  - specification 18, 36
- freeness constraint 18, 37
- function
  - partial 13, 47
  - persistent 19, 100
  - total 13, 23
- generated
  - datatype 33
  - specification 33
- generic
  - component 100
  - application 101
  - specification 15, 18, 30, 77
  - composition 84
  - view 91
- global variable declaration 27
- HETS 132
- hiding
  - specification 17, 71
  - with revealing 72
- HOL-CASL 138
- homomorphism 16
- idempotence attribute 29
- identity renaming 71
- ill-formed instantiation 80, 90
- implementation 93, 140
- implicit fitting 80
- implies annotation 28
- imports 19, 88
- inconsistent 12
- independence, institution 127, 137
- induction 34
  - proof 139
  - scheme 35
- inductive definition 38
  - style 44
- informal requirements 156
- inherited 58
- initial model 16, 38
- input format 27
- instantiation 15, 19, 30, 77, 78
  - ill-formed 80, 90
- institution 12, 127, 128
  - independence 127, 137
- interpretation 11, 12
- isomorphism 16
- junk 15
- label annotation 25
- libraries of basic datatypes 117, 119, 143
- library 111
  - current version 121
  - distributed 111, 116
  - item 114
  - local 111, 112
  - registered 117
  - self-contained 111, 112
  - under development 120
  - validation 117
  - version 112, 121
- linear visibility 26, 96, 113
- literal syntax annotation 118
- local
  - library 111, 112
  - specification 73
  - unit definition 109
- logic 15, *see* institution
- loose specification 24
- many-sorted signature 12
- map, symbol 70
- mathematical theory 126
- MAYA 138
- meaning 12
- mixfix notation 24
- model
  - class 12
  - free 18
  - of specification 12
  - semantics 127
- morphism
  - fitting 19
  - signature 16
  - specification 16
- multi-line comment 25
- mutual recursion 36, 113
- name, path 117
- named
  - specification 15, 24, 75

- view 19
- natural semantics 127
- negation as failure 38
- 'no junk, no confusion' principle 37
- 'no junk' principle 33
- non-isomorphic models 25
- non-linear visibility 36, 63
- notation, decimal 118
- number 144
- observer 43
  - definition style 44
- open system 106
- operation
  - auxiliary 71
  - boolean-valued v. predicate 13, 18
  - constant 13
  - definition 26
  - symbol 13
    - compound 87
- overloaded
  - constant 31
  - symbol 31
    - renaming 70
- overloading 14, 61
  - compatibility
    - with embedding 62
    - with renaming 70
- parameter 30
  - specification 18, 77, 78
- parsing 133, 137, 139
  - annotation 115, 119
- partial
  - constructor 55
  - function 13, 47
    - declaration 47
  - selector 54
- path name 117
- persistent function 19, 100
- place-holder 24
- precedence annotation 115
- predicate
  - auxiliary 71
  - definition 28
  - symbol 13, 24
    - compound 87
  - v. boolean-valued operation 13, 18
- principle
  - 'no junk, no confusion' 37
  - 'no junk' 33
  - 'same name, same thing' 17, 31, 113
- process of formalization 156
- profile 13
- programming language 140
- proof
  - calculus 128
  - induction 139
  - obligation 28, 134
- propagate undefinedness 48
- quantification
  - existential 14, 25
  - unique-existential 14
  - universal 14, 25
- recursion, mutual 36, 113
- reduct 16
- reference to name 15
- refinement 156
- registered library 117
- relative precedence annotation 115
- renaming
  - compatibility with overloading 70
  - identity 71
  - overloaded symbol 70
  - specification 69
- requirements
  - informal 156
  - specification 156
- result sort 13
- reuse of specification 111
- revealing with hiding 72
- rewriting 140
- 'same name, same thing' principle 17, 31, 113
- satisfaction relation 12
- scheme, induction 35
- selector 53
  - partial 54
  - total 54
- self-contained
  - library 111, 112
  - specification *see* closed
- semantic domain 127
- semantics 12, 126
  - model 127
  - natural 127

- static 127
- set
  - carrier 12
  - inclusion 13
- signature 12
  - algebraic 12
  - many-sorted 12
  - morphism 16
  - subsorted 12
- software modules 94
- sort 12, 24
  - compound 12, 85
  - generation constraint 15, 33
  - result 13
  - union 59
- soundness 128
- specification
  - architectural 93, 186
  - argument 19, 30
    - fitting 77
  - basic 11, 194
  - body 18
  - closed 114
  - downloaded 112, 118
  - evolution 112
  - extension 18, 68
  - fitting argument 77
  - formal 156
  - free 18, 36
  - generated 33
  - generic 15, 18, 30, 77
  - hiding 17, 71
  - instantiation 77, 78
  - local 73
  - loose 24
  - model of 12
  - morphism 16
  - named 15, 24, 75
  - of component 96
    - named 102
  - parameter 18, 77, 78
  - reference 15
  - renaming 69
  - requirements 156
  - reuse 111
  - structured 67
  - subsorted 57
  - translation 17
  - union 17, 68
- unit 98
- static
  - analysis 127, 133
  - semantics 127
- strong equation 15, 48, 55
- structured
  - datatype 147
  - specification 67
- subsort 13, 57
  - declaration 57, 59
  - definition 60
  - membership assertion 59
- subsorted
  - signature 12
  - specification 57
- supersort 13, 57
- symbol 12
  - auxiliary 73
  - compound 85
  - constant 29
  - declaration 11
  - exported 16, 25
  - map 70
    - fitting 81
  - operation 13
  - overloaded 31
  - predicate 13, 24
- syntax 126
  - directed editor 139
- system 106
- term, unit 100
- theorem prover 138
- total
  - function 13, 23
    - declaration 26
  - selector 54
- translation 17
- two-valued logic 48
- undefined value 48
- union
  - sort 59
  - specification 17, 68
- unique-existential quantification 14
- unit
  - attribute 29
  - declaration 96
  - definition 109

- expression 106
- local definition 109
- specification 98
- term 100
- universal quantification 14, 25
- validation 156
  - of library 117
- variable 14
  - declaration 25
- version
  - control 112, 120
  - current 121
  - library 121
  - number 121
  - of library 112
- view 90
  - generic 91
  - named 19
- visibility
  - component name 97
  - linear 26, 96, 113